

Mark3 Realtime Kernel

Generated by Doxygen 1.8.8

Fri Jul 7 2017 01:40:39

Contents

1	The Mark3 Realtime Kernel	1
2	License	3
2.1	License	3
3	Configuring The Mark3 Kernel	5
3.1	Overview	5
3.2	Timer Options	5
3.3	Blocking Objects	6
3.4	Inter-process/thread Communication	7
3.5	Debug Features	7
3.6	Enhancements, Security, Miscellaneous	7
4	Building Mark3	9
4.1	Source Layout	9
4.2	Toolchain Integration	9
4.3	Installing Dependencies	9
4.4	Building Mark3 Kernel and Libraries	10
4.5	Exporting the kernel source	10
5	Getting Started With The Mark3 API	13
5.1	Kernel Setup	13
5.2	Threads	14
5.2.1	Thread Setup	14
5.2.2	Entry Functions	15
5.3	Timers	15
5.4	Semaphores	16
5.5	Mutexes	17
5.6	Event Flags	17
5.7	Messages	18
5.7.1	Message Objects	18
5.7.2	Global Message Pool	19
5.7.3	Message Queues	19

5.7.4	Messaging Example	19
5.8	Mailboxes	20
5.8.1	Mailbox Example	20
5.9	Notification Objects	21
5.9.1	Notification Example	21
5.10	Sleep	21
5.11	Round-Robin Quantum	22
6	Why Mark3?	23
7	When should you use an RTOS?	25
7.1	The reality of system code	25
7.2	Superloops, and their limitations	26
7.2.1	Intro to Superloops	26
7.2.2	The simplest loop	26
7.2.3	Interrupt-Driven Super-loop	27
7.2.4	Cooperative multi-tasking	28
7.2.5	Hybrid cooperative/preemptive multi-tasking	29
7.3	Problems with superloops	30
7.3.1	Hidden Costs	30
7.3.2	Tightly-coupled code	31
7.3.3	No blocking Calls	31
7.3.4	Difficult to guarantee responsiveness	31
7.3.5	Limited preemption capability	31
8	Can you afford an RTOS?	33
8.1	Intro	33
8.2	Application description	34
8.3	Runtime Overhead	35
8.4	Analysis	36
9	Mark3 Design Goals	37
9.1	Overview	37
9.1.1	Services Provided by an RTOS Kernel	37
9.1.2	Guiding Principles of Mark3	37
9.1.3	Be feature competitive	37
9.1.4	Be highly configuration	37
9.1.5	No external dependencies, no new language features	38
9.1.6	Target the most popular hobbyist platforms available	38
9.1.7	Maximize determinism – but be pragmatic	38
9.1.8	Apply engineering principles – and that means discipline, measurement and verification	38

9.1.9 Use Virtualization For Verification	38
10 Mark3 Kernel Architecture	41
10.1 Overview	41
10.2 Threads and Scheduling	43
10.2.1 A Bit About Threads	43
10.2.2 Thread States and ThreadLists	44
10.2.3 Blocking and Unblocking	44
10.2.4 Blocking Objects	45
10.3 Inside the Mark3 Scheduler	46
10.3.1 Considerations for Round-Robin Scheduling	47
10.3.2 Context Switching	47
10.3.3 Putting It All Together	48
10.4 Timers	48
10.4.1 Tick-based Timers	50
10.4.2 Tickless Timers	50
10.4.3 Timer Processing Algorithm	50
10.5 Synchronization and IPC	51
10.6 Blocking Objects	51
10.6.1 Semaphores	52
10.6.2 Mutex	52
10.6.3 Event Flags	52
10.6.4 Notification Objects	52
10.7 Messages and Global Message Queue	53
10.7.1 Messages	53
10.7.2 Message Objects	53
10.7.3 Global Message Pool	53
10.7.4 Message Queues	54
10.7.5 Mailboxes	54
10.7.6 Atomic Operations	54
10.7.7 Drivers	55
10.8 Kernel Proper and Porting	57
11 Mark3C - C-language API bindings for the Mark3 Kernel.	65
11.1 API Conventions	65
11.2 Allocating Objects	66
11.3 Drivers in Mark3C	66
12 Release Notes	69
12.1 R6 Release	69
12.2 R5 Release	69

12.3 R4 Release	70
12.4 R3 Release	70
12.5 R2	71
12.6 R1 - 2nd Release Candidate	71
12.7 R1 - 1st Release Candidate	71
13 Profiling Results	73
13.1 Date Performed	73
13.2 Compiler Information	73
13.3 Profiling Results	73
14 Code Size Profiling	75
14.1 Information	75
14.2 Compiler Version	75
14.3 Profiling Results	75
15 Hierarchical Index	77
15.1 Class Hierarchy	77
16 Class Index	79
16.1 Class List	79
17 File Index	81
17.1 File List	81
18 Class Documentation	85
18.1 BlockingObject Class Reference	85
18.1.1 Detailed Description	86
18.1.2 Member Function Documentation	86
18.1.2.1 Block	86
18.1.2.2 BlockPriority	86
18.1.2.3 IsInitialized	86
18.1.2.4 Unblock	86
18.2 CircularLinkedList Class Reference	87
18.2.1 Detailed Description	87
18.2.2 Member Function Documentation	87
18.2.2.1 Add	87
18.2.2.2 InsertNodeBefore	88
18.2.2.3 PivotBackward	88
18.2.2.4 PivotForward	88
18.2.2.5 Remove	88
18.3 DoubleLinkedList Class Reference	88

18.3.1 Detailed Description	89
18.3.2 Constructor & Destructor Documentation	89
18.3.2.1 DoubleLinkedList	89
18.3.3 Member Function Documentation	89
18.3.3.1 Add	89
18.3.3.2 Remove	89
18.4 Driver Class Reference	90
18.4.1 Detailed Description	91
18.4.2 Member Function Documentation	91
18.4.2.1 Close	91
18.4.2.2 Control	91
18.4.2.3 GetPath	91
18.4.2.4 Init	91
18.4.2.5 Open	92
18.4.2.6 Read	92
18.4.2.7 SetName	92
18.4.2.8 Write	92
18.5 DriverList Class Reference	93
18.5.1 Detailed Description	93
18.5.2 Member Function Documentation	93
18.5.2.1 Add	93
18.5.2.2 FindByPath	93
18.5.2.3 Init	94
18.5.2.4 Remove	94
18.6 EventFlag Class Reference	94
18.6.1 Detailed Description	95
18.6.2 Member Function Documentation	95
18.6.2.1 Clear	95
18.6.2.2 GetMask	95
18.6.2.3 Set	95
18.6.2.4 Wait	96
18.6.2.5 Wait	96
18.6.2.6 Wait_i	96
18.6.2.7 WakeMe	97
18.7 FakeThread_t Struct Reference	97
18.7.1 Detailed Description	97
18.8 GlobalMessagePool Class Reference	98
18.8.1 Detailed Description	98
18.8.2 Member Function Documentation	98
18.8.2.1 GetHead	98

18.8.2.2	GetPool	99
18.8.2.3	Init	99
18.8.2.4	Pop	99
18.8.2.5	Push	99
18.9	Kernel Class Reference	99
18.9.1	Detailed Description	101
18.9.2	Member Function Documentation	101
18.9.2.1	GetIdleThread	101
18.9.2.2	GetThreadContextSwitchCallout	101
18.9.2.3	GetThreadCreateCallout	101
18.9.2.4	GetThreadExitCallout	101
18.9.2.5	Init	102
18.9.2.6	IsPanic	102
18.9.2.7	IsStarted	102
18.9.2.8	Panic	102
18.9.2.9	SetIdleFunc	102
18.9.2.10	SetPanic	103
18.9.2.11	SetThreadContextSwitchCallout	104
18.9.2.12	SetThreadCreateCallout	104
18.9.2.13	SetThreadExitCallout	104
18.9.2.14	Start	105
18.10	KernelAware Class Reference	105
18.10.1	Detailed Description	106
18.10.2	Member Function Documentation	106
18.10.2.1	ExitSimulator	106
18.10.2.2	IsSimulatorAware	106
18.10.2.3	Print	106
18.10.2.4	ProfileInit	107
18.10.2.5	ProfileReport	107
18.10.2.6	ProfileStart	107
18.10.2.7	ProfileStop	107
18.10.2.8	Trace	107
18.10.2.9	Trace	107
18.10.2.10	Trace	108
18.10.2.11	Trace_i	108
18.11	KernelSWI Class Reference	108
18.11.1	Detailed Description	109
18.11.2	Member Function Documentation	109
18.11.2.1	Clear	109
18.11.2.2	Config	109

18.11.2.3 DI	109
18.11.2.4 RI	109
18.11.2.5 Start	109
18.11.2.6 Stop	110
18.11.2.7 Trigger	110
18.12KernelTimer Class Reference	110
18.12.1 Detailed Description	111
18.12.2 Member Function Documentation	111
18.12.2.1 ClearExpiry	111
18.12.2.2 Config	111
18.12.2.3 DI	111
18.12.2.4 EI	111
18.12.2.5 GetOvertime	111
18.12.2.6 Read	111
18.12.2.7 RI	112
18.12.2.8 SetExpiry	112
18.12.2.9 Start	112
18.12.2.10Stop	112
18.12.2.11SubtractExpiry	112
18.12.2.12TimeToExpiry	113
18.13LinkedList Class Reference	113
18.13.1 Detailed Description	114
18.13.2 Member Function Documentation	114
18.13.2.1 GetHead	114
18.13.2.2 GetTail	114
18.13.2.3 Init	114
18.14LinkedListNode Class Reference	114
18.14.1 Detailed Description	115
18.14.2 Member Function Documentation	115
18.14.2.1 ClearNode	115
18.14.2.2 GetNext	116
18.14.2.3 GetPrev	116
18.15Mailbox Class Reference	116
18.15.1 Detailed Description	117
18.15.2 Member Function Documentation	118
18.15.2.1 CopyData	118
18.15.2.2 GetHeadPointer	118
18.15.2.3 GetTailPointer	118
18.15.2.4 Init	118
18.15.2.5 MoveHeadBackward	119

18.15.2.6 MoveHeadForward	119
18.15.2.7 MoveTailBackward	119
18.15.2.8 MoveTailForward	119
18.15.2.9 Receive	119
18.15.2.10Receive	119
18.15.2.11Receive_i	120
18.15.2.12ReceiveTail	120
18.15.2.13ReceiveTail	120
18.15.2.14Send	120
18.15.2.15Send	122
18.15.2.16Send_i	122
18.15.2.17SendTail	122
18.15.2.18SendTail	123
18.15.3 Member Data Documentation	123
18.15.3.1 m_clSendSem	123
18.16Message Class Reference	123
18.16.1 Detailed Description	124
18.16.2 Member Function Documentation	124
18.16.2.1 GetCode	124
18.16.2.2 GetData	124
18.16.2.3 Init	125
18.16.2.4 SetCode	125
18.16.2.5 SetData	125
18.17MessagePool Class Reference	125
18.17.1 Detailed Description	126
18.17.2 Member Function Documentation	126
18.17.2.1 GetHead	126
18.17.2.2 Init	126
18.17.2.3 Pop	126
18.17.2.4 Push	126
18.18MessageQueue Class Reference	127
18.18.1 Detailed Description	127
18.18.2 Member Function Documentation	128
18.18.2.1 GetCount	128
18.18.2.2 Init	128
18.18.2.3 Receive	128
18.18.2.4 Receive	128
18.18.2.5 Receive_i	128
18.18.2.6 Send	129
18.19Mutex Class Reference	129

18.19.1 Detailed Description	130
18.19.2 Member Function Documentation	130
18.19.2.1 Claim	130
18.19.2.2 Claim	130
18.19.2.3 Claim_i	131
18.19.2.4 Init	131
18.19.2.5 Release	131
18.19.2.6 WakeMe	132
18.19.2.7 WakeNext	133
18.20Notify Class Reference	133
18.20.1 Detailed Description	133
18.20.2 Member Function Documentation	134
18.20.2.1 Init	134
18.20.2.2 Signal	134
18.20.2.3 Wait	134
18.20.2.4 Wait	134
18.20.2.5 WakeMe	134
18.21PriorityMap Class Reference	135
18.21.1 Detailed Description	135
18.21.2 Constructor & Destructor Documentation	135
18.21.2.1 PriorityMap	135
18.21.3 Member Function Documentation	135
18.21.3.1 Clear	135
18.21.3.2 HighestPriority	136
18.21.3.3 Set	136
18.22ProfileTimer Class Reference	136
18.22.1 Detailed Description	137
18.22.2 Member Function Documentation	137
18.22.2.1 ComputeCurrentTicks	137
18.22.2.2 GetAverage	137
18.22.2.3 GetCurrent	138
18.22.2.4 Init	138
18.22.2.5 Start	138
18.22.2.6 Stop	138
18.23Quantum Class Reference	138
18.23.1 Detailed Description	139
18.23.2 Member Function Documentation	139
18.23.2.1 AddThread	139
18.23.2.2 ClearInTimer	139
18.23.2.3 RemoveThread	139

18.23.2.4 SetInTimer	139
18.23.2.5 SetTimer	139
18.23.2.6 UpdateTimer	139
18.24 Scheduler Class Reference	140
18.24.1 Detailed Description	140
18.24.2 Member Function Documentation	141
18.24.2.1 Add	141
18.24.2.2 GetCurrentThread	141
18.24.2.3 GetNextThread	141
18.24.2.4 GetStopList	141
18.24.2.5 GetThreadList	141
18.24.2.6 Init	142
18.24.2.7 IsEnabled	142
18.24.2.8 QueueScheduler	142
18.24.2.9 Remove	142
18.24.2.10 Schedule	142
18.24.2.11 SetScheduler	143
18.25 Semaphore Class Reference	143
18.25.1 Detailed Description	144
18.25.2 Member Function Documentation	144
18.25.2.1 GetCount	144
18.25.2.2 Init	144
18.25.2.3 Pend	145
18.25.2.4 Pend	145
18.25.2.5 Pend_i	145
18.25.2.6 Post	145
18.25.2.7 WakeMe	146
18.25.2.8 WakeNext	146
18.26 Thread Class Reference	146
18.26.1 Detailed Description	149
18.26.2 Member Function Documentation	149
18.26.2.1 ContextSwitchSWI	149
18.26.2.2 Exit	149
18.26.2.3 GetCurPriority	149
18.26.2.4 GetCurrent	149
18.26.2.5 GetEventFlagMask	150
18.26.2.6 GetEventFlagMode	150
18.26.2.7 GetExpired	150
18.26.2.8 GetExtendedContext	150
18.26.2.9 GetID	150

18.26.2.10	GetOwner	151
18.26.2.11	GetPriority	151
18.26.2.12	GetQuantum	151
18.26.2.13	GetStack	151
18.26.2.14	GetStackSize	151
18.26.2.15	GetStackSlack	152
18.26.2.16	GetState	152
18.26.2.17	InheritPriority	152
18.26.2.18	nit	152
18.26.2.19	nitIdle	153
18.26.2.20	SetCurrent	153
18.26.2.21	SetEventFlagMask	153
18.26.2.22	SetEventFlagMode	153
18.26.2.23	SetExpired	153
18.26.2.24	SetExtendedContext	154
18.26.2.25	SetID	154
18.26.2.26	SetOwner	154
18.26.2.27	SetPriority	154
18.26.2.28	SetPriorityBase	154
18.26.2.29	SetQuantum	155
18.26.2.30	SetState	155
18.26.2.31	Sleep	155
18.26.2.32	Start	155
18.26.2.33	Stop	156
18.26.2.34	USleep	156
18.26.2.35	Yield	156
18.27	ThreadList Class Reference	156
18.27.1	Detailed Description	157
18.27.2	Constructor & Destructor Documentation	157
18.27.2.1	ThreadList	157
18.27.3	Member Function Documentation	157
18.27.3.1	Add	157
18.27.3.2	Add	158
18.27.3.3	AddPriority	158
18.27.3.4	HighestWaiter	158
18.27.3.5	Remove	158
18.27.3.6	SetMapPointer	158
18.27.3.7	SetPriority	159
18.28	ThreadPort Class Reference	159
18.28.1	Detailed Description	159

18.28.2 Member Function Documentation	160
18.28.2.1 InitStack	160
18.28.2.2 StartThreads	160
18.29 Timer Class Reference	160
18.29.1 Detailed Description	162
18.29.2 Constructor & Destructor Documentation	162
18.29.2.1 Timer	162
18.29.3 Member Function Documentation	162
18.29.3.1 GetInterval	162
18.29.3.2 Init	162
18.29.3.3 IsInitialized	162
18.29.3.4 SetCallback	162
18.29.3.5 SetData	163
18.29.3.6 SetFlags	163
18.29.3.7 SetIntervalMSeconds	163
18.29.3.8 SetIntervalSeconds	163
18.29.3.9 SetIntervalTicks	163
18.29.3.10 SetIntervalUSeconds	164
18.29.3.11 SetOwner	164
18.29.3.12 SetTolerance	164
18.29.3.13 Start	164
18.29.3.14 Start	164
18.29.3.15 Start	166
18.29.3.16 Stop	166
18.30 TimerList Class Reference	166
18.30.1 Detailed Description	167
18.30.2 Member Function Documentation	167
18.30.2.1 Add	167
18.30.2.2 Init	167
18.30.2.3 Process	167
18.30.2.4 Remove	167
18.31 TimerScheduler Class Reference	168
18.31.1 Detailed Description	168
18.31.2 Member Function Documentation	168
18.31.2.1 Add	168
18.31.2.2 Init	168
18.31.2.3 Process	168
18.31.2.4 Remove	169

19.1	/media/usb/project/github/Mark3/kernel/atomic.cpp File Reference	171
19.1.1	Detailed Description	171
19.2	atomic.cpp	171
19.3	/media/usb/project/github/Mark3/kernel/autoalloc.cpp File Reference	173
19.3.1	Detailed Description	173
19.4	autoalloc.cpp	173
19.5	/media/usb/project/github/Mark3/kernel/blocking.cpp File Reference	175
19.5.1	Detailed Description	175
19.6	blocking.cpp	175
19.7	/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kernelprofile.cpp File Reference	177
19.7.1	Detailed Description	177
19.8	kernelprofile.cpp	177
19.9	/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kernelswi.cpp File Reference	178
19.9.1	Detailed Description	178
19.10	kernelswi.cpp	178
19.11	/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kerneltimer.cpp File Reference	179
19.11.1	Detailed Description	180
19.12	kerneltimer.cpp	180
19.13	/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kernelprofile.h File Reference	182
19.13.1	Detailed Description	182
19.14	kernelprofile.h	183
19.15	/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kernelswi.h File Reference	183
19.15.1	Detailed Description	183
19.16	kernelswi.h	184
19.17	/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kerneltimer.h File Reference	184
19.17.1	Detailed Description	184
19.18	kerneltimer.h	184
19.19	/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/portcfg.h File Reference	185
19.19.1	Detailed Description	186
19.19.2	Macro Definition Documentation	186
19.19.2.1	AVR	186
19.19.2.2	K_WORD	186
19.19.2.3	PORT_PRIO_TYPE	186
19.19.2.4	PORT_SYSTEM_FREQ	187
19.19.2.5	PORT_TIMER_COUNT_TYPE	187
19.19.2.6	PORT_TIMER_FREQ	187
19.20	portcfg.h	187
19.21	/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/threadport.h File Reference	188
19.21.1	Detailed Description	188
19.21.2	Macro Definition Documentation	189

19.21.2.1 CS_ENTER	189
19.21.3 Function Documentation	189
19.21.3.1 __mark3_clz8	189
19.22threadport.h	189
19.23/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/threadport.cpp File Reference . .	191
19.23.1 Detailed Description	192
19.24threadport.cpp	192
19.25/media/usb/project/github/Mark3/kernel/driver.cpp File Reference	194
19.25.1 Detailed Description	194
19.26driver.cpp	194
19.27/media/usb/project/github/Mark3/kernel/eventflag.cpp File Reference	195
19.27.1 Detailed Description	195
19.28eventflag.cpp	196
19.29/media/usb/project/github/Mark3/kernel/kernel.cpp File Reference	200
19.29.1 Detailed Description	200
19.30kernel.cpp	200
19.31/media/usb/project/github/Mark3/kernel/kernelaware.cpp File Reference	201
19.31.1 Detailed Description	202
19.32kernelaware.cpp	202
19.33/media/usb/project/github/Mark3/kernel/ksemaphore.cpp File Reference	203
19.33.1 Detailed Description	204
19.34ksemaphore.cpp	204
19.35/media/usb/project/github/Mark3/kernel/ll.cpp File Reference	207
19.35.1 Detailed Description	207
19.36ll.cpp	207
19.37/media/usb/project/github/Mark3/kernel/mailbox.cpp File Reference	209
19.37.1 Detailed Description	209
19.38mailbox.cpp	210
19.39/media/usb/project/github/Mark3/kernel/message.cpp File Reference	213
19.39.1 Detailed Description	213
19.40message.cpp	213
19.41/media/usb/project/github/Mark3/kernel/mutex.cpp File Reference	216
19.41.1 Detailed Description	216
19.42mutex.cpp	216
19.43/media/usb/project/github/Mark3/kernel/notify.cpp File Reference	219
19.43.1 Detailed Description	220
19.44notify.cpp	220
19.45/media/usb/project/github/Mark3/kernel/priomap.cpp File Reference	222
19.45.1 Detailed Description	222
19.46priomap.cpp	222

19.47/media/usb/project/github/Mark3/kernel/profile.cpp File Reference	224
19.47.1 Detailed Description	224
19.48profile.cpp	224
19.49/media/usb/project/github/Mark3/kernel/public/atomic.h File Reference	225
19.49.1 Detailed Description	226
19.50atomic.h	226
19.51/media/usb/project/github/Mark3/kernel/public/autoalloc.h File Reference	226
19.51.1 Detailed Description	226
19.52autoalloc.h	227
19.53/media/usb/project/github/Mark3/kernel/public/blocking.h File Reference	228
19.53.1 Detailed Description	228
19.54blocking.h	229
19.55/media/usb/project/github/Mark3/kernel/public/buffalogger.h File Reference	229
19.55.1 Detailed Description	229
19.56buffalogger.h	230
19.57/media/usb/project/github/Mark3/kernel/public/driver.h File Reference	230
19.57.1 Detailed Description	230
19.57.2 Intro	231
19.57.3 Driver Design	231
19.57.4 Driver API	231
19.58driver.h	231
19.59/media/usb/project/github/Mark3/kernel/public/eventflag.h File Reference	232
19.59.1 Detailed Description	233
19.60eventflag.h	233
19.61/media/usb/project/github/Mark3/kernel/public/kernel.h File Reference	234
19.61.1 Detailed Description	234
19.62kernel.h	234
19.63/media/usb/project/github/Mark3/kernel/public/kernelaware.h File Reference	235
19.63.1 Detailed Description	236
19.63.2 Enumeration Type Documentation	236
19.63.2.1 KernelAwareCommand_t	236
19.64kernelaware.h	236
19.65/media/usb/project/github/Mark3/kernel/public/kerneldebug.h File Reference	237
19.65.1 Detailed Description	238
19.66kerneldebug.h	238
19.67/media/usb/project/github/Mark3/kernel/public/kerneltypes.h File Reference	243
19.67.1 Detailed Description	243
19.67.2 Enumeration Type Documentation	243
19.67.2.1 EventFlagOperation_t	243
19.68kerneltypes.h	244

19.69/media/usb/project/github/Mark3/kernel/public/ksemaphore.h File Reference	244
19.69.1 Detailed Description	245
19.70ksemaphore.h	245
19.71/media/usb/project/github/Mark3/kernel/public/ll.h File Reference	246
19.71.1 Detailed Description	246
19.72ll.h	246
19.73/media/usb/project/github/Mark3/kernel/public/mailbox.h File Reference	248
19.73.1 Detailed Description	248
19.74mailbox.h	248
19.75/media/usb/project/github/Mark3/kernel/public/manual.h File Reference	250
19.75.1 Detailed Description	250
19.76manual.h	250
19.77/media/usb/project/github/Mark3/kernel/public/mark3.h File Reference	251
19.77.1 Detailed Description	251
19.78mark3.h	251
19.79/media/usb/project/github/Mark3/kernel/public/mark3cfg.h File Reference	252
19.79.1 Detailed Description	253
19.79.2 Macro Definition Documentation	254
19.79.2.1 GLOBAL_MESSAGE_POOL_SIZE	254
19.79.2.2 KERNEL_AWARE_SIMULATION	254
19.79.2.3 KERNEL_EXTRA_CHECKS	254
19.79.2.4 KERNEL_NUM_PRIORITIES	254
19.79.2.5 KERNEL_TIMERS_MINIMUM_DELAY_US	254
19.79.2.6 KERNEL_TIMERS_THREADED	254
19.79.2.7 KERNEL_TIMERS_TICKLESS	255
19.79.2.8 KERNEL_USE_ATOMIC	255
19.79.2.9 KERNEL_USE_AUTO_ALLOC	255
19.79.2.10KERNEL_USE_DYNAMIC_THREADS	255
19.79.2.11KERNEL_USE_EVENTFLAG	255
19.79.2.12KERNEL_USE_IDLE_FUNC	255
19.79.2.13KERNEL_USE_MAILBOX	256
19.79.2.14KERNEL_USE_MESSAGE	256
19.79.2.15KERNEL_USE_PROFILER	256
19.79.2.16KERNEL_USE_QUANTUM	256
19.79.2.17KERNEL_USE_SEMAPHORE	256
19.79.2.18KERNEL_USE_STACK_GUARD	256
19.79.2.19KERNEL_USE_THREAD_CALLOUTS	257
19.79.2.20KERNEL_USE_THREADNAME	257
19.79.2.21KERNEL_USE_TIMEOUTS	257
19.79.2.22KERNEL_USE_TIMERS	257

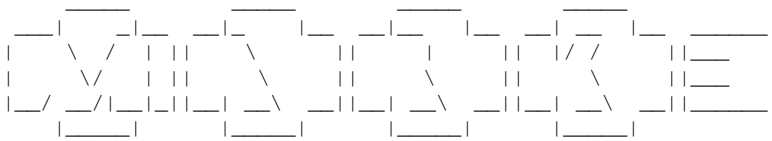
19.79.2.23SAFE_UNLINK	257
19.79.2.24THREAD_QUANTUM_DEFAULT	257
19.80mark3cfg.h	258
19.81/media/usb/project/github/Mark3/kernel/public/message.h File Reference	259
19.81.1 Detailed Description	260
19.81.2 using Messages, Queues, and the Global Message Pool	260
19.82message.h	260
19.83/media/usb/project/github/Mark3/kernel/public/mutex.h File Reference	262
19.83.1 Detailed Description	262
19.83.2 Initializing	262
19.83.3 Resource protection example	262
19.84mutex.h	263
19.85/media/usb/project/github/Mark3/kernel/public/notify.h File Reference	263
19.85.1 Detailed Description	264
19.86notify.h	264
19.87/media/usb/project/github/Mark3/kernel/public/paniccodes.h File Reference	264
19.87.1 Detailed Description	265
19.88paniccodes.h	265
19.89/media/usb/project/github/Mark3/kernel/public/priomap.h File Reference	265
19.89.1 Detailed Description	265
19.90priomap.h	265
19.91/media/usb/project/github/Mark3/kernel/public/profile.h File Reference	266
19.91.1 Detailed Description	267
19.92profile.h	267
19.93/media/usb/project/github/Mark3/kernel/public/quantum.h File Reference	268
19.93.1 Detailed Description	268
19.94quantum.h	268
19.95/media/usb/project/github/Mark3/kernel/public/scheduler.h File Reference	269
19.95.1 Detailed Description	269
19.96scheduler.h	270
19.97/media/usb/project/github/Mark3/kernel/public/thread.h File Reference	271
19.97.1 Detailed Description	271
19.98thread.h	271
19.99/media/usb/project/github/Mark3/kernel/public/threadlist.h File Reference	274
19.99.1 Detailed Description	274
19.100threadlist.h	274
19.101/media/usb/project/github/Mark3/kernel/public/timer.h File Reference	275
19.101.1 Detailed Description	276
19.101.2 Macro Definition Documentation	276
19.101.2.1TIMERLIST_FLAG_EXPIRED	276

19.101.3 Typedef Documentation	276
19.101.3.1 TimerCallback_t	276
19.102 mer.h	276
19.103 media/usb/project/github/Mark3/kernel/public/timerlist.h File Reference	278
19.103.1 Detailed Description	278
19.104 merlist.h	278
19.105 media/usb/project/github/Mark3/kernel/public/timerscheduler.h File Reference	279
19.105.1 Detailed Description	279
19.106 merscheduler.h	280
19.107 media/usb/project/github/Mark3/kernel/public/tracebuffer.h File Reference	280
19.107.1 Detailed Description	280
19.108 acebuffer.h	280
19.109 media/usb/project/github/Mark3/kernel/quantum.cpp File Reference	281
19.109.1 Detailed Description	281
19.110 quantum.cpp	281
19.111 media/usb/project/github/Mark3/kernel/scheduler.cpp File Reference	283
19.111.1 Detailed Description	283
19.112 scheduler.cpp	284
19.113 media/usb/project/github/Mark3/kernel/thread.cpp File Reference	285
19.113.1 Detailed Description	285
19.114 hread.cpp	285
19.115 media/usb/project/github/Mark3/kernel/threadlist.cpp File Reference	292
19.115.1 Detailed Description	292
19.116 hreadlist.cpp	292
19.117 media/usb/project/github/Mark3/kernel/timer.cpp File Reference	294
19.117.1 Detailed Description	294
19.118 mer.cpp	294
19.119 media/usb/project/github/Mark3/kernel/timerlist.cpp File Reference	297
19.119.1 Detailed Description	297
19.120 merlist.cpp	297
19.121 media/usb/project/github/Mark3/kernel/tracebuffer.cpp File Reference	300
19.121.1 Detailed Description	300
19.122 acebuffer.cpp	301
19.123 media/usb/project/github/Mark3/libs/mark3c/public/fake_types.h File Reference	302
19.123.1 Detailed Description	302
19.124 ake_types.h	302
19.125 media/usb/project/github/Mark3/libs/mark3c/public/mark3c.h File Reference	304
19.125.1 Detailed Description	305
19.125.2 Function Documentation	306
19.125.2.1 Kernel_Init	306

19.125.2.2	Kernel_IsPanic	306
19.125.2.3	Kernel_IsStarted	306
19.125.2.4	Kernel_Panic	306
19.125.2.5	Kernel_SetPanic	306
19.125.2.6	Kernel_Start	307
19.125.2.7	Scheduler_Enable	307
19.125.2.8	Scheduler_GetCurrentThread	307
19.125.2.9	Scheduler_IsEnabled	307
19.125.2.10	Thread_GetCurPriority	307
19.125.2.11	Thread_GetID	308
19.125.2.12	Thread_GetPriority	308
19.125.2.13	Thread_GetStackSlack	308
19.125.2.14	Thread_GetState	309
19.125.2.15	Thread_Init	309
19.125.2.16	Thread_SetID	309
19.125.2.17	Thread_SetPriority	309
19.125.2.18	Thread_Start	310
19.125.2.19	Thread_Stop	310
19.125.2.20	Thread_Yield	310
19.126	mark3c.h	310
20	Example Documentation	317
20.1	buffalogger/main.cpp	317
20.2	lab10_notifications/main.cpp	319
20.3	lab11_mailboxes/main.cpp	320
20.4	lab1_kernel_setup/main.cpp	322
20.5	lab2_idle_function/main.cpp	324
20.6	lab3_round_robin/main.cpp	325
20.7	lab4_semaphores/main.cpp	326
20.8	lab5_mutexes/main.cpp	328
20.9	lab6_timers/main.cpp	330
20.10	lab7_events/main.cpp	331
20.11	lab8_messages/main.cpp	333
20.12	lab9_dynamic_threads/main.cpp	335
	Index	339

Chapter 1

The Mark3 Realtime Kernel



```
--[Mark3 Realtime Platform]-----
```

```
Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.  
See license for more information
```

The Mark3 Realtime [Kernel](#) is a completely free, open-source, real-time operating system aimed at bringing powerful, easy-to-use multitasking to microcontroller systems without MMUs.

It uses modern programming languages and concepts to minimize code duplication, and its object-oriented design enhances readability. The API is simple – in six function calls, you can set up the kernel, initialize two threads, and start the scheduler.

The source is fully-documented with example code provided to illustrate concepts. The result is a performant RTOS, which is easy to read, easy to understand, and easy to extend to fit your needs.

But Mark3 is bigger than just a real-time kernel, it also contains a number of class-leading features:

- Native implementation in C++, with C-language bindings.
- Device driver HAL which provides a meaningful abstraction around device-specific peripherals.
- CMake-based build system which can be used to build all libraries, examples, tests, documentation, and user-projects for any number of targets from the command-line.
- Graphics and UI code designed to simplify the implementation of systems using displays, keypads, joysticks, and touchscreens
- Robust and deterministic dynamic memory management libraries
- A Variety of general-purpose libraries to speed up embedded app development
- Emulator-aware debugging via the fIAVR AVR emulator
- A bulletproof, well-documented bootloader for AVR microcontrollers Support for kernel-aware simulators, including Funkenstein's own fIAVR.

Chapter 2

License

2.1 License

Copyright (c) 2012 - 2017, Funkenstein Software Consulting
All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
- Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
- Neither the name of Funkenstein Software Consulting, nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL FUNKENSTEIN SOFTWARE AND/OR ITS CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Chapter 3

Configuring The Mark3 Kernel

3.1 Overview

The Mark3 [Kernel](#) features a large number of compile-time options that can be set by the user. In this way, the user can build a custom OS kernel that provides only the necessary feature set required by the application, and reduce the code and data requirements of the kernel.

Care has been taken to ensure that all valid combinations of features can be enabled or disabled, barring direct dependencies.

When Mark3 is built, the various compile-time definitions are used to alter how the kernel is compiled, and include or exclude various bits and pieces in order to satisfy the requirements of the selected features. As a result, the kernel must be rebuilt whenever changes are made to the configuration header.

Note that not all demos, libraries, and tests will build successfully if the prerequisite features are not included.

[Kernel](#) options are set by modifying [mark3cfg.h](#), located within the /kernel/public folder.

In the following sections, we will discuss the various configuration options, grouped by functionality.

3.2 Timer Options

KERNEL_USE_TIMERS

This option is related to all kernel time-tracking:

- Timers provide a way for events to be periodically triggered in a lightweight manner. These can be periodic, or one-shot.
- [Thread Quantum](#) (used for round-robin scheduling) is dependent on this module, as is [Thread Sleep](#) functionality.

Setting this option to 0 disables all timer-based functionality within the kernel.

KERNEL_TIMERS_TICKLESS

If you've opted to use the kernel timers module, you have an option as to which timer implementation to use: Tick-based or Tick-less.

Tick-based timers provide a "traditional" RTOS timer implementation based on a fixed-frequency timer interrupt. While this provides very accurate, reliable timing, it also means that the CPU is being interrupted far more often than may be necessary (as not all timer ticks result in "real work" being done).

Tick-less timers still rely on a hardware timer interrupt, but uses a dynamic expiry interval to ensure that the interrupt is only called when the next timer expires. This increases the complexity of the timer interrupt handler, but reduces the number and frequency.

Note that the CPU port ([kerneltimer.cpp](#)) must be implemented for the particular timer variant desired.

Set this option to 1 to use the tickless timer implementation, 0 to use the traditional tick-based approach. Tickless timers are a bit more heavy weight (larger code footprint), but can yield significant power savings as the CPU does not need to wake up at a fixed, high frequency.

KERNEL_USE_TIMEOUTS

By default, if you opt to enable kernel timers, you also get timeout- enabled versions of the blocking object APIs along with it. This support comes at a small cost to code size, but a slightly larger cost to realtime performance - as checking for the use of timers in the underlying internal code costs some cycles.

As a result, the option is given to the user here to manually disable these timeout-based APIs if desired by the user for performance and code-size reasons.

Set this option to 1 to enable timeout-based APIs for blocking calls.

KERNEL_USE_QUANTUM

Do you want to enable time quanta? This is useful when you want to have tasks in the same priority group share time in a controlled way. This allows equal tasks to use unequal amounts of the CPU, which is a great way to set up CPU budgets per thread in a round-robin scheduling system. If enabled, you can specify a number of ticks that serves as the default time period (quantum). Unless otherwise specified, every thread in a priority will get the default quantum.

Set this option to 1 to enable round-robin scheduling.

THREAD_QUANTUM_DEFAULT

This value defines the default thread quantum when `KERNEL_USE_QUANTUM` is enabled. The value defined is a time in milliseconds.

KERNEL_USE_SLEEP

This define enables the [Thread::Sleep\(\)](#) API, which allows a thread to suspend its operation for a defined length of time, specified in ms.

3.3 Blocking Objects

KERNEL_USE_NOTIFY

This is a simple blocking object, where a thread (or threads) are guaranteed to block until an asynchronous event signals the object.

KERNEL_USE_SEMAPHORE

Do you want the ability to use counting/binary semaphores for thread synchronization? Enabling this features provides fully-blocking semaphores and enables all API functions declared in `semaphore.h`. If you have to pick one blocking mechanism, this is the one to choose.

Note that all IPC mechanisms (mailboxes, messages) rely on semaphores, so keep in mind that this is a prerequisite for many other features in the kernel.

KERNEL_USE_MUTEX

Do you want the ability to use mutual exclusion semaphores (mutex) for resource/block protection? Enabling this feature provides mutexes, with priority inheritance, as declared in [mutex.h](#).

KERNEL_USE_EVENTFLAG

Provides additional event-flag based blocking. This relies on an additional per-thread flag-mask to be allocated, which adds 2 bytes to the size of each thread object.

3.4 Inter-process/thread Communication

KERNEL_USE_MESSAGE

Enable inter-thread messaging using message queues. This is the preferred mechanism for IPC for serious multi-threaded communications; generally anywhere a semaphore or event-flag is insufficient.

GLOBAL_MESSAGE_POOL_SIZE

If Messages are enabled, define the size of the default kernel message pool. Messages can be manually added to the message pool, but this mechanism is more convenient and automatic. All message queues can share their message objects from this global pool to maximize efficiency and simplify data management.

KERNEL_USE_MAILBOX

Enable inter-thread messaging using mailboxes. A mailbox manages a blob of data provided by the user, that is partitioned into fixed-size blocks called envelopes. The size of an envelope is set by the user when the mailbox is initialized. Any number of threads can read-from and write-to the mailbox. Envelopes can be sent-to or received-from the mailbox at the head or tail. In this way, mailboxes essentially act as a circular buffer that can be used as a blocking FIFO or LIFO queue.

3.5 Debug Features

KERNEL_USE_THREADNAME

Provide [Thread](#) method to allow the user to set a name for each thread in the system. Adds a const char* pointer to the size of the thread object.

KERNEL_USE_DEBUG

Provides extra logic for kernel debugging, and instruments the kernel with extra asserts, and kernel trace functionality.

KERNEL_ENABLE_LOGGING

Set this to 1 to enable very chatty kernel logging. Since most important things in the kernel emit logs, a large log-buffer and fast output are required in order to keep up. This is a pretty advanced power-user type feature, so it's disabled by default.

KERNEL_ENABLE_USER_LOGGING

This enables a set of logging macros similar to the kernel-logging macros; however, these can be enabled or disabled independently. This allows for user-code to benefit from the built-in kernel logging macros without having to account for the super-high-volume of logs generated by kernel code.

KERNEL_EXTRA_CHECKS

This option provides extra safety checks within the kernel APIs in order to minimize the potential for unsafe operations. This is especially helpful during development, and can help catch problems at development time, instead of in the field.

KERNEL_USE_STACK_GUARD

This feature, when enabled, tells the kernel to check whether any [Thread](#)'s stack has been exhausted (or slack falls below a certain safety threshold) before executing each context switch. Enabling this is the most effective means to guard against stack corruption and stack overflow in the kernel, at the cost of increased context

3.6 Enhancements, Security, Miscellaneous

KERNEL_USE_DRIVER

Enabling device drivers provides a posix-like filesystem interface for peripheral device drivers.

KERNEL_USE_DYNAMIC_THREADS

Provide extra [Thread](#) methods to allow the application to create (and more importantly destroy) threads at runtime. useful for designs implementing worker threads, or threads that can be restarted after encountering error conditions.

KERNEL_USE_PROFILER

Provides extra classes for profiling the performance of code. useful for debugging and development, but uses an additional hardware timer.

KERNEL_USE_ATOMIC

Provides support for atomic operations, including addition, subtraction, set, and test-and-set. Add/Sub/Set contain 8, 16, and 32-bit variants.

SAFE_UNLINK

"Safe unlinking" performs extra checks on data to make sure that there are no consistencies when performing operations on linked lists. This goes beyond pointer checks, adding a layer of structural and metadata validation to help detect system corruption early.

KERNEL_AWARE_SIMULATION

Include support for kernel-aware simulation. Enabling this feature adds advanced profiling, trace, and environment-aware debugging and diagnostic functionality when Mark3-based applications are run on the fIAVR AVR simulator.

KERNEL_USE_IDLE_FUNC

Enabling this feature removes the necessity for the user to dedicate a complete thread for idle functionality. This saves a full thread stack, but also requires a bit extra static data. This also adds a slight overhead to the context switch and scheduler, as a special case has to be taken into account.

KERNEL_USE_AUTO_ALLOC

This feature enables an additional set of APIs that allow for objects to be created on-the-fly out of a special heap, without having to explicitly allocate them (from stack, heap, or static memory). Note that auto-alloc memory cannot be reclaimed.

AUTO_ALLOC_SIZE

Size (in bytes) of the static pool of memory reserved from RAM for use by the auto allocator (if enabled).

KERNEL_USE_THREAD_CALLOUTS

This feature provides additional kernel APIs to register callout functions that are activated when threads are created or exited. This is useful for implementing low-level instrumentation based on information held in the threads.

KERNEL_USE_EXTENDED_CONTEXT

Allocate an extra pointer's worth of storage within a [Thread](#) object (and corresponding accessor methods) to provide the user with a means to implement arbitrary Thread-local storage.

Chapter 4

Building Mark3

4.1 Source Layout

One key aspect of Mark3 is that system features are organized into their own separate modules. These modules are further grouped together into folders based on the type of features represented:

Root	Base folder, contains license info and build system configuration
arduino	Arduino-specific headers and API documentation files
bootloader	Mark3 Bootloader code for AVR microcontrollers
build	Device-specific toolchain configuraton files for various platforms
docs	Documentation (pdf + html)
drivers	Device driver code for various supported devices
example	Example applications
export	Platform specific output folder, used when running export.sh
fonts	Bitmap fonts converted from TTF, used by Mark3 graphics library
kbuild	Build output directory
kernel	Basic Mark3 Components (the focus of this manual)
cpu	CPU-specific porting code
scripts	Scripts used to simplify build, documentation, and profiling
libs	Utility code and services, extended system features
tests	Unit tests, written as C/C++ applications
util	Host utilities - including a ttf font converter and device programmer

4.2 Toolchain Integration

Mark3 supports a variety of GCC ports out of the box - however, depending on your host OS and target processor, there may be some effort required to tie the toolchain into the build system.

After installing your toolchain of choice, you must make sure that the main toolchain binary paths are set in your systems PATH environment variable, ensuring that they are accessible directly from the command-line. Without this step, the build configuration step (cmake) will inevitably fail.

Depending on your toolchain, you may also be required to add toolchain-specific include directories to the build flags. These flags can be added to the cmake variables defined in `/build/<cpu>/<variant>/<toolchain>/platform.cmake` for your target architecture.

4.3 Installing Dependencies

The Mark3 build system uses CMake (3.4.2 or above) for configuration management, and Ninja to execute the build steps. The combination of these two tools results in exceptionally fast builds - so fast that the previous makefile build system was scrapped in its favor.

These tools are readily available for most common host operating systems.

CMake can be found here: <https://cmake.org> Ninja can be found here: <https://ninja-build.org>

4.4 Building Mark3 Kernel and Libraries

Once a sane environment has been created, the kernel, libraries, examples and tests can be built by running `./scripts/build.sh` from the root directory. By default, Mark3 builds for the `atmega328p` target, but the target can be selected by manually configuring the above environment variables, or by running the included `./scripts/set_target.sh` script as follows:

```
./scripts/set_target.sh <architecture> <variant> <toolchain>
```

Where:

```
<architecture> is the target CPU architecture(i.e. avr, msp430, cm0, cm3, cm4f)
<variant>      is the part name (i.e. atmega328p, msp430f2274, generic)
<toolchain>    is the build toolchain (i.e. gcc)
```

This script is a thin wrapper for the cmake configuration commands, and clears the `kbuild` output directory before re-initializing cmake for the selected target.

To build the Mark3 kernel and middleware libraries for a generic ARM Cortex-M0 using a pre-configured `arm-none-eabi-gcc` toolchain, one would run the following commands:

```
./scripts/set_target.sh cm0 generic gcc
./scripts/build.sh
```

To perform an incremental build, go into the cmake build directory (`kbuild`) and simply run `'ninja'`.

Note that not all libraries/tests/examples will build in all kernel configurations. The default kernel configuration may need adjustment/tweaking to support a specific part. See `CMakeLists.txt` and [mark3cfg.h](#) respectively for more information

4.5 Exporting the kernel source

While the build system is flexible enough to adapt to any toolchain, it may be desirable to integrate the Mark3 kernel and associated drivers/libraries into another build system.

Mark3 provides a script (the aptly-named `export.sh`) which allow for the source for any supported port to be exported for this purpose. This script will also generate appropriate doxygen documentation, and package the whole of it together in a zip file. The files in the archive are placed in a "flat" hierarchy, and do not require any specific path structure to be maintained when imported into another build system.

As a special feature, if the "arduino" AVR target is specified, additional pre-processing is done on the source to turn the standard Mark3 kernel into a library that can be imported directly into Arduino IDE. This is also how the official Mark3 arduino-compatible releases are generated (hosted on [mark3os.com](#) and [sourceforge.net](#))

To exercise the build system, type the following from the main mark3 embedded source directory:

```
> ./scripts/export.sh <target>
```

Where:

Target is one of the following:

```
atmega328p
atmega644
atmega1280
atmega2560
atmega1284p
```



```
atxmega256a3  
arduino  
arduino2560  
cortex_m0  
cortex_m3  
cortex_m4f  
msp430f2274
```

If successful, the generated artifacts will be placed in an output folder under the `./export` directory.

Additionally, if `doxygen` is found on the host system's `PATH`, a copy of the manual (using the specific port's source code) will be generated and archived with the source release. If `pdflatex` is also found on the host's `PATH`, a PDF copy of the manual will be generated, tailored to the selected target.

Chapter 5

Getting Started With The Mark3 API

5.1 Kernel Setup

This section details the process of defining threads, initializing the kernel, and adding threads to the scheduler.

If you're at all familiar with real-time operating systems, then these setup and initialization steps should be familiar. I've tried very hard to ensure that as much of the heavy lifting is hidden from the user, so that only the bare minimum of calls are required to get things started.

The examples presented in this chapter are real, working examples taken from the ATmega328p port.

First, you'll need to create the necessary data structures and functions for the threads:

1. Create a [Thread](#) object for all of the "root" or "initial" tasks.
2. Allocate stacks for each of the Threads
3. Define an entry-point function for each [Thread](#)

This is shown in the example code below:

```
//-----  
#include "thread.h"  
#include "kernel.h"  
  
//1) Create a thread object for all of the "root" or "initial" tasks  
static Thread AppThread;  
static Thread IdleThread;  
  
//2) Allocate stacks for each thread  
#define STACK_SIZE_APP      (192)  
#define STACK_SIZE_IDLE     (128)  
  
static uint8_t aucAppStack[STACK_SIZE_APP];  
static uint8_t aucIdleStack[STACK_SIZE_IDLE];  
  
//3) Define entry point functions for each thread  
void AppThread(void);  
void IdleThread(void);
```

Next, we'll need to add the required kernel initialization code to main. This consists of running the [Kernel's](#) init routine, initializing all of the threads we defined, adding the threads to the scheduler, and finally calling [Kernel::Start\(\)](#), which transfers control of the system to the RTOS.

These steps are illustrated in the following example.

```
int main(void)  
{  
    //1) Initialize the kernel prior to use  
    Kernel::Init();           // MUST be before other kernel ops  
  
    //2) Initialize all of the threads we've defined
```

```

AppThread.Init( aucAppStack,      // Pointer to the stack
                STACK_SIZE_APP,   // Size of the stack
                1,                // Thread priority
                (void*)AppEntry,   // Entry function
                NULL );           // Entry function argument

IdleThread.Init( aucIdleStack,    // Pointer to the stack
                 STACK_SIZE_IDLE, // Size of the stack
                 0,               // Thread priority
                 (void*)IdleEntry, // Entry function
                 NULL );          // Entry function argument

//3) Add the threads to the scheduler
AppThread.Start();           // Actively schedule the threads
IdleThread.Start();

//4) Give control of the system to the kernel
Kernel::Start();             // Start the kernel!
}

```

Not much to it, is there? There are a few noteworthy points in this code, though.

In order for the kernel to work properly, a system must always contain an idle thread; that is, a thread at priority level 0 that never blocks. This thread is responsible for performing any of the low-level power management on the CPU in order to maximize battery life in an embedded device. The idle thread must also never block, and it must never exit. Either of these operations will cause undefined behavior in the system.

The App thread is at a priority level greater-than 0. This ensures that as long as the App thread has something useful to do, it will be given control of the CPU. In this case, if the app thread blocks, control will be given back to the Idle thread, which will put the CPU into a power-saving mode until an interrupt occurs.

Stack sizes must be large enough to accommodate not only the requirements of the threads, but also the requirements of interrupts - up to the maximum interrupt-nesting level used. Stack overflows are super-easy to run into in an embedded system; if you encounter strange and unexplained behavior in your code, chances are good that one of your threads is blowing its stack.

5.2 Threads

Mark3 Threads act as independent tasks in the system. While they share the same address-space, global data, device-drivers, and system peripherals, each thread has its own set of CPU registers and stack, collectively known as the thread's **context**. The context is what allows the RTOS kernel to rapidly switch between threads at a high rate, giving the illusion that multiple things are happening in a system, when really, only one thread is executing at a time.

5.2.1 Thread Setup

Each instance of the [Thread](#) class represents a thread, its stack, its CPU context, and all of the state and metadata maintained by the kernel. Before a [Thread](#) will be scheduled to run, it must first be initialized with the necessary configuration data.

The Init function gives the user the opportunity to set the stack, stack size, thread priority, entry-point function, entry-function argument, and round-robin time quantum:

[Thread](#) stacks are pointers to blobs of memory (usually char arrays) carved out of the system's address space. Each thread must have a stack defined that's large enough to handle not only the requirements of local variables in the thread's code path, but also the maximum depth of the ISR stack.

Priorities should be chosen carefully such that the shortest tasks with the most strict determinism requirements are executed first - and are thus located in the highest priorities. Tasks that take the longest to execute (and require the least degree of responsiveness) must occupy the lower thread priorities. The idle thread must be the only thread occupying the lowest priority level.

The thread quantum only applies when there are multiple threads in the ready queue at the same priority level. This interval is used to kick-off a timer that will cycle execution between the threads in the priority list so that they each get a fair chance to execute.

The entry function is the function that the kernel calls first when the thread instance is first started. Entry functions have at most one argument - a pointer to a data-object specified by the user during initialization.

An example thread initialization is shown below:

```
Thread clMyThread;
uint8_t aucStack[192];

void AppEntry(void)
{
    while(1)
    {
        // Do something!
    }
}

...
{
    clMyThread.Init(aucStack,    // Pointer to the stack to use by this thread
                    192,        // Size of the stack in bytes
                    1,          // Thread priority (0 = idle, 7 = max)
                    (void*)AppEntry, // Function where the thread starts executing
                    NULL );      // Argument passed into the entry function
}
```

Once a thread has been initialized, it can be added to the scheduler by calling:

```
clMyThread.Start();
```

The thread will be placed into the [Scheduler](#)'s queue at the designated priority, where it will wait its turn for execution.

5.2.2 Entry Functions

Mark3 Threads should not run-to-completion - they should execute as infinite loops that perform a series of tasks, appropriately partitioned to provide the responsiveness characteristics desired in the system.

The most basic [Thread](#) loop is shown below:

```
void Thread( void *param )
{
    while(1)
    {
        // Do Something
    }
}
```

Threads can interact with eachother in the system by means of synchronization objects ([Semaphore](#)), mutual-exclusion objects ([Mutex](#)), Inter-process messaging ([MessageQueue](#)), and timers ([Timer](#)).

Threads can suspend their own execution for a predetermined period of time by using the static [Thread::Sleep\(\)](#) method. Calling this will block the [Thread](#)'s executin until the amount of time specified has ellapsed. Upon expiry, the thread will be placed back into the ready queue for its priority level, where it awaits its next turn to run.

5.3 Timers

[Timer](#) objects are used to trigger callback events periodic or on a one-shot (alarm) basis.

While extremely simple to use, they provide one of the most powerful execution contexts in the system. The timer callbacks execute from within the timer callback ISR in an interrupt-enabled context. As such, timer callbacks are considered higher-priority than any thread in the system, but lower priority than other interrupts. Care must be taken to ensure that timer callbacks execute as quickly as possible to minimize the impact of processing on the throughput of tasks in the system. Wherever possible, heavy-lifting should be deferred to the threads by way of semaphores or messages.

Below is an example showing how to start a periodic system timer which will trigger every second:

```

{
    Timer clTimer;
    clTimer.Init();

    clTimer.Start( 1000,
                  1,
                  MyCallback,
                  (void*)&my_data );

    ... // Keep doing work in the thread
}

// Callback function, executed from the timer-expiry context.
void MyCallback( Thread *pclOwner_, void *pvData_ )
{
    LED.Flash(); // Flash an LED.
}

```

5.4 Semaphores

Semaphores are used to synchronized execution of threads based on the availability (and quantity) of application-specific resources in the system. They are extremely useful for solving producer-consumer problems, and are the method-of-choice for creating efficient, low latency systems, where ISRs post semaphores that are handled from within the context of individual threads. (Yes, Semaphores can be posted - but not pended - from the interrupt context).

The following is an example of the producer-consumer usage of a binary semaphore:

```

Semaphore clSemaphore; // Declare a semaphore shared between a producer and a consumer thread.

void Producer()
{
    clSemaphore.Init(0, 1);
    while(1)
    {
        // Do some work, create something to be consumed

        // Post a semaphore, allowing another thread to consume the data
        clSemaphore.Post();
    }
}

void Consumer()
{
    // Assumes semaphore initialized before use...
    While(1)
    {
        // Wait for new data from the producer thread
        clSemaphore.Pend();

        // Consume the data!
    }
}

```

And an example of using semaphores from the ISR context to perform event- driven processing.

```

Semaphore clSemaphore;

__interrupt__ MyISR()
{
    clSemaphore.Post(); // Post the interrupt. Lightweight when uncontested.
}

void MyThread()
{
    clSemaphore.Init(0, 1); // Ensure this is initialized before the MyISR interrupt is enabled.
    while(1)
    {
        // Wait until we get notification from the interrupt
        clSemaphore.Pend();

        // Interrupt has fired, do the necessary work in this thread's context
        HeavyLifting();
    }
}

```

5.5 Mutexes

Mutexes (Mutual exclusion objects) are provided as a means of creating "protected sections" around a particular resource, allowing for access of these objects to be serialized. Only one thread can hold the mutex at a time - other threads have to wait until the region is released by the owner thread before they can take their turn operating on the protected resource. Note that mutexes can only be owned by threads - they are not available to other contexts (i.e. interrupts). Calling the mutex APIs from an interrupt will cause catastrophic system failures.

Note that these objects are also not recursive- that is, the owner thread can not attempt to claim a mutex more than once.

Priority inheritance is provided with these objects as a means to avoid priority inversions. Whenever a thread at a priority than the mutex owner blocks on a mutex, the priority of the current thread is boosted to the highest-priority waiter to ensure that other tasks at intermediate priorities cannot artificially prevent progress from being made.

Mutex objects are very easy to use, as there are only three operations supported: Initialize, Claim and Release. An example is shown below.

```

Mutex clMutex; // Create a mutex globally.

void Init()
{
    // Initialize the mutex before use.
    clMutex.Init();
}

// Some function called from a thread
void Thread1Function()
{
    clMutex.Claim();

    // Once the mutex is owned, no other thread can
    // enter a block protect by the same mutex

    my_protected_resource.do_something();
    my_protected_resource.do_something_else();

    clMutex.Release();
}

// Some function called from another thread
void Thread2Function()
{
    clMutex.Claim();

    // Once the mutex is owned, no other thread can
    // enter a block protect by the same mutex

    my_protected_resource.do_something();
    my_protected_resource.do_different_things();

    clMutex.Release();
}

```

5.6 Event Flags

Event Flags are another synchronization object, conceptually similar to a semaphore.

Unlike a semaphore, however, the condition on which threads are unblocked is determined by a more complex set of rules. Each Event Flag object contains a 16-bit field, and threads block, waiting for combinations of bits within this field to become set.

A thread can wait on any pattern of bits from this field to be set, and any number of threads can wait on any number of different patterns. Threads can wait on a single bit, multiple bits, or bits from within a subset of bits within the field.

As a result, setting a single value in the flag can result in any number of threads becoming unblocked simultaneously. This mechanism is extremely powerful, allowing for all sorts of complex, yet efficient, thread synchronization schemes that can be created using a single shared object.

Note that Event Flags can be set from interrupts, but you cannot wait on an event flag from within an interrupt.

Examples demonstrating the use of event flags are shown below.

```
// Simple example showing a thread blocking on a multiple bits in the
// fields within an event flag.

EventFlag clEventFlag;

int main()
{
    ...
    clEventFlag.Init(); // Initialize event flag prior to use
    ...
}

void MyInterrupt()
{
    // Some interrupt corresponds to event 0x0020
    clEventFlag.Set(0x0020);
}

void MyThreadFunc()
{
    ...
    while(1)
    {
        ...
        uint16_t ul6WakeCondition;

        // Allow this thread to block on multiple flags
        ul6WakeCondition = clEventFlag.Wait(0x00FF, EVENT_FLAG_ANY);

        // Clear the event condition that caused the thread to wake (in this case,
        // ul6WakeCondition will equal 0x20 when triggered from the interrupt above)
        clEventFlag.Clear(ul6WakeCondition);

        // <do something>
    }
}
```

5.7 Messages

Sending messages between threads is the key means of synchronizing access to data, and the primary mechanism to perform asynchronous data processing operations.

Sending a message consists of the following operations:

- Obtain a [Message](#) object from the global message pool
- Set the message data and event fields
- Send the message to the destination message queue

While receiving a message consists of the following steps:

- Wait for a messages in the destination message queue
- Process the message data
- Return the message back to the global message pool

These operations, and the various data objects involved are discussed in more detail in the following section.

5.7.1 Message Objects

[Message](#) objects are used to communicate arbitrary data between threads in a safe and synchronous way.

The message object consists of an event code field and a data field. The event code is used to provide context to the message object, while the data field (essentially a void * data pointer) is used to provide a payload of data corresponding to the particular event.

Access to these fields is marshalled by accessors - the transmitting thread uses the `SetData()` and `SetCode()` methods to seed the data, while the receiving thread uses the `GetData()` and `GetCode()` methods to retrieve it.

By providing the data as a void data pointer instead of a fixed-size message, we achieve an unprecedented measure of simplicity and flexibility. Data can be either statically or dynamically allocated, and sized appropriately for the event without having to format and reformat data by both sending and receiving threads. The choices here are left to the user - and the kernel doesn't get in the way of efficiency.

It is worth noting that you can send messages to message queues from within ISR context. This helps maintain consistency, since the same APIs can be used to provide event-driven programming facilities throughout the whole of the OS.

5.7.2 Global Message Pool

To maintain efficiency in the messaging system (and to prevent over-allocation of data), a global pool of message objects is provided. The size of this message pool is specified in the implementation, and can be adjusted depending on the requirements of the target application as a compile-time option.

Allocating a message from the message pool is as simple as calling the `GlobalMessagePool::Pop()` Method.

Messages are returned back to the `GlobalMessagePool::Push()` method once the message contents are no longer required.

One must be careful to ensure that discarded messages always are returned to the pool, otherwise a resource leak can occur, which may cripple the operating system's ability to pass data between threads.

5.7.3 Message Queues

`Message` objects specify data with context, but do not specify where the messages will be sent. For this purpose we have a `MessageQueue` object. Sending an object to a message queue involves calling the `MessageQueue::Send()` method, passing in a pointer to the `Message` object as an argument.

When a message is sent to the queue, the first thread blocked on the queue (as a result of calling the `MessageQueue::Receive()` method) will wake up, with a pointer to the `Message` object returned.

It's worth noting that multiple threads can block on the same message queue, providing a means for multiple threads to share work in parallel.

5.7.4 Messaging Example

```
// Message queue object shared between threads
MessageQueue cMsgQ;

// Function that initializes the shared message queue
void MsgQInit()
{
    cMsgQ.Init();
}

// Function called by one thread to send message data to
// another
void TxMessage()
{
    // Get a message, initialize its data
    Message *pclMsg = GlobalMessagePool::Pop();

    pclMsg->SetCode(0xAB);
    pclMsg->SetData((void*)some_data);

    // Send the data to the message queue
    cMsgQ.Send(pclMsg);
}

// Function called in the other thread to block until
// a message is received in the message queue.
void RxMessage()
{
    Message *pclMsg;
```

```

// Block until we have a message in the queue
pclMsg = clMsgQ.Receive();

// Do something with the data once the message is received
pclMsg->GetCode();

// Free the message once we're done with it.
GlobalMessagePool::Push(pclMsg);
}

```

5.8 Mailboxes

Another form of IPC is provided by Mark3, in the form of Mailboxes and Envelopes.

Mailboxes are similar to message queues in that they provide a synchronized interface by which data can be transmitted between threads.

Where [Message](#) Queues rely on linked lists of lightweight message objects (containing only message code and a void* data-pointer), which are inherently abstract, Mailboxes use a dedicated blob of memory, which is carved up into fixed-size chunks called Envelopes (defined by the user), which are sent and received. Unlike message queues, mailbox data is copied to and from the mailboxes dedicated pool.

Mailboxes also differ in that they provide not only a blocking "receive" call, but also a blocking "send" call, providing the opportunity for threads to block on "mailbox full" as well as "mailbox empty" conditions.

All send/receive APIs support an optional timeout parameter if the `KERNEL_USE_TIMEOUTS` option has been configured in [mark3cfg.h](#)

5.8.1 Mailbox Example

```

// Create a mailbox object, and define a buffer that will be used to store the
// mailbox' envelopes.
static Mailbox clMbox;
static uint8_t aucMBoxBuffer[128];

...
void InitMailbox(void)
{
    // Initialize our mailbox, telling it to use our defined buffer for envelope
    // storage. Pass in the size of the buffer, and set the size of each
    // envelope to 16 bytes. This gives u16 a mailbox capacity of (128 / 16) = 8
    // envelopes.
    clMbox.Init((void*)aucMBoxBuffer, 128, 16);
}

...
void SendThread(void)
{
    // Define a buffer that we'll eventually send to the
    // mailbox. Note the size is the same as that of an
    // envelope.
    uint8_t aucTxBuf[16];

    while(1)
    {
        // Copy some data into aucTxBuf, a 16-byte buffer, the
        // same size as a mailbox envelope.
        ...

        // Deliver the envelope (our buffer) into the mailbox
        clMbox.Send((void*)aucTxBuf);
    }
}

...
void RecvThred(void)
{
    uint8_t aucRxBuf[16];

    while(1)
    {
        // Wait until there's a message in our mailbox. Once
        // there is a message, read it into our local buffer.
        cmMbox.Receive((void*)aucRxBuf);
    }
}

```

```

        // Do something with the contents of aucRxBuf, which now
        // contains an envelope of data read from the mailbox.
        ...
    }
}

```

5.9 Notification Objects

Notification objects are the most lightweight of all blocking objects supplied by Mark3.

using this blocking primitive, one or more threads wait for the notification object to be signalled by code elsewhere in the system (i.e. another thread or interrupt). Once the the notification has been signalled, all threads currently blocked on the object become unblocked.

5.9.1 Notification Example

```

static Notify clNotifier;

...
void MyThread(void *unused_)
{
    // Initialize our notification object before use
    clNotifier.Init();

    while (1)
    {
        // Wait until our thread has been notified that it
        // can wake up.
        clNotifier.Wait();

        ...
        // Thread has woken up now -- do something!
    }
}

...
void SignalCallback(void)
{
    // Something in the system (interrupt, thread event, IPC,
    // etc.,) has called this function. As a result, we need
    // our other thread to wake up. Call the Notify object's
    // Signal() method to wake the thread up. Note that this
    // will have no effect if the thread is not presently
    // blocked.

    clNotifier.Signal();
}

```

5.10 Sleep

There are instances where it may be necessary for a thread to poll a resource, or wait a specific amount of time before proceeding to operate on a peripheral or volatile piece of data.

While the [Timer](#) object is generally a better choice for performing time-sensitive operations (and certainly a better choice for periodic operations), the [Thread::Sleep\(\)](#) method provides a convenient (and efficient) mechanism that allows for a thread to suspend its execution for a specified interval.

Note that when a thread is sleeping it is blocked, during which other threads can operate, or the system can enter its idle state.

```

int GetPeripheralData()
{
    int value;
    // The hardware manual for a peripheral specifies that
    // the "foo()" method will result in data being generated
    // that can be captured using the "bar()" method.
    // However, the value only becomes valid after 10ms

    peripheral.foo();
    Thread::Sleep(10); // Wait 10ms for data to become valid
    value = peripheral.bar();
}

```

```
    return value;  
}
```

5.11 Round-Robin Quantum

Threads at the same thread priority are scheduled using a round-robin scheme. Each thread is given a timeslice (which can be configured) of which it shares time amongst ready threads in the group. Once a thread's timeslice has expired, the next thread in the priority group is chosen to run until its quantum has expired - the cycle continues over and over so long as each thread has work to be done.

By default, the round-robin interval is set at 4ms.

This value can be overridden by calling the thread's `SetQuantum()` with a new interval specified in milliseconds.

Chapter 6

Why Mark3?

My first job after graduating from university in 2005 was with a small company that had a very old-school, low-budget philosophy when it came to software development.

Every make-or-buy decision ended with "make" when it came to tools. It was the kind of environment where vendors cost us money, but manpower was free. In retrospect, we didn't have a ton of business during the time that I worked there, and that may have had something to do with the fact that we were constantly short on ready cash for things we could code ourselves.

Early on, I asked why we didn't use industry-standard tools - like JTAG debuggers or IDEs. One senior engineer scoffed that debuggers were tools for wimps - and something that a good programmer should be able to do without. After all - we had serial ports, GPIOs, and a bi-color LED on our boards. Since these were built into the hardware, they didn't cost us a thing. We also had a single software "build" server that took 5 minutes to build a 32k binary on its best days, so when we had to debug code, it was a painful process of trial and error, with lots of Youtube between iterations. We complained that tens of thousands of dollars of productivity was being flushed away that could have been solved by implementing a proper build server - and while we eventually got our wish, it took far more time than it should have.

Needless to say, software development was painful at that company. We made life hard on ourselves purely out of pride, and for the right to say that we walked "up-hills both ways through 3 feet of snow, everyday". Our code was tied ever-so-tightly to our hardware platform, and the system code was indistinguishable from the application. While we didn't use an RTOS, we had effectively implemented a 3-priority threading scheme using a carefully designed interrupt nesting scheme with event flags and a while(1) superloop running as a background thread. Nothing was abstracted, and the code was always optimized for the platform, presumably in an effort to save on code size and wasted cycles. I asked why we didn't use an RTOS in any of our systems and received dismissive scoffs - the overhead from thread switching and maintaining multiple threads could not be tolerated in our systems according to our chief engineers. In retrospect, our ad-hoc system was likely as large as my smallest kernel, and had just as much context switching (although it was hidden by the compiler).

And every time a new iteration of our product was developed, the firmware took far too long to bring up, because the algorithms and data structures had to be re-tooled to work with the peripherals and sensors attached to the new boards. We worked very hard in an attempt to reinvent the wheel, all in the name of producing "efficient" code.

Regardless, I learned a lot about embedded software development.

Most important, I learned that good design is the key to good software; and good design doesn't have to come at a price. In all but the smallest of projects, the well-designed, well-abstracted code is not only more portable, but it's usually smaller, easier to read, and easier to reuse.

Also, since we had all the time in the world to invest in developing our own tools, I gained a lot of experience building them, and making use of good, free PC tools that could be used to develop and debug a large portion of our code. I ended up writing PC-based device and peripheral simulators, state-machine frameworks, and abstractions for our horrible ad-hoc system code. At the end of the day, I had developed enough tools that I could solve a lot of our development problems without having to re-inventing the wheel at each turn. Gaining a background in how these tools worked gave me a better understanding of how to use them - making me more productive at the jobs that I've had since.

I am convinced that designing good software takes honest effort up-front, and that good application code cannot be written unless it is based on a solid framework. Just as the wise man builds his house on rocks, and not on sand, wise developers write applications based on a well-defined platforms. And while you can probably build a house using nothing but a hammer and sheer will, you can certainly build one a lot faster with all the right tools.

This conviction lead me to development my first RTOS kernel in 2009 - FunkOS. It is a small, yet surprisingly full-featured kernel. It has all the basics (semaphores, mutexes, round-robin and preemptive scheduling), and some pretty advanced features as well (device drivers and other middleware). However, it had two major problems - it doesn't scale well, and it doesn't support many devices.

While I had modest success with this kernel (it has been featured on some blogs, and still gets around 125 downloads a month), it was nothing like the success of other RTOS kernels like uC/OS-II and FreeRTOS. To be honest, as a one-man show, I just don't have the resources to support all of the devices, toolchains, and evaluation boards that a real vendor can. I had never expected my kernel to compete with the likes of them, and I don't expect Mark3 to change the embedded landscape either.

My main goal with Mark3 was to solve the technical shortfalls in the FunkOS kernel by applying my experience in kernel development. As a result, Mark3 is better than FunkOS in almost every way; it scales better, has lower interrupt latency, and is generally more thoughtfully designed (all at a small cost to code size).

Another goal I had was to create something easy to understand, that could be documented and serve as a good introduction to RTOS kernel design. The end result of these goals is the kernel as presented in this book - a full source listing of a working OS kernel, with each module completely documented and explained in detail.

Finally, I wanted to prove that a kernel written entirely in C++ could perform just as well as one written in C. Mark3 is fully benchmarked and profiled – you can see exactly how much it costs to call certain APIs or include various features in the kernel.

And in addition, the code is more readable and easier to understand as a result of making use of object-oriented concepts provided by C++. Applications are easier to write because common concepts are encapsulated into objects (Threads, Semaphores, Mutexes, etc.) with their own methods and data, as opposed to APIs which rely on lots of explicit pointer or handle-passing, type casting, and other operations that are typically considered "unsafe" or "advanced" topics in C.

Chapter 7

When should you use an RTOS?

7.1 The reality of system code

System code can be defined as the program logic required to manage, synchronize, and schedule all of the resources (CPU time, memory, peripherals, etc.) used by the application running on the CPU. And it's true that a significant portion of the code running on an embedded system will be system code. No matter how simple a system is, whether or not this logic is embedded directly into the application (bare-metal system), or included as part of a well-defined stack on which an application is written (RTOS-based); system code is still present, and it comes with a cost.

As an embedded systems is being designed, engineers have to decide which approach to take: Bare-metal, or RTOS. There are advantages and disadvantages to each – and a reasonable engineer should always perform a thorough analysis of the pros and cons of each - in the context of the given application - before choosing a path.

The following figure demonstrates the differences between the architecture of a bare-metal system and RTOS based system at a high level:




Figure 7.1: Arch

As can be seen, the RTOS (And associated middleware + libraries) captures a certain fixed size.

As a generalization, bare-metal systems typically have the advantage in that the system code overhead is small to start – but grows significantly as the application grows in complexity. At a certain point, it becomes extremely difficult and error-prone to add more functionality to an application running on such a system. There's a tipping point, where the cost of the code used to work-around the limitations of a bare-metal system outweigh the cost of a capable RTOS. Bare-metal systems also generally take longer to implement, because the system code has to be written from scratch (or derived from existing code) for the application. The resulting code also tend to be less portable, as it takes serious discipline to keep the system-specific elements of the code separated – in an RTOS-based system, once the kernel and drivers are ported, the application code is generally platform agnostic.

Conversely, an RTOS-based system incurs a slightly higher fixed cost up-front, but scales infinitely better than a bare-metal system as application's complexity increases. Using an RTOS for simple systems reduces application development time, but may cause an application not to fit into some extremely size-constrained microcontroller. An

RTOS can also cause the size of an application to grow more slowly relative to a bare-metal system – especially as a result of applying synchronization mechanisms and judicious IPC. As a result, an RTOS makes it significantly easier to “go agile” with an application – iteratively adding features and functionality, without having to consider refactoring the underlying system at each turn.

Some of these factors may be more important than others. Requirements, specifications, schedules, chip-selection, and volume projections for a project should all be used to feed into the discussions to decide whether or to go bare-metal or RTOS as a result.

Consider the following questions when making that decision:

- What is the application?
- How efficient is efficient enough?
- How fast is fast enough?
- How small is small enough?
- How responsive is responsive enough?
- How much code space/RAM/etc is available on the target system?
- How much code space/RAM do I need for an RTOS?
- How much code space/RAM do I think I'll need for my application?
- How much time do I have to deliver my system?
- How many units do we plan to sell?

7.2 Superloops, and their limitations

7.2.1 Intro to Superloops

Before we start taking a look at designing a real-time operating system, it's worthwhile taking a look through one of the most-common design patterns that developers use to manage task execution in bare-metal embedded systems - Superloops.

Systems based on superloops favor the system control logic baked directly into the application code, usually under the guise of simplicity, or memory (code and RAM) efficiency. For simple systems, superloops can definitely get the job done. However, they have some serious limitations, and are not suitable for every kind of project. In a lot of cases you can squeak by using superloops - especially in extremely constrained systems, but in general they are not a solid basis for reusable, portable code.

Nonetheless, a variety of examples are presented here- from the extremely simple, to cooperative and limited-preemptive multitasking systems, all of which are examples are representative of real-world systems that I've either written the firmware for, or have seen in my experience.

7.2.2 The simplest loop

Let's start with the simplest embedded system design possible - an infinite loop that performs a single task repeatedly:

```
int main()
{
    while(1)
    {
        Do_Something();
    }
}
```


Here, the code inside the loop will run a single function forever and ever. Not much to it, is there? But you might be surprised at just how much embedded system firmware is implemented using essentially the same mechanism - there isn't anything wrong with that, but it's just not that interesting.

Despite its simplicity we can see the beginnings of some core OS concepts. Here, the `while(1)` statement can be logically seen as the operating system kernel - this one control statement determines what tasks can run in the system, and defines the constraints that could modify their execution. But at the end of the day, that's a big part of what a kernel is - a mechanism that controls the execution of application code.

The second concept here is the task. This is application code provided by the user to perform some useful purpose in a system. In this case `Do_something()` represents that task - it could be monitoring blood pressure, reading a sensor and writing its data to a terminal, or playing an MP3; anything you can think of for an embedded system to do. A simple round-robin multi-tasking system can be built off of this example by simply adding additional tasks in sequence in the main `while`-loop. Note that in this example the CPU is always busy running tasks - at no time is the CPU idle, meaning that it is likely burning a lot of power.

While we conceptually have two separate pieces of code involved here (an operating system kernel and a set of running tasks), they are not logically separate. The OS code is indistinguishable from the application. It's like a single-celled organism - everything is crammed together within the walls of an indivisible unit; and specialized to perform its given function relying solely on instinct.

7.2.3 Interrupt-Driven Super-loop

In the previous example, we had a system without any way to control the execution of the task- it just runs forever. There's no way to control when the task can (or more importantly can't) run, which greatly limits the usefulness of the system. Say you only want your task to run every 100 milliseconds - in the previous code, you have to add a hard-coded delay at the end of your task's execution to ensure your code runs only when it should.

Fortunately, there is a much more elegant way to do this. In this example, we introduce the concept of the synchronization object. A Synchronization object is some data structure which works within the bounds of the operating system to tell tasks when they can run, and in many cases includes special data unique to the synchronization event.

There are a whole family of synchronization objects, which we'll get into later. In this example, we make use of the simplest synchronization primitive

- the global flag.

With the addition of synchronization brings the addition of event-driven systems. If you're programming a microcontroller system, you generally have scores of peripherals available to you - timers, GPIOs, ADCs, UARTs, ethernet, USB, etc. All of which can be configured to provide a stimulus to your system by means of interrupts. This stimulus gives us the ability not only to program our micros to `do_something()`, but to `do_something()` if-and-only-if a corresponding trigger has occurred.

The following concepts are shown in the example below:

```
volatile K_BOOL something_to_do = false;

__interrupt__ My_Interrupt_Source(void)
{
    something_to_do = true;
}

int main()
{
    while (1)
    {
        if (something_to_do)
        {
            Do_something();
            something_to_do = false;
        }
        else
        {
            Idle();
        }
    }
}
```

So there you have it - an event driven system which uses a global variable to synchronize the execution of our task based on the occurrence of an interrupt. It's still just a bare-metal, OS-baked-into-the-application system, but it's introduced a whole bunch of added complexity (and control!) into the system.

The first thing to notice in the source is that the global variable, `something_to_do`, is used as a synchronization object. When an interrupt occurs from some external event, triggering the `My_Interrupt_Source()` ISR, program flow in `main()` is interrupted, the interrupt handler is run, and `something_to_do` is set to true, letting us know that when we get back to `main()`, that we should run our `Do_something()` task.

Another new concept at play here is that of the idle function. In general, when running an event driven system, there are times when the CPU has no application tasks to run. In order to minimize power consumption, CPUs usually contain instructions or registers that can be set up to disable non-essential subsets of the system when there's nothing to do. In general, the sleeping system can be re-activated quickly as a result of an interrupt or other external stimulus, allowing normal processing to resume.

Now, we could just call `Do_something()` from the interrupt itself - but that's generally not a great solution. In general, the more time we spend inside an interrupt, the more time we spend with at least some interrupts disabled. As a result, we end up with interrupt latency. Now, in this system, with only one interrupt source and only one task this might not be a big deal, but say that `Do_something()` takes several seconds to complete, and in that time several other interrupts occur from other sources. While executing in our long-running interrupt, no other interrupts can be processed - in many cases, if two interrupts of the same type occur before the first is processed, one of these interrupt events will be lost. This can be utterly disastrous in a real-time system and should be avoided at all costs. As a result, it's generally preferable to use synchronization objects whenever possible to defer processing outside of the ISR.

Another OS concept that is implicitly introduced in this example is that of task priority. When an interrupt occurs, the normal execution of code in `main()` is preempted: control is swapped over to the ISR (which runs to completion), and then control is given back to `main()` where it left off. The very fact that interrupts take precedence over what's running shows that `main` is conceptually a "low-priority" task, and that all ISRs are "high-priority" tasks. In this example, our "high-priority" task is setting a variable to tell our "low-priority" task that it can do something useful. We will investigate the concept of task priority further in the next example.

Preemption is another key principle in embedded systems. This is the notion that whatever the CPU is doing when an interrupt occurs, it should stop, cache its current state (referred to as its context), and allow the high-priority event to be processed. The context of the previous task is then restored its state before the interrupt, and resumes processing. We'll come back to preemption frequently, since the concept comes up frequently in RTOS-based systems.

7.2.4 Cooperative multi-tasking

Our next example takes the previous example one step further by introducing cooperative multi-tasking:

```
// Bitfield values used to represent three distinct tasks
#define TASK_1_EVENT (0x01)
#define TASK_2_EVENT (0x02)
#define TASK_3_EVENT (0x04)

volatile K_UCHAR event_flags = 0;

// Interrupt sources used to trigger event execution

__interrupt__ My_Interrupt_1(void)
{
    event_flags |= TASK_1_EVENT;
}

__interrupt__ My_Interrupt_2(void)
{
    event_flags |= TASK_2_EVENT;
}

__interrupt__ My_Interrupt_3(void)
{
    event_flags |= TASK_3_EVENT;
}

// Main tasks
int main(void)
{

```

```

while(1)
{
    while(event_flags)
    {
        if( event_flags & TASK_1_EVENT)
        {
            Do_Task_1();
            event_flags &= ~TASK_1_EVENT;
        } else if( event_flags & TASK_2_EVENT) {
            Do_Task_2();
            event_flags &= ~TASK_2_EVENT;
        } else if( event_flags & TASK_3_EVENT) {
            Do_Task_3();
            event_flags &= ~TASK_3_EVENT;
        }
    }
    Idle();
}
}

```

This system is very similar to what we had before - however the differences are worth discussing. First, we have stimulus from multiple interrupt sources: each ISR is responsible for setting a single bit in our global event flag, which is then used to control execution of individual tasks from within main().

Next, we can see that tasks are explicitly given priorities inside the main loop based on the logic of the if/else if structure. As long as there is something set in the event flag, we will always try to execute Task1 first, and only when Task1 isn't set will we attempt to execute Task2, and then Task3. This added logic provides the notion of priority. However, because each of these tasks exist within the same context (they're just different functions called from our main control loop), we don't have the same notion of preemption that we have when dealing with interrupts.

That means that even through we may be running Task2 and an event flag for Task1 is set by an interrupt, the CPU still has to finish processing Task2 to completion before Task1 can be run. And that's why this kind of scheduling is referred to as cooperative multitasking: we can have as many tasks as we want, but unless they cooperate by means of returning back to main, the system can end up with high-priority tasks getting starved for CPU time by lower-priority, long-running tasks.

This is one of the more popular Os-baked-into-the-application approaches, and is widely used in a variety of real-time embedded systems.

7.2.5 Hybrid cooperative/preemptive multi-tasking

The final variation on the superloop design utilizes software-triggered interrupts to simulate a hybrid cooperative/preemptive multitasking system. Consider the example code below.

```

// Bitfields used to represent high-priority tasks. Tasks in this group
// can preempt tasks in the group below - but not eachother.
#define HP_TASK_1(0x01)
#define HP_TASK_2(0x02)

volatile K_UCHAR hp_tasks = 0;

// Bitfields used to represent low-priority tasks.
#define LP_TASK_1(0x01)
#define LP_TASK_2(0x02)

volatile K_UCHAR lp_tasks = 0;

// Interrupt sources, used to trigger both high and low priority tasks.
__interrupt__ System_Interrupt_1(void)
{
    // Set any of the other tasks from here...
    hp_tasks |= HP_TASK_1;
    // Trigger the SWI that calls the High_Priority_Tasks interrupt handler
    SWI();
}

__interrupt__ System_Interrupt_n...(void)
{
    // Set any of the other tasks from here...
}

// Interrupt handler that is used to implement the high-priority event context
__interrupt__ High_Priority_Tasks(void)
{

```

```

// Enabled every interrupt except this one
Disable_My_Interrupt();
Enable_Interrupts();
while( hp_tasks)
{
    if( hp_tasks & HP_TASK_1)
    {
        HP_Task1();
        hp_tasks &= ~HP_TASK_1;
    }
    else if (hp_tasks & HP_TASK_2)
    {
        HP_Task2();
        hp_tasks &= ~HP_TASK_2;
    }
}
Restore_Interrupts();
Enable_My_Interrupt();
}

// Main loop, used to implement the low-priority events
int main(void)
{
    // Set the function to run when a SWI is triggered
    Set_SWI(High_Priority_Tasks);

    // Run our super-loop
    while(1)
    {
        while (lp_tasks)
        {
            if (lp_tasks & LP_TASK_1)
            {
                LP_Task1();
                lp_tasks &= ~LP_TASK_1;
            }
            else if (lp_tasks & LP_TASK_2)
            {
                LP_Task2();
                lp_tasks &= ~LP_TASK_2;
            }
        }
        Idle();
    }
}

```

In this example, `High_Priority_Tasks()` can be triggered at any time as a result of a software interrupt (SWI),. When a high-priority event is set, the code that sets the event calls the SWI as well, which instantly preempts whatever is happening in main, switching to the high-priority interrupt handler. If the CPU is executing in an interrupt handler already, the current ISR completes, at which point control is given to the high priority interrupt handler.

Once inside the HP ISR, all interrupts (except the software interrupt) are re-enabled, which allows this interrupt to be preempted by other interrupt sources, which is called interrupt nesting. As a result, we end up with two distinct execution contexts (main and `HighPriorityTasks()`), in which all tasks in the high-priority group are guaranteed to preempt main() tasks, and will run to completion before returning control back to tasks in main(). This is a very basic preemptive multitasking scenario, approximating a "real" RTOS system with two threads of different priorities.

7.3 Problems with superloops

As mentioned earlier, a lot of real-world systems are implemented using a superloop design; and while they are simple to understand due to the limited and obvious control logic involved, they are not without their problems.

7.3.1 Hidden Costs

It's difficult to calculate the overhead of the superloop and the code required to implement workarounds for blocking calls, scheduling, and preemption. There's a cost in both the logic used to implement workarounds (usually involving state machines), as well as a cost to maintainability that comes with breaking up into chunks based on execution time instead of logical operations. In moderate firmware systems, this size cost can exceed the overhead of a reasonably well-featured RTOS, and the deficit in maintainability is something that is measurable in terms of lost productivity through debugging and profiling.

7.3.2 Tightly-coupled code

Because the control logic is integrated so closely with the application logic, a lot of care must be taken not to compromise the separation between application and system code. The timing loops, state machines, and architecture-specific control mechanisms used to avoid (or simulate) preemption can all contribute to the problem. As a result, a lot of superloop code ends up being difficult to port without effectively simulating or replicating the underlying system for which the application was written. Abstraction layers can mitigate the risks, but a lot of care should be taken to fully decouple the application code from the system code.

7.3.3 No blocking Calls

In a super-loop environment, there's no such thing as a blocking call or blocking objects. Tasks cannot stop mid-execution for event-driven I/O from other contexts - they must always run to completion. If busy-waiting and polling are used as a substitute, it increases latency and wastes cycles. As a result, extra code complexity is often times necessary to work-around this lack of blocking objects, often times through implementing additional state machines. In a large enough system, the added overhead in code size and cycles can add up.

7.3.4 Difficult to guarantee responsiveness

Without multiple levels of priority, it may be difficult to guarantee a certain degree of real-time responsiveness without added profiling and tweaking. The latency of a given task in a priority-based cooperative multitasking system is the length of the longest task. Care must be taken to break tasks up into appropriate sized chunks in order to ensure that higher- priority tasks can run in a timely fashion - a manual process that must be repeated as new tasks are added in the system. Once again, this adds extra complexity that makes code larger, more difficult to understand and maintain due to the artificial subdivision of tasks into time-based components.

7.3.5 Limited preemption capability

As shown in the example code, the way to gain preemption in a superloop is through the use of nested interrupts. While this isn't unwieldy for two levels of priority, adding more levels beyond this becomes complicated. In this case, it becomes necessary to track interrupt nesting manually, and separate sets of tasks that can run within given priority loops - and deadlock becomes more difficult to avoid.

Chapter 8

Can you afford an RTOS?

8.1 Intro

Of course, since you're reading the manual for an RTOS that I've been developing over the course of the several years, you can guess that the conclusion that I draw.

If your code is of any sort of non-trivial complexity (say, at least a few- thousand lines), then a more appropriate question would be "can you afford not* to use an RTOS in your system?".

In short, there are simply too many benefits of an RTOS to ignore, the most important being:

Threading, along with priority and time-based scheduling Sophisticated synchronization objects and IPC Flexible, powerful Software Timers Ability to write more portable, decoupled code

Sure, these features have a cost in code space and RAM, but from my experience the cost of trying to code around a lack of these features will cost you as much - if not more. The results are often far less maintainable, error prone, and complex. And that simply adds time and cost. Real developers ship, and the RTOS is quickly becoming one of the standard tools that help keep developers shipping.

One of the main arguments against using an RTOS in an embedded project is that the overhead incurred is too great to be justified. Concerns over "wasted" RAM caused by using multiple stacks, added CPU utilization, and the "large" code footprint from the kernel cause a large number of developers to shun using a preemptive RTOS, instead favoring a non-preemptive, application-specific solution.

I believe that not only is the impact negligible in most cases, but that the benefits of writing an application with an RTOS can lead to savings around the board (code size, quality, reliability, and development time). While these other benefits provide the most compelling case for using an RTOS, they are far more challenging to demonstrate in a quantitative way, and are clearly documented in numerous industry-based case studies.

While there is some overhead associated with an RTOS, the typical arguments are largely unfounded when an RTOS is correctly implemented in a system. By measuring the true overhead of a preemptive RTOS in a typical application, we will demonstrate that the impact to code space, RAM, and CPU usage is minimal, and indeed acceptable for a wide range of CPU targets.

To illustrate just how little an RTOS impacts the size of an embedded software design we will look at a typical microcontroller project and analyze the various types of overhead associated with using a pre-emptive realtime kernel versus a similar non-preemptive event-based framework.

RTOS overhead can be broken into three distinct areas:

- Code space: The amount of code space eaten up by the kernel (static)
- Memory overhead: The RAM associated with running the kernel and application threads.
- Runtime overhead: The CPU cycles required for the kernel's functionality (primarily scheduling and thread switching)

While there are other notable reasons to include or avoid the use of an RTOS in certain applications (determinism,

responsiveness, and interrupt latency among others), these are not considered in this discussion - as they are difficult to consider for the scope of our "canned" application.

8.2 Application description

For the purpose of this comparison, we first create an application using the standard preemptive Mark3 kernel with 2 system threads running: A foreground thread and a background thread. This gives three total priority levels in the system - the interrupt level (high), and two application priority threads (medium and low), which is quite a common paradigm for microcontroller firmware designs. The foreground thread processes a variety of time-critical events at a fixed frequency, while the background thread processes lower priority, aperiodic events. When there are no background thread events to process, the processor enters its low-power mode until the next interrupt is acknowledged.

The contents of the threads themselves are unimportant for this comparison, but we can assume they perform a variety of realtime I/O functions. As a result, a number of device drivers are also implemented.

Code Space and Memory Overhead:

The application is compiled for an ATmega328p processor which contains 32kB of code space in flash, and 2kB of RAM, which is a lower-mid-range microcontroller in Atmel's 8-bit AVR line of microcontrollers. Using the AVR GCC compiler with -Os level optimizations, an executable is produced with the following code/RAM utilization:

```
Program: 27914 bytes
Data:    1313 bytes
```

An alternate version of this project is created using a custom "super-loop" kernel, which uses a single application thread and provides 2 levels of priority (interrupt and application). In this case, the event handler processes the different priority application events to completion from highest to lowest priority.

This approach leaves the application itself largely unchanged. Using the same optimization levels as the preemptive kernel, the code compiles as follows:

```
Program: 24886 bytes
Data:    750 bytes
```

At first glance, the difference in RAM utilization seems quite a lot higher for the preemptive mode version of the application, but the raw numbers don't tell the whole story.

The first issue is that the cooperative-mode total does not take into account the system stack - whereas these values are included in the totals for RTOS version of the project. As a result, some further analysis is required to determine how the stack sizes truly compare.

In cooperative mode, there is only one thread of execution - so considering that multiple event handlers are executed in turn, the stack requirements for cooperative mode is simply determined by those of the most stack-intensive event handler (ignoring stack use contributions due to interrupts).

In contrast, the preemptive kernel requires a separate stack for each active thread, and as a result the stack usage of the system is the sum of the stacks for all threads.

Since the application and idle events are the same for both preemptive and cooperative mode, we know that their (independent) stack requirements will be the same in both cases.

For cooperative mode, we see that the idle thread stack utilization is lower than that of the application thread, and so the application thread's determines the stack size requirement. Again, with the preemptive kernel the stack utilization is the sum of the stacks defined for both threads.

As a result, the difference in overhead between the two cases becomes the extra stack required for the idle thread - which in our case is (a somewhat generous) 128 bytes.

The numbers still don't add up completely, but looking into the linker output we see that the rest of the difference comes from the extra data structures used to manage the kernel in preemptive mode, and the kernel data itself.

Fixed kernel data costs:


```
--- 134 Bytes Kernel data
--- 26 Bytes Kernel Vtables
```

Application (Variable) data costs:

```
--- 24 Bytes Driver Vtables
--- 123 Bytes - statically-allocated kernel objects (semaphores, timers, etc.)
```

With this taken into account, the true memory cost of a 2-thread system ends up being around 428 bytes of RAM - which is about 20% of the total memory available on this particular microcontroller. Whether or not this is reasonable certainly depends on the application, but more importantly, it is not so unreasonable as to eliminate an RTOS-based solution from being considered. Also note that by using the “simulated idle” feature provided in Mark3 R3 and onward, the idle thread (and its associated stack) can be eliminated altogether to reduce the cost in constrained devices.

The difference in code space overhead between the preemptive and cooperative mode solutions is less of an issue. Part of this reason is that both the preemptive and cooperative kernels are relatively small, and even an average target device (like the Atmega328 we’ve chosen) has plenty of room.

Mark3 can be configured so that only features necessary for the application are included in the RTOS - you only pay for the parts of the system that you use. In this way, we can measure the overhead on a feature-by-feature basis, which is shown below for the kernel as configured for this application:

```
Kernel ..... 2563 Bytes
Synchronization Objects. 644 Bytes
Port ..... 974 Bytes
Features ..... 871 Bytes
```

The configuration tested in this comparison uses the thread/port module with timers, drivers, and semaphores, and mutexes, for a total kernel size of 5052 Bytes, with the rest of the code space occupied by the application.

As can be seen from the compiler’s output, the difference in code space between the two versions of the application is 3028 bytes - or about 9% of the available code space on the selected processor. While nearly all of this comes from the added overhead of the kernel, the rest of the difference comes the changes to the application necessary to facilitate the different frameworks. This also demonstrates that the system-software code size in the cooperative case is about 2024 bytes.

8.3 Runtime Overhead

On the cooperative kernel, the overhead associated with running the thread is the time it takes the kernel to notice a pending event flag and launch the appropriate event handler, plus the timer interrupt execution time.

Similarly, on the preemptive kernel, the overhead is the time it takes to switch contexts to the application thread, plus the timer interrupt execution time.

The timer interrupt overhead is similar for both cases, so the overhead then becomes the difference between the following:

Preemptive mode:

- Posting the semaphore that wakes the high-priority thread
- Performing a context switch to the high-priority thread

Cooperative mode:

- Setting the event flag from the timer interrupt
- Acknowledging the event from the event loop

coop – 438 cycles preempt – 764 cycles

Using a cycle-accurate AVR simulator (fIAVR) running with a simulated speed of 16MHz, we find the end-to-end event sequence time to be 27us for the cooperative mode scheduler and 48us for the preemptive, and a raw difference of 20us.

With a fixed high-priority event frequency of 30Hz, we achieve a runtime overhead of 611us per second, or 0.06% of the total available CPU time. Now, obviously this value would expand at higher event frequencies and/or slower CPU frequencies, but for this typical application we find the difference in runtime overhead to be negligible for a preemptive system.

8.4 Analysis

For the selected test application and platform, including a preemptive RTOS is entirely reasonable, as the costs are low relative to a non-preemptive kernel solution. But these costs scale relative to the speed, memory and code space of the target processor. Because of these variables, there is no "magic bullet" environment suitable for every application, but Mark3 attempts to provide a framework suitable for a wide range of targets.

On the one hand, if these tests had been performed on a higher-end microcontroller such as the ATmega1284p (containing 128kB of code space and 16kB of RAM), the overhead would be in the noise. For this type of resource-rich microcontroller, there would be no reason to avoid using the Mark3 preemptive kernel.

Conversely, using a lower-end microcontroller like an ATmega88pa (which has only 8kB of code space and 1kB of RAM), the added overhead would likely be prohibitive for including a preemptive kernel. In this case, the cooperative-mode kernel would be a better choice.

As a rule of thumb, if one budgets 25% of a microcontroller's code space/RAM for system code, you should only require at minimum a microcontroller with 16kB of code space and 2kB of RAM as a base platform for an RTOS. Unless there are serious constraints on the system that require much better latency or responsiveness than can be achieved with RTOS overhead, almost any modern platform is sufficient for hosting a kernel. In the event you find yourself with a microprocessor with external memory, there should be no reason to avoid using an RTOS at all.

Chapter 9

Mark3 Design Goals

9.1 Overview

9.1.1 Services Provided by an RTOS Kernel

At its lowest-levels, an operating system kernel is responsible for managing and scheduling resources within a system according to the application. In a typical thread-based RTOS, the resources involved is CPU time, and the kernel manages this by scheduling threads and timers. But capable RTOS kernels provide much more than just threading and timers.

In the following section, we discuss the Mark3 kernel architecture, all of its features, and a thorough discussion of how the pieces all work together to make an awesome RTOS kernel.

9.1.2 Guiding Principles of Mark3

Mark3 was designed with a number of over-arching principles, coming from years of experience designing, implementing, refining, and experimenting with RTOS kernels. Through that process I not only discovered what features I wanted in an RTOS, but how I wanted to build those features to look, work, and “feel”. With that understanding, I started with a clean slate and began designing a new RTOS. Mark3 is the result of that process, and its design goals can be summarized in the following guiding principles.

9.1.3 Be feature competitive

To truly be taken seriously as more than just a toy or educational tool, an RTOS needs to have a certain feature suite. While Mark3 isn't a clone of any existing RTOS, it should at least attempt parity with the most common software in its class.

Looking at its competitors, Mark3 as a kernel supports most, if not all of the compelling features found in modern RTOS kernels, including dynamic threads, tickless timers, efficient message passing, and multiple types of synchronization primitives.

9.1.4 Be highly configuration

Mark3 isn't a one-size-fits-all kernel – and as a result, it provides the means to build a custom kernel to suit your needs. By configuring the kernel at compile-time, Mark3 can be built to contain the optimal feature set for a given application. And since features can be configured individually, you only pay the code/RAM footprint for the features you actually use.

9.1.5 No external dependencies, no new language features

To maximize portability and promote adoption to new platforms, Mark3 is written in a widely supported subset of C++ that lends itself to embedded applications. It avoids RTTI, exceptions, templates, and libraries (C standard, STL, etc.), with all fundamental data structures and types implemented completely for use by the kernel. As a result, the portable parts of Mark3 should compile for any capable C++ toolchain.

9.1.6 Target the most popular hobbyist platforms available

Realistically, this means supporting the various Arduino-compatible target CPUs, including AVR and ARM Cortex-M series microcontrollers. As a result, the current default target for Mark3 is the atmega328p, which has 32KB of flash and 2KB of RAM. All decisions regarding default features, code size, and performance need to take that target system into account.

Mark3 integrates cleanly as a library into the Arduino IDE to support atmega328-based targets. Other AVR and Cortex-M targets can be supported using the port code provided in the source package.

9.1.7 Maximize determinism – but be pragmatic

Guaranteeing deterministic and predictable behavior is tough to do in an embedded system, and often comes with a heavy price tag in either RAM or code-space. With Mark3, we strive to keep the core kernel APIs and features as lightweight as possible, while avoiding algorithms that don't scale to large numbers of threads. We also achieve minimal latency by keeping interrupts enabled (operating out of the critical section) wherever possible.

In Mark3, the most important parts of the kernel are fixed-time, including thread scheduling and context switching. Operations that are not fixed time can be characterized as a function of their dependent data. For instances, the [Mutex](#) and [Semaphore](#) APIs operate in fixed time in the uncontested case, and execute in linear time for the contested case – where the speed of execution is dependent on the number of threads currently waiting on that object.

The caveat here is that while we want to minimize latency and time spent in critical sections, that has to be balanced against increases in code size, and uncontested-case performance.

9.1.8 Apply engineering principles – and that means discipline, measurement and verification

My previous RTOS, FunkOS, was designed to be very ad-hoc. The usage instructions were along the lines of “drag and drop the source files into your IDE and compile”. There was no regression/unit testing, no code size/speed profiling, and all documentation was done manually. It worked, but the process was a bit of a mess, and resulted in a lot of re-spins of the software, and a lot of time spent stepping through emulators to measure parameters.

We take a different approach in Mark3. Here, we've designed not only the kernel-code, but the build system, unit tests, profiling code, documentation and reporting that supports the kernel. Each release is built and tested using automation in order to ensure quality and correctness, with supporting documentation containing all critical metrics. Only code that passes testing is submitted to the repos and public forums for distribution. These metrics can be traced from build-to-build to ensure that performance remains consistent from one drop to the next, and that no regressions are introduced by new/refactored code.

And while the kernel code can still be exported into an IDE directly, that takes place with the knowledge that the kernel code has already been rigorously tested and profiled. Exporting source in Mark3 is also supported by scripting to ensure reliable, reproducible results without the possibility for human-error.

9.1.9 Use Virtualization For Verification

Mark3 was designed to work with automated simulation tools as the primary means to validate changes to the kernel, due to the power and flexibility of automatic tests on virtual hardware. I was also intrigued by the thought of extending the virtual target to support functionality useful for a kernel, but not found on real hardware.

When the project was started, simavr was the tool of choice- however, its simulation was found to be incorrect compared to execution on a real MCU, and it did not provide the degree of extension that I desired for use with kernel development.

The flAVR AVR simulator was written to replace the dependency on that tool, and overcome those limitations. It also provides a GDB interface, as well as its own built-in debugger, profilers, and trace tools.

The example and test code relies heavily on flAVR kernel aware messaging, so it is recommended that you familiarize yourself with that tool if you intend to do any sort of customizations or extensions to the kernel.

flAVR is hosted on sourceforge at <http://www.sourceforge.net/projects/flavr/> . In its basic configuration, it builds with minimal external dependencies.

- On linux, it requires only pthreads.
- On Windows, it requires pthreads and ws2_32, both satisfied via MinGW.
- Optional SDL builds for both targets (featuring graphics and simulated joystick input) can be built, and rely on libSDL.

Chapter 10

Mark3 Kernel Architecture

10.1 Overview

At a high level, the Mark3 RTOS is organized into the following features, and layered as shown below:

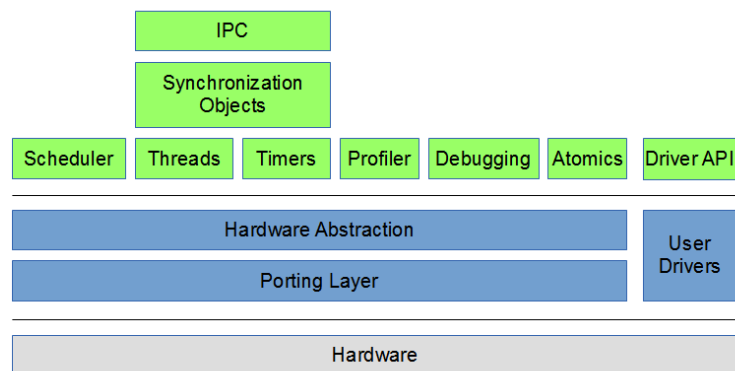


Figure 10.1: Overview

Everything in the “green” layer represents the Mark3 public API and classes, beneath which lives all hardware abstraction and CPU-specific porting and driver code, which runs on a given target CPU.

The features and concepts introduced in this diagram can be described as follows:

Threads: The ability to multiplex the CPU between multiple tasks to give the perception that multiple programs are running simultaneously. Each thread runs in its own context with its own stack.

Scheduler: Algorithm which determines the thread that gets to run on the CPU at any given time. This algorithm takes into account the priorities (and other execution parameters) associated with the threads in the system.

IPC: Inter-process-communications. Message-passing and [Mailbox](#) interfaces used to communicate between threads synchronously or asynchronously.

Synchronization Objects: Ability to schedule thread execution relative to system conditions and events, allowing for sharing global data and resources safely and effectively.

Timers: High-resolution software timers that allow for actions to be triggered on a periodic or one-shot basis.

Profiler: Special timer used to measure the performance of arbitrary blocks of code.

Debugging: Realtime logging and trace functionality, facilitating simplified debugging of systems using the OS.

Atomsics: Support for UN-interruptible arithmetic operations.

Driver API: Hardware abstraction interface allowing for device drivers to be written in a consistent, portable manner.

Hardware Abstraction Layer: Class interface definitions to represent threading, context-switching, and timers in a generic, abstracted manner.

Porting Layer: Class interface implementation to support threading, context-switching, and timers for a given CPU.

User Drivers: Code written by the user to implement device-specific peripheral drivers, built to make use of the Mark3 driver API.

Each of these features will be described in more detail in the following sections of this chapter.

The concepts introduced in the above architecture are implemented in a variety of source modules, which are logically broken down into classes (or in some cases, groups of functions/macros). The relationship between objects in the Mark3 kernel is shown below:

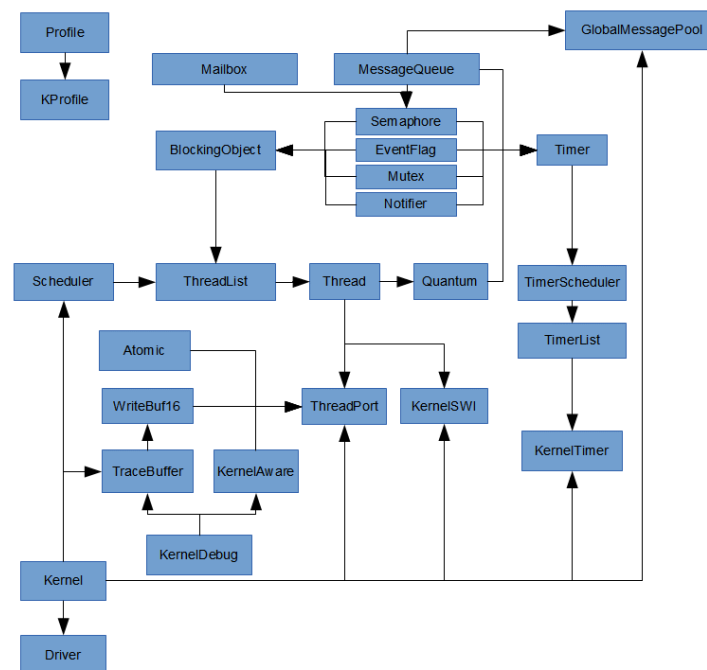


Figure 10.2: Overview

The objects shown in the preceding table can be grouped together by feature. In the table below, we group each feature by object, referencing the source module in which they can be found in the Mark3 source tree.

Feature	Kernel Object	Source Files
Profiling	ProfileTimer	profile.cpp/.h
Threads + Scheduling	Thread	thread.cpp/.h
	Scheduler	scheduler.cpp/.h
	PriorityMap	priomap.cpp/.h
	Quantum	quantum.cpp/.h
	ThreadPort	threadport.cpp/.h **
	KernelSWI	kernelswi.cpp/.h **
Timers	Timer	timer.h/timer.cpp
	TimerScheduler	timerscheduler.h
	TimerList	timerlist.h/cpp

Let's look at a very basic CPU without any sort of special multi-threading hardware, and without interrupts. When the CPU is powered up, the program counter is loaded with some default location, at which point the processor core will start executing instructions sequentially - running forever and ever according to whatever has been loaded into program memory. This single instance of a simple program sequence is the only thing that runs on the processor, and the execution of the program can be predicted entirely by looking at the CPU's current register state, its program, and any affected system memory (the CPU's "context").

It's simple enough, and that's exactly the definition we have for a thread in an RTOS.

Each thread contains an instance of a CPU's register context, its own stack, and any other bookkeeping information necessary to define the minimum unique execution state of a system at runtime. It is the job of a RTOS to multiplex the execution of multiple threads on a single physical CPU, thereby creating the illusion that many programs are being executed simultaneously. In reality there can only ever be one thread truly executing at any given moment on a CPU core, so it's up to the scheduler to set and enforce rules about what thread gets to run when, for how long, and under what conditions. As mentioned earlier, any system without an RTOS executes as a single thread, so at least two threads are required for an RTOS to serve any useful purpose.

Note that all of this information is common to pretty well every RTOS in existence - the implementation details, including the scheduler rules, are all part of what differentiates one RTOS from another.

10.2.2 Thread States and ThreadLists

Since only one thread can run on a CPU at a time, the scheduler relies on thread information to make its decisions. Mark3's scheduler relies on a variety of such information, including:

- The thread's current priority
- Round-Robin execution quanta
- Whether or not the thread is blocked on a synchronization object, such as a mutex or semaphore
- Whether or not the thread is currently suspended

The scheduler further uses this information to logically place each thread into 1 of 4 possible states:

- Ready - The thread is currently running
- Running - The thread is able to run
- Blocked - The thread cannot run until a system condition is met
- Stopped - The thread cannot run because its execution has been suspended

In order to determine a thread's state, threads are placed in "buckets" corresponding to these states. Ready and running threads exist in the scheduler's buckets, blocked threads exist in a bucket belonging to the object they're blocked on, and stopped threads exist in a separate bucket containing all stopped threads.

In reality, the various buckets are just doubly-linked lists of [Thread](#) objects - implemented in something called the [ThreadList](#) class. To facilitate this, the [Thread](#) class directly inherits from a [LinkListNode](#) class, which contains the node pointers required to implement a doubly-linked list. As a result, Threads may be effortlessly moved from one state to another using efficient linked-list operations built into the [ThreadList](#) class.

10.2.3 Blocking and Unblocking

While many developers new to the concept of an RTOS assume that all threads in a system are entirely separate from each other, the reality is that practical systems typically involve multiple threads working together, or at the very least sharing resources. In order to synchronize the execution of threads for that purpose, a number of synchronization primitives (blocking objects) are implemented to create specific sets of conditions under which threads can continue execution. The concept of "blocking" a thread until a specific condition is met is fundamental to understanding RTOS applications design, as well as any highly-multithreaded applications.

10.2.4 Blocking Objects

Blocking objects and primitives provided by Mark3 include:

- Semaphores (binary and counting)
- Mutexes
- Event Flags
- Thread Notification Objects
- Thread Sleep
- Message Queues
- Mailboxes

The relationship between these objects in the system are shown below:

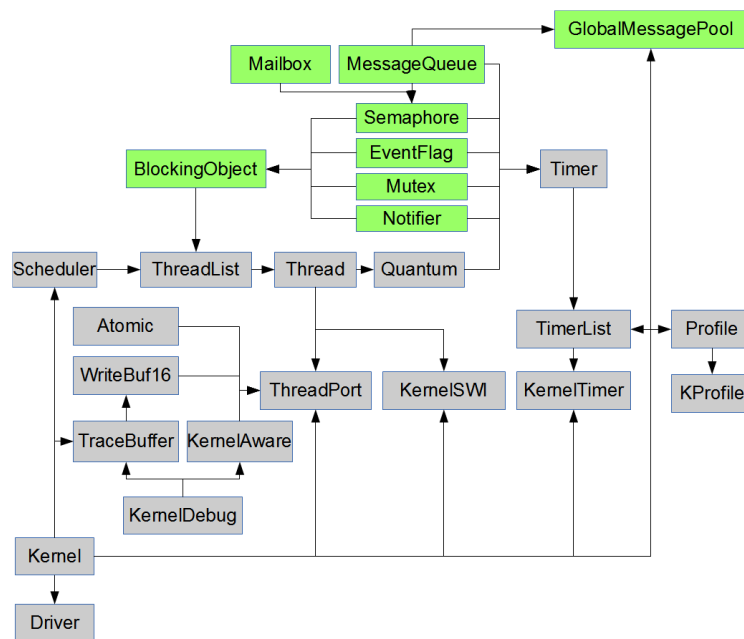


Figure 10.4: Blocking Objects

Each of these objects inherit from the [BlockingObject](#) class, which itself contains a [ThreadList](#) object. This class contains methods to `Block()` a thread (remove it from the [Scheduler](#)'s "Ready" or "Running" ThreadLists), as well as `UnBlock()` a thread (move a thread back to the "Ready" lists). This object handles transitioning threads from list-to-list (and state-to-state), as well as taking care of any other [Scheduler](#) bookkeeping required in the process. While each of the Blocking types implement a different condition, they are effectively variations on the same theme. Many simple Blocking objects are also used to build complex blocking objects - for instance, the [Thread](#) Sleep mechanism is essentially a binary semaphore and a timer object, while a message queue is a linked-list of message objects combined with a semaphore.

10.3 Inside the Mark3 Scheduler

At this point we've covered the following concepts:

- Threads
- [Thread](#) States and [Thread](#) Lists
- Blocking and Un-Blocking Threads

Thankfully, this is all the background required to understand how the Mark3 [Scheduler](#) works. In technical terms, Mark3 implements "strict priority scheduling, with round-robin scheduling among threads in each priority group". In plain English, this boils down to a scheduler which follows a few simple rules:

```
Find the highest-priority "Ready" list that has at least one Threads.
If the first thread in that bucket is not the current thread, select it
to run next. Otherwise, rotate the linked list, and choose the next
thread in the list to run
```

Since context switching is one of the most common and frequent operation performed by an RTOS, this needs to be as fast and deterministic as possible. While the logic is simple, a lot of care must be put into optimizing the scheduler to achieve those goals. In the section below we discuss the optimization approaches taken in Mark3.

There are a number of ways to find the highest-priority thread. The naive approach would be to simply iterate through the scheduler's array of ThreadLists from highest to lowest, stopping when the first non-empty list is found, such as in the following block of code:

```
for (prio = num_prio - 1; prio >= 0; prio--)
{
    if (thread_list[prio].get_head() != NULL)
    {
        break;
    }
}
```

While that would certainly work and be sufficient for a variety of systems, it's a non-deterministic approach (complexity $O(n)$) whose cost varies substantially based on how many priorities have to be evaluated. It's simple to read and understand, but it's non-optimal.

Fortunately, a functionally-equivalent and more deterministic approach can be implemented with a few tricks.

In addition to maintaining an array of ThreadLists, Mark3 also maintains a bitmap (one bit per priority level) that indicates which thread lists have ready threads. This bitmap is maintained automatically by the [ThreadList](#) class, and is updated every time a thread is moved to/from the [Scheduler's](#) ready lists.

By inspecting this bitmap using a technique to count the leading zero bits in the bitmap, we determine which threadlist to choose in fixed time.

Now, to implement the leading-zeros check, this can once again be performed iteratively using bitshifts and compares (which isn't any more efficient than the raw list traversal), but it can also be evaluated using either a lookup table, or via a special CPU instruction to count the leading zeros in a value. In Mark3, we opt for the lookup-table approach since we have a limited number of priorities and not all supported CPU architectures support a count leading zero instruction. To achieve a balance between performance and memory use, we use a 4-bit lookup table (costing 16 bytes) to perform the lookup.

(As a sidenote - this is actually a very common approach in OS schedulers. It's actually part of the reason why modern ARM cores implement a dedicated count-leading-zeros [CLZ] instruction!)

With a 4-bit lookup table and an 8-bit priority-level bitmap, the priority check algorithm looks something like this:

```
// Check the highest 4 priority levels, represented in the
// upper 4 bits in the bitmap
priority = priority_lookup_table[(priority_bitmap >> 4)];

// priority is non-zero if we found something there
if( priority )
```

```

{
    // Add 4 because we were looking at the higher levels
    priority += 4;
}
else
{
    // Nothing in the upper 4, look at the lowest 4 priority levels
    // represented by the lowest 4 bits in the bitmap
    priority = priority_lookup_table[priority_bitmap & 0x0F];
}

```

Deconstructing this algorithm, you can see that the priority lookup will have an $O(1)$ complexity - and is extremely low-cost.

This operation is thus fully deterministic and time bound - no matter how many threads are scheduled, the operation will always be time-bound to the most expensive of these two code paths. Even with only 8 priority levels, this is still much faster than iteratively checking the thread lists manually, compared with the previous example implementation.

Once the priority level has been found, selecting the next thread to run is trivial, consisting of something like this:

```
next_thread = thread_list[prio].get_head();
```

In the case of the `get_head()` calls, this evaluates to an inline-load of the "head" pointer in the particular thread list.

One important thing to take away from this analysis is that the scheduler is only responsible for selecting the next-to-run thread. In fact, these two operations are totally decoupled - no context switching is performed by the scheduler, and the scheduler isn't called from the context switch. The scheduler simply produces new "next thread" values that are consumed from within the context switch code.

10.3.1 Considerations for Round-Robin Scheduling

One thing that isn't considered directly from the scheduler algorithm is the problem of dealing with multiple threads within a single priority group; all of the algorithms that have been explored above simply look at the first [Thread](#) in each group.

Mark3 addresses this issue indirectly, using a software timer to manage round-robin scheduling, as follows.

In some instances where the scheduler is run by the kernel directly (typically as a result of calling [Thread::Yield\(\)](#)), the kernel will perform an additional check after running the [Scheduler](#) to determine whether or there are multiple ready [Threads](#) in the priority of the next ready thread.

If there are multiple threads within that priority, the kernel adds a one-shot software timer which is programmed to expire at the next [Thread](#)'s configured quantum. When this timer expires, the timer's callback function executes to perform two simple operations:

"Pivot" the current [Thread](#)'s priority list. Set a flag telling the kernel to trigger a Yield after exiting the main [Timer](#)↔[Scheduler](#) processing loop

Pivoting the thread list basically moves the head of a circular-linked-list to its next value, which in our case ensures that a new thread will be chosen the next time the scheduler is run (the scheduler only looks at the head node of the priority lists). And by calling Yield, the system forces the scheduler to run, a new round-robin software timer to be installed (if necessary), and triggers a context switch SWI to load the newly-chosen thread. Note that if the thread attached to the round-robin timer is pre-empted, the kernel will take steps to abort and invalidate that round-robin software timer, installing a new one tied to the next thread to run if necessary.

Because the round-robin software timer is dynamically installed when there are multiple ready threads at the highest ready priority level, there is no CPU overhead with this feature unless that condition is met. The cost of round-robin scheduling is also fixed - no matter how many threads there are, and the cost is identical to any other one-shot software timer in the system.

10.3.2 Context Switching

There's really not much to say about the actual context switch operation at a high level. Context switches are triggered whenever it has been determined that a new thread needs to be swapped into the CPU core when the scheduler is run. Mark3 implements also context switches as a call to a software interrupt - on AVR platforms, we

typically use INT0 or INT2 for this (although any pin-change GPIO interrupt can be used), and on ARM we achieve this by triggering a PendSV exception.

However, regardless of the architecture, the context-switch ISR will perform the following three operations:

Save the current [Thread](#)'s context to the current [Thread](#) stack
Make the "next to run" thread the "currently running" thread
Restore the context of the next [Thread](#) from the [Thread](#) stack

The code to implement the context switch is entirely architecture-specific, so it won't be discussed in detail here. It's almost always gory inline-assembly which is used to load and store various CPU registers, and is highly-optimized for speed. We dive into an example implementation for the ARM Cortex-M0 microcontroller in a later section of this book.

10.3.3 Putting It All Together

In short, we can say that the Mark3 scheduler works as follows:

- The scheduler is run whenever a [Thread::Yield\(\)](#) is called by a user, as part of blocking calls, or whenever a new thread is started
- The Mark3 scheduler is deterministic, selecting the next thread to run in fixed-time
- The scheduler only chooses the next thread to run, the context switch SWI consumes that information to get that thread running
- Where there are multiple ready threads in the highest populated priority level, a software timer is used to manage round-robin scheduling

While we've covered a lot of ground in this section, there's not a whole lot of code involved. However, the code that performs these operations is nuanced and subtle. If you're interested in seeing how this all works in practice, I suggest reading through the Mark3 source code (which is heavily annotated), and stepping through the code with a simulator/emulator.

10.4 Timers

Mark3 implements one-shot and periodic software-timers via the [Timer](#) class. The user configures the timer for duration, repetition, and action, at which point the timer can be activated. When an active timer expires, the kernel calls a user-specified callback function, and then reloads the timer in the case of periodic timers. The same timer objects exposed to the user are also used within the kernel to implement round-robin scheduling, and timeout-based APIs for semaphores, mutexes, events, and messages.

Timers are implemented using the following components in the Mark3 [Kernel](#):

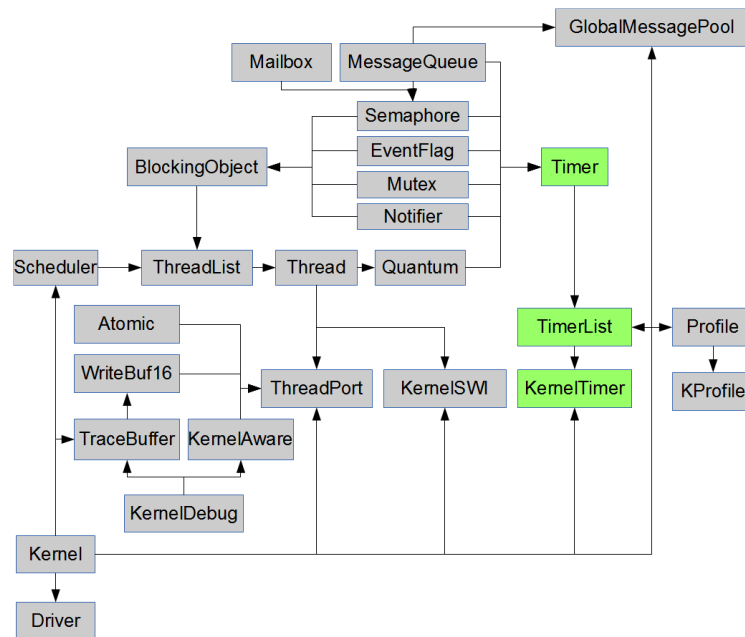


Figure 10.5: Timers

The [Timer](#) class provides the basic periodic and one-shot timer functionality used by application code, blocking objects, and IPC.

The [TimerList](#) class implements a doubly-linked list of [Timer](#) objects, and the logic required to implement a timer tick (tick-based kernel) or timer expiry (tickless kernel) event.

The [TimerScheduler](#) class contains a single [TimerList](#) object, implementing a single, system-wide list of [Timer](#) objects within the kernel. It also provides hooks for the hardware timer, such that when a timer tick or expiry event occurs, the [TimerList](#) expiry handler is run.

The [KernelTimer](#) class ([kerneltimer.cpp/.h](#)) implements the CPU specific hardware timer driver that is used by the kernel and the [TimerScheduler](#) to implement software timers.

While extremely simple to use, they provide one of the most powerful execution contexts in the system.

The software timers implemented in Mark3 use interrupt-nesting within the kernel timer's interrupt handler. This context is considered higher-priority than the highest priority user thread, but lower-priority than other interrupts in the system. As a result, this minimizes critical interrupt latency in the system, albeit at the expense of responsiveness of the user-threads.

For this reason, it's critical to ensure that all timer callback events are kept as short as possible to prevent adding thread-level latency. All heavy-lifting should be left to the threads, so the callback should only implement signalling via IPC or synchronization object.

The time spent in this interrupt context is also dependent on the number of active timers at any given time. However, Mark3 also can be used to minimize the frequency of these interrupts wakeups, by using an optional "tolerance" parameter in the timer API calls. In this way, periodic tasks that have less rigorous real-time constraints can all be grouped together – executing as a group instead of one-after-another.

Mark3 also contains two different timer implementations that can be configured at build-time, each with their own advantages.

10.4.1 Tick-based Timers

In a tick-based timing scheme, the kernel relies on a system-timer interrupt to fire at a relatively-high frequency, on which all kernel timer events are derived. On modern CPUs and microcontrollers, a 1kHz system tick is common, although quite often lower frequencies such as 60Hz, 100Hz, or 120Hz are used. The resolution of this timer also defines the maximum resolution of timer objects as a result. That is, if the timer frequency is 1kHz, a user cannot specify a timer resolution lower than 1ms.

The advantage of a tick-based timer is its sheer simplicity. It typically doesn't take much to set up a timer to trigger an interrupt at a fixed-interval, at which point, all system timer intervals are decremented by 1 count. When each system timer interval reaches zero, a callback is called for the event, and the events are either reset and restarted (repeated timers) or cleared (1-shot).

Unfortunately, that simplicity comes at a cost of increased interrupt count, which cause frequent CPU wakeups and utilization, and power consumption.

10.4.2 Tickless Timers

In a tickless system, the kernel timer only runs when there are active timers pending expiry, and even then, the timer module only generates interrupts when a timer expires, or a timer reaches its maximum count value. Additionally, when there are no active timer objects, the timer can be completely disabled – saving even more cycles, power, and CPU wakeups. These factors make the tickless timer approach a highly-optimal solution, suitable for a wide array of low-power applications.

Also, since tickless timers do not rely on a fixed, periodic clock, they can potentially be higher resolution. The only limitation in timer resolution is the precision of the underlying hardware timer as configured. For example, if a 32kHz hardware timer is being used to drive the timer scheduler, the resolution of timer objects would be in the $\sim 33\mu\text{s}$ range.

The only downside of the tickless timer system is an added complexity to the timer code, requiring more code space, and slightly longer execution of the timer routines when the timer interrupt is executed.

10.4.3 Timer Processing Algorithm

Timer interrupts occur at either a fixed-frequency (tick-based), or at the next timer expiry interval (tickless), at which point the timer processing algorithm runs. While the timer count is reset by the timer-interrupt, it is still allowed to accumulate ticks while this algorithm is executed in order to ensure that timer-accuracy is kept in real-time. It is also important to note that round-robin scheduling changes are disabled during the execution of this algorithm to prevent race conditions, as the round-robin code also relies on timer objects.

All active timer objects are stored in a doubly-linked list within the timer-scheduler, and this list is processed in two passes by the algorithm which runs from the timer-interrupt (with interrupt nesting enabled). The first pass determines which timers have expired and the next timer interval, while the second pass deals with executing the timer callbacks themselves. Both phases are discussed in more detail below.

In the first pass, the active timers are decremented by either 1 tick (tick-based), or by the duration of the last elapsed timer interval (tickless). Timers that have zero (or less-than-zero) time remaining have a "callback" flag set, telling the algorithm to call the timer's callback function in the second pass of the loop. In the event of a periodic timer, the timer's interval is reset to its starting value.

For the tickless case, the next timer interval is also computed in the first-pass by looking for the active timer with the least amount of time remaining in its interval. Note that this calculation is irrelevant in the tick-based timer code, as the timer interrupt fires at a fixed-frequency.

In the second pass, the algorithm loops through the list of active timers, looking for those with their "callback" flag set in the first pass. The callback function is then executed for each expired timer, and the "callback" flag cleared. In the event that a non-periodic (one-shot) timer expires, the timer is also removed from the timer scheduler at this time.

In a tickless system, once the second pass of the loop has been completed, the hardware timer is checked to see if the next timer interval has expired while processing the expired timer callbacks. In that event, the complete

algorithm is re-run to ensure that no expired timers are missed. Once the algorithm has completed without the next timer expiring during processing, the expiry time is programmed into the hardware timer. Round-robin scheduling is re-enabled, and if a new thread has been scheduled as a result of action taken during a timer callback, a context switch takes place on return from the timer interrupt.

10.5 Synchronization and IPC

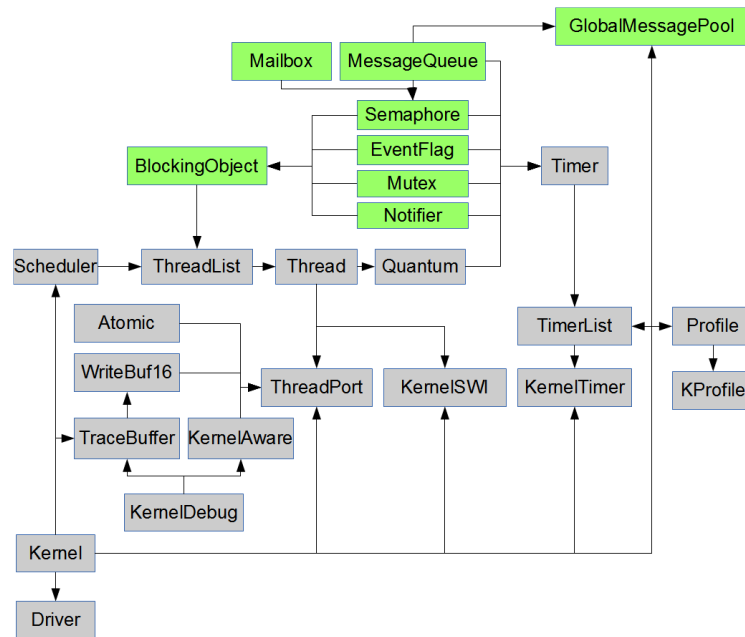


Figure 10.6: Synchronization and IPC

10.6 Blocking Objects

A Blocking object in Mark3 is essentially a thread list. Any blocking object implementation (being a semaphore, mutex, event flag, etc.) can be built on top of this class, utilizing the provided functions to manipulate thread location within the [Kernel](#).

Blocking a thread results in that thread becoming de-scheduled, placed in the blocking object's own private list of threads which are waiting on the object.

Unblocking a thread results in the reverse: The thread is moved back to its original location from the blocking list.

The only difference between a blocking object based on this class is the logic used to determine what constitutes a Block or Unblock condition.

For instance, a semaphore Pend operation may result in a call to the Block() method with the currently-executing thread in order to make that thread wait for a semaphore Post. That operation would then invoke the Unblock() method, removing the blocking thread from the semaphore's list, and back into the appropriate thread inside the scheduler.

Care must be taken when implementing blocking objects to ensure that critical sections are used judiciously, otherwise asynchronous events like timers and interrupts could result in non-deterministic and often catastrophic behavior.

Mark3 implements a variety of blocking objects including semaphores, mutexes, event flags, and IPC mechanisms that all inherit from the basic Blocking-object class found in [blocking.h/cpp](#), ensuring consistency and a high degree of code-reuse between components.

10.6.1 Semaphores

Semaphores are used to synchronized execution of threads based on the availability (and quantity) of application-specific resources in the system. They are extremely useful for solving producer-consumer problems, and are the method-of-choice for creating efficient, low latency systems, where ISRs post semaphores that are handled from within the context of individual threads. Semaphores can also be posted (but not pended) from within the interrupt context.

10.6.2 Mutex

Mutexes (Mutual exclusion objects) are provided as a means of creating "protected sections" around a particular resource, allowing for access of these objects to be serialized. Only one thread can hold the mutex at a time

- other threads have to wait until the region is released by the owner thread before they can take their turn operating on the protected resource. Note that mutexes can only be owned by threads - they are not available to other contexts (i.e. interrupts). Calling the mutex APIs from an interrupt will cause catastrophic system failures.

Note that these objects are recursive in Mark3 - that is, the owner thread can claim a mutex more than once. The caveat here is that a recursively-held mutex will not be released until a matching "release" call is made for each "claim" call.

Priority inheritance is provided with these objects as a means to avoid priority inversions. Whenever a thread at a priority than the mutex owner blocks on a mutex, the priority of the current thread is boosted to the highest-priority waiter to ensure that other tasks at intermediate priorities cannot artificially prevent progress from being made.

10.6.3 Event Flags

Event Flags are another synchronization object, conceptually similar to a semaphore.

Unlike a semaphore, however, the condition on which threads are unblocked is determined by a more complex set of rules. Each Event Flag object contains a 16-bit field, and threads block, waiting for combinations of bits within this field to become set.

A thread can wait on any pattern of bits from this field to be set, and any number of threads can wait on any number of different patterns. Threads can wait on a single bit, multiple bits, or bits from within a subset of bits within the field.

As a result, setting a single value in the flag can result in any number of threads becoming unblocked simultaneously. This mechanism is extremely powerful, allowing for all sorts of complex, yet efficient, thread synchronization schemes that can be created using a single shared object.

Note that Event Flags can be set from interrupts, but you cannot wait on an event flag from within an interrupt.

10.6.4 Notification Objects

Notification objects are the most lightweight of all blocking objects supplied by Mark3.

using this blocking primitive, one or more threads wait for the notification object to be signalled by code elsewhere in the system (i.e. another thread or interrupt). Once the notification has been signalled, all threads currently blocked on the object become unblocked and moved into the ready list.

Signalling a notification object that has no actively-waiting threads has no effect.

10.7 Messages and Global Message Queue

10.7.1 Messages

Sending messages between threads is the key means of synchronizing access to data, and the primary mechanism to perform asynchronous data processing operations.

Sending a message consists of the following operations:

- Obtain a [Message](#) object from the global message pool
- Set the message data and event fields
- Send the message to the destination message queue

While receiving a message consists of the following steps:

- Wait for a messages in the destination message queue
- Process the message data
- Return the message back to the global message pool

These operations, and the various data objects involved are discussed in more detail in the following section.

10.7.2 Message Objects

[Message](#) objects are used to communicate arbitrary data between threads in a safe and synchronous way.

The message object consists of an event code field and a data field. The event code is used to provide context to the message object, while the data field (essentially a void * data pointer) is used to provide a payload of data corresponding to the particular event.

Access to these fields is marshalled by accessors - the transmitting thread uses the `SetData()` and `SetCode()` methods to seed the data, while the receiving thread uses the `GetData()` and `GetCode()` methods to retrieve it.

By providing the data as a void data pointer instead of a fixed-size message, we achieve an unprecedented measure of simplicity and flexibility. Data can be either statically or dynamically allocated, and sized appropriately for the event without having to format and reformat data by both sending and receiving threads. The choices here are left to the user - and the kernel doesn't get in the way of efficiency.

It is worth noting that you can send messages to message queues from within ISR context. This helps maintain consistency, since the same APIs can be used to provide event-driven programming facilities throughout the whole of the OS.

10.7.3 Global Message Pool

To maintain efficiency in the messaging system (and to prevent over-allocation of data), a global pool of message objects is provided. The size of this message pool is specified in the implementation, and can be adjusted depending on the requirements of the target application as a compile-time option.

Allocating a message from the message pool is as simple as calling the

[GlobalMessagePool::Pop\(\)](#) Method.

Messages are returned back to the [GlobalMessagePool::Push\(\)](#) method once the message contents are no longer required.

One must be careful to ensure that discarded messages always are returned to the pool, otherwise a resource leak will occur, which may cripple the operating system's ability to pass data between threads.

10.7.4 Message Queues

[Message](#) objects specify data with context, but do not specify where the messages will be sent. For this purpose we have a [MessageQueue](#) object. Sending an object to a message queue involves calling the [MessageQueue::Send\(\)](#) method, passing in a pointer to the [Message](#) object as an argument.

When a message is sent to the queue, the first thread blocked on the queue (as a result of calling the [MessageQueue Receive\(\)](#) method) will wake up, with a pointer to the [Message](#) object returned.

It's worth noting that multiple threads can block on the same message queue, providing a means for multiple threads to share work in parallel.

10.7.5 Mailboxes

Another form of IPC is provided by Mark3, in the form of Mailboxes and Envelopes. Mailboxes are similar to message queues in that they provide a synchronized interface by which data can be transmitted between threads.

Where [Message](#) Queues rely on linked lists of lightweight message objects (containing only message code and a void* data-pointer), which are inherently abstract, Mailboxes use a dedicated blob of memory, which is carved up into fixed-size chunks called Envelopes (defined by the user), which are sent and received. Unlike message queues, mailbox data is copied to and from the mailboxes dedicated pool.

Mailboxes also differ in that they provide not only a blocking "receive" call, but also a blocking "send" call, providing the opportunity for threads to block on "mailbox full" as well as "mailbox empty" conditions.

All send/receive APIs support an optional timeout parameter if the `KERNEL_USE_TIMEOUTS` option has been configured in [mark3cfg.h](#)

10.7.6 Atomic Operations

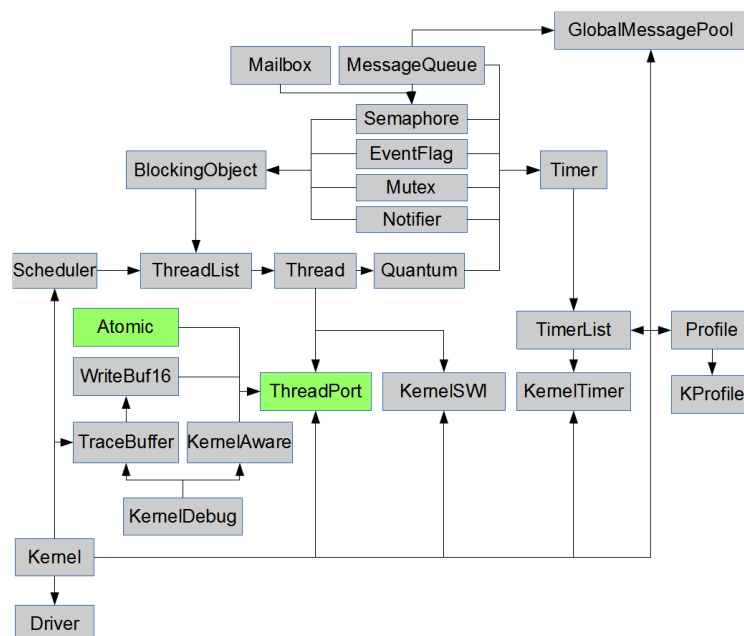


Figure 10.7: Atomic operations

This utility class provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Delete for 8, 16, and 32-bit integer types, as well as an atomic test-and-set.

10.7.7 Drivers

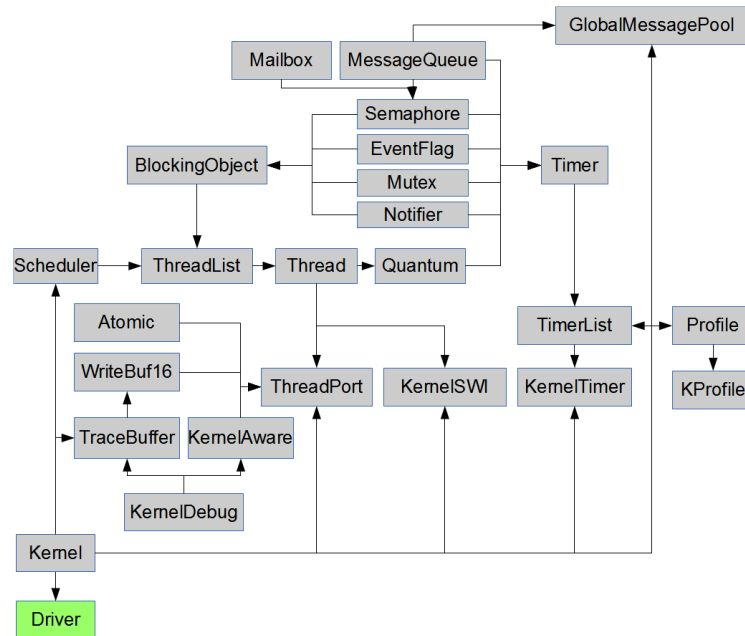


Figure 10.8: Drivers

This is the basis of the driver framework. In the context of Mark3, drivers don't necessarily have to be based on physical hardware peripherals. They can be used to represent algorithms (such as random number generators), files, or protocol stacks. Unlike FunkOS, where driver IO is protected automatically by a mutex, we do not use this kind of protection - we leave it up to the driver implementor to do what's right in its own context. This also frees up the driver to implement all sorts of other neat stuff, like sending messages to threads associated with the driver. Drivers are implemented as character devices, with the standard array of posix-style accessor methods for reading, writing, and general driver control.

A global driver list is provided as a convenient and minimal "filesystem" structure, in which devices can be accessed by name.

Driver Design

A device driver needs to be able to perform the following operations:

- Initialize a peripheral
- Start/stop a peripheral
- Handle I/O control operations
- Perform various read/write operations

At the end of the day, that's pretty much all a device driver has to do, and all of the functionality that needs to be presented to the developer.

We abstract all device drivers using a base-class which implements the following methods:

- Start/Open
- Stop/Close
- Control
- Read
- Write

A basic driver framework and API can thus be implemented in five function calls - that's it! You could even reduce that further by handling the initialize, start, and stop operations inside the "control" operation.

Driver API

In C++, we can implement this as a class to abstract these event handlers, with virtual void functions in the base class overridden by the inherited objects.

To add and remove device drivers from the global table, we use the following methods:

```
void DriverList::Add( Driver *pclDriver_ );
void DriverList::Remove( Driver *pclDriver_ );
```

`DriverList::Add()/Remove()` takes a single argument - the pointer to the object to operate on.

Once a driver has been added to the table, drivers are opened by NAME using `DriverList::FindByName("/dev/name")`. This function returns a pointer to the specified driver if successful, or to a built in /dev/null device if the path name is invalid. After a driver is open, that pointer is used for all other driver access functions.

This abstraction is incredibly useful - any peripheral or service can be accessed through a consistent set of APIs, that make it easy to substitute implementations from one platform to another. Portability is ensured, the overhead is negligible, and it emphasizes the reuse of both driver and application code as separate entities.

Consider a system with drivers for I2C, SPI, and UART peripherals - under our driver framework, an application can initialize these peripherals and write a greeting to each using the same simple API functions for all drivers:

```
pclI2C = DriverList::FindByName("/dev/i2c");
pclUART = DriverList::FindByName("/dev/tty0");
pclSPI = DriverList::FindByName("/dev/spi");

pclI2C->Write(12, "Hello World!");
pclUART->Write(12, "Hello World!");
pclSPI->Write(12, "Hello World!");
```

10.8 Kernel Proper and Porting

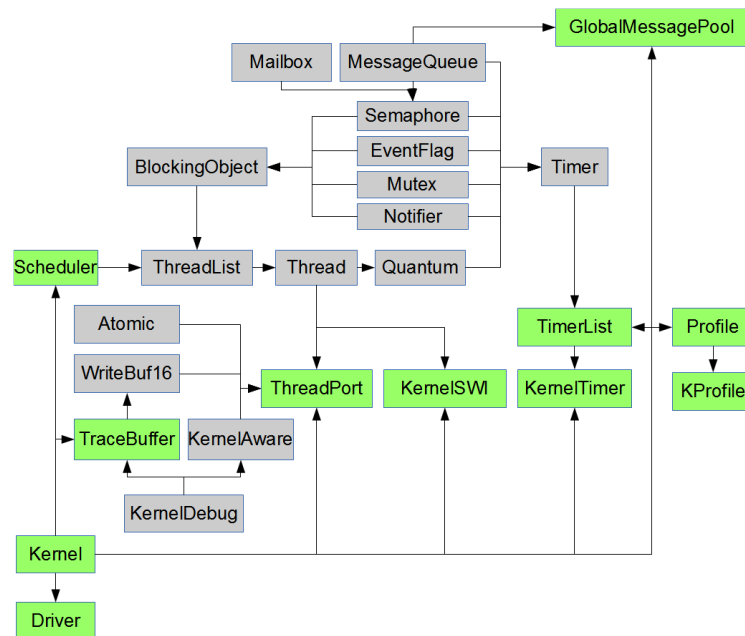


Figure 10.9: Kernel Proper and Porting

The `Kernel` class is a static class with methods to handle the initialization and startup of the RTOS, manage errors, and provide user-hooks for fatal error handling (functions called when `Kernel::Panic()` conditions are encountered), or when the Idle function is run.

Internally, `Kernel::Init()` calls the initialization routines for various kernel objects, providing a single interface by which all RTOS-related system initialization takes place.

`Kernel::Start()` is called to begin running OS functionality, and does not return. Control of the CPU is handed over to the scheduler, and the highest-priority ready thread begins execution in the RTOS environment.

Hardware Abstraction Layer

Almost all of the Mark3 kernel (and middleware) is completely platform independent, and should compile cleanly on any platform with a modern C++ compiler. However, there are a few areas within Mark3 that can only be implemented by touching hardware directly.

These interfaces generally cover four features:

- `Thread` initialization and context-switching logic
- Software interrupt control (used to generate context switches)
- Hardware timer control (support for time-based functionality, such as `Sleep()`)
- Code-execution profiling timer (not necessary to port if code-profiling is not compiled into the kernel)

The hardware abstraction layer in Mark3 provides a consistent interface for each of these four features. Mark3 is ported to new target architectures by providing an implementation for all of the interfaces declared in the abstraction layer. In the following section, we will explore how this was used to port the kernel to ARM Cortex-M0.

Real-world Porting Example – Cortex M0

This section serves as a real-world example of how Mark3 can be ported to new architectures, how the Mark3 abstraction layer works, and as a practical reference for using the RTOS support functionality baked in modern ARM Cortex-M series microcontrollers. Most of this documentation here is taken directly from the source code found in the `/kernel/cpu/cm0/ports` directory, with additional annotations to explain the port in more detail. Note that a familiarity with Cortex-M series parts will go a long way to understanding the subject matter presented, especially a basic understanding of the ARM CPU registers, exception models, and OS support features (PendSV, SysTick and SVC). If you're unfamiliar with ARM architecture, pay attention to the comments more than the source itself to illustrate the concepts.

Porting Mark3 to a new architecture consists of a few basic pieces; for developers familiar with the target architecture and the porting process, it's not a tremendously onerous endeavour to get Mark3 up-and-running somewhere new. For starters, all non-portable components are completely isolated in the source-tree under:

```
/embedded/kernel/CPU/VARIANT/TOOLCHAIN/
```

where CPU is the architecture, VARIANT is the vendor/part, and TOOLCHAIN is the compiler tool suite used to build the code.

From within the specific port folder, a developer needs only implement a few classes and headers that define the port-specific behavior of Mark3:

- [KernelSWI](#) ([kernelswi.cpp](#)/[kernelswi.h](#)) - Provides a maskable software-triggered interrupt used to perform context switching.
- [KernelTimer](#) ([kerneltimer.cpp](#)/[kerneltimer.h](#)) - Provides either a fixed-frequency or programmable-interval timer, which triggers an interrupt on expiry. This is used for implementing round-robin scheduling, thread-sleeps, and generic software timers.
- [Profiler](#) ([kprofile.cpp](#)/[kprofile.h](#)) - Contains code for runtime code-profiling. This is optional and may be stubbed out if left unimplemented (we won't cover profiling timers here).
- [ThreadPort](#) ([threadport.cpp](#)/[threadport.h](#)) - The meat-and-potatoes of the port code lives here. This class contains architecture/part-specific code used to initialize threads, implement critical-sections, perform context-switching, and start the kernel. Most of the time spent in this article focuses on the code found here.

Summarizing the above, these modules provide the following list of functionality:

```
- Thread stack initialization
- Kernel startup and first thread entry
- Context switch and SWI
- Kernel timers
- Critical Sections
.
```

The implementation of each of these pieces will be analyzed in detail in the sections that follow.

Thread Stack Initialization

Before a thread can be used, its stack must first be initialized to its default state. This default state ensures that when the thread is scheduled for the first time and its context restored, that it will cause the CPU to jump to the user's specified entry-point function.

All of the platform independent thread setup is handled by the generic kernel code. However, since every CPU architecture has its own register set, and stacks different information as part of an interrupt/exception, we have to implement this thread setup code for each platform we want the kernel to support (Combination of Architecture + Variant + Toolchain).

In the ARM Cortex-M0 architecture, the stack frame consists of the following information:

a) Exception Stack Frame

Contains the 8 registers which the ARM Cortex-M0 CPU automatically pushes to the stack when entering an exception. The following registers are included (in stack'd order):

```
[ XPSR ] <-- Highest address in context
[ PC   ]
```



```
[ LR  ]
[ R12 ]
[ R3  ]
[ R2  ]
[ R1  ]
[ R0  ]
```

XPSR – This is the CPU's status register. We need to set this to 0x01000000 (the "T" bit), which indicates that the CPU is executing in "thumb" mode. Note that ARMv6m and ARMv7m processors only run thumb2 instructions, so an exception is liable to occur if this bit ever gets cleared.

PC – Program Counter. This should be set with the initial entry point (function pointer) for that the user wishes to start executing with this thread.

LR - The base link register. Normally, this register contains the return address of the calling function, which is where the CPU jumps when a function returns. However, our threads generally don't return (and if they do, they're placed into the stop state). As a result we can leave this as 0.

The other registers in the stack frame are generic working registers, and have no special meaning, with the exception that R0 will hold the user's argument value passed into the entrypoint.

b) Complimentary CPU Register Context

```
[ R11 ]
...
[ R4  ] <-- Lowest address in context
```

These are the other general-purpose CPU registers that need to be backed up/ restored on a context switch, but aren't stacked by default on a Cortex-M0 exception. If there were any additional hardware registers to back up, then we'd also have to include them in this part of the context as well.

As a result, these registers all need to be manually pushed to the stack on stack creation, and will need to be explicitly pushed and pop as part of a normal context switch.

With this default exception state in mind, the following code is used to initialize a thread's stack for a Cortex-M0.

```
void ThreadPort::InitStack(Thread *pclThread_)
{
    K_ULONG *pulStack;
    K_ULONG *pulTemp;
    K_ULONG ulAddr;
    K_USHORT i;

    // Get the entrypoint for the thread
    ulAddr = (K_ULONG)(pclThread_>m_pfEntryPoint);

    // Get the top-of-stack pointer for the thread
    pulStack = (K_ULONG*)pclThread_>m_pwStackTop;

    // Initialize the stack to all FF's to aid in stack depth checking
    pulTemp = (K_ULONG*)pclThread_>m_pwStack;
    for (i = 0; i < pclThread_>m_usStackSize / sizeof(K_ULONG); i++)
    {
        pulTemp[i] = 0xFFFFFFFF;
    }

    PUSH_TO_STACK(pulStack, 0); // Apply one word of padding

    //-- Simulated Exception Stack Frame --
    PUSH_TO_STACK(pulStack, 0x01000000); // XPSR;set "T" bit for thumb-mode
    PUSH_TO_STACK(pulStack, ulAddr); // PC
    PUSH_TO_STACK(pulStack, 0); // LR
    PUSH_TO_STACK(pulStack, 0x12);
    PUSH_TO_STACK(pulStack, 0x3);
    PUSH_TO_STACK(pulStack, 0x2);
    PUSH_TO_STACK(pulStack, 0x1);
    PUSH_TO_STACK(pulStack, (K_ULONG)pclThread_>m_pvArg); // R0 = argument

    //-- Simulated Manually-Stacked Registers --
    PUSH_TO_STACK(pulStack, 0x11);
    PUSH_TO_STACK(pulStack, 0x10);
    PUSH_TO_STACK(pulStack, 0x09);
    PUSH_TO_STACK(pulStack, 0x08);
    PUSH_TO_STACK(pulStack, 0x07);
    PUSH_TO_STACK(pulStack, 0x06);
    PUSH_TO_STACK(pulStack, 0x05);
```

```

    PUSH_TO_STACK(pulStack, 0x04);
    pulStack++;

    pclThread->m_pwStackTop = pulStack;
}

```

Kernel Startup

The same general process applies to starting the kernel on an ARM Cortex-M0 as on other platforms. Here, we initialize and start the platform specific timer and software-interrupt modules, find the first thread to run, and then jump to that first thread.

Now, to perform that last step, we have two options:

1) Simulate a return from an exception manually to start the first thread, or.. 2) Use a software interrupt to trigger the first "Context Restore/Return from Interrupt"

For 1), we basically have to restore the whole stack manually, not relying on the CPU to do any of this for us. That's certainly doable, but not all Cortex parts support this (other members of the family support privileged modes, etc.). That, and the code required to do this is generally more complex due to all of the exception-state simulation. So, we will opt for the second option instead.

To implement a software to start our first thread, we will use the SVC instruction to generate an exception. From that exception, we can then restore the context from our first thread, set the CPU up to use the right "process" stack, and return-from-exception back to our first thread. We'll explore the code for that later.

But, before we can call the SVC exception, we're going to do a couple of things.

First, we're going to reset the default MSP stack pointer to its original top-of-stack value. The rationale here is that we no longer care about the data on the MSP stack, since calling the SVC instruction triggers a chain of events from which we never return. The MSP is also used by all exception-handling, so regaining a few words of stack here can be useful. We'll also enable all maskable exceptions at this point, since this code results in the kernel being started with the CPU executing the RTOS threads, at which point a user would expect interrupts to be enabled.

Note, the default stack pointer location is stored at address 0x00000000 on all ARM Cortex M0 parts. That explains the code below...

```

void ThreadPort_StartFirstThread( void )
{
    asm(
        " ldr r1, [r0] \n" // Reset the MSP to the default base address
        " msr msp, r1 \n"
        " cpsie i \n"      // Enable interrupts
        " svc 0 \n"        // Jump to SVC Call
    );
}

```

First Thread Entry

This handler has the job of taking the first thread object's stack, and restoring the default state data in a way that ensures that the thread starts executing when returning from the call.

We also keep in mind that there's an 8-byte offset from the beginning of the thread object to the location of the thread stack pointer. This offset is a result of the thread object inheriting from the linked-list node class, which has 8-bytes of data. This is stored first in the object, before the first element of the class, which is the "stack top" pointer.

The following assembly code shows how the SVC call is implemented in Mark3 for the purpose of starting the first thread.

```

get_thread_stack:
; Get the stack pointer for the current thread
ldr r0, g_pstCurrent
ldr r1, [r0]
add r1, #8
ldr r2, [r1]      ; r2 contains the current stack-top

load_manually_placed_context_r11_r8:
; Handle the bottom 32-bytes of the stack frame
; Start with r11-r8, because only r0-r7 can be used
; with ldmia on CM0.
add r2, #16
ldmia r2!, {r4-r7}

```

```

    mov r11, r7
    mov r10, r6
    mov r9, r5
    mov r8, r4

set_psp:
    ; Since r2 is coincidentally back to where the stack pointer should be,
    ; Set the program stack pointer such that returning from the exception handler
    msr psp, r2

load_manually_placed_context_r7_r4:
    ; Get back to the bottom of the manually stacked registers and pop.
    sub r2, #32
    ldmia r2!, {r4-r7} ; Register r4-r11 are restored.

set_thread_and_privilege_modes:
    ; Also modify the control register to force use of thread mode as well
    ; For CM3 forward-compatibility, also set user mode.
    mrs r0, control
    mov r1, #0x03
    orr r0, r1
    control, r0

set_lr:
    ; Set up the link register such that on return, the code operates
    ; in thread mode using the PSP. To do this, we or 0x0D to the value stored
    ; in the lr by the exception hardware EXC_RETURN. Alternately, we could
    ; just force lr to be 0xFFFFFFF0 (we know that's what we want from the
    ; hardware, anyway)
    mov r0, #0x0D
    mov r1, lr
    orr r0, r1

exit_exception:
    ; Return from the exception handler.
    ; The CPU will automatically unstack R0-R3, R12, PC, LR, and xPSR
    ; for us. If all goes well, our thread will start execution at the
    ; entrypoint, with the us-specified argument.
    bx r0

```

On ARM Cortex parts, there's dedicated hardware that's used primarily to support RTOS (or RTOS-like) functionality. This functionality includes the SysTick timer, and the PendSV Exception. SysTick is used for a tick-based kernel timer, while the PendSV exception is used for performing context switches. In reality, it's a "special SVC" call that's designed to be lower-overhead, in that it isn't mux'd with a bunch of other system or application functionality.

So how do we go about actually implementing a context switch here? There are a lot of different parts involved, but it essentially comes down to 3 steps:

1) Saving the context.

Thread's top-of-stack value is stored, all registers are stacked. We're good to go!

2) Swap threads

We swap the Scheduler's "next" thread with the "current" thread.

3) Restore Context

This is more or less identical to what we did when restoring the first context. Some operations may be optimized for data already stored in registers.

The code used to implement these steps on Cortex-M0 is presented below:

```

void PendSV_Handler(void)
{
    ASM(
        // Thread_SaveContext()
        " ldr r1, CURR_ \n"
        " ldr r1, [r1] \n "
        " mov r3, r1 \n "
        " add r3, #8 \n "

        // Grab the psp and adjust it by 32 based on extra registers we're going
        // to be manually stacking.
        " mrs r2, psp \n "
    )
}

```

```

" sub r2, #32 \n "

// While we're here, store the new top-of-stack value
" str r2, [r3] \n "

// And, while r2 is at the bottom of the stack frame, stack r7-r4
" stmia r2!, {r4-r7} \n "

// Stack r11-r8
" mov r7, r11 \n "
" mov r6, r10 \n "
" mov r5, r9 \n "
" mov r4, r8 \n "
" stmia r2!, {r4-r7} \n "

// Equivalent of Thread_Swap() - performs g_pstCurrent = g_pstNext
" ldr r1, CURR_ \n"
" ldr r0, NEXT_ \n"
" ldr r0, [r0] \n"
" str r0, [r1] \n"

// Thread_RestoreContext()
// Get the pointer to the next thread's stack
" add r0, #8 \n "
" ldr r2, [r0] \n "

// Stack pointer is in r2, start loading registers from
// the "manually-stacked" set
// Start with r11-r8, since these can't be accessed directly.
" add r2, #16 \n "
" ldmbia r2!, {r4-r7} \n "
" mov r11, r7 \n "
" mov r10, r6 \n "
" mov r9, r5 \n "
" mov r8, r4 \n "

// After subbing R2 #16 manually, and #16 through ldmbia, our PSP is where it
// needs to be when we return from the exception handler
" msr psp, r2 \n "

// Pop manually-stacked R4-R7
" sub r2, #32 \n "
" ldmbia r2!, {r4-r7} \n "

// lr contains the proper EXC_RETURN value
// we're done with the exception, so return back to newly-chosen thread
" bx lr \n "
" nop \n "

// Must be 4-byte aligned.
" NEXT_: .word g_pstNext \n"
" CURR_: .word g_pstCurrent \n"
);
}

```

Kernel Timers

ARM Cortex-M series microcontrollers each contain a SysTick timer, which was designed to facilitate a fixed-interval RTOS timer-tick. This timer is a precise 24-bit down-count timer, run at the main CPU clock frequency, that can be programmed to trigger an exception when the timer expires. The handler for this exception can thus be used to drive software timers throughout the system on a fixed interval.

Unfortunately, this hardware is extremely simple, and does not offer the flexibility of other timer hardware commonly implemented by MCU vendors - specifically a suitable timer prescaler that can be used to generate efficient, long-counting intervals. As a result, while the "generic" port of Mark3 for Cortex-M0 leverages the common SysTick timer interface, it only supports the tick-based version of the kernel's timer (note that specific Cortex-M0 ports such as the Atmel SAMD20 do have tickless timers).

Setting up a tick-based [KernelTimer](#) class to use the SysTick timer is, however, extremely easy, as is illustrated below:

```

void KernelTimer::Start(void)
{
    SysTick_Config(PORT_SYSTEM_FREQ / 1000); // 1KHz fixed clock...
    NVIC_EnableIRQ(SysTick_IRQn);
}

```

In [this](#) instance, the call to SysTick_Config() generates a 1kHz system-tick signal, and the NVIC_EnableIRQ() call ensures that a SysTick exception is generated for each tick. All other functions in the Cortex version of the [KernelTimer](#) class are essentially stubbed out (see the source for more details).

Note that the functions used in this call are part of the ARM Cortex Microcontroller Software Interface Standard (cmsis), and are supplied by all parts vendors selling Cortex hardware. This greatly simplifies the design of our port-code, since we can be reasonably assured that these APIs will work the same on all devices.

The handler code called when a SysTick exception occurs is basically the same as on other platforms (such as AVR), except that we explicitly clear the "exception pending" bit before returning. This is implemented in the following code:

```
\code{.cpp}
void SysTick_Handler(void)
{
    #if KERNEL_USE_TIMERS
        TimerScheduler::Process();
    #endif
    #if KERNEL_USE_QUANTUM
        Quantum::UpdateTimer();
    #endif

    // Clear the systick interrupt pending bit.
    SCB->ICSR |= SCB_ICSR_PENDSTCLR_Msk;
}
```

Critical Sections

A "critical section" is a block of code whose execution cannot be interrupted by means of context switches or an interrupt. In a traditional single-core operating system, it is typically implemented as a block of code where the interrupts are disabled - this is also the approach taken by Mark3. Given that every CPU has its own means of disabling/enabling interrupts, the implementation of the critical section APIs is also non-portable.

In the Cortex-M0 port, we implement the two critical section APIs ([CS_ENTER\(\)](#) and [CS_EXIT\(\)](#)) as function-like macros containing inline assembly. All uses of these calls are called in pairs within a function and must take place at the same level-of-scope. Also, as nesting may occur (critical section within a critical section), this must be taken into account in the code.

In general, [CS_ENTER\(\)](#) performs the following tasks:

- Cache the current interrupt-enabled state within a local variable in the thread's state
- Disable interrupts
- .

Conversely, [CS_EXIT\(\)](#) performs the following tasks:

- Read the original interrupt-enabled state from the cached value
- Restore interrupts to the original value
- .

On Cortex-M series microcontrollers, the PRIMASK special register contains a single status bit which can be used to enable/disable all maskable interrupts at once. This register can be read directly to examine or modify its state. For convenience, ARMv6m provides two instructions to enable/disable interrupts

- cpsid (disable interrupts) and cpsie (enable interrupts). Mark3 Implements these steps according to the following code:

```
//-----
#define CS_ENTER() \
{ \
    K_ULONG __ulRegState; \
    asm ( \
        " mrs r0, PRIMASK \n" \
        " mov %[STATUS], r0 \n" \
        " cpsid i \n" \
        : [STATUS] "=r" (__ulRegState) \
        ); \
}

//-----
#define CS_EXIT() \
asm ( \
    " mov r0, %[STATUS] \n" \
    " msr primask, r0 \n" \
    )
```

```

: \
: [STATUS] "r" (__ulRegState) \
); \
}

```

Summary

In this section we have investigated how the main non-portable areas of the Mark3 RTOS are implemented on a Cortex-M0 microcontroller. Mark3 leverages all of the hardware blocks designed to enable RTOS functionality on ARM Cortex-M series microcontrollers: the SVC call provides the mechanism by which we start the kernel, the PendSV exception provides the necessary software interrupt, and the SysTick timer provides an RTOS tick. As a result, Mark3 is a perfect fit for these devices - and as a result of this approach, the same RTOS port code should work with little to no modification on all ARM Cortex-M parts.

We have discussed what functionality in the RTOS is not portable, and what interfaces must be implemented in order to complete a fully-functional port. The five specific areas which are non-portable (stack initialization, kernel startup/entry, kernel timers, context switching, and critical sections) have been discussed in detail, with the platform-specific source provided as a practical reference to ARM-specific OS features, as well as Mark3's porting infrastructure. From this example (and the accompanying source), it should be possible for an experienced developers to create a port Mark3 to other microcontroller targets.

Chapter 11

Mark3C - C-language API bindings for the Mark3 Kernel.

Mark3 now includes an optional additional library with C language bindings for all core kernel APIs, known as Mark3C.

This library allows applications to be written in C, while still enjoying all of the benefits of the clean, modular design of the core RTOS kernel.

The C-language Mark3C APIs map directly to their Mark3 counterparts using a simple set of conventions, documented below. As a result, explicit API documentation for Mark3C is not necessary, as the functions map 1-1 to their C++ counterparts.

11.1 API Conventions

1) Static Methods:

<code><ClassName>::<MethodName>()</code>	Becomes	<code><ClassName>_<MethodhName>()</code>
i.e. <code>Kernel::Start()</code>	Becomes	<code>Kernel_Start()</code>

2) [Kernel](#) Object Methods:

In short, any class instance is represented using an object handle, and is always passed into the relevant APIs as the first argument. Further, any method that returns a pointer to an object in the C++ implementation now returns a handle to that object.

<code><Object>.<MethodName>(<args>)</code>	Becomes	<code><ClassName>_<MethodhName>(<ObjectHandle>, <args>)</code>
i.e. <code>clAppThread.Start()</code>	Becomes	<code>Thread_Start(hAppThread)</code>

3) Overloaded Methods:

a) Methods overloaded with a Timeout parameter:

<code><Object>.<MethodName>(<args>)</code>	Becomes	<code><ClassName>_Timed<MethodhName>(<ObjectHandle>, <args>)</code>
i.e. <code>clSemaphore.Wait(1000)</code>	Becomes	<code>Semaphore_Wait(hSemaphore, 1000)</code>

b) Methods overloaded based on number of arguments:

<code><Object>.<MethodName>()</code>	Becomes	<code><ClassName>_<MethodhName>(<ObjectHandle>)</code>
<code><Object>.<MethodName>(<arg1>)</code>	Becomes	<code><ClassName>_<MethodhName>1(<ObjectHandle>, <arg1>)</code>
<code><Object>.<MethodName>(<arg1>, <arg2>)</code>	Becomes	<code><ClassName>_<MethodhName>2(<ObjectHandle>, <arg1>, <arg2>)</code>

<ClassName>::<MethodName>()	Becomes	<ClassName>_<MethodName>(<ObjectHandle>)
<ClassName>::<MethodName>(<arg1>)	Becomes	<ClassName>_<MethodName>1(<ObjectHandle>, <arg1>)
<ClassName>::<MethodName>(<arg1>, <arg2>)	Becomes	<ClassName>_<MethodName>2(<ObjectHandle>, <arg1>, <arg2>)

c) Methods overloaded base on parameter types:

<Object>.<MethodName>(<arg type_a>)	Becomes	<ClassName>_<MethodName><type_a>(<ObjectHandle>, <arg type_a>)
<Object>.<MethodName>(<arg type_b>)	Becomes	<ClassName>_<MethodName><type_b>(<ObjectHandle>, <arg type_b>)
<ClassName>::<MethodName>(<arg type_a>)	Becomes	<ClassName>_<MethodName><type_a>(<arg type_a>)
<ClassName>::<MethodName>(<arg type_b>)	Becomes	<ClassName>_<MethodName><type_b>(<arg type_b>)

d) Allocate-once memory allocation APIs

AutoAlloc::New<ObjectName>	Becomes	Alloc_<ObjectName>
AutoAlloc::Allocate(uint16_t u16Size_)	Becomes	AutoAlloc(uint16_t u16Size_)

11.2 Allocating Objects

Aside from the API name translations, the object allocation scheme is the major different between Mark3C and Mark3. Instead of instantiating objects of the various kernel types, kernel objects must be declared using Declaration macros, which serve the purpose of reserving memory for the kernel object, and provide an opaque handle to that object memory. This is the case for statically-allocated objects, and objects allocated on the stack.

Example: Declaring a thread

```
#include "mark3c.h"

// Statically-allocated
DECLARE_THREAD(hMyThread1);
...

// On stack
int main(void)
{
    DECLARE_THREAD(hMyThread2);
    ...
}
```

Where:

hMyThread1 - is a handle to a statically-allocated thread
hMyThread2 - is a handle to a thread allocated from the main stack.

Alternatively, the AutoAlloc APIs can be used to dynamically allocate objects, as demonstrated in the following example.

```
void Allocate_Example(void)
{
    Thread_t hMyThread = AutoAlloc_Thread();

    Thread_Init(hMyThread, awMyStack, sizeof(awMyStack), 1, MyFunction, 0);
}
```

Note that the relevant kernel-object Init() function *must* be called prior to using any kernel object, whether or not they have been allocated statically, or dynamically.

11.3 Drivers in Mark3C

Because the Mark3 driver framework makes extensive use of inheritance and virtual functions in C++, it is difficult to wrap for use with C. In addition, all derived drivers types would still need to have their custom interfaces wrapped

with C-language bindings in order to be accessible from C, which is cumbersome and inelegant, and duplicates large portions of code. As a result, it's probably less work to write a Mark3C specific driver module with a similar interface to Mark3, on which drivers can be ported where necessary, or implemented directly on for efficiency. The APIs presented in [driver3c.h](#) provide such an interface for use in Mark3c.

Chapter 12

Release Notes

12.1 R6 Release

- New: Replace recursive-make build system with CMake and Ninja
- New: Transitioned version control to Git from Subversion.
- New: Socket library, implementing named "domain-socket" style IPC
- New: State Machine framework library
- New: Software I2C library completed, with demo app
- New: [Kernel Timer](#) loop can optionally be run within its own thread instead of a nested interrupt
- New: UART drivers are all now abstracted through UartDriver base class for portability
- Experimental: Process library, allowing for the creation of resource-isolated processes
- Removed: Bare-metal support for Atmel SAMD20 (generic port still works)
- Cleanup all compiler warnings on atmega328p
- Various Bugfixes and optimizations
- Various Script changes related to automating the build + release process

12.2 R5 Release

- New: Shell library for creating responsive CLIs for embedded applications (M3Shell)
- New: Stream library for creating thread-safe buffered streams (streamer)
- New: Blocking UART implementation for AVR (drvUARTplus)
- New: "Extended context" kernel feature, which is used to implement thread-local storage
- New: "Extra Checks" kernel feature, which enforces safe API usage under pain of [Kernel Panic](#)
- New: Realtime clock library
- New: Example application + bsp for the open-hardware Mark3no development board (mark3no)
- New: [Kernel](#) objects descoped/destroyed while still in active use will now cause kernel panic
- New: [Kernel](#) callouts for thread creation/destruction/context switching, used for time tracking
- New: Simple power management class

- New: WIP software-based I2C + SPI drivers
- Optimized thread scheduling via target-optimized "count-leading-zero" macros
- Expanded memutil library
- Various optimizations of ARM Cortex-M assembly code
- Various bugfixes to [Timer](#) code
- Improved stack overflow checking + warning (stack guard kernel feature)
- AVR bootloader now supports targets with more than 64K of flash
- Moved some port configuration out of platform.mak into header files in the kernel port code
- The usual minor bugfixes and "gentle refactoring"

12.3 R4 Release

- New: C-language bindings for Mark3 kernel (mark3c library)
- New: Support for ARM Cortex-M3 and Cortex-M4 (floating point) targets
- New: Support for Atmel AVR atmega2560 and arduino pro mega
- New: Full-featured, lightweight heap implementation
- New: [Mailbox](#) IPC class
- New: Notification object class
- New: lightweight tracelogger/instrumentation implementation (buffalogger), with sample parser
- New: High-performance AVR Software UART implementation
- New: Allocate-once "AutoAlloc" memory allocator
- New: Fixed-time blocking/unblocking operations added to ThreadList/Blocking class
- Placement-new supported for all kernel objects
- [Scheduler](#) now supports up to 1024 levels of thread priority, up from 8 (configurable at build-time)
- [Kernel](#) now uses stdint.h types for standard integers (instead of K_CHAR, K_ULONG, etc.)
- Greatly expanded documentation, with many new examples covering all key kernel features
- Expanded unit test coverage on AVR
- Updated build system and scripts for easier kernel configuration
- Updated builds to only attempt to build tests for supported platforms

12.4 R3 Release

- New: Added support for MSP430 microcontrollers
- New: Added [Kernel](#) Idle-Function hook to eliminate the need for a dedicated idle-thread (where supported)
- New: Support for kernel-aware simulation and testing via fIAVR AVR simulator
- Updated AVR driver selection
- General bugfixes and maintenance
- Expanded documentation and test coverage

12.5 R2

- Experimental release, using a "kernel transaction queue" for serializing kernel calls
- Works as a proof-of-concept, but abandoned due to overhead of the transaction mechanism in the general case.

12.6 R1 - 2nd Release Candidate

- New: Added support for ARM Cortex-M0 targets
- New: Added support for various AVR targets
- New: Timers now support a "tolerance" parameter for grouping timers with close expiry times
- Expanded scripts and automation used in build/test
- Updated and expanded graphics APIs
- Large number of bugfixes

12.7 R1 - 1st Release Candidate

- Initial release, with support for AVR microcontrollers

Chapter 13

Profiling Results

The following profiling results were obtained using an ATmega328p @ 16MHz.

The test cases are designed to make use of the kernel profiler, which accurately measures the performance of the fundamental system APIs, in order to provide information for user comparison, as well as to ensure that regressions are not being introduced into the system.

13.1 Date Performed

Fri 7 Jul 01:40:14 UTC 2017

13.2 Compiler Information

The kernel and test code used in these results were built using the following compiler:

```
Using built-in specs.
COLLECT_GCC=avr-gcc
COLLECT_LTO_WRAPPER=/usr/lib/gcc/avr/4.8.1/lto-wrapper
Target: avr
Configured with: ../src/configure -v --enable-languages=c,c++ --prefix=/usr/lib --infodir=/usr/share/info -
    -mandir=/usr/share/man --bindir=/usr/bin --libexecdir=/usr/lib --libdir=/usr/lib --enable-shared --with-
    system-zlib --enable-long-long --enable-nls --without-included-gettext --disable-libssp --build=arm-linux-
    gnueabi --host=arm-linux-gnueabi --target=avr
Thread model: single
gcc version 4.8.1 (GCC)
```

13.3 Profiling Results

```
- Semaphore Initialization: 40 cycles (averaged over 100 iterations)
- Semaphore Post (uncontested): 120 cycles (averaged over 100 iterations)
- Semaphore Pend (uncontested): 32 cycles (averaged over 100 iterations)
- Semaphore Flyback Time (Contested Pend): 2272 cycles (averaged over 100 iterations)
- Mutex Init: 000004294967272 cycles (averaged over 100 iterations)
- Mutex Claim: 216 cycles (averaged over 100 iterations)
- Mutex Release: 120 cycles (averaged over 100 iterations)
- Thread Initialize: 8312 cycles (averaged over 100 iterations)
- Thread Start: 1096 cycles (averaged over 100 iterations)
- Context Switch: 168 cycles (averaged over 100 iterations)
- Thread Schedule: 32 cycles (averaged over 100 iterations)
```


Chapter 14

Code Size Profiling

The following report details the size of each module compiled into the kernel.

The size of each component is dependent on the flags specified in [mark3cfg.h](#) at compile time. Note that these sizes represent the maximum size of each module before dead code elimination and any additional link-time optimization, and represent the maximum possible size that any module can take.

The results below are for profiling on Atmel AVR atmega328p-based targets using gcc. Results are not necessarily indicative of relative or absolute performance on other platforms or toolchains.

14.1 Information

Date Profiled: Fri 7 Jul 01:40:15 UTC 2017

14.2 Compiler Version

avr-gcc (GCC) 4.8.1 Copyright (C) 2013 Free Software Foundation, Inc. This is free software; see the source for copying conditions. There is NO warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

14.3 Profiling Results

Mark3 Module Size Report:

```
- Atomic Operations..... : 0 Bytes
- Allocate-once Heap..... : 0 Bytes
- Synchronization Objects - Base Class..... : 126 Bytes
- Device Driver Framework (including /dev/null)... : 212 Bytes
- Synchronization Object - Event Flag..... : 754 Bytes
- Mark3 Kernel Base Class..... : 163 Bytes
- Atmel AVR - Kernel Aware Simulation Support..... : 190 Bytes
- Semaphore (Synchronization Object)..... : 530 Bytes
- Fundamental Kernel Linked-List Classes..... : 458 Bytes
- Mailbox IPC Support..... : 862 Bytes
- Message-based IPC..... : 384 Bytes
- Mutex (Synchronization Object)..... : 712 Bytes
- Notification Blocking Object..... : 580 Bytes
- 2D Priority Map Object - Scheduler..... : 116 Bytes
- Performance-profiling timers..... : 474 Bytes
- Round-Robin Scheduling Support..... : 261 Bytes
- Thread Scheduling..... : 318 Bytes
- Thread Implementation..... : 1665 Bytes
- Fundamental Kernel Thread-list Data Structures.. : 250 Bytes
- Software Timer Kernel Object..... : 508 Bytes
- Software Timer Management..... : 645 Bytes
- Runtime Kernel Trace Implementation..... : 0 Bytes
```

- Atmel AVR - Profiling Timer Implementation..... : 216 Bytes
- Atmel AVR - Kernel Interrupt Implemenation..... : 56 Bytes
- Atmel AVR - Kernel Timer Implementation..... : 382 Bytes
- Atmel AVR - Basic Threading Support..... : 526 Bytes

Mark3 Kernel Size Summary:

- Kernel : 3096 Bytes
- Synchronization Objects : 2380 Bytes
- Port : 2812 Bytes
- Features : 2100 Bytes
- Total Size : 10388 Bytes

Chapter 15

Hierarchical Index

15.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

BlockingObject	85
EventFlag	94
Mutex	129
Notify	133
Semaphore	143
DriverList	93
FakeThread_t	97
GlobalMessagePool	98
Kernel	99
KernelAware	105
KernelSWI	108
KernelTimer	110
LinkList	113
CircularLinkList	87
ThreadList	156
DoubleLinkList	88
TimerList	166
LinkListNode	114
Driver	90
Message	123
Thread	146
Timer	160
Mailbox	116
MessagePool	125
MessageQueue	127
PriorityMap	135
ProfileTimer	136
Quantum	138
Scheduler	140
ThreadPort	159
TimerScheduler	168

Chapter 16

Class Index

16.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

BlockingObject	Class implementing thread-blocking primitives	85
CircularLinkedList	Circular-linked-list data type, inherited from the base LinkedList type	87
DoubleLinkedList	Doubly-linked-list data type, inherited from the base LinkedList type	88
Driver	Base device-driver class used in hardware abstraction	90
DriverList	List of Driver objects used to keep track of all device drivers in the system	93
EventFlag	Blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system	94
FakeThread_t	If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system	97
GlobalMessagePool	Implements a list of message objects shared between all threads	98
Kernel	Class that encapsulates all of the kernel startup functions	99
KernelAware	The KernelAware class	105
KernelSWI	Class providing the software-interrupt required for context-switching in the kernel	108
KernelTimer	Hardware timer interface, used by all scheduling/timer subsystems	110
LinkedList	Abstract-data-type from which all other linked-lists are derived	113
LinkedListNode	Basic linked-list node data structure	114
Mailbox	Implements an IPC mechanism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user	116
Message	Class to provide message-based IPC services in the kernel	123
MessagePool	Implements a list of message objects	125

MessageQueue	List of messages, used as the channel for sending and receiving messages between threads	127
Mutex	Mutual-exclusion locks, based on BlockingObject	129
Notify	Blocking object type, that allows one or more threads to wait for an event to occur before resuming operation	133
PriorityMap	The PriorityMap class	135
ProfileTimer	Profiling timer	136
Quantum	Static-class used to implement Thread quantum functionality, which is a key part of round-robin scheduling	138
Scheduler	Priority-based round-robin Thread scheduling, using ThreadLists for housekeeping	140
Semaphore	Binary & Counting semaphores, based on BlockingObject base class	143
Thread	Object providing fundamental multitasking support in the kernel	146
ThreadList	This class is used for building thread-management facilities, such as schedulers, and blocking objects	156
ThreadPort	Class defining the architecture specific functions required by the kernel	159
Timer	Kernel-managed software timers	160
TimerList	TimerList class - a doubly-linked-list of timer objects	166
TimerScheduler	"Static" Class used to interface a global TimerList with the rest of the kernel	168

Chapter 17

File Index

17.1 File List

Here is a list of all documented files with brief descriptions:

/media/usb/project/github/Mark3/kernel/ atomic.cpp	
Basic Atomic Operations	171
/media/usb/project/github/Mark3/kernel/ autoalloc.cpp	
Automatic memory allocation for kernel objects	173
/media/usb/project/github/Mark3/kernel/ blocking.cpp	
Implementation of base class for blocking objects	175
/media/usb/project/github/Mark3/kernel/ driver.cpp	
Device driver/hardware abstraction layer	194
/media/usb/project/github/Mark3/kernel/ eventflag.cpp	
Event Flag Blocking Object/IPC-Object implementation	195
/media/usb/project/github/Mark3/kernel/ kernel.cpp	
Kernel initialization and startup code	200
/media/usb/project/github/Mark3/kernel/ kernelaware.cpp	
Kernel aware simulation support	201
/media/usb/project/github/Mark3/kernel/ ksemaphore.cpp	
Semaphore Blocking-Object Implementation	203
/media/usb/project/github/Mark3/kernel/ ll.cpp	
Core Linked-List implementation, from which all kernel objects are derived	207
/media/usb/project/github/Mark3/kernel/ mailbox.cpp	
Mailbox + Envelope IPC mechanism	209
/media/usb/project/github/Mark3/kernel/ message.cpp	
Inter-thread communications via message passing	213
/media/usb/project/github/Mark3/kernel/ mutex.cpp	
Mutual-exclusion object	216
/media/usb/project/github/Mark3/kernel/ notify.cpp	
Lightweight thread notification - blocking object	219
/media/usb/project/github/Mark3/kernel/ priomap.cpp	
Priority map data structure	222
/media/usb/project/github/Mark3/kernel/ profile.cpp	
Code profiling utilities	224
/media/usb/project/github/Mark3/kernel/ quantum.cpp	
Thread Quantum Implementation for Round-Robin Scheduling	281
/media/usb/project/github/Mark3/kernel/ scheduler.cpp	
Strict-Priority + Round-Robin thread scheduler implementation	283
/media/usb/project/github/Mark3/kernel/ thread.cpp	
Platform-Independent thread class Definition	285
/media/usb/project/github/Mark3/kernel/ threadlist.cpp	
Thread linked-list definitions	292

/media/usb/project/github/Mark3/kernel/timer.cpp	
Timer implementations	294
/media/usb/project/github/Mark3/kernel/timerlist.cpp	
Implements timer list processing algorithms, responsible for all timer tick and expiry logic	297
/media/usb/project/github/Mark3/kernel/tracebuffer.cpp	
Kernel trace buffer class definition	300
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kernelprofile.cpp	
ATMega328p Profiling timer implementation	177
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kernelswi.cpp	
Kernel Software interrupt implementation for ATMega328p	178
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kerneltimer.cpp	
Kernel Timer Implementation for ATMega328p	179
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/threadport.cpp	
ATMega328p Multithreading	191
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kernelprofile.h	
Profiling timer hardware interface	182
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kernelswi.h	
Kernel Software interrupt declarations	183
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kerneltimer.h	
Kernel Timer Class declaration	184
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/portcfg.h	
Mark3 Port Configuration	185
/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/threadport.h	
ATMega328p Multithreading support	188
/media/usb/project/github/Mark3/kernel/public/atomic.h	
Basic Atomic Operations	225
/media/usb/project/github/Mark3/kernel/public/autoalloc.h	
Automatic memory allocation for kernel objects	226
/media/usb/project/github/Mark3/kernel/public/blocking.h	
Blocking object base class declarations	228
/media/usb/project/github/Mark3/kernel/public/buffalogger.h	
Super-efficient, super-secure logging routines	229
/media/usb/project/github/Mark3/kernel/public/dbg_file_list.h	
??	
/media/usb/project/github/Mark3/kernel/public/driver.h	
Driver abstraction framework	230
/media/usb/project/github/Mark3/kernel/public/eventflag.h	
Event Flag Blocking Object/IPC-Object definition	232
/media/usb/project/github/Mark3/kernel/public/kernel.h	
Kernel initialization and startup class	234
/media/usb/project/github/Mark3/kernel/public/kernelaware.h	
Kernel aware simulation support	235
/media/usb/project/github/Mark3/kernel/public/kerneldebug.h	
Macros and functions used for assertions, kernel traces, etc	237
/media/usb/project/github/Mark3/kernel/public/kerneltypes.h	
Basic data type primitives used throughout the OS	243
/media/usb/project/github/Mark3/kernel/public/ksemaphore.h	
Semaphore Blocking Object class declarations	244
/media/usb/project/github/Mark3/kernel/public/ll.h	
Core linked-list declarations, used by all kernel list types	246
/media/usb/project/github/Mark3/kernel/public/mailbox.h	
Mailbox + Envelope IPC Mechanism	248
/media/usb/project/github/Mark3/kernel/public/manual.h	
/brief Ascii-format documentation, used by doxygen to create various printable and viewable forms	250
/media/usb/project/github/Mark3/kernel/public/mark3.h	
Single include file given to users of the Mark3 Kernel API	251
/media/usb/project/github/Mark3/kernel/public/mark3cfg.h	
Mark3 Kernel Configuration	252

/media/usb/project/github/Mark3/kernel/public/ message.h	
Inter-thread communication via message-passing	259
/media/usb/project/github/Mark3/kernel/public/ mutex.h	
Mutual exclusion class declaration	262
/media/usb/project/github/Mark3/kernel/public/ notify.h	
Lightweight thread notification - blocking object	263
/media/usb/project/github/Mark3/kernel/public/ paniccodes.h	
Defines the reason codes thrown when a kernel panic occurs	264
/media/usb/project/github/Mark3/kernel/public/ priomap.h	
Priority map data structure	265
/media/usb/project/github/Mark3/kernel/public/ profile.h	
High-precision profiling timers	266
/media/usb/project/github/Mark3/kernel/public/ profiling_results.h	??
/media/usb/project/github/Mark3/kernel/public/ quantum.h	
Thread Quantum declarations for Round-Robin Scheduling	268
/media/usb/project/github/Mark3/kernel/public/ scheduler.h	
Thread scheduler function declarations	269
/media/usb/project/github/Mark3/kernel/public/ sizeprofile.h	??
/media/usb/project/github/Mark3/kernel/public/ thread.h	
Platform independent thread class declarations	271
/media/usb/project/github/Mark3/kernel/public/ threadlist.h	
Thread linked-list declarations	274
/media/usb/project/github/Mark3/kernel/public/ timer.h	
Timer object declarations	275
/media/usb/project/github/Mark3/kernel/public/ timerlist.h	
Timer list declarations	278
/media/usb/project/github/Mark3/kernel/public/ timerscheduler.h	
Timer scheduler declarations	279
/media/usb/project/github/Mark3/kernel/public/ tracebuffer.h	
Kernel trace buffer class declaration	280
/media/usb/project/github/Mark3/libs/mark3c/public/ driver3c.h	??
/media/usb/project/github/Mark3/libs/mark3c/public/ fake_types.h	
C-struct definitions that mirror	302
/media/usb/project/github/Mark3/libs/mark3c/public/ mark3c.h	
Header providing C-language API bindings for the Mark3 kernel	304

Chapter 18


Class Documentation

18.1 BlockingObject Class Reference

Class implementing thread-blocking primitives.

```
#include <blocking.h>
```

Inheritance diagram for BlockingObject:



Protected Member Functions

- void **Block** (Thread *pclThread_)
Block.
- void **BlockPriority** (Thread *pclThread_)
BlockPriority.
- void **UnBlock** (Thread *pclThread_)
UnBlock.
- void **SetInitialized** (void)
SetInitialized.
- bool **IsInitialized** (void)
IsInitialized.

Protected Attributes

- ThreadList **m_clBlockList**
ThreadList which is used to hold the list of threads blocked on a given object.
- uint8_t **m_u8Initialized**
Token used to check whether or not the object has been initialized prior to use.

18.1.1 Detailed Description

Class implementing thread-blocking primitives.

used for implementing things like semaphores, mutexes, message queues, or anything else that could cause a thread to suspend execution on some external stimulus.

Definition at line 71 of file [blocking.h](#).

18.1.2 Member Function Documentation

18.1.2.1 void BlockingObject::Block (Thread * *pcIThread_*) [protected]

Block.

Blocks a thread on this object. This is the fundamental operation performed by any sort of blocking operation in the operating system. All semaphores/mutexes/sleeping/messaging/etc ends up going through the blocking code at some point as part of the code that manages a transition from an "active" or "waiting" thread to a "blocked" thread.

The steps involved in blocking a thread (which are performed in the function itself) are as follows;

1) Remove the specified thread from the current owner's list (which is likely one of the scheduler's thread lists) 2) Add the thread to this object's thread list 3) Setting the thread's "current thread-list" point to reference this object's threadlist.

Parameters

<i>pcIThread_</i>	Pointer to the thread object that will be blocked.
-------------------	--

18.1.2.2 void BlockingObject::BlockPriority (Thread * *pcIThread_*) [protected]

BlockPriority.

Same as [Block\(\)](#), but ensures that threads are added to the block-list in priority-order, which optimizes the unblock procedure.

Parameters

<i>pcIThread_</i>	Pointer to the Thread to Block.
-------------------	---

18.1.2.3 bool BlockingObject::IsInitialized (void) [inline],[protected]

IsInitialized.

Returns

Definition at line 147 of file [blocking.h](#).

18.1.2.4 void BlockingObject::UnBlock (Thread * *pcIThread_*) [protected]

UnBlock.

Unblock a thread that is already blocked on this object, returning it to the "ready" state by performing the following steps:

Parameters

<code>pcThread_</code>	Pointer to the thread to unblock.
------------------------	-----------------------------------

1) Removing the thread from this object's threadlist 2) Restoring the thread to its "original" owner's list

The documentation for this class was generated from the following file:


- [/media/usb/project/github/Mark3/kernel/public/blocking.h](#)

18.2 CircularLinkedList Class Reference

Circular-linked-list data type, inherited from the base [LinkedList](#) type.

```
#include <ll.h>
```

Inheritance diagram for CircularLinkedList:



Public Member Functions

- void [Add](#) ([LinkedListNode](#) *node_)
Add the linked list node to this linked list.
- void [Remove](#) ([LinkedListNode](#) *node_)
Remove.
- void [PivotForward](#) ()
PivotForward.
- void [PivotBackward](#) ()
PivotBackward.
- void [InsertNodeBefore](#) ([LinkedListNode](#) *node_, [LinkedListNode](#) *insert_)
InsertNodeBefore.

Additional Inherited Members

18.2.1 Detailed Description

Circular-linked-list data type, inherited from the base [LinkedList](#) type.

Definition at line [187](#) of file [ll.h](#).

18.2.2 Member Function Documentation

18.2.2.1 void CircularLinkedList::Add ([LinkedListNode](#) * node_)

Add the linked list node to this linked list.

Parameters

<i>node_</i>	Pointer to the node to add
--------------	----------------------------

Definition at line 98 of file [ll.cpp](#).

18.2.2.2 void CircularLinkedList::InsertNodeBefore (LinkListNode * *node_*, LinkListNode * *insert_*)

InsertNodeBefore.

Insert a linked-list node into the list before the specified insertion point.

Parameters

<i>node_</i>	Node to insert into the list
<i>insert_</i>	Insert point.

Definition at line 172 of file [ll.cpp](#).

18.2.2.3 void CircularLinkedList::PivotBackward ()

PivotBackward.

Pivot the head of the circularly linked list backward (Head = Head->prev, Tail = Tail->prev)

Definition at line 163 of file [ll.cpp](#).

18.2.2.4 void CircularLinkedList::PivotForward ()

PivotForward.

Pivot the head of the circularly linked list forward (Head = Head->next, Tail = Tail->next)

Definition at line 154 of file [ll.cpp](#).

18.2.2.5 void CircularLinkedList::Remove (LinkListNode * *node_*)

Remove.

Add the linked list node to this linked list

Parameters

<i>node_</i>	Pointer to the node to remove
--------------	-------------------------------

Definition at line 120 of file [ll.cpp](#).

The documentation for this class was generated from the following files:


- [/media/usb/project/github/Mark3/kernel/public/ll.h](#)
- [/media/usb/project/github/Mark3/kernel/ll.cpp](#)

18.3 DoubleLinkedList Class Reference

Doubly-linked-list data type, inherited from the base [LinkList](#) type.

```
#include <ll.h>
```

Inheritance diagram for DoubleLinkedList:



Public Member Functions

- [DoubleLinkedList](#) ()
DoubleLinkedList.
- void [Add](#) ([LinkedListNode](#) *node_)
Add.
- void [Remove](#) ([LinkedListNode](#) *node_)
Remove.

Additional Inherited Members

18.3.1 Detailed Description

Doubly-linked-list data type, inherited from the base [LinkedList](#) type.

Definition at line 149 of file [ll.h](#).

18.3.2 Constructor & Destructor Documentation

18.3.2.1 [DoubleLinkedList::DoubleLinkedList](#) () [[inline](#)]

[DoubleLinkedList](#).

Default constructor - initializes the head/tail nodes to NULL

Definition at line 158 of file [ll.h](#).

18.3.3 Member Function Documentation

18.3.3.1 void [DoubleLinkedList::Add](#) ([LinkedListNode](#) * node_)

Add.

Add the linked list node to this linked list

Parameters

<i>node_</i>	Pointer to the node to add
--------------	----------------------------

Definition at line 47 of file [ll.cpp](#).

18.3.3.2 void [DoubleLinkedList::Remove](#) ([LinkedListNode](#) * node_)

Remove.

Add the linked list node to this linked list

Parameters

<i>node_</i>	Pointer to the node to remove
--------------	-------------------------------

Definition at line 68 of file [ll.cpp](#).

The documentation for this class was generated from the following files:


- [/media/usb/project/github/Mark3/kernel/public/ll.h](#)
- [/media/usb/project/github/Mark3/kernel/ll.cpp](#)

18.4 Driver Class Reference

Base device-driver class used in hardware abstraction.

```
#include <driver.h>
```

Inheritance diagram for Driver:



Public Member Functions

- virtual void [Init](#) ()=0
Init.
- virtual uint8_t [Open](#) ()=0
Open.
- virtual uint8_t [Close](#) ()=0
Close.
- virtual uint16_t [Read](#) (uint16_t u16Bytes_, uint8_t *pu8Data_)=0
Read.
- virtual uint16_t [Write](#) (uint16_t u16Bytes_, uint8_t *pu8Data_)=0
Write.
- virtual uint16_t [Control](#) (uint16_t u16Event_, void *pvDataIn_, uint16_t u16SizeIn_, void *pvDataOut_↔, uint16_t u16SizeOut_)=0
Control.
- void [SetName](#) (const char *pcName_)
SetName.
- const char * [GetPath](#) ()
GetPath.

Private Attributes

- const char * [m_pcPath](#)
string pointer that holds the driver path (name)

Additional Inherited Members

18.4.1 Detailed Description

Base device-driver class used in hardware abstraction.

All other device drivers inherit from this class

Definition at line 121 of file [driver.h](#).

18.4.2 Member Function Documentation

18.4.2.1 `virtual uint8_t Driver::Close () [pure virtual]`

Close.

Close a previously-opened device driver.

Returns

Driver-specific return code, 0 = OK, non-0 = error

18.4.2.2 `virtual uint16_t Driver::Control (uint16_t u16Event_, void * pvDataIn_, uint16_t u16SizeIn_, void * pvDataOut_, uint16_t u16SizeOut_) [pure virtual]`

Control.

This is the main entry-point for device-specific io and control operations. This is used for implementing all "side-channel" communications with a device, and any device-specific IO operations that do not conform to the typical POSIX read/write paradigm. use of this function is analogous to the non-POSIX (yet still common) `devctl()` or `ioctl()`.

Parameters

<i>u16Event_</i>	Code defining the io event (driver-specific)
<i>pvDataIn_</i>	Pointer to the input data
<i>u16SizeIn_</i>	Size of the input data (in bytes)
<i>pvDataOut_</i>	Pointer to the output data
<i>u16SizeOut_</i>	Size of the output data (in bytes)

Returns

Driver-specific return code, 0 = OK, non-0 = error

18.4.2.3 `const char* Driver::GetPath () [inline]`

GetPath.

Returns a string containing the device path.

Returns

pcName_ Return the string constant representing the device path

Definition at line 221 of file [driver.h](#).

18.4.2.4 `virtual void Driver::Init () [pure virtual]`

Init.

Initialize a driver, must be called prior to use

18.4.2.5 `virtual uint8_t Driver::Open () [pure virtual]`

Open.

Open a device driver prior to use.

Returns

Driver-specific return code, 0 = OK, non-0 = error

18.4.2.6 `virtual uint16_t Driver::Read (uint16_t u16Bytes_, uint8_t* pu8Data_) [pure virtual]`

Read.

Read a specified number of bytes from the device into a specific buffer. Depending on the driver-specific implementation, this may be a number less than the requested number of bytes read, indicating that there was less input than desired, or that as a result of buffering, the data may not be available.

Parameters

<i>u16Bytes_</i>	Number of bytes to read (<= size of the buffer)
<i>pu8Data_</i>	Pointer to a data buffer receiving the read data

Returns

Number of bytes actually read

18.4.2.7 `void Driver::SetName (const char * pcName_) [inline]`

SetName.

Set the path for the driver. Name must be set prior to access (since driver access is name-based).

Parameters

<i>pcName_</i>	String constant containing the device path
----------------	--

Definition at line 213 of file [driver.h](#).

18.4.2.8 `virtual uint16_t Driver::Write (uint16_t u16Bytes_, uint8_t* pu8Data_) [pure virtual]`

Write.

Write a payload of data of a given length to the device. Depending on the implementation of the driver, the amount of data written to the device may be less than the requested number of bytes. A result less than the requested size may indicate that the device buffer is full, indicating that the user must retry the write at a later point with the remaining data.

Parameters

<i>u16Bytes_</i>	Number of bytes to write (<= size of the buffer)
<i>pu8Data_</i>	Pointer to a data buffer containing the data to write

Returns

Number of bytes actually written

The documentation for this class was generated from the following file:

- [/media/usb/project/github/Mark3/kernel/public/driver.h](#)

18.5 DriverList Class Reference

List of [Driver](#) objects used to keep track of all device drivers in the system.

```
#include <driver.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init.
- static void [Add](#) ([Driver](#) *pclDriver_)
Add.
- static void [Remove](#) ([Driver](#) *pclDriver_)
Remove.
- static [Driver](#) * [FindByPath](#) (const char *m_pcPath)
FindByPath.

Static Private Attributes

- static [DoubleLinkedList](#) [m_clDriverList](#)
LinkedList object used to implementing the driver object management.

18.5.1 Detailed Description

List of [Driver](#) objects used to keep track of all device drivers in the system.

By default, the list contains a single entity, "/dev/null".

Definition at line 232 of file [driver.h](#).

18.5.2 Member Function Documentation

18.5.2.1 static void [DriverList::Add](#) ([Driver](#) * [pclDriver_](#)) [inline], [static]

Add.

Add a [Driver](#) object to the managed global driver-list.

Parameters

pclDriver_	pointer to the driver object to add to the global driver list.
----------------------------	--

Examples:

[buffalogger/main.cpp](#).

Definition at line 252 of file [driver.h](#).

18.5.2.2 static [Driver](#)* [DriverList::FindByPath](#) (const char * [m_pcPath](#)) [static]

FindByPath.

Look-up a driver in the global driver-list based on its path. In the event that the driver is not found in the list, a pointer to the default "/dev/null" object is returned. In this way, unimplemented drivers are automatically stubbed out.

18.5.2.3 static void DriverList::Init () [static]

Init.

Initialize the list of drivers. Must be called prior to using the device driver library.

18.5.2.4 static void DriverList::Remove (Driver * *pclDriver_*) [inline],[static]

Remove.

Remove a driver from the global driver list.

Parameters

<i>pclDriver_</i>	Pointer to the driver object to remove from the global table
-------------------	--

Definition at line 261 of file [driver.h](#).

The documentation for this class was generated from the following file:


- [/media/usb/project/github/Mark3/kernel/public/driver.h](#)

18.6 EventFlag Class Reference

The [EventFlag](#) class is a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system.

```
#include <eventflag.h>
```

Inheritance diagram for EventFlag:



Public Member Functions

- void [Init](#) ()
Init Initializes the [EventFlag](#) object prior to use.
- uint16_t [Wait](#) (uint16_t u16Mask_, [EventFlagOperation_t](#) eMode_)
Wait - Block a thread on the specific flags in this event flag group.
- uint16_t [Wait](#) (uint16_t u16Mask_, [EventFlagOperation_t](#) eMode_, uint32_t u32TimeMS_)
Wait - Block a thread on the specific flags in this event flag group.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)
WakeMe.
- void [Set](#) (uint16_t u16Mask_)
Set - Set additional flags in this object (logical OR).
- void [Clear](#) (uint16_t u16Mask_)
ClearFlags - Clear a specific set of flags within this object, specific by bitmask.
- uint16_t [GetMask](#) ()
GetMask Returns the state of the 16-bit bitmask within this object.

Private Member Functions

- uint16_t [Wait_i](#) (uint16_t u16Mask_, [EventFlagOperation_t](#) eMode_, uint32_t u32TimeMS_)
Wait_i.

Private Attributes

- uint16_t [m_u16SetMask](#)
Event flags currently set in this object.

Additional Inherited Members

18.6.1 Detailed Description

The [EventFlag](#) class is a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system.

Each [EventFlag](#) object contains a 16-bit bitmask, which is used to trigger events on associated threads. Threads wishing to block, waiting for a specific event to occur can wait on any pattern within this 16-bit bitmask to be set. Here, we provide the ability for a thread to block, waiting for ANY bits in a specified mask to be set, or for ALL bits within a specific mask to be set. Depending on how the object is configured, the bits that triggered the wakeup can be automatically cleared once a match has occurred.

Examples:

[lab7_events/main.cpp](#).

Definition at line 46 of file [eventflag.h](#).

18.6.2 Member Function Documentation

18.6.2.1 void EventFlag::Clear (uint16_t u16Mask_)

ClearFlags - Clear a specific set of flags within this object, specific by bitmask.

Parameters

u16Mask_	- Bitmask of flags to clear
--------------------------	-----------------------------

Examples:

[lab7_events/main.cpp](#).

18.6.2.2 uint16_t EventFlag::GetMask ()

GetMask Returns the state of the 16-bit bitmask within this object.

Returns

The state of the 16-bit bitmask

18.6.2.3 void EventFlag::Set (uint16_t u16Mask_)

Set - Set additional flags in this object (logical OR).

This API can potentially result in threads blocked on [Wait\(\)](#) to be unblocked.

Parameters

<i>u16Mask_</i>	- Bitmask of flags to set.
-----------------	----------------------------

Examples:

[lab7_events/main.cpp](#).

18.6.2.4 `uint16_t EventFlag::Wait (uint16_t u16Mask_, EventFlagOperation_t eMode_)`

Wait - Block a thread on the specific flags in this event flag group.

Parameters

<i>u16Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EVENT_FLAG_ANY: Thread will block on any of the bits in the mask <ul style="list-style-type: none"> • EVENT_FLAG_ALL: Thread will block on all of the bits in the mask

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

Examples:

[lab7_events/main.cpp](#).

18.6.2.5 `uint16_t EventFlag::Wait (uint16_t u16Mask_, EventFlagOperation_t eMode_, uint32_t u32TimeMS_)`

Wait - Block a thread on the specific flags in this event flag group.

Parameters

<i>u16Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EVENT_FLAG_ANY: Thread will block on any of the bits in the mask <ul style="list-style-type: none"> • EVENT_FLAG_ALL: Thread will block on all of the bits in the mask
<i>u32TimeMS_</i>	- Time to block (in ms)

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

18.6.2.6 `uint16_t EventFlag::Wait_i (uint16_t u16Mask_, EventFlagOperation_t eMode_, uint32_t u32TimeMS_)`
[private]

Wait_i.

Internal abstraction used to manage both timed and untimed wait operations

Parameters

--	--

<i>u16Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EVENT_FLAG_ANY: Thread will block on any of the bits in the mask • EVENT_FLAG_ALL: Thread will block on all of the bits in the mask
<i>u32TimeMS_</i>	- Time to block (in ms)

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

18.6.2.7 void EventFlag::WakeMe (Thread * *pclChosenOne_*)

WakeMe.

Wake the given thread, currently blocking on this object

Parameters

<i>pclOwner_</i>	Pointer to the owner thread to unblock.
------------------	---

The documentation for this class was generated from the following file:

- [/media/usb/project/github/Mark3/kernel/public/eventflag.h](#)

18.7 FakeThread_t Struct Reference

If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system.

```
#include <thread.h>
```

Public Attributes

- [K_WORD * m_pwStackTop](#)
Pointer to the top of the thread's stack.
- [K_WORD * m_pwStack](#)
Pointer to the thread's stack.
- [uint8_t m_u8ThreadID](#)
Thread ID.
- [PORT_PRIO_TYPE m_uXPriority](#)
Default priority of the thread.
- [PORT_PRIO_TYPE m_uXCurPriority](#)
Current priority of the thread (priority inheritance)
- [ThreadState_t m_eState](#)
Enum indicating the thread's current state.
- [void * m_pvExtendedContext](#)
Pointer provided to a [Thread](#) to implement thread-local storage.

18.7.1 Detailed Description

If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system.

When cast to a [Thread](#), this data structure will still result in `GetPriority()` calls being valid, which is all that is needed to support the tick-based/tickless times – while saving a fairly decent chunk of RAM on a small micro.

Note that this struct must have the same memory layout as the [Thread](#) class up to the last item.

Definition at line 533 of file [thread.h](#).

The documentation for this struct was generated from the following file:

- [/media/usb/project/github/Mark3/kernel/public/thread.h](#)

18.8 GlobalMessagePool Class Reference

Implements a list of message objects shared between all threads.

```
#include <message.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init.
- static void [Push](#) ([Message](#) *pclMessage_)
Push.
- static [Message](#) * [Pop](#) ()
Pop.
- static [Message](#) * [GetHead](#) ()
GetHead.
- static [MessagePool](#) * [GetPool](#) ()
GetPool.

Static Private Attributes

- static [Message](#) [m_aclMessagePool](#) [[GLOBAL_MESSAGE_POOL_SIZE](#)]
Array of message objects that make up the message pool.

18.8.1 Detailed Description

Implements a list of message objects shared between all threads.

Definition at line 208 of file [message.h](#).

18.8.2 Member Function Documentation

18.8.2.1 static [Message](#)* [GlobalMessagePool::GetHead](#) () [static]

[GetHead](#).

Return a pointer to the first element in the message list

Returns

Pointer to head message element, or NULL if empty

18.8.2.2 `static MessagePool* GlobalMessagePool::GetPool () [static]`

GetPool.

Get the pointer to the underlying message pool object

Returns

Pointer to message pool.

18.8.2.3 `static void GlobalMessagePool::Init () [static]`

Init.

Initialize the message queue prior to use

18.8.2.4 `static Message* GlobalMessagePool::Pop () [static]`

Pop.

Pop a message from the global queue, returning it to the user to be populated before sending by a transmitter.

Returns

Pointer to a [Message](#) object

Examples:

[lab8_messages/main.cpp](#).

18.8.2.5 `static void GlobalMessagePool::Push (Message * pciMessage_) [static]`

Push.

Return a previously-claimed message object back to the global queue. used once the message has been processed by a receiver.

Parameters

<i>pciMessage_</i>	Pointer to the Message object to return back to the global queue
--------------------	--

Examples:

[lab8_messages/main.cpp](#).

The documentation for this class was generated from the following file:

- [/media/usb/project/github/Mark3/kernel/public/message.h](#)

18.9 Kernel Class Reference

Class that encapsulates all of the kernel startup functions.

```
#include <kernel.h>
```

Static Public Member Functions

- static void [Init](#) (void)
Kernel Initialization Function, call before any other OS function.
- static void [Start](#) (void)
Start the operating system kernel - the current execution context is cancelled, all kernel services are started, and the processor resumes execution at the entrypoint for the highest-priority thread.
- static bool [IsStarted](#) ()
IsStarted.
- static void [SetPanic](#) ([PanicFunc_t](#) pfPanic_)
SetPanic Set a function to be called when a kernel panic occurs, giving the user to determine the behavior when a catastrophic failure is observed.
- static bool [IsPanic](#) ()
IsPanic Returns whether or not the kernel is in a panic state.
- static void [Panic](#) (uint16_t u16Cause_)
Panic Cause the kernel to enter its panic state.
- static void [SetIdleFunc](#) ([IdleFunc_t](#) pfIdle_)
SetIdleFunc Set the function to be called when no active threads are available to be scheduled by the scheduler.
- static void [IdleFunc](#) (void)
IdleFunc Call the low-priority idle function when no active threads are available to be scheduled.
- static [Thread *](#) [GetIdleThread](#) (void)
GetIdleThread Return a pointer to the Kernel's idle thread object to the user.
- static void [SetThreadCreateCallout](#) ([ThreadCreateCallout_t](#) pfCreate_)
SetThreadCreateCallout.
- static void [SetThreadExitCallout](#) ([ThreadExitCallout_t](#) pfExit_)
SetThreadExitCallout.
- static void [SetThreadContextSwitchCallout](#) ([ThreadContextCallout_t](#) pfContext_)
SetThreadContextSwitchCallout.
- static [ThreadCreateCallout_t](#) [GetThreadCreateCallout](#) (void)
GetThreadCreateCallout.
- static [ThreadExitCallout_t](#) [GetThreadExitCallout](#) (void)
GetThreadExitCallout.
- static [ThreadContextCallout_t](#) [GetThreadContextSwitchCallout](#) (void)
GetThreadContextSwitchCallout.

Static Private Attributes

- static bool [m_bIsStarted](#)
true if kernel is running, false otherwise
- static bool [m_bIsPanic](#)
true if kernel is in panic state, false otherwise
- static [PanicFunc_t](#) [m_pfPanic](#)
set panic function
- static [IdleFunc_t](#) [m_pfIdle](#)
set idle function
- static [FakeThread_t](#) [m_clIdle](#)
Idle thread object (note: not a real thread)
- static [ThreadCreateCallout_t](#) [m_pfThreadCreateCallout](#)
Function to call on thread creation.
- static [ThreadExitCallout_t](#) [m_pfThreadExitCallout](#)
Function to call on thread exit.
- static [ThreadContextCallout_t](#) [m_pfThreadContextCallout](#)
Function to call on context switch.

18.9.1 Detailed Description

Class that encapsulates all of the kernel startup functions.

Definition at line 44 of file [kernel.h](#).

18.9.2 Member Function Documentation

18.9.2.1 `static Thread* Kernel::GetIdleThread (void) [inline],[static]`

GetIdleThread Return a pointer to the [Kernel](#)'s idle thread object to the user.

Note that the [Thread](#) object involved is to be used for comparisons only – the thread itself is "virtual", and doesn't represent a unique execution context with its own stack.

Returns

Pointer to the [Kernel](#)'s idle thread object

Definition at line 122 of file [kernel.h](#).

18.9.2.2 `static ThreadContextCallout_t Kernel::GetThreadContextSwitchCallout (void) [inline],[static]`

GetThreadContextSwitchCallout.

Return the current function called on every [Thread::ContextSwitchSWI\(\)](#)

Returns

Pointer to the currently-installed callout function, or NULL if not set.

Definition at line 190 of file [kernel.h](#).

18.9.2.3 `static ThreadCreateCallout_t Kernel::GetThreadCreateCallout (void) [inline],[static]`

GetThreadCreateCallout.

Return the current function called on every [Thread::Init\(\)](#);

Returns

Pointer to the currently-installed callout function, or NULL if not set.

Definition at line 172 of file [kernel.h](#).

18.9.2.4 `static ThreadExitCallout_t Kernel::GetThreadExitCallout (void) [inline],[static]`

GetThreadExitCallout.

Return the current function called on every [Thread::Exit\(\)](#);

Returns

Pointer to the currently-installed callout function, or NULL if not set.

Definition at line 181 of file [kernel.h](#).

18.9.2.5 void Kernel::Init(void) [static]

[Kernel](#) Initialization Function, call before any other OS function.

Initializes all global resources used by the operating system. This must be called before any other kernel function is invoked.

Examples:

[buffalogger/main.cpp](#), [lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 67 of file [kernel.cpp](#).

18.9.2.6 static bool Kernel::IsPanic() [inline],[static]

IsPanic Returns whether or not the kernel is in a panic state.

Returns

Whether or not the kernel is in a panic state

Definition at line 90 of file [kernel.h](#).

18.9.2.7 static bool Kernel::IsStarted() [inline],[static]

IsStarted.

Returns

Whether or not the kernel has started - true = running, false = not started

Definition at line 77 of file [kernel.h](#).

18.9.2.8 void Kernel::Panic(uint16_t u16Cause_) [static]

Panic Cause the kernel to enter its panic state.

Parameters

<i>u16Cause_</i>	Reason for the kernel panic
------------------	-----------------------------

Definition at line 110 of file [kernel.cpp](#).

18.9.2.9 static void Kernel::SetIdleFunc(IdleFunc_t pIdle_) [inline],[static]

SetIdleFunc Set the function to be called when no active threads are available to be scheduled by the scheduler.

Parameters

<i>pIdle_</i>	Pointer to the idle function
---------------	------------------------------

Examples:

[lab2_idle_function/main.cpp](#).

Definition at line 103 of file [kernel.h](#).

18.9.2.10 `static void Kernel::SetPanic (PanicFunc_t pfPanic_) [inline],[static]`

SetPanic Set a function to be called when a kernel panic occurs, giving the user to determine the behavior when a catastrophic failure is observed.

Parameters

<i>pfPanic_</i>	Panic function pointer
-----------------	------------------------

Definition at line 85 of file [kernel.h](#).

```
18.9.2.11 static void Kernel::SetThreadContextSwitchCallout ( ThreadContextCallout_t pfContext_ ) [inline],
[static]
```

SetThreadContextSwitchCallout.

Set a function to be called on each context switch.

A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pfContext_</i>	Pointer to a function to call on context switch
-------------------	---

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 159 of file [kernel.h](#).

```
18.9.2.12 static void Kernel::SetThreadCreateCallout ( ThreadCreateCallout_t pfCreate_ ) [inline],[static]
```

SetThreadCreateCallout.

Set a function to be called on creation of a new thread. This callout is executed on the successful completion of a [Thread::Init\(\)](#) call. A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pfCreate_</i>	Pointer to a function to call on thread creation
------------------	--

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 136 of file [kernel.h](#).

```
18.9.2.13 static void Kernel::SetThreadExitCallout ( ThreadExitCallout_t pfExit_ ) [inline],[static]
```

SetThreadExitCallout.

Set a function to be called on thread exit. This callout is executed from the beginning of [Thread::Exit\(\)](#).

A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pfCreate_</i>	Pointer to a function to call on thread exit
------------------	--

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 148 of file [kernel.h](#).

18.9.2.14 void Kernel::Start (void) [static]

Start the operating system kernel - the current execution context is cancelled, all kernel services are started, and the processor resumes execution at the entrypoint for the highest-priority thread.

You must have at least one thread added to the kernel before calling this function, otherwise the behavior is undefined. The exception to this is if the system is configured to use the threadless idle hook, in which case the kernel is allowed to run without any ready threads.

Examples:

[buffalogger/main.cpp](#), [lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 101 of file [kernel.cpp](#).

The documentation for this class was generated from the following files:

- [/media/usb/project/github/Mark3/kernel/public/kernel.h](#)
- [/media/usb/project/github/Mark3/kernel/kernel.cpp](#)

18.10 KernelAware Class Reference

The [KernelAware](#) class.

```
#include <kernelaware.h>
```

Static Public Member Functions

- static void [ProfileInit](#) (const char *szStr_)
ProfileInit.
- static void [ProfileStart](#) (void)
ProfileStart.
- static void [ProfileStop](#) (void)
ProfileStop.
- static void [ProfileReport](#) (void)
ProfileReport.
- static void [ExitSimulator](#) (void)
ExitSimulator.
- static void [Print](#) (const char *szStr_)
Print.
- static void [Trace](#) (uint16_t u16File_, uint16_t u16Line_)
Trace.
- static void [Trace](#) (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_)
Trace.
- static void [Trace](#) (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t u16Arg2_)
Trace.
- static bool [IsSimulatorAware](#) (void)
IsSimulatorAware.

Static Private Member Functions

- static void [Trace_i](#) (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t u16Arg2_, [KernelAwareCommand_t](#) eCmd_)

Trace_i.

18.10.1 Detailed Description

The [KernelAware](#) class.

This class contains functions that are used to trigger kernel-aware functionality within a supported simulation environment (i.e. fIAVR).

These static methods operate on a singleton set of global variables, which are monitored for changes from within the simulator. The simulator hooks into these variables by looking for the correctly-named symbols in an elf-formatted binary being run and registering callbacks that are called whenever the variables are changed. On each change of the command variable, the kernel-aware data is analyzed and interpreted appropriately.

If these methods are run in an unsupported simulator or on actual hardware the commands generally have no effect (except for the exit-on-reset command, which will result in a jump-to-0 reset).

Definition at line 64 of file [kernelaware.h](#).

18.10.2 Member Function Documentation

18.10.2.1 static void KernelAware::ExitSimulator (void) [static]

ExitSimulator.

Instruct the kernel-aware simulator to terminate (destroying the virtual CPU).

18.10.2.2 static bool KernelAware::IsSimulatorAware (void) [static]

IsSimulatorAware.

use this function to determine whether or not the code is running on a simulator that is aware of the kernel.

Returns

true - the application is being run in a kernel-aware simulator. false - otherwise.

18.10.2.3 static void KernelAware::Print (const char * szStr_) [static]

Print.

Instruct the kernel-aware simulator to print a char string

Parameters

<i>szStr_</i>	
---------------	--

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

18.10.2.4 static void KernelAware::ProfileInit (const char * *szStr_*) [static]

ProfileInit.

Initializes the kernel-aware profiler. This function instructs the kernel-aware simulator to reset its accounting variables, and prepare to start counting profiling data tagged to the given string. How this is handled is the responsibility of the simulator.

Parameters

<i>szStr_</i>	String to use as a tag for the profiling session.
---------------	---

18.10.2.5 static void KernelAware::ProfileReport (void) [static]

ProfileReport.

Instruct the kernel-aware simulator to print a report for its current profiling data.

18.10.2.6 static void KernelAware::ProfileStart (void) [static]

ProfileStart.

Instruct the kernel-aware simulator to begin counting cycles towards the current profiling counter.

18.10.2.7 static void KernelAware::ProfileStop (void) [static]

ProfileStop.

Instruct the kernel-aware simulator to end counting cycles relative to the current profiling counter's iteration.

18.10.2.8 static void KernelAware::Trace (uint16_t *u16File_*, uint16_t *u16Line_*) [static]

Trace.

Insert a kernel trace statement into the kernel-aware simulator's debug data stream.

Parameters

<i>u16File_</i>	16-bit code representing the file
<i>u16Line_</i>	16-bit code representing the line in the file

Examples:

[lab11_mailboxes/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

18.10.2.9 static void KernelAware::Trace (uint16_t *u16File_*, uint16_t *u16Line_*, uint16_t *u16Arg1_*) [static]

Trace.

Insert a kernel trace statement into the kernel-aware simulator's debug data stream.

Parameters

<i>u16File_</i>	16-bit code representing the file
<i>u16Line_</i>	16-bit code representing the line in the file

<i>u16Arg1_</i>	16-bit argument to the format string.
-----------------	---------------------------------------

18.10.2.10 `static void KernelAware::Trace (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t u16Arg2_)`
`[static]`

Trace.

Insert a kernel trace statement into the kernel-aware simulator's debug data stream.

Parameters

<i>u16File_</i>	16-bit code representing the file
<i>u16Line_</i>	16-bit code representing the line in the file
<i>u16Arg1_</i>	16-bit argument to the format string.
<i>u16Arg2_</i>	16-bit argument to the format string.

18.10.2.11 `static void KernelAware::Trace_i (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t u16Arg2_, KernelAwareCommand_t eCmd_)` `[static], [private]`

Trace_i.

Private function by which the class's [Trace\(\)](#) methods are reflected, which allows u16 to realize a modest code saving.

Parameters

<i>u16File_</i>	16-bit code representing the file
<i>u16Line_</i>	16-bit code representing the line in the file
<i>u16Arg1_</i>	16-bit argument to the format string.
<i>u16Arg2_</i>	16-bit argument to the format string.
<i>eCmd_</i>	Code indicating the number of arguments to emit.

The documentation for this class was generated from the following file:

- </media/usb/project/github/Mark3/kernel/public/kernelaware.h>

18.11 KernelSWI Class Reference

Class providing the software-interrupt required for context-switching in the kernel.

```
#include <kernelswi.h>
```

Static Public Member Functions

- static void [Config](#) (void)

Config.

- static void [Start](#) (void)

Start.

- static void [Stop](#) (void)

Stop.

- static void [Clear](#) (void)

Clear.

- static void [Trigger](#) (void)

Trigger.

- static uint8_t [DI](#) ()
DI.
- static void [RI](#) (bool bEnable_)
RI.

18.11.1 Detailed Description

Class providing the software-interrupt required for context-switching in the kernel.

Definition at line 31 of file [kernelswi.h](#).

18.11.2 Member Function Documentation

18.11.2.1 void KernelSWI::Clear (void) [static]

Clear.

Clear the software interrupt

Definition at line 68 of file [kernelswi.cpp](#).

18.11.2.2 void KernelSWI::Config (void) [static]

Config.

Configure the software interrupt - must be called before any other software interrupt functions are called.

Definition at line 29 of file [kernelswi.cpp](#).

18.11.2.3 uint8_t KernelSWI::DI () [static]

DI.

Disable the SWI flag itself

Returns

previous status of the SWI, prior to the DI call

Definition at line 50 of file [kernelswi.cpp](#).

18.11.2.4 void KernelSWI::RI (bool bEnable_) [static]

RI.

Restore the state of the SWI to the value specified

Parameters

<i>bEnable_</i>	true - enable the SWI, false - disable SWI
-----------------	--

Definition at line 58 of file [kernelswi.cpp](#).

18.11.2.5 void KernelSWI::Start (void) [static]

Start.

Enable ("Start") the software interrupt functionality

Definition at line 37 of file [kernelswi.cpp](#).

18.11.2.6 void KernelSWI::Stop (void) [static]

Stop.

Disable the software interrupt functionality

Definition at line 44 of file [kernelswi.cpp](#).

18.11.2.7 void KernelSWI::Trigger (void) [static]

Trigger.

Call the software interrupt

Definition at line 74 of file [kernelswi.cpp](#).

The documentation for this class was generated from the following files:

- /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/[kernelswi.h](#)
- /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/[kernelswi.cpp](#)

18.12 KernelTimer Class Reference

Hardware timer interface, used by all scheduling/timer subsystems.

```
#include <kerneltimer.h>
```

Static Public Member Functions

- static void [Config](#) (void)
Config.
- static void [Start](#) (void)
Start.
- static void [Stop](#) (void)
Stop.
- static uint8_t [DI](#) (void)
DI.
- static void [RI](#) (bool bEnable_)
RI.
- static void [EI](#) (void)
EI.
- static [PORT_TIMER_COUNT_TYPE](#) [SubtractExpiry](#) ([PORT_TIMER_COUNT_TYPE](#) uInterval_)
SubtractExpiry.
- static [PORT_TIMER_COUNT_TYPE](#) [TimeToExpiry](#) (void)
TimeToExpiry.
- static [PORT_TIMER_COUNT_TYPE](#) [SetExpiry](#) (uint32_t u32Interval_)
SetExpiry.
- static [PORT_TIMER_COUNT_TYPE](#) [GetOvertime](#) (void)
GetOvertime.
- static void [ClearExpiry](#) (void)
ClearExpiry.
- static [PORT_TIMER_COUNT_TYPE](#) [Read](#) (void)
Read.

18.12.1 Detailed Description

Hardware timer interface, used by all scheduling/timer subsystems.

Definition at line 31 of file [kerneltimer.h](#).

18.12.2 Member Function Documentation

18.12.2.1 void KernelTimer::ClearExpiry (void) [static]

ClearExpiry.

Clear the hardware timer expiry register

Definition at line 196 of file [kerneltimer.cpp](#).

18.12.2.2 void KernelTimer::Config (void) [static]

Config.

Initializes the kernel timer before use

Definition at line 83 of file [kerneltimer.cpp](#).

18.12.2.3 uint8_t KernelTimer::DI (void) [static]

DI.

Disable the kernel timer's expiry interrupt

Definition at line 204 of file [kerneltimer.cpp](#).

18.12.2.4 void KernelTimer::EI (void) [static]

EI.

Enable the kernel timer's expiry interrupt

Definition at line 217 of file [kerneltimer.cpp](#).

18.12.2.5 PORT_TIMER_COUNT_TYPE KernelTimer::GetOvertime (void) [static]

GetOvertime.

Return the number of ticks that have elapsed since the last expiry.

Returns

Number of ticks that have elapsed after last timer expiration

Definition at line 172 of file [kerneltimer.cpp](#).

18.12.2.6 PORT_TIMER_COUNT_TYPE KernelTimer::Read (void) [static]

Read.

Safely read the current value in the timer register

Returns

Value held in the timer register

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 126 of file [kerneltimer.cpp](#).

18.12.2.7 `void KernelTimer::RI (bool bEnable_) [static]`

RI.

Restore the state of the kernel timer's expiry interrupt.

Parameters

<i>bEnable_</i>	1 enable, 0 disable
-----------------	---------------------

Definition at line 223 of file [kerneltimer.cpp](#).

18.12.2.8 `PORT_TIMER_COUNT_TYPE KernelTimer::SetExpiry (uint32_t u32Interval_) [static]`

SetExpiry.

Resets the kernel timer's expiry interval to the specified value

Parameters

<i>u32Interval_</i>	Desired interval in ticks to set the timer for
---------------------	--

Returns

Actual number of ticks set (may be less than desired)

Definition at line 178 of file [kerneltimer.cpp](#).

18.12.2.9 `void KernelTimer::Start (void) [static]`

Start.

Starts the kernel time (must be configured first)

Definition at line 99 of file [kerneltimer.cpp](#).

18.12.2.10 `void KernelTimer::Stop (void) [static]`

Stop.

Shut down the kernel timer, used when no timers are scheduled

Definition at line 114 of file [kerneltimer.cpp](#).

18.12.2.11 `PORT_TIMER_COUNT_TYPE KernelTimer::SubtractExpiry (PORT_TIMER_COUNT_TYPE uInterval_) [static]`

SubtractExpiry.

Subtract the specified number of ticks from the timer's expiry count register. Returns the new expiry value stored in the register.

Parameters

<i>u32Interval_</i>	Time (in HW-specific) ticks to subtract
---------------------	---

Returns

Value in ticks stored in the timer's expiry register

Definition at line 144 of file [kerneltimer.cpp](#).

18.12.2.12 PORT_TIMER_COUNT_TYPE KernelTimer::TimeToExpiry (void) [static]

TimeToExpiry.

Returns the number of ticks remaining before the next timer expiry.

Returns

Time before next expiry in platform-specific ticks

Definition at line 155 of file [kerneltimer.cpp](#).

The documentation for this class was generated from the following files:


- [/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kerneltimer.h](#)
- [/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kerneltimer.cpp](#)

18.13 LinkedList Class Reference

Abstract-data-type from which all other linked-lists are derived.

```
#include <ll.h>
```

Inheritance diagram for LinkedList:



Public Member Functions

- void [Init](#) ()
Init.
- [LinkedListNode](#) * [GetHead](#) ()
GetHead.
- [LinkedListNode](#) * [GetTail](#) ()
GetTail.

Protected Attributes

- [LinkedListNode](#) * [m_pstHead](#)
Pointer to the head node in the list.
- [LinkedListNode](#) * [m_pstTail](#)
Pointer to the tail node in the list.

18.13.1 Detailed Description

Abstract-data-type from which all other linked-lists are derived.

Definition at line 109 of file [ll.h](#).

18.13.2 Member Function Documentation

18.13.2.1 [LinkedListNode](#)* [LinkedList::GetHead](#) () [\[inline\]](#)

GetHead.

Get the head node in the linked list

Returns

Pointer to the head node in the list

Definition at line 134 of file [ll.h](#).

18.13.2.2 [LinkedListNode](#)* [LinkedList::GetTail](#) () [\[inline\]](#)

GetTail.

Get the tail node of the linked list

Returns

Pointer to the tail node in the list

Definition at line 142 of file [ll.h](#).

18.13.2.3 void [LinkedList::Init](#) (void) [\[inline\]](#)

Init.

Clear the linked list.

Definition at line 121 of file [ll.h](#).

The documentation for this class was generated from the following file:


- [/media/usb/project/github/Mark3/kernel/public/ll.h](#)

18.14 LinkedListNode Class Reference

Basic linked-list node data structure.

```
#include <ll.h>
```

Inheritance diagram for [LinkedListNode](#):



Public Member Functions

- [LinkedListNode * GetNext](#) (void)
GetNext.
- [LinkedListNode * GetPrev](#) (void)
GetPrev.

Protected Member Functions

- void [ClearNode](#) ()
ClearNode.

Protected Attributes

- [LinkedListNode * next](#)
Pointer to the next node in the list.
- [LinkedListNode * prev](#)
Pointer to the previous node in the list.

Friends

- class **LinkedList**
- class **DoubleLinkedList**
- class **CircularLinkedList**
- class **ThreadList**

18.14.1 Detailed Description

Basic linked-list node data structure.

This data is managed by the linked-list class types, and can be used transparently between them.

Definition at line 68 of file [ll.h](#).

18.14.2 Member Function Documentation

18.14.2.1 void LinkedListNode::ClearNode () [protected]

ClearNode.

Initialize the linked list node, clearing its next and previous node.

Definition at line 40 of file [ll.cpp](#).

18.14.2.2 `LinkedListNode* LinkedListNode::GetNext (void)` `[inline]`

`GetNext`.

Returns a pointer to the next node in the list.

Returns

a pointer to the next node in the list.

Definition at line 90 of file [ll.h](#).

18.14.2.3 `LinkedListNode* LinkedListNode::GetPrev (void)` `[inline]`

`GetPrev`.

Returns a pointer to the previous node in the list.

Returns

a pointer to the previous node in the list.

Definition at line 98 of file [ll.h](#).

The documentation for this class was generated from the following files:

- [/media/usb/project/github/Mark3/kernel/public/ll.h](#)
- [/media/usb/project/github/Mark3/kernel/ll.cpp](#)

18.15 Mailbox Class Reference

The [Mailbox](#) class implements an IPC mechnism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user.

```
#include <mailbox.h>
```

Public Member Functions

- void [Init](#) (void *pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_)
Init.
- bool [Send](#) (void *pvData_)
Send.
- bool [SendTail](#) (void *pvData_)
SendTail.
- bool [Send](#) (void *pvData_, uint32_t u32TimeoutMS_)
Send.
- bool [SendTail](#) (void *pvData_, uint32_t u32TimeoutMS_)
SendTail.
- void [Receive](#) (void *pvData_)
Receive.
- void [ReceiveTail](#) (void *pvData_)
ReceiveTail.
- bool [Receive](#) (void *pvData_, uint32_t u32TimeoutMS_)
Receive.
- bool [ReceiveTail](#) (void *pvData_, uint32_t u32TimeoutMS_)
ReceiveTail.

Private Member Functions

- void * [GetHeadPointer](#) (void)
GetHeadPointer.
- void * [GetTailPointer](#) (void)
GetTailPointer.
- void [CopyData](#) (const void *src_, const void *dst_, uint16_t len_)
CopyData.
- void [MoveTailForward](#) (void)
MoveTailForward.
- void [MoveHeadForward](#) (void)
MoveHeadForward.
- void [MoveTailBackward](#) (void)
MoveTailBackward.
- void [MoveHeadBackward](#) (void)
MoveHeadBackward.
- bool [Send_i](#) (const void *pvData_, bool bTail_, uint32_t u32TimeoutMS_)
Send_i.
- bool [Receive_i](#) (const void *pvData_, bool bTail_, uint32_t u32WaitTimeMS_)
Receive_i.

Private Attributes

- uint16_t [m_u16Head](#)
Current head index.
- uint16_t [m_u16Tail](#)
Current tail index.
- uint16_t [m_u16Count](#)
Count of items in the mailbox.
- volatile uint16_t [m_u16Free](#)
Current number of free slots in the mailbox.
- uint16_t [m_u16ElementSize](#)
Size of the objects tracked in this mailbox.
- const void * [m_pvBuffer](#)
Pointer to the data-buffer managed by this mailbox.
- [Semaphore m_clRecvSem](#)
Counting semaphore used to synchronize threads on the object.
- [Semaphore m_clSendSem](#)
Binary semaphore for send-blocked threads.

18.15.1 Detailed Description

The [Mailbox](#) class implements an IPC mechanism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user.

Examples:

[lab11_mailboxes/main.cpp](#).

Definition at line 36 of file [mailbox.h](#).

18.15.2 Member Function Documentation

18.15.2.1 `void Mailbox::CopyData (const void * src_, const void * dst_, uint16_t len_)` `[inline],[private]`

CopyData.

Perform a direct byte-copy from a source to a destination object.

Parameters

<i>src_</i>	Pointer to an object to read from
<i>dst_</i>	Pointer to an object to write to
<i>len_</i>	Length to copy (in bytes)

Definition at line 239 of file [mailbox.h](#).

18.15.2.2 `void* Mailbox::GetHeadPointer (void)` `[inline],[private]`

GetHeadPointer.

Return a pointer to the current head of the mailbox's internal circular buffer.

Returns

pointer to the head element in the mailbox

Definition at line 208 of file [mailbox.h](#).

18.15.2.3 `void* Mailbox::GetTailPointer (void)` `[inline],[private]`

GetTailPointer.

Return a pointer to the current tail of the mailbox's internal circular buffer.

Returns

pointer to the tail element in the mailbox

Definition at line 223 of file [mailbox.h](#).

18.15.2.4 `void Mailbox::Init (void * pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_)`

Init.

Initialize the mailbox object prior to its use. This must be called before any calls can be made to the object.

Parameters

<i>pvBuffer_</i>	Pointer to the static buffer to use for the mailbox
<i>u16BufferSize_</i>	Size of the mailbox buffer, in bytes
<i>u16ElementSize_</i>	Size of each envelope, in bytes

Examples:

[lab11_mailboxes/main.cpp](#).

18.15.2.5 `void Mailbox::MoveHeadBackward (void) [inline],[private]`

MoveHeadBackward.

Move the head index backward one element

Definition at line 292 of file [mailbox.h](#).

18.15.2.6 `void Mailbox::MoveHeadForward (void) [inline],[private]`

MoveHeadForward.

Move the head index forward one element

Definition at line 266 of file [mailbox.h](#).

18.15.2.7 `void Mailbox::MoveTailBackward (void) [inline],[private]`

MoveTailBackward.

Move the tail index backward one element

Definition at line 279 of file [mailbox.h](#).

18.15.2.8 `void Mailbox::MoveTailForward (void) [inline],[private]`

MoveTailForward.

Move the tail index forward one element

Definition at line 253 of file [mailbox.h](#).

18.15.2.9 `void Mailbox::Receive (void * pvData_)`

Receive.

Read one envelope from the head of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
----------------	---

Examples:

[lab11_mailboxes/main.cpp](#).

18.15.2.10 `bool Mailbox::Receive (void * pvData_, uint32_t u32TimeoutMS_)`

Receive.

Read one envelope from the head of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered, or the specified time has elapsed without delivery.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
----------------	---

<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.
----------------------	------------------------------------

Returns

true - envelope was delivered, false - delivery timed out.

18.15.2.11 **bool Mailbox::Receive_i (const void * *pvData_*, bool *bTail_*, uint32_t *u32WaitTimeMS_*)** [private]

Receive_i.

Internal method which implements all Read() methods in the class.

Parameters

<i>pvData_</i>	Pointer to the envelope data
<i>bTail_</i>	true - read from tail, false - read from head
<i>u32WaitTimeMS_</i>	Time to wait before timeout (in ms).

Returns

true - read successfully, false - timeout.

18.15.2.12 **void Mailbox::ReceiveTail (void * *pvData_*)**

ReceiveTail.

Read one envelope from the tail of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
----------------	---

18.15.2.13 **bool Mailbox::ReceiveTail (void * *pvData_*, uint32_t *u32TimeoutMS_*)**

ReceiveTail.

Read one envelope from the tail of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered, or the specified time has elapsed without delivery.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

18.15.2.14 **bool Mailbox::Send (void * *pvData_*)**

Send.

Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the head of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
----------------	--

Returns

true - envelope was delivered, false - mailbox is full.

Examples:

[lab11_mailboxes/main.cpp](#).

18.15.2.15 `bool Mailbox::Send (void * pvData_, uint32_t u32TimeoutMS_)`

Send.

Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the head of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

18.15.2.16 `bool Mailbox::Send_i (const void * pvData_, bool bTail_, uint32_t u32TimeoutMS_) [private]`

Send_i.

Internal method which implements all [Send\(\)](#) methods in the class.

Parameters

<i>pvData_</i>	Pointer to the envelope data
<i>bTail_</i>	true - write to tail, false - write to head
<i>u32WaitTimeMS_</i>	Time to wait before timeout (in ms).

Returns

true - data successfully written, false - buffer full

18.15.2.17 `bool Mailbox::SendTail (void * pvData_)`

SendTail.

Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the tail of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
----------------	--

Returns

true - envelope was delivered, false - mailbox is full.

18.15.2.18 `bool Mailbox::SendTail (void * pvData_, uint32_t u32TimeoutMS_)`

SendTail.

Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the tail of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

18.15.3 Member Data Documentation

18.15.3.1 Semaphore `Mailbox::m_clSendSem` `[private]`

Binary semaphore for send-blocked threads.

Definition at line 361 of file [mailbox.h](#).

The documentation for this class was generated from the following file:


- [/media/usb/project/github/Mark3/kernel/public/mailbox.h](#)

18.16 Message Class Reference

Class to provide message-based IPC services in the kernel.

```
#include <message.h>
```

Inheritance diagram for Message:

**Public Member Functions**

- void [Init](#) ()

- Init.*
- void [SetData](#) (void *pvData_)
- SetData.*
- void * [GetData](#) ()
- GetData.*
- void [SetCode](#) (uint16_t u16Code_)
- SetCode.*
- uint16_t [GetCode](#) ()
- GetCode.*

Private Attributes

- void * [m_pvData](#)
- Pointer to the message data.*
- uint16_t [m_u16Code](#)
- Message code, providing context for the message.*

Additional Inherited Members

18.16.1 Detailed Description

Class to provide message-based IPC services in the kernel.

Examples:

[lab8_messages/main.cpp](#).

Definition at line 99 of file [message.h](#).

18.16.2 Member Function Documentation

18.16.2.1 uint16_t Message::GetCode () [inline]

GetCode.

Return the code set in the message upon receipt

Returns

user code set in the object

Examples:

[lab8_messages/main.cpp](#).

Definition at line 146 of file [message.h](#).

18.16.2.2 void* Message::GetData () [inline]

GetData.

Get the data pointer stored in the message upon receipt

Returns

Pointer to the data set in the message object

Examples:

[lab8_messages/main.cpp](#).

Definition at line 130 of file [message.h](#).

18.16.2.3 `void Message::Init(void) [inline]`

Init.

Initialize the data and code in the message.

Definition at line 108 of file [message.h](#).

18.16.2.4 `void Message::SetCode(uint16_t u16Code_) [inline]`

SetCode.

Set the code in the message before transmission

Parameters

<i>u16Code_</i>	Data code to set in the object
-----------------	--------------------------------

Examples:

[lab8_messages/main.cpp](#).

Definition at line 138 of file [message.h](#).

18.16.2.5 `void Message::SetData(void * pvData_) [inline]`

SetData.

Set the data pointer for the message before transmission.

Parameters

<i>pvData_</i>	Pointer to the data object to send in the message
----------------	---

Examples:

[lab8_messages/main.cpp](#).

Definition at line 122 of file [message.h](#).

The documentation for this class was generated from the following file:

- [/media/usb/project/github/Mark3/kernel/public/message.h](#)

18.17 MessagePool Class Reference

Implements a list of message objects.

```
#include <message.h>
```

Public Member Functions

- void [Init](#) ()
Init.
- void [Push](#) ([Message](#) *pclMessage_)
Push.
- [Message](#) * [Pop](#) ()
Pop.
- [Message](#) * [GetHead](#) ()
GetHead.

Private Attributes

- [DoubleLinkedList](#) m_clList
Linked list used to manage the [Message](#) objects.

18.17.1 Detailed Description

Implements a list of message objects.

Definition at line 159 of file [message.h](#).

18.17.2 Member Function Documentation

18.17.2.1 [Message](#)* [MessagePool](#)::[GetHead](#) ()

[GetHead](#).

Return a pointer to the first element in the message list

Returns

18.17.2.2 void [MessagePool](#)::[Init](#) ()

[Init](#).

Initialize the message queue prior to use

18.17.2.3 [Message](#)* [MessagePool](#)::[Pop](#) ()

[Pop](#).

Pop a message from the queue, returning it to the user to be popu32ated before sending by a transmitter.

Returns

Pointer to a [Message](#) object

18.17.2.4 void [MessagePool](#)::[Push](#) ([Message](#) * [pclMessage_](#))

[Push](#).

Return a previously-claimed message object back to the queue. used once the message has been processed by a receiver.

Parameters

<code>pclMessage_</code>	Pointer to the Message object to return back to the queue
--------------------------	---

The documentation for this class was generated from the following file:

- </media/usb/project/github/Mark3/kernel/public/message.h>

18.18 MessageQueue Class Reference

List of messages, used as the channel for sending and receiving messages between threads.

```
#include <message.h>
```

Public Member Functions

- void [Init](#) ()
Init.
- [Message](#) * [Receive](#) ()
Receive.
- [Message](#) * [Receive](#) (uint32_t u32TimeWaitMS_)
Receive.
- void [Send](#) ([Message](#) *pclSrc_)
Send.
- uint16_t [GetCount](#) ()
GetCount.

Private Member Functions

- [Message](#) * [Receive_i](#) (uint32_t u32TimeWaitMS_)
Receive_i.

Private Attributes

- [Semaphore](#) [m_clSemaphore](#)
Counting semaphore used to manage thread blocking.
- [DoubleLinkedList](#) [m_clLinkList](#)
List object used to store messages.

18.18.1 Detailed Description

List of messages, used as the channel for sending and receiving messages between threads.

Examples:

[lab8_messages/main.cpp](#).

Definition at line 269 of file [message.h](#).

18.18.2 Member Function Documentation

18.18.2.1 uint16_t MessageQueue::GetCount ()

GetCount.

Return the number of messages pending in the "receive" queue.

Returns

Count of pending messages in the queue.

18.18.2.2 void MessageQueue::Init ()

Init.

Initialize the message queue prior to use.

Examples:

[lab8_messages/main.cpp](#).

18.18.2.3 Message* MessageQueue::Receive ()

Receive.

Receive a message from the message queue. If the message queue is empty, the thread will block until a message is available.

Returns

Pointer to a message object at the head of the queue

Examples:

[lab8_messages/main.cpp](#).

18.18.2.4 Message* MessageQueue::Receive (uint32_t u32TimeWaitMS_)

Receive.

Receive a message from the message queue. If the message queue is empty, the thread will block until a message is available for the duration specified. If no message arrives within that duration, the call will return with NULL.

Parameters

<i>u32TimeWaitMS_</i>	The amount of time in ms to wait for a message before timing out and unblocking the waiting thread.
-----------------------	---

Returns

Pointer to a message object at the head of the queue or NULL on timeout.

18.18.2.5 Message* MessageQueue::Receive_i (uint32_t u32TimeWaitMS_) [private]

Receive_i.

Internal function used to abstract timed and un-timed Receive calls.

Parameters

<code>u32WaitMS_</code>	Time (in ms) to block, 0 for un-timed call.
-------------------------	---

Returns

Pointer to a message, or 0 on timeout.

18.18.2.6 void MessageQueue::Send (Message * pSrc_)

Send.

Send a message object into this message queue. Will un-block the first waiting thread blocked on this queue if that occurs.

Parameters

<code>pSrc_</code>	Pointer to the message object to add to the queue
--------------------	---

Examples:

[lab8_messages/main.cpp](#).

The documentation for this class was generated from the following file:


- [/media/usb/project/github/Mark3/kernel/public/message.h](#)

18.19 Mutex Class Reference

Mutual-exclusion locks, based on [BlockingObject](#).

```
#include <mutex.h>
```

Inheritance diagram for Mutex:



Public Member Functions

- void [Init](#) ()
Init.
- void [Claim](#) ()
Claim.
- bool [Claim](#) (uint32_t u32WaitTimeMS_)
Claim.
- void [WakeMe](#) (Thread *pOwner_)
WakeMe.
- void [Release](#) ()
Release.

Private Member Functions

- `uint8_t WakeNext ()`
WakeNext.
- `bool Claim_i (uint32_t u32WaitTimeMS_)`
Claim_i.

Private Attributes

- `uint8_t m_u8Recurse`
The recursive lock-count when a mutex is claimed multiple times by the same owner.
- `bool m_bReady`
State of the mutex - true = ready, false = claimed.
- `uint8_t m_u8MaxPri`
Maximum priority of thread in queue, used for priority inheritance.
- `Thread * m_pclOwner`
Pointer to the thread that owns the mutex (when claimed)

Additional Inherited Members

18.19.1 Detailed Description

Mutual-exclusion locks, based on [BlockingObject](#).

Examples:

[lab5_mutexes/main.cpp](#).

Definition at line 64 of file [mutex.h](#).

18.19.2 Member Function Documentation

18.19.2.1 void Mutex::Claim ()

Claim.

Claim the mutex. When the mutex is claimed, no other thread can claim a region protected by the object. If another [Thread](#) currently holds the [Mutex](#) when the Claim method is called, that [Thread](#) will block until the current owner of the mutex releases the [Mutex](#).

If the calling [Thread](#)'s priority is lower than that of a [Thread](#) that currently owns the [Mutex](#) object, then the priority of that [Thread](#) will be elevated to that of the highest-priority calling object until the [Mutex](#) is released. This property is known as "Priority Inheritance"

Note: A single thread can recursively claim a mutex up to a count of

1. Attempting to claim a mutex beyond that will cause a kernel panic.

Examples:

[lab5_mutexes/main.cpp](#).

18.19.2.2 bool Mutex::Claim (uint32_t u32WaitTimeMS_)

Claim.

Claim a mutex, with timeout.

Parameters

<code>u32WaitTimeMS_</code>	
-----------------------------	--

Returns

true - mutex was claimed within the time period specified
false - mutex operation timed-out before the claim operation.

18.19.2.3 `bool Mutex::Claim_i(uint32_t u32WaitTimeMS_) [private]`

Claim_i.

Abstracts out timed/non-timed mutex claim operations.

Parameters

<code>u32WaitTimeMS_</code>	Time in MS to wait, 0 for infinite
-----------------------------	------------------------------------

Returns

true on successful claim, false otherwise

18.19.2.4 `void Mutex::Init ()`

Init.

Initialize a mutex object for use - must call this function before using the object.

Examples:

[lab5_mutexes/main.cpp](#).

18.19.2.5 `void Mutex::Release ()`

Release.

Release the mutex. When the mutex is released, another object can enter the mutex-protected region.

If there are Threads waiting for the [Mutex](#) to become available, then the highest priority [Thread](#) will be unblocked at this time and will claim the [Mutex](#) lock immediately - this may result in an immediate context switch, depending on relative priorities.

If the calling [Thread](#)'s priority was boosted as a result of priority inheritance, the [Thread](#)'s previous priority will also be restored at this time.

Note that if a [Mutex](#) is held recursively, it must be Release'd the same number of times that it was Claim'd before it will be available for use by another [Thread](#).

Examples:

[lab5_mutexes/main.cpp](#).

18.19.2.6 void Mutex::WakeMe (Thread * *pclOwner_*)

WakeMe.

Wake a thread blocked on the mutex. This is an internal function used for implementing timed mutexes relying on timer callbacks. Since these do not have access to the private data of the mutex and its base classes, we have to wrap this as a public method - do not use this for any other purposes.

Parameters

<code>pclOwner_</code>	Thread to unblock from this object.
------------------------	---

18.19.2.7 `uint8_t Mutex::WakeNext () [private]`

WakeNext.

Wake the next thread waiting on the [Mutex](#).

The documentation for this class was generated from the following file:


- </media/usb/project/github/Mark3/kernel/public/mutex.h>

18.20 Notify Class Reference

The [Notify](#) class is a blocking object type, that allows one or more threads to wait for an event to occur before resuming operation.

```
#include <notify.h>
```

Inheritance diagram for Notify:



Public Member Functions

- void [Init](#) (void)
Init.
- void [Signal](#) (void)
Signal.
- void [Wait](#) (bool *pbFlag_)
Wait.
- bool [Wait](#) (uint32_t u32WaitTimeMS_, bool *pbFlag_)
Wait.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)
WakeMe.

Additional Inherited Members

18.20.1 Detailed Description

The [Notify](#) class is a blocking object type, that allows one or more threads to wait for an event to occur before resuming operation.

Examples:

[lab10_notifications/main.cpp](#).

Definition at line 33 of file [notify.h](#).

18.20.2 Member Function Documentation

18.20.2.1 void Notify::Init (void)

Init.

Initialize the Notification object prior to use.

Examples:

[lab10_notifications/main.cpp](#).

18.20.2.2 void Notify::Signal (void)

Signal.

Signal the notification object. This will cause the highest priority thread currently blocking on the object to wake. If no threads are currently blocked on the object, the call has no effect.

Examples:

[lab10_notifications/main.cpp](#).

18.20.2.3 void Notify::Wait (bool * *pbFlag_*)

Wait.

Block the current thread, waiting for a signal on the object.

Parameters

<i>pbFlag_</i>	Flag set to false on block, and true upon wakeup.
----------------	---

Examples:

[lab10_notifications/main.cpp](#).

18.20.2.4 bool Notify::Wait (uint32_t *u32WaitTimeMS_*, bool * *pbFlag_*)

Wait.

Block the current thread, waiting for a signal on the object.

Parameters

<i>u32WaitTimeMS_</i>	Time to wait for the notification event.
<i>pbFlag_</i>	Flag set to false on block, and true upon wakeup.

Returns

true on notification, false on timeout

18.20.2.5 void Notify::WakeMe (Thread * *pclChosenOne_*)

WakeMe.

Wake the specified thread from its current blocking queue. Note that this is only public in order to be accessible from a timer callack.

Parameters

<i>pclChosenOne</i> ↔	Thread to wake up
—	

The documentation for this class was generated from the following file:

- </media/usb/project/github/Mark3/kernel/public/notify.h>

18.21 PriorityMap Class Reference

The [PriorityMap](#) class.

```
#include <priomap.h>
```

Public Member Functions

- [PriorityMap](#) ()
PriorityMap.
- void [Set](#) ([PORT_PRIO_TYPE](#) [uXPrio_](#))
Set Set the priority map bitmap data, at all levels, for the given priority.
- void [Clear](#) ([PORT_PRIO_TYPE](#) [uXPrio_](#))
Clear Clear the priority map bitmap data, at all levels, for the given priority.
- [PORT_PRIO_TYPE](#) [HighestPriority](#) (void)
HighestPriority.

18.21.1 Detailed Description

The [PriorityMap](#) class.

Definition at line 70 of file [priomap.h](#).

18.21.2 Constructor & Destructor Documentation

18.21.2.1 [PriorityMap::PriorityMap](#) ()

[PriorityMap.](#)

Initialize the priority map object, clearing the bitamp data to all 0's.

Definition at line 51 of file [priomap.cpp](#).

18.21.3 Member Function Documentation

18.21.3.1 void [PriorityMap::Clear](#) ([PORT_PRIO_TYPE](#) [uXPrio_](#))

Clear Clear the priority map bitmap data, at all levels, for the given priority.

Parameters

<i>uXPrio_</i>	Priority level to clear the bitmap data for.
----------------	--

Definition at line 78 of file [priomap.cpp](#).

18.21.3.2 `PORT_PRIO_TYPE` `PriorityMap::HighestPriority (void)`

HighestPriority.

Computes the numeric priority of the highest-priority thread represented in the priority map.

Returns

Highest priority ready-thread's number.

Definition at line 94 of file [priomap.cpp](#).

18.21.3.3 `void` `PriorityMap::Set (PORT_PRIO_TYPE uXPrio_)`

Set Set the priority map bitmap data, at all levels, for the given priority.

Parameters

<code>uXPrio_</code>	Priority level to set the bitmap data for.
----------------------	--

Definition at line 64 of file [priomap.cpp](#).

The documentation for this class was generated from the following files:

- [/media/usb/project/github/Mark3/kernel/public/priomap.h](#)
- [/media/usb/project/github/Mark3/kernel/priomap.cpp](#)

18.22 ProfileTimer Class Reference

Profiling timer.

```
#include <profile.h>
```

Public Member Functions

- `void` [Init](#) ()
Init.
- `void` [Start](#) ()
Start.
- `void` [Stop](#) ()
Stop.
- `uint32_t` [GetAverage](#) ()
GetAverage.
- `uint32_t` [GetCurrent](#) ()
GetCurrent.

Private Member Functions

- `uint32_t` [ComputeCurrentTicks](#) (`uint16_t` `u16Current_`, `uint32_t` `u32Epoch_`)
ComputeCurrentTicks.

Private Attributes

- uint32_t [m_u32Cumulative](#)
Cumulative tick-count for this timer.
- uint32_t [m_u32CurrentIteration](#)
Tick-count for the current iteration.
- uint16_t [m_u16Initial](#)
Initial count.
- uint32_t [m_u32InitialEpoch](#)
Initial Epoch.
- uint16_t [m_u16Iterations](#)
Number of iterations executed for this profiling timer.
- bool [m_bActive](#)
Whether or not the timer is active or stopped.

18.22.1 Detailed Description

Profiling timer.

This class is used to perform high-performance profiling of code to see how int32_t certain operations take. useful in instrumenting the performance of key algorithms and time-critical operations to ensure real-timer behavior.

Definition at line 69 of file [profile.h](#).

18.22.2 Member Function Documentation

18.22.2.1 uint32_t ProfileTimer::ComputeCurrentTicks (uint16_t *u16Current_*, uint32_t *u32Epoch_*) [private]

ComputeCurrentTicks.

Figure out how many ticks have elapsed in this iteration

Parameters

<i>u16Count_</i>	Current timer count
<i>u32Epoch_</i>	Current timer epoch

Returns

Current tick count

18.22.2.2 uint32_t ProfileTimer::GetAverage ()

GetAverage.

Get the average time associated with this operation.

Returns

Average tick count normalized over all iterations

18.22.2.3 uint32_t ProfileTimer::GetCurrent ()

GetCurrent.

Return the current tick count held by the profiler. Valid for both active and stopped timers.

Returns

The currently held tick count.

18.22.2.4 void ProfileTimer::Init ()

Init.

Initialize the profiling timer prior to use. Can also be used to reset a timer that's been used previously.

18.22.2.5 void ProfileTimer::Start ()

Start.

Start a profiling session, if the timer is not already active. Has no effect if the timer is already active.

18.22.2.6 void ProfileTimer::Stop ()

Stop.

Stop the current profiling session, adding to the cumulative time for this timer, and the total iteration count.

The documentation for this class was generated from the following file:

- /media/usb/project/github/Mark3/kernel/public/[profile.h](#)

18.23 Quantum Class Reference

Static-class used to implement [Thread](#) quantum functionality, which is a key part of round-robin scheduling.

```
#include <quantum.h>
```

Static Public Member Functions

- static void [UpdateTimer](#) ()
UpdateTimer.
- static void [AddThread](#) ([Thread](#) *pclThread_)
AddThread.
- static void [RemoveThread](#) ()
RemoveThread.
- static void [SetInTimer](#) (void)
SetInTimer.
- static void [ClearInTimer](#) (void)
ClearInTimer.

Static Private Member Functions

- static void [SetTimer](#) ([Thread](#) *pclThread_)
SetTimer.

18.23.1 Detailed Description

Static-class used to implement [Thread](#) quantum functionality, which is a key part of round-robin scheduling.

Definition at line 41 of file [quantum.h](#).

18.23.2 Member Function Documentation

18.23.2.1 `static void Quantum::AddThread (Thread * pclThread_) [static]`

AddThread.

Add the thread to the quantum timer. Only one thread can own the quantum, since only one thread can be running on a core at a time.

18.23.2.2 `static void Quantum::ClearInTimer (void) [inline],[static]`

ClearInTimer.

Clear the flag once the timer callback function has been completed.

Definition at line 83 of file [quantum.h](#).

18.23.2.3 `static void Quantum::RemoveThread () [static]`

RemoveThread.

Remove the thread from the quantum timer. This will cancel the timer.

18.23.2.4 `static void Quantum::SetInTimer (void) [inline],[static]`

SetInTimer.

Set a flag to indicate that the CPU is currently running within the timer-callback routine. This prevents the [Quantum](#) timer from being updated in the middle of a callback cycle, potentially resulting in the kernel timer becoming disabled.

Definition at line 77 of file [quantum.h](#).

18.23.2.5 `static void Quantum::SetTimer (Thread * pclThread_) [static],[private]`

SetTimer.

Set up the quantum timer in the timer scheduler. This creates a one-shot timer, which calls a static callback in [quantum.cpp](#) that on expiry will pivot the head of the threadlist for the thread's priority. This is the mechanism that provides round-robin scheduling in the system.

Parameters

<i>pclThread_</i>	Pointer to the thread to set the Quantum timer on
-------------------	---

18.23.2.6 `static void Quantum::UpdateTimer () [static]`

UpdateTimer.

This function is called to update the thread quantum timer whenever something in the scheduler has changed. This can result in the timer being re-loaded or started. The timer is never stopped, but it may be ignored on expiry.

The documentation for this class was generated from the following file:

- [/media/usb/project/github/Mark3/kernel/public/quantum.h](#)

18.24 Scheduler Class Reference

Priority-based round-robin [Thread](#) scheduling, using ThreadLists for housekeeping.

```
#include <scheduler.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init.
- static void [Schedule](#) ()
Schedule.
- static void [Add](#) ([Thread](#) *pclThread_)
Add.
- static void [Remove](#) ([Thread](#) *pclThread_)
Remove.
- static bool [SetScheduler](#) (bool bEnable_)
SetScheduler.
- static [Thread](#) * [GetCurrentThread](#) ()
GetCurrentThread.
- static volatile [Thread](#) * [GetNextThread](#) ()
GetNextThread.
- static [ThreadList](#) * [GetThreadList](#) ([PORT_PRIO_TYPE](#) uXPriority_)
GetThreadList.
- static [ThreadList](#) * [GetStopList](#) ()
GetStopList.
- static bool [IsEnabled](#) ()
IsEnabled.
- static void [QueueScheduler](#) ()
QueueScheduler.

Static Private Attributes

- static bool [m_bEnabled](#)
[Scheduler](#)'s state - enabled or disabled.
- static bool [m_bQueuedSchedule](#)
Variable representing whether or not there's a queued scheduler operation.
- static [ThreadList](#) [m_clStopList](#)
[ThreadList](#) for all stopped threads.
- static [ThreadList](#) [m_aclPriorities](#) [[KERNEL_NUM_PRIORITIES](#)]
[ThreadLists](#) for all threads at all priorities.
- static [PriorityMap](#) [m_clPrioMap](#)
Priority bitmap lookup structure, 1-bit per thread priority.

18.24.1 Detailed Description

Priority-based round-robin [Thread](#) scheduling, using ThreadLists for housekeeping.

Definition at line 62 of file [scheduler.h](#).

18.24.2 Member Function Documentation

18.24.2.1 void Scheduler::Add (Thread * *pclThread_*) [static]

Add.

Add a thread to the scheduler at its current priority level.

Parameters

<i>pclThread_</i>	Pointer to the thread to add to the scheduler
-------------------	---

Definition at line 89 of file [scheduler.cpp](#).

18.24.2.2 static Thread* Scheduler::GetCurrentThread () [inline],[static]

GetCurrentThread.

Return the pointer to the currently-running thread.

Returns

Pointer to the currently-running thread

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 121 of file [scheduler.h](#).

18.24.2.3 static volatile Thread* Scheduler::GetNextThread () [inline],[static]

GetNextThread.

Return the pointer to the thread that should run next, according to the last run of the scheduler.

Returns

Pointer to the next-running thread

Definition at line 130 of file [scheduler.h](#).

18.24.2.4 static ThreadList* Scheduler::GetStopList () [inline],[static]

GetStopList.

Return the pointer to the list of threads that are in the scheduler's stopped state.

Returns

Pointer to the [ThreadList](#) containing the stopped threads

Definition at line 150 of file [scheduler.h](#).

18.24.2.5 static ThreadList* Scheduler::GetThreadList (PORT_PRIO_TYPE *uXPriority_*) [inline],[static]

GetThreadList.

Return the pointer to the active list of threads that are at the given priority level in the scheduler.

Parameters

<i>uXPriority_</i>	Priority level of the threadlist
--------------------	----------------------------------

Returns

Pointer to the [ThreadList](#) for the given priority level

Definition at line 141 of file [scheduler.h](#).

18.24.2.6 void Scheduler::Init (void) [static]

Init.

Intialize the scheduler, must be called before use.

Definition at line 54 of file [scheduler.cpp](#).

18.24.2.7 static bool Scheduler::IsEnabled () [inline],[static]

IsEnabled.

Return the current state of the scheduler - whether or not scheudling is enabled or disabled.

Returns

true - scheduler enabled, false - disabled

Definition at line 159 of file [scheduler.h](#).

18.24.2.8 static void Scheduler::QueueScheduler () [inline],[static]

QueueScheduler.

Tell the kernel to perform a scheduling operation as soon as the scheduler is re-enabled.

Definition at line 166 of file [scheduler.h](#).

18.24.2.9 void Scheduler::Remove (Thread * *pciThread_*) [static]

Remove.

Remove a thread from the scheduler at its current priority level.

Parameters

<i>pciThread_</i>	Pointer to the thread to be removed from the scheduler
-------------------	--

Definition at line 95 of file [scheduler.cpp](#).

18.24.2.10 void Scheduler::Schedule () [static]

Schedule.

Run the scheduler, determines the next thread to run based on the current state of the threads. Note that the next-thread chosen from this function is only valid while in a critical section.

Definition at line 63 of file [scheduler.cpp](#).

18.24.2.11 `bool Scheduler::SetScheduler (bool bEnable_) [static]`

SetScheduler.

Set the active state of the scheduler. When the scheduler is disabled, the *next thread* is never set; the currently running thread will run forever until the scheduler is enabled again. Care must be taken to ensure that we don't end up trying to block while the scheduler is disabled, otherwise the system ends up in an unusable state.

Parameters

<i>bEnable_</i>	true to enable, false to disable the scheduler
-----------------	--

Definition at line 101 of file [scheduler.cpp](#).

The documentation for this class was generated from the following files:


- [/media/usb/project/github/Mark3/kernel/public/scheduler.h](#)
- [/media/usb/project/github/Mark3/kernel/scheduler.cpp](#)

18.25 Semaphore Class Reference

Binary & Counting semaphores, based on [BlockingObject](#) base class.

```
#include <ksemaphore.h>
```

Inheritance diagram for Semaphore:



Public Member Functions

- void [Init](#) (uint16_t u16InitVal_, uint16_t u16MaxVal_)
Initialize a semaphore before use.
- bool [Post](#) ()
Increment the semaphore count.
- void [Pend](#) ()
Decrement the semaphore count.
- uint16_t [GetCount](#) ()
Return the current semaphore counter.
- bool [Pend](#) (uint32_t u32WaitTimeMS_)
Decrement the semaphore count.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)
Wake a thread blocked on the semaphore.

Private Member Functions

- uint8_t [WakeNext](#) ()
Wake the next thread waiting on the semaphore.
- bool [Pend_i](#) (uint32_t u32WaitTimeMS_)
Pend_i.

Private Attributes

- `uint16_t m_u16Value`
Current count held by the semaphore.
- `uint16_t m_u16MaxValue`
Maximum count that can be held by this semaphore.

Additional Inherited Members

18.25.1 Detailed Description

Binary & Counting semaphores, based on [BlockingObject](#) base class.

Examples:

[buffalogger/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab6_timers/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).↔

Definition at line 37 of file [ksemaphore.h](#).

18.25.2 Member Function Documentation

18.25.2.1 `uint16_t Semaphore::GetCount ()`

Return the current semaphore counter.

This can be used by a thread to bypass blocking on a semaphore - allowing it to do other things until a non-zero count is returned, instead of blocking until the semaphore is posted.

Returns

The current semaphore counter value.

18.25.2.2 `void Semaphore::Init (uint16_t u16InitVal_, uint16_t u16MaxVal_)`

Initialize a semaphore before use.

Must be called before attempting post/pend operations on the object.

This initialization is required to configure the behavior of the semaphore with regards to the initial and maximum values held by the semaphore. By providing access to the raw initial and maximum count elements of the semaphore, these objects are able to be used as either counting or binary semaphores.

To configure a semaphore object for use as a binary semaphore, set values of 0 and 1 respectively for the initial/maximum value parameters.

Any other combination of values can be used to implement a counting semaphore.

Parameters

<code>u16InitVal_</code>	Initial value held by the semaphore
<code>u16MaxVal_</code>	Maximum value for the semaphore. Must be nonzero.

Examples:

[buffalogger/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab6_timers/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).↔

18.25.2.3 void Semaphore::Pend ()

Decrement the semaphore count.

If the count is zero, the calling [Thread](#) will block until the semaphore is posted, and the [Thread](#)'s priority is higher than that of any other [Thread](#) blocked on the object.

Examples:

[buffalogger/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab6_timers/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

18.25.2.4 bool Semaphore::Pend (uint32_t u32WaitTimeMS_)

Decrement the semaphore count.

If the count is zero, the thread will block until the semaphore is pended. If the specified interval expires before the thread is unblocked, then the status is returned back to the user.

Returns

true - semaphore was acquired before the timeout false - timeout occurred before the semaphore was claimed.

18.25.2.5 bool Semaphore::Pend_i (uint32_t u32WaitTimeMS_) [private]

Pend_i.

Internal function used to abstract timed and untimed semaphore pend operations.

Parameters

u32WaitTimeMS_ S_	Time in MS to wait
--------------------------------------	--------------------

Returns

true on success, false on failure.

18.25.2.6 bool Semaphore::Post ()

Increment the semaphore count.

If the semaphore count is zero at the time this is called, and there are threads blocked on the object, this will immediately unblock the highest-priority blocked [Thread](#).

Note that if the priority of that [Thread](#) is higher than the current thread's priority, a context switch will occur and control will be relinquished to that [Thread](#).

Returns

true if the semaphore was posted, false if the count is already maxed out.

Examples:

[buffalogger/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab6_timers/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

18.25.2.7 void Semaphore::WakeMe (Thread * pChosenOne_)

Wake a thread blocked on the semaphore.

This is an internal function used for implementing timed semaphores relying on timer callbacks. Since these do not have access to the private data of the semaphore and its base classes, we have to wrap this as a public method - do not use this for any other purposes.

18.25.2.8 uint8_t Semaphore::WakeNext () [private]

Wake the next thread waiting on the semaphore.

Used internally.

The documentation for this class was generated from the following file:


- </media/usb/project/github/Mark3/kernel/public/ksemaphore.h>

18.26 Thread Class Reference

Object providing fundamental multitasking support in the kernel.

```
#include <thread.h>
```

Inheritance diagram for Thread:



Public Member Functions

- void **Init** (K_WORD *pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_, ThreadEntry_↔ t pfEntryPoint_, void *pvArg_)
Init.
- void **Start** ()
Start.
- void **Stop** ()
Stop.
- ThreadList * **GetOwner** (void)
GetOwner.
- ThreadList * **GetCurrent** (void)
GetCurrent.
- PORT_PRIO_TYPE **GetPriority** (void)
GetPriority.
- PORT_PRIO_TYPE **GetCurPriority** (void)
GetCurPriority.
- void **SetQuantum** (uint16_t u16Quantum_)
SetQuantum.
- uint16_t **GetQuantum** (void)
GetQuantum.
- void **SetCurrent** (ThreadList *pClNewList_)

- SetCurrent.*
- void [SetOwner](#) ([ThreadList](#) *pclNewList_)
- SetOwner.*
- void [SetPriority](#) ([PORT_PRIO_TYPE](#) uXPriority_)
- SetPriority.*
- void [InheritPriority](#) ([PORT_PRIO_TYPE](#) uXPriority_)
- InheritPriority.*
- void [Exit](#) ()
- Exit.*
- void [SetID](#) ([uint8_t](#) u8ID_)
- SetID.*
- [uint8_t](#) [GetID](#) ()
- GetID.*
- [uint16_t](#) [GetStackSlack](#) ()
- GetStackSlack.*
- [uint16_t](#) [GetEventFlagMask](#) ()
- GetEventFlagMask* returns the thread's current event-flag mask, which is used in conjunction with the [EventFlag](#) blocking object type.
- void [SetEventFlagMask](#) ([uint16_t](#) u16Mask_)
- SetEventFlagMask* Sets the active event flag bitfield mask.
- void [SetEventFlagMode](#) ([EventFlagOperation_t](#) eMode_)
- SetEventFlagMode* Sets the active event flag operation mode.
- [EventFlagOperation_t](#) [GetEventFlagMode](#) ()
- GetEventFlagMode* Returns the thread's event flag's operating mode.
- [Timer](#) * [GetTimer](#) ()
- Return a pointer to the thread's timer object.*
- void [SetExpired](#) (bool bExpired_)
- SetExpired.*
- bool [GetExpired](#) ()
- GetExpired.*
- void [InitIdle](#) ()
- InitIdle* Initialize this [Thread](#) object as the [Kernel](#)'s idle thread.
- void * [GetExtendedContext](#) ()
- GetExtendedContext.*
- void [SetExtendedContext](#) (void *pvData_)
- SetExtendedContext.*
- [ThreadState_t](#) [GetState](#) ()
- GetState* Returns the current state of the thread to the caller.
- void [SetState](#) ([ThreadState_t](#) eState_)
- SetState* Set the thread's state to a new value.
- [K_WORD](#) * [GetStack](#) ()
- GetStack.*
- [uint16_t](#) [GetStackSize](#) ()
- GetStackSize.*

Static Public Member Functions

- static void [Sleep](#) ([uint32_t](#) u32TimeMs_)
- Sleep.*
- static void [USleep](#) ([uint32_t](#) u32TimeUs_)
- USleep.*
- static void [Yield](#) (void)
- Yield.*

Private Member Functions

- void [SetPriorityBase](#) ([PORT_PRIO_TYPE](#) uXPriority_)
SetPriorityBase.

Static Private Member Functions

- static void [ContextSwitchSWI](#) (void)
ContextSwitchSWI.

Private Attributes

- [K_WORD](#) * [m_pwStackTop](#)
Pointer to the top of the thread's stack.
- [K_WORD](#) * [m_pwStack](#)
Pointer to the thread's stack.
- [uint8_t](#) [m_u8ThreadID](#)
Thread ID.
- [PORT_PRIO_TYPE](#) [m_uXPriority](#)
Default priority of the thread.
- [PORT_PRIO_TYPE](#) [m_uXCurPriority](#)
Current priority of the thread (priority inheritance)
- [ThreadState_t](#) [m_eState](#)
Enum indicating the thread's current state.
- void * [m_pvExtendedContext](#)
Pointer provided to a [Thread](#) to implement thread-local storage.
- [uint16_t](#) [m_u16StackSize](#)
Size of the stack (in bytes)
- [ThreadList](#) * [m_pclCurrent](#)
Pointer to the thread-list where the thread currently resides.
- [ThreadList](#) * [m_pclOwner](#)
Pointer to the thread-list where the thread resides when active.
- [ThreadEntry_t](#) [m_pfEntryPoint](#)
The entry-point function called when the thread starts.
- void * [m_pvArg](#)
Pointer to the argument passed into the thread's entypoint.
- [uint16_t](#) [m_u16Quantum](#)
Thread quantum (in milliseconds)
- [uint16_t](#) [m_u16FlagMask](#)
Event-flag mask.
- [EventFlagOperation_t](#) [m_eFlagMode](#)
Event-flag mode.
- [Timer](#) [m_clTimer](#)
Timer used for blocking-object timeouts.
- bool [m_bExpired](#)
Indicate whether or not a blocking-object timeout has occurred.

Friends

- class [ThreadPort](#)

Additional Inherited Members

18.26.1 Detailed Description

Object providing fundamental multitasking support in the kernel.

Examples:

[buffalogger/main.cpp](#), [lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 60 of file [thread.h](#).

18.26.2 Member Function Documentation

18.26.2.1 void Thread::ContextSwitchSWI(void) [static], [private]

ContextSwitchSWI.

This code is used to trigger the context switch interrupt. Called whenever the kernel decides that it is necessary to swap out the current thread for the "next" thread.

Definition at line 493 of file [thread.cpp](#).

18.26.2.2 void Thread::Exit()

Exit.

Remove the thread from being scheduled again. The thread is effectively destroyed when this occurs. This is extremely useful for cases where a thread encounters an unrecoverable error and needs to be restarted, or in the context of systems where threads need to be created and destroyed dynamically.

This must not be called on the idle thread.

Examples:

[lab9_dynamic_threads/main.cpp](#).

18.26.2.3 PORT_PRIO_TYPE Thread::GetCurPriority(void) [inline]

GetCurPriority.

Return the priority of the current thread

Returns

Priority of the current thread

Definition at line 182 of file [thread.h](#).

18.26.2.4 ThreadList* Thread::GetCurrent(void) [inline]

GetCurrent.

Return the [ThreadList](#) where the thread is currently located

Returns

Pointer to the thread's current list

Definition at line 165 of file [thread.h](#).

18.26.2.5 uint16_t Thread::GetEventFlagMask () [inline]

GetEventFlagMask returns the thread's current event-flag mask, which is used in conjunction with the [EventFlag](#) blocking object type.

Returns

A copy of the thread's event flag mask

Definition at line 327 of file [thread.h](#).

18.26.2.6 EventFlagOperation_t Thread::GetEventFlagMode () [inline]

GetEventFlagMode Returns the thread's event flag's operating mode.

Returns

The thread's event flag mode.

Definition at line 343 of file [thread.h](#).

18.26.2.7 bool Thread::GetExpired ()

GetExpired.

Return the status of the most-recent blocking call on the thread.

Returns

true - call expired, false - call did not expire

18.26.2.8 void* Thread::GetExtendedContext () [inline]

GetExtendedContext.

Return the [Thread](#) object's extended-context data pointer. Used by code implementing a user-defined thread-local storage model. Pointer exists only for the lifespan of the [Thread](#).

Returns

[Thread](#)'s extended context data pointer.

Definition at line 391 of file [thread.h](#).

18.26.2.9 uint8_t Thread::GetID () [inline]

GetID.

Return the 8-bit ID corresponding to this thread.

Returns

[Thread](#)'s 8-bit ID, set by the user

Definition at line 304 of file [thread.h](#).

18.26.2.10 ThreadList* Thread::GetOwner (void) [inline]

GetOwner.

Return the [ThreadList](#) where the thread belongs when it's in the active/ready state in the scheduler.

Returns

Pointer to the [Thread](#)'s owner list

Definition at line 157 of file [thread.h](#).

18.26.2.11 PORT_PRIO_TYPE Thread::GetPriority (void) [inline]

GetPriority.

Return the priority of the current thread

Returns

Priority of the current thread

Definition at line 174 of file [thread.h](#).

18.26.2.12 uint16_t Thread::GetQuantum (void) [inline]

GetQuantum.

Get the thread's round-robin execution quantum.

Returns

The thread's quantum

Definition at line 199 of file [thread.h](#).

18.26.2.13 K_WORD* Thread::GetStack () [inline]

GetStack.

Returns

Pointer to the blob of memory used as the thread's stack

Definition at line 427 of file [thread.h](#).

18.26.2.14 uint16_t Thread::GetStackSize () [inline]

GetStackSize.

Returns

Size of the thread's stack in bytes

Definition at line 433 of file [thread.h](#).

18.26.2.15 `uint16_t Thread::GetStackSlack ()`

GetStackSlack.

Performs a (somewhat lengthy) check on the thread stack to check the amount of stack margin (or "slack") remaining on the stack. If you're having problems with blowing your stack, you can run this function at points in your code during development to see what operations cause problems. Also useful during development as a tool to optimally size thread stacks.

Returns

The amount of slack (unused bytes) on the stack

ToDo : Reverse the logic for MCUs where stack grows UP instead of down

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 352 of file [thread.cpp](#).

18.26.2.16 `ThreadState_t Thread::GetState ()` `[inline]`

GetState Returns the current state of the thread to the caller.

Can be used to determine whether or not a thread is ready (or running), stopped, or terminated/exit'd.

Returns

ThreadState_t representing the thread's current state

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 413 of file [thread.h](#).

18.26.2.17 `void Thread::InheritPriority (PORT_PRIO_TYPE uXPriority_)`

InheritPriority.

Allow the thread to run at a different priority level (temporarily) for the purpose of avoiding priority inversions. This should only be called from within the implementation of blocking-objects.

Parameters

<i>uXPriority_</i>	New Priority to boost to.
--------------------	---------------------------

Definition at line 482 of file [thread.cpp](#).

18.26.2.18 `void Thread::Init (K_WORD * pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_, ThreadEntry_t pfEntryPoint_, void * pvArg_)`

Init.

Initialize a thread prior to its use. Initialized threads are placed in the stopped state, and are not scheduled until the thread's start method has been invoked first.

Parameters

<i>pwStack_</i>	Pointer to the stack to use for the thread
<i>u16StackSize_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint_</i>	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entrypoint function.

Examples:

[buffalogger/main.cpp](#), [lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 70 of file [thread.cpp](#).

18.26.2.19 void Thread::InitIdle ()

InitIdle Initialize this [Thread](#) object as the [Kernel](#)'s idle thread.

There should only be one of these, maximum, in a given system.

18.26.2.20 void Thread::SetCurrent (ThreadList * *pcNewList_*) [inline]

SetCurrent.

Set the thread's current to the specified thread list

Parameters

<i>pcNewList_</i>	Pointer to the threadlist to apply thread ownership
-------------------	---

Definition at line 209 of file [thread.h](#).

18.26.2.21 void Thread::SetEventFlagMask (uint16_t *u16Mask_*) [inline]

SetEventFlagMask Sets the active event flag bitfield mask.

Parameters

<i>u16Mask_</i>	
-----------------	--

Definition at line 332 of file [thread.h](#).

18.26.2.22 void Thread::SetEventFlagMode (EventFlagOperation_t *eMode_*) [inline]

SetEventFlagMode Sets the active event flag operation mode.

Parameters

<i>eMode_</i>	Event flag operation mode, defines the logical operator to apply to the event flag.
---------------	---

Definition at line 338 of file [thread.h](#).

18.26.2.23 void Thread::SetExpired (bool *bExpired_*)

SetExpired.

Set the status of the current blocking call on the thread.

Parameters

<i>bExpired_</i>	true - call expired, false - call did not expire
------------------	--

18.26.2.24 void Thread::SetExtendedContext (void * *pvData_*) [inline]

SetExtendedContext.

Assign the [Thread](#) object's extended-context data pointer. Used by code implementing a user-defined thread-local storage model.

Object assigned to the context pointer should persist for the duration of the [Thread](#).

Parameters

<i>pvData_</i>	Object to assign to the extended data pointer.+
----------------	---

Definition at line 404 of file [thread.h](#).

18.26.2.25 void Thread::SetID (uint8_t *u8ID_*) [inline]

SetID.

Set an 8-bit ID to uniquely identify this thread.

Parameters

<i>u8ID_</i>	8-bit Thread ID, set by the user
--------------	--

Definition at line 296 of file [thread.h](#).

18.26.2.26 void Thread::SetOwner (ThreadList * *pclNewList_*) [inline]

SetOwner.

Set the thread's owner to the specified thread list

Parameters

<i>pclNewList_</i>	Pointer to the threadlist to apply thread ownership
--------------------	---

Definition at line 217 of file [thread.h](#).

18.26.2.27 void Thread::SetPriority (PORT_PRIO_TYPE *uXPriority_*)

SetPriority.

Set the priority of the [Thread](#) (running or otherwise) to a different level. This activity involves re-scheduling, and must be done so with due caution, as it may effect the determinism of the system.

This should *always* be called from within a critical section to prevent system issues.

Parameters

<i>uXPriority_</i>	New priority of the thread
--------------------	----------------------------

Definition at line 433 of file [thread.cpp](#).

18.26.2.28 void Thread::SetPriorityBase (PORT_PRIO_TYPE *uXPriority_*) [private]

SetPriorityBase.

Parameters

<i>uXPriority_</i>	
--------------------	--

Definition at line 419 of file [thread.cpp](#).

18.26.2.29 void Thread::SetQuantum (uint16_t *u16Quantum_*) [inline]

SetQuantum.

Set the thread's round-robin execution quantum.

Parameters

<i>u16Quantum_</i>	Thread's execution quantum (in milliseconds)
--------------------	--

Examples:

[lab3_round_robin/main.cpp](#).

Definition at line 191 of file [thread.h](#).

18.26.2.30 void Thread::SetState (ThreadState_t *eState_*) [inline]

SetState Set the thread's state to a new value.

This is only to be used by code within the kernel, and is not intended for use by an end-user.

Parameters

<i>eState_</i>	New thread state to set.
----------------	--------------------------

Definition at line 421 of file [thread.h](#).

18.26.2.31 static void Thread::Sleep (uint32_t *u32TimeMs_*) [static]

Sleep.

Put the thread to sleep for the specified time (in milliseconds). Actual time slept may be longer (but not less than) the interval specified.

Parameters

<i>u32TimeMs_</i>	Time to sleep (in ms)
-------------------	-----------------------

Examples:

[buffalogger/main.cpp](#), [lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

18.26.2.32 void Thread::Start (void)

Start.

Start the thread - remove it from the stopped list, add it to the scheduler's list of threads (at the thread's set priority), and continue along.

Examples:

[buffalogger/main.cpp](#), [lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), and [lab6_threads/main.cpp](#).

[cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 146 of file [thread.cpp](#).

18.26.2.33 void Thread::Stop ()

Stop.

Stop a thread that's actively scheduled without destroying its stacks. Stopped threads can be restarted using the [Start\(\)](#) API.

Definition at line 188 of file [thread.cpp](#).

18.26.2.34 static void Thread::USleep (uint32_t u32TimeUs) [static]

USleep.

Put the thread to sleep for the specified time (in microseconds). Actual time slept may be longer (but not less than) the interval specified.

Parameters

<i>u32TimeUs_</i>	Time to sleep (in microseconds)
-------------------	---------------------------------

18.26.2.35 void Thread::Yield (void) [static]

Yield.

Yield the thread - this forces the system to call the scheduler and determine what thread should run next. This is typically used when threads are moved in and out of the scheduler.

Definition at line 387 of file [thread.cpp](#).

The documentation for this class was generated from the following files:


- [/media/usb/project/github/Mark3/kernel/public/thread.h](#)
- [/media/usb/project/github/Mark3/kernel/thread.cpp](#)

18.27 ThreadList Class Reference

This class is used for building thread-management facilities, such as schedulers, and blocking objects.

```
#include <threadlist.h>
```

Inheritance diagram for ThreadList:



Public Member Functions

- [ThreadList](#) ()
ThreadList.
- void [SetPriority](#) ([PORT_PRIO_TYPE](#) uXPriority_)
SetPriority.
- void [SetMapPointer](#) ([PriorityMap](#) *pclMap_)
SetMapPointer.
- void [Add](#) ([LinkListNode](#) *node_)
Add.
- void [Add](#) ([LinkListNode](#) *node_, [PriorityMap](#) *pclMap_, [PORT_PRIO_TYPE](#) uXPriority_)
Add.
- void [AddPriority](#) ([LinkListNode](#) *node_)
AddPriority.
- void [Remove](#) ([LinkListNode](#) *node_)
Remove.
- [Thread](#) * [HighestWaiter](#) ()
HighestWaiter.

Private Attributes

- [PORT_PRIO_TYPE](#) m_uXPriority
Priority of the threadlist.
- [PriorityMap](#) * m_pclMap
Pointer to the bitmap/flag to set when used for scheduling.

Additional Inherited Members

18.27.1 Detailed Description

This class is used for building thread-management facilities, such as schedulers, and blocking objects.

Definition at line 35 of file [threadlist.h](#).

18.27.2 Constructor & Destructor Documentation

18.27.2.1 ThreadList::ThreadList () [inline]

[ThreadList](#).

Default constructor - zero-initializes the data.

Definition at line 44 of file [threadlist.h](#).

18.27.3 Member Function Documentation

18.27.3.1 void ThreadList::Add (LinkListNode * node_)

Add.

Add a thread to the threadlist.

Parameters

<i>node_</i>	Pointer to the thread (link list node) to add to the list
--------------	---

Definition at line 52 of file [threadlist.cpp](#).

18.27.3.2 void ThreadList::Add (LinkListNode * *node_*, PriorityMap * *pclMap_*, PORT_PRIO_TYPE *uXPriority_*)

Add.

Add a thread to the threadlist, specifying the flag and priority at the same time.

Parameters

<i>node_</i>	Pointer to the thread to add (link list node)
<i>pclMap_</i>	Pointer to the bitmap flag to set (if used in a scheduler context), or NULL for non-scheduler.
<i>uXPriority_</i>	Priority of the threadlist

Definition at line 101 of file [threadlist.cpp](#).

18.27.3.3 void ThreadList::AddPriority (LinkListNode * *node_*)

AddPriority.

Add a thread to the list such that threads are ordered from highest to lowest priority from the head of the list.

Parameters

<i>node_</i>	Pointer to a thread to add to the list.
--------------	---

Definition at line 65 of file [threadlist.cpp](#).

18.27.3.4 Thread * ThreadList::HighestWaiter ()

HighestWaiter.

Return a pointer to the highest-priority thread in the thread-list.

Returns

Pointer to the highest-priority thread

Definition at line 124 of file [threadlist.cpp](#).

18.27.3.5 void ThreadList::Remove (LinkListNode * *node_*)

Remove.

Remove the specified thread from the threadlist

Parameters

<i>node_</i>	Pointer to the thread to remove
--------------	---------------------------------

Definition at line 111 of file [threadlist.cpp](#).

18.27.3.6 void ThreadList::SetMapPointer (PriorityMap * *pclMap_*)

SetMapPointer.

Set the pointer to a bitmap to use for this threadlist. Once again, only needed when the threadlist is being used for scheduling purposes.

Parameters

<i>pclMap_</i>	Pointer to the priority map object used to track this thread.
----------------	---

Definition at line 46 of file [threadlist.cpp](#).

18.27.3.7 void ThreadList::SetPriority (PORT_PRIO_TYPE uXPriority_)

SetPriority.

Set the priority of this threadlist (if used for a scheduler).

Parameters

<i>uXPriority_</i>	Priority level of the thread list
--------------------	-----------------------------------

Definition at line 40 of file [threadlist.cpp](#).

The documentation for this class was generated from the following files:

- [/media/usb/project/github/Mark3/kernel/public/threadlist.h](#)
- [/media/usb/project/github/Mark3/kernel/threadlist.cpp](#)

18.28 ThreadPort Class Reference

Class defining the architecture specific functions required by the kernel.

```
#include <threadport.h>
```

Static Public Member Functions

- static void [StartThreads](#) ()
StartThreads.

Static Private Member Functions

- static void [InitStack](#) ([Thread](#) *pstThread_)
InitStack.

Friends

- class [Thread](#)

18.28.1 Detailed Description

Class defining the architecture specific functions required by the kernel.

This is limited (at this point) to a function to start the scheduler, and a function to initialize the default stack-frame for a thread.

Definition at line 186 of file [threadport.h](#).

18.28.2 Member Function Documentation

18.28.2.1 `void ThreadPort::InitStack (Thread * pstThread_) [static], [private]`

InitStack.

Initialize the thread's stack.

Parameters

<i>pstThread_</i>	Pointer to the thread to initialize
-------------------	-------------------------------------

Definition at line 39 of file [threadport.cpp](#).

18.28.2.2 `void ThreadPort::StartThreads () [static]`

StartThreads.

Function to start the scheduler, initial threads, etc.

Definition at line 130 of file [threadport.cpp](#).

The documentation for this class was generated from the following files:


- [/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/threadport.h](#)
- [/media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/threadport.cpp](#)

18.29 Timer Class Reference

Kernel-managed software timers.

```
#include <timer.h>
```

Inheritance diagram for Timer:



Public Member Functions

- [Timer](#) ()
Timer.
- void [Init](#) ()
Init.
- void [Start](#) (bool bRepeat_, uint32_t u32IntervalMs_, [TimerCallback_t](#) pfCallback_, void *pvData_)
Start.
- void [Start](#) (bool bRepeat_, uint32_t u32IntervalMs_, uint32_t u32ToleranceMs_, [TimerCallback_t](#) pfCallback_, void *pvData_)
Start.
- void [Start](#) ()
Start.
- void [Stop](#) ()
Stop.

- void [SetFlags](#) (uint8_t u8Flags_)
SetFlags.
- void [SetCallback](#) (TimerCallback_t pfCallback_)
SetCallback.
- void [SetData](#) (void *pvData_)
SetData.
- void [SetOwner](#) (Thread *pclOwner_)
SetOwner.
- void [SetIntervalTicks](#) (uint32_t u32Ticks_)
SetIntervalTicks.
- void [SetIntervalSeconds](#) (uint32_t u32Seconds_)
SetIntervalSeconds.
- uint32_t [GetInterval](#) ()
GetInterval.
- void [SetIntervalMSeconds](#) (uint32_t u32MSeconds_)
SetIntervalMSeconds.
- void [SetIntervalUSeconds](#) (uint32_t u32USeconds_)
SetIntervalUSeconds.
- void [SetTolerance](#) (uint32_t u32Ticks_)
SetTolerance.

Private Member Functions

- void [SetInitialized](#) ()
SetInitialized.
- bool [IsInitialized](#) (void)
IsInitialized.

Private Attributes

- uint8_t [m_u8Initialized](#)
Cookie used to determine whether or not the timer is initialized.
- uint8_t [m_u8Flags](#)
Flags for the timer, defining if the timer is one-shot or repeated.
- [TimerCallback_t](#) [m_pfCallback](#)
Pointer to the callback function.
- uint32_t [m_u32Interval](#)
Interval of the timer in timer ticks.
- uint32_t [m_u32TimeLeft](#)
Time remaining on the timer.
- uint32_t [m_u32TimerTolerance](#)
Maximum tolerance (used for timer harmonization)
- [Thread](#) * [m_pclOwner](#)
Pointer to the owner thread.
- void * [m_pvData](#)
Pointer to the callback data.

Friends

- class [TimerList](#)

Additional Inherited Members

18.29.1 Detailed Description

Kernel-managed software timers.

Kernel-managed timers, used to provide high-precision high-resolution delays. Functionality is useful to both user-code, and is used extensively within the kernel and its blocking objects to implement round-robin scheduling, thread sleep, and timeouts. Relies on a single hardware timer, which is multiplexed through the kernel.

Examples:

[lab6_timers/main.cpp](#).

Definition at line 112 of file [timer.h](#).

18.29.2 Constructor & Destructor Documentation

18.29.2.1 Timer::Timer ()

[Timer](#).

Default Constructor - Do nothing. Allow the init call to perform the necessary object initialization prior to use.

18.29.3 Member Function Documentation

18.29.3.1 uint32_t Timer::GetInterval () [inline]

GetInterval.

Return the timer's configured interval in ticks

Returns

[Timer](#) interval in ticks.

Definition at line 238 of file [timer.h](#).

18.29.3.2 void Timer::Init ()

Init.

Re-initialize the [Timer](#) to default values.

18.29.3.3 bool Timer::IsInitialized (void) [inline],[private]

IsInitialized.

Returns

Definition at line 280 of file [timer.h](#).

18.29.3.4 void Timer::SetCallback (TimerCallback_t pfCallback_) [inline]

SetCallback.

Define the callback function to be executed on expiry of the timer

Parameters

<i>pfCallback_</i>	Pointer to the callback function to call
--------------------	--

Definition at line 195 of file [timer.h](#).

18.29.3.5 void Timer::SetData (void * *pvData_*) [inline]

SetData.

Define a pointer to be sent to the timer callbacak on timer expiry

Parameters

<i>pvData_</i>	Pointer to data to pass as argument into the callback
----------------	---

Definition at line 203 of file [timer.h](#).

18.29.3.6 void Timer::SetFlags (uint8_t *u8Flags_*) [inline]

SetFlags.

Set the timer's flags based on the bits in the *u8Flags_* argument

Parameters

<i>u8Flags_</i>	Flags to assign to the timer object. TIMERLIST_FLAG_ONE_SHOT for a one-shot timer, 0 for a continuous timer.
-----------------	--

Definition at line 187 of file [timer.h](#).

18.29.3.7 void Timer::SetIntervalMSeconds (uint32_t *u32MSeconds_*)

SetIntervalMSeconds.

Set the timer expiry interval in milliseconds (platform agnostic)

Parameters

<i>u32MSeconds_</i>	Time in milliseconds
---------------------	----------------------

18.29.3.8 void Timer::SetIntervalSeconds (uint32_t *u32Seconds_*)

SetIntervalSeconds.

Set the timer expiry interval in seconds (platform agnostic)

Parameters

<i>u32Seconds_</i>	Time in seconds
--------------------	-----------------

18.29.3.9 void Timer::SetIntervalTicks (uint32_t *u32Ticks_*)

SetIntervalTicks.

Set the timer expiry in system-ticks (platform specific!)

Parameters

<i>u32Ticks_</i>	Time in ticks
------------------	---------------

18.29.3.10 void Timer::SetIntervalUSeconds (uint32_t *u32USeconds_*)

SetIntervalUSeconds.

Set the timer expiry interval in microseconds (platform agnostic)

Parameters

<i>u32USeconds_</i>	Time in microseconds
---------------------	----------------------

18.29.3.11 void Timer::SetOwner (Thread * *pclOwner_*) [inline]

SetOwner.

Set the owner-thread of this timer object (all timers must be owned by a thread).

Parameters

<i>pclOwner_</i>	Owner thread of this timer object
------------------	-----------------------------------

Definition at line 212 of file [timer.h](#).

18.29.3.12 void Timer::SetTolerance (uint32_t *u32Ticks_*)

SetTolerance.

Set the timer's maximum tolerance in order to synchronize timer processing with other timers in the system.

Parameters

<i>u32Ticks_</i>	Maximum tolerance in ticks
------------------	----------------------------

18.29.3.13 void Timer::Start (bool *bRepeat_*, uint32_t *u32IntervalMs_*, TimerCallback_t *pfCallback_*, void * *pvData_*)

Start.

Start a timer using default ownership, using repeats as an option, and millisecond resolution.

Parameters

<i>bRepeat_</i>	0 - timer is one-shot. 1 - timer is repeating.
<i>u32IntervalMs_</i>	- Interval of the timer in milliseconds
<i>pfCallback_</i>	- Function to call on timer expiry
<i>pvData_</i>	- Data to pass into the callback function

Examples:

[lab6_timers/main.cpp](#).

18.29.3.14 void Timer::Start (bool *bRepeat_*, uint32_t *u32IntervalMs_*, uint32_t *u32ToleranceMs_*, TimerCallback_t *pfCallback_*, void * *pvData_*)

Start.

Start a timer using default ownership, using repeats as an option, and millisecond resolution.

Parameters

<i>bRepeat_</i>	0 - timer is one-shot. 1 - timer is repeating.
<i>u32IntervalMs_</i>	- Interval of the timer in milliseconds
<i>u32ToleranceMs_</i>	- Allow the timer expiry to be delayed by an additional maximum time, in order to have as many timers expire at the same time as possible.
<i>pfCallback_</i>	- Function to call on timer expiry
<i>pvData_</i>	- Data to pass into the callback function

18.29.3.15 void Timer::Start ()

Start.

Start or restart a timer using parameters previously configured via calls to Start(<with args>), or via the a-la-carte parameter setter methods. This is especially useful for retriggering one-shot timers that have previously expired, using the timer's previous configuration.

18.29.3.16 void Timer::Stop ()

Stop.

Stop a timer already in progress. Has no effect on timers that have already been stopped.

The documentation for this class was generated from the following file:


- </media/usb/project/github/Mark3/kernel/public/timer.h>

18.30 TimerList Class Reference

[TimerList](#) class - a doubly-linked-list of timer objects.

```
#include <timerlist.h>
```

Inheritance diagram for TimerList:



Public Member Functions

- void [Init](#) ()
Init.
- void [Add](#) (Timer *pclListNode_)
Add.
- void [Remove](#) (Timer *pclLinkedListNode_)
Remove.
- void [Process](#) ()
Process.

Private Attributes

- uint32_t [m_u32NextWakeup](#)
The time (in system clock ticks) of the next wakeup event.
- bool [m_bTimerActive](#)
Whether or not the timer is active.

Additional Inherited Members

18.30.1 Detailed Description

[TimerList](#) class - a doubly-linked-list of timer objects.

Definition at line 39 of file [timerlist.h](#).

18.30.2 Member Function Documentation

18.30.2.1 void TimerList::Add (Timer * *pcListNode_*)

Add.

Add a timer to the [TimerList](#).

Parameters

<i>pcListNode_</i>	Pointer to the Timer to Add
--------------------	---

18.30.2.2 void TimerList::Init ()

Init.

Initialize the [TimerList](#) object. Must be called before using the object.

18.30.2.3 void TimerList::Process ()

Process.

Process all timers in the timerlist as a result of the timer expiring. This will select a new timer epoch based on the next timer to expire.

18.30.2.4 void TimerList::Remove (Timer * *pcLinkListNode_*)

Remove.

Remove a timer from the [TimerList](#), cancelling its expiry.

Parameters

<i>pcListNode_</i>	Pointer to the Timer to remove
--------------------	--

The documentation for this class was generated from the following file:

- [/media/usb/project/github/Mark3/kernel/public/timerlist.h](#)

18.31 TimerScheduler Class Reference

"Static" Class used to interface a global [TimerList](#) with the rest of the kernel.

```
#include <timerscheduler.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init.
- static void [Add](#) ([Timer](#) *pclListNode_)
Add.
- static void [Remove](#) ([Timer](#) *pclListNode_)
Remove.
- static void [Process](#) ()
Process.

Static Private Attributes

- static [TimerList](#) [m_clTimerList](#)
[TimerList](#) object manipulated by the [Timer Scheduler](#).

18.31.1 Detailed Description

"Static" Class used to interface a global [TimerList](#) with the rest of the kernel.

Definition at line 38 of file [timerscheduler.h](#).

18.31.2 Member Function Documentation

18.31.2.1 static void [TimerScheduler::Add](#) ([Timer](#) * [pclListNode_](#)) [inline],[static]

Add.

Add a timer to the timer scheduler. Adding a timer implicitly starts the timer as well.

Parameters

pclListNode_	Pointer to the timer list node to add
------------------------------	---------------------------------------

Definition at line 56 of file [timerscheduler.h](#).

18.31.2.2 static void [TimerScheduler::Init](#) (void) [inline],[static]

Init.

Initialize the timer scheduler. Must be called before any timer, or timer-derived functions are used.

Definition at line 47 of file [timerscheduler.h](#).

18.31.2.3 static void [TimerScheduler::Process](#) () [inline],[static]

Process.

This function must be called on timer expiry (from the timer's ISR context). This will result in all timers being updated based on the epoch that just elapsed. The next timer epoch is set based on the next [Timer](#) object to expire.

Definition at line 74 of file [timerscheduler.h](#).

18.31.2.4 static void TimerScheduler::Remove (Timer * *pcListNode_*) [inline],[static]

Remove.

Remove a timer from the timer scheduler. May implicitly stop the timer if this is the only active timer scheduled.

Parameters

<i>pcListNode_</i>	Pointer to the timer list node to remove
--------------------	--

Definition at line 65 of file [timerscheduler.h](#).

The documentation for this class was generated from the following file:

- [/media/usb/project/github/Mark3/kernel/public/timerscheduler.h](#)

File Documentation

Basic Atomic Operations.

```
00001 /*-----
00002
00003 |_____|_____|_____|_____|_____|_____|_____|_____|_____|_____|
00004 | \   / | \   / | \   / | \   / | \   / | \   / | \   / |
00005 |  V   |  V   |  V   |  V   |  V   |  V   |  V   |  V   |
00006 |_/___\|_/___\|_/___\|_/___\|_/___\|_/___\|_/___\|_/___\|
00007 |_____||_____||_____||_____||_____||_____||_____||
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023 #include "atomic.h"
00024 #include "threadport.h"
00025
00026 #define _CAN_HAS_DEBUG
00027 //--[Autogenerated - Do Not Modify]-----
00028 #include "dbg_file_list.h"
00029 #include "buffalogger.h"
00030 #if defined(DBG_FILE)
00031 #error "Debug logging file token already defined! Bailing."
00032 #else
00033 #define DBG_FILE _DBG__KERNEL_ATOMIC_CPP
00034 #endif
00035 //--[End Autogenerated content]-----
00036
00037 #if KERNEL_USE_ATOMIC
```

```

00038
00039 //-----
00040 uint8_t Atomic::Set(uint8_t* pu8Source_, uint8_t u8Val_)
00041 {
00042     uint8_t u8Ret;
00043     CS_ENTER();
00044     u8Ret = *pu8Source_;
00045     *pu8Source_ = u8Val_;
00046     CS_EXIT();
00047     return u8Ret;
00048 }
00049 //-----
00050 uint16_t Atomic::Set(uint16_t* pu16Source_, uint16_t u16Val_)
00051 {
00052     uint16_t u16Ret;
00053     CS_ENTER();
00054     u16Ret = *pu16Source_;
00055     *pu16Source_ = u16Val_;
00056     CS_EXIT();
00057     return u16Ret;
00058 }
00059 //-----
00060 uint32_t Atomic::Set(uint32_t* pu32Source_, uint32_t u32Val_)
00061 {
00062     uint32_t u32Ret;
00063     CS_ENTER();
00064     u32Ret = *pu32Source_;
00065     *pu32Source_ = u32Val_;
00066     CS_EXIT();
00067     return u32Ret;
00068 }
00069
00070 //-----
00071 uint8_t Atomic::Add(uint8_t* pu8Source_, uint8_t u8Val_)
00072 {
00073     uint8_t u8Ret;
00074     CS_ENTER();
00075     u8Ret = *pu8Source_;
00076     *pu8Source_ += u8Val_;
00077     CS_EXIT();
00078     return u8Ret;
00079 }
00080
00081 //-----
00082 uint16_t Atomic::Add(uint16_t* pu16Source_, uint16_t u16Val_)
00083 {
00084     uint16_t u16Ret;
00085     CS_ENTER();
00086     u16Ret = *pu16Source_;
00087     *pu16Source_ += u16Val_;
00088     CS_EXIT();
00089     return u16Ret;
00090 }
00091
00092 //-----
00093 uint32_t Atomic::Add(uint32_t* pu32Source_, uint32_t u32Val_)
00094 {
00095     uint32_t u32Ret;
00096     CS_ENTER();
00097     u32Ret = *pu32Source_;
00098     *pu32Source_ += u32Val_;
00099     CS_EXIT();
00100     return u32Ret;
00101 }
00102
00103 //-----
00104 uint8_t Atomic::Sub(uint8_t* pu8Source_, uint8_t u8Val_)
00105 {
00106     uint8_t u8Ret;
00107     CS_ENTER();
00108     u8Ret = *pu8Source_;
00109     *pu8Source_ -= u8Val_;
00110     CS_EXIT();
00111     return u8Ret;
00112 }
00113
00114 //-----
00115 uint16_t Atomic::Sub(uint16_t* pu16Source_, uint16_t u16Val_)
00116 {
00117     uint16_t u16Ret;
00118     CS_ENTER();
00119     u16Ret = *pu16Source_;
00120     *pu16Source_ -= u16Val_;
00121     CS_EXIT();
00122     return u16Ret;
00123 }
00124

```

```
00125 //-----
00126 uint32_t Atomic::Sub(uint32_t* pu32Source_, uint32_t u32Val_)
00127 {
00128     uint32_t u32Ret;
00129     CS_ENTER();
00130     u32Ret = *pu32Source_;
00131     *pu32Source_ -= u32Val_;
00132     CS_EXIT();
00133     return u32Ret;
00134 }
00135
00136 //-----
00137 bool Atomic::TestAndSet(bool* pbLock_)
00138 {
00139     uint8_t u8Ret;
00140     CS_ENTER();
00141     u8Ret = *pbLock_;
00142     if (!u8Ret) {
00143         *pbLock_ = 1;
00144     }
00145     CS_EXIT();
00146     return u8Ret;
00147 }
00148
00149 #endif // KERNEL_USE_ATOMIC
```

19.3 /media/usb/project/github/Mark3/kernel/autoalloc.cpp File Reference

Automatic memory allocation for kernel objects.

```
#include "mark3cfg.h"
#include "mark3.h"
#include "autoalloc.h"
#include "threadport.h"
#include "kernel.h"
```

19.3.1 Detailed Description

Automatic memory allocation for kernel objects.

Definition in file [autoalloc.cpp](#).

19.4 autoalloc.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00020 #include "mark3cfg.h"
00021 #include "mark3.h"
00022 #include "autoalloc.h"
00023 #include "threadport.h"
00024 #include "kernel.h"
00025
00026 #if KERNEL_USE_AUTO_ALLOC
00027
00028 // Align to nearest word boundary
00029 #define ALLOC_ALIGN(x) (((x) + (sizeof(K_ADDR) - 1)) & (sizeof(K_ADDR) - 1))
00030
00031 //-----
00032 uint8_t AutoAlloc::m_au8AutoHeap[AUTO_ALLOC_SIZE];
00033 K_ADDR AutoAlloc::m_aHeapTop;

```

```

00034
00035 //-----
00036 void AutoAlloc::Init(void)
00037 {
00038     m_aHeapTop = (K_ADDR) (m_au8AutoHeap);
00039 }
00040
00041 //-----
00042 void* AutoAlloc::Allocate(uint16_t u16Size_)
00043 {
00044     void* pvRet = 0;
00045
00046     CS_ENTER();
00047     uint16_t u16AllocSize = ALLOC_ALIGN(u16Size_);
00048     if (((K_ADDR)m_aHeapTop - (K_ADDR)&m_au8AutoHeap[0]) + u16AllocSize) < AUTO_ALLOC_SIZE) {
00049         pvRet = (void*)m_aHeapTop;
00050         m_aHeapTop += u16AllocSize;
00051     }
00052     CS_EXIT();
00053
00054     if (!pvRet) {
00055         Kernel::Panic(PANIC_AUTO_HEAP_EXHAUSTED);
00056     }
00057
00058     return pvRet;
00059 }
00060
00061 #if KERNEL_USE_SEMAPHORE
00062 //-----
00063 Semaphore* AutoAlloc::NewSemaphore(void)
00064 {
00065     void* pvObj = Allocate(sizeof(Semaphore));
00066     if (pvObj) {
00067         return new (pvObj) Semaphore();
00068     }
00069     return 0;
00070 }
00071 #endif
00072
00073 #if KERNEL_USE_MUTEX
00074 //-----
00075 Mutex* AutoAlloc::NewMutex(void)
00076 {
00077     void* pvObj = Allocate(sizeof(Mutex));
00078     if (pvObj) {
00079         return new (pvObj) Mutex();
00080     }
00081     return 0;
00082 }
00083 #endif
00084
00085 #if KERNEL_USE_EVENTFLAG
00086 //-----
00087 EventFlag* AutoAlloc::NewEventFlag(void)
00088 {
00089     void* pvObj = Allocate(sizeof(EventFlag));
00090     if (pvObj) {
00091         return new (pvObj) EventFlag();
00092     }
00093     return 0;
00094 }
00095 #endif
00096
00097 #if KERNEL_USE_MESSAGE
00098 //-----
00099 Message* AutoAlloc::NewMessage(void)
00100 {
00101     void* pvObj = Allocate(sizeof(Message));
00102     if (pvObj) {
00103         return new (pvObj) Message();
00104     }
00105     return 0;
00106 }
00107 //-----
00108 MessageQueue* AutoAlloc::NewMessageQueue(void)
00109 {
00110     void* pvObj = Allocate(sizeof(MessageQueue));
00111     if (pvObj) {
00112         return new (pvObj) MessageQueue();
00113     }
00114     return 0;
00115 }
00116
00117 #endif
00118
00119 #if KERNEL_USE_NOTIFY
00120 //-----

```

```
00121 Notify* AutoAlloc::NewNotify(void)
00122 {
00123     void* pvObj = Allocate(sizeof(Notify));
00124     if (pvObj) {
00125         return new (pvObj) Notify();
00126     }
00127     return 0;
00128 }
00129 #endif
00130
00131 #if KERNEL_USE_MAILBOX
00132 //-----
00133 Mailbox* AutoAlloc::NewMailbox(void)
00134 {
00135     void* pvObj = Allocate(sizeof(Mailbox));
00136     if (pvObj) {
00137         return new (pvObj) Mailbox();
00138     }
00139     return 0;
00140 }
00141 #endif
00142
00143 //-----
00144 Thread* AutoAlloc::NewThread(void)
00145 {
00146     void* pvObj = Allocate(sizeof(Thread));
00147     if (pvObj) {
00148         return new (pvObj) Thread();
00149     }
00150     return 0;
00151 }
00152
00153 #if KERNEL_USE_TIMERS
00154 //-----
00155 Timer* AutoAlloc::NewTimer(void)
00156 {
00157     void* pvObj = Allocate(sizeof(Timer));
00158     if (pvObj) {
00159         return new (pvObj) Timer();
00160     }
00161     return 0;
00162 }
00163 #endif
00164
00165 #endif
```

19.5 /media/usb/project/github/Mark3/kernel/blocking.cpp File Reference

Implementation of base class for blocking objects.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "blocking.h"
#include "thread.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

19.5.1 Detailed Description

Implementation of base class for blocking objects.

Definition in file [blocking.cpp](#).

19.6 blocking.cpp

[illegible]

19.7 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kernelprofile.cpp File Reference

ATMega328p Profiling timer implementation.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "profile.h"
#include "kernelprofile.h"
#include "threadport.h"
#include <avr/io.h>
#include <avr/interrupt.h>
```

19.7.1 Detailed Description

ATMega328p Profiling timer implementation.

Definition in file [kernelprofile.cpp](#).

19.8 kernelprofile.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #include "kerneltypes.h"
00021 #include "mark3cfg.h"
00022 #include "profile.h"
00023 #include "kernelprofile.h"
00024 #include "threadport.h"
00025 #include <avr/io.h>
00026 #include <avr/interrupt.h>
00027
00028 #if KERNEL_USE_PROFILER
00029 uint32_t Profiler::m_u32Epoch;
00030
00031 //-----
00032 void Profiler::Init()
00033 {
00034     TCCR0A = 0;
00035     TCCR0B = 0;
00036     TIFR0 = 0;
00037     TIMSK0 = 0;
00038     m_u32Epoch = 0;
00039 }
00040
00041 //-----
00042 void Profiler::Start()
00043 {
00044     TIFR0 = 0;
00045     TCNT0 = 0;
00046     TCCR0B |= (1 << CS01);
00047     TIMSK0 |= (1 << TOIE0);
00048 }
00049
00050 //-----
00051 void Profiler::Stop()
00052 {
00053     TIFR0 = 0;
00054     TCCR0B &= ~(1 << CS01);
00055     TIMSK0 &= ~(1 << TOIE0);
00056 }
00057 //-----
```

```

00058 PORT_TIMER_COUNT_TYPE Profiler::Read()
00059 {
00060     uint16_t u16Ret;
00061     CS_ENTER();
00062     TCCR0B &= ~(1 << CS01);
00063     u16Ret = TCNT0;
00064     TCCR0B |= (1 << CS01);
00065     CS_EXIT();
00066     return u16Ret;
00067 }
00068
00069 //-----
00070 void Profiler::Process()
00071 {
00072     CS_ENTER();
00073     m_u32Epoch++;
00074     CS_EXIT();
00075 }
00076
00077 //-----
00078 ISR(TIMER0_OVF_vect)
00079 {
00080     Profiler::Process();
00081 }
00082
00083 #endif

```

19.9 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kernelswi.cpp File Reference

Kernel Software interrupt implementation for ATmega328p.

```

#include "kerneltypes.h"
#include "kernelswi.h"
#include <avr/io.h>
#include <avr/interrupt.h>

```

19.9.1 Detailed Description

Kernel Software interrupt implementation for ATmega328p.

Definition in file [kernelswi.cpp](#).

19.10 kernelswi.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "kernelswi.h"
00024
00025 #include <avr/io.h>
00026 #include <avr/interrupt.h>
00027
00028 //-----
00029 void KernelSWI::Config(void)
00030 {
00031     PORTD &= ~0x04; // Clear INT0
00032     DDRD |= 0x04; // Set PortD, bit 2 (INT0) As Output
00033     EICRA |= (1 << ISC00) | (1 << ISC01); // Rising edge on INT0

```



```

00034 }
00035
00036 //-----
00037 void KernelSWI::Start(void)
00038 {
00039     EIFR &= ~(1 << INTF0); // Clear any pending interrupts on INT0
00040     EIMSK |= (1 << INT0); // Enable INT0 interrupt (as int32_t as I-bit is set)
00041 }
00042
00043 //-----
00044 void KernelSWI::Stop(void)
00045 {
00046     EIMSK &= ~(1 << INT0); // Disable INT0 interrupts
00047 }
00048
00049 //-----
00050 uint8_t KernelSWI::DI()
00051 {
00052     bool bEnabled = ((EIMSK & (1 << INT0)) != 0);
00053     EIMSK &= ~(1 << INT0);
00054     return bEnabled;
00055 }
00056
00057 //-----
00058 void KernelSWI::RI(bool bEnable_)
00059 {
00060     if (bEnable_) {
00061         EIMSK |= (1 << INT0);
00062     } else {
00063         EIMSK &= ~(1 << INT0);
00064     }
00065 }
00066
00067 //-----
00068 void KernelSWI::Clear(void)
00069 {
00070     EIFR &= ~(1 << INTF0); // Clear the interrupt flag for INT0
00071 }
00072
00073 //-----
00074 void KernelSWI::Trigger(void)
00075 {
00076     // if(Thread_IsSchedulerEnabled())
00077     {
00078         PORTD &= ~0x04;
00079         PORTD |= 0x04;
00080     }
00081 }

```

19.11 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/kerneltimer.cpp File Reference

[Kernel Timer](#) Implementation for ATmega328p.

```

#include "kerneltypes.h"
#include "kerneltimer.h"
#include "mark3cfg.h"
#include "ksemaphore.h"
#include "thread.h"
#include <avr/common.h>
#include <avr/io.h>
#include <avr/interrupt.h>

```

Functions

- [ISR \(TIMER1_COMPA_vect\)](#)

[ISR\(TIMER1_COMPA_vect\)](#) *Timer interrupt ISR - service the timer thread.*


```

00086 #if KERNEL_TIMERS_THREADED
00087     s_clTimerSemaphore.Init(0, 1);
00088     s_clTimerThread.Init(s_clTimerThreadStack,
00089                          sizeof(s_clTimerThreadStack) / sizeof(K_WORD),
00090                          KERNEL_TIMERS_THREAD_PRIORITY,
00091                          KernelTimer_Task,
00092                          0);
00093     Quantum::SetTimerThread(&s_clTimerThread);
00094     s_clTimerThread.Start();
00095 #endif
00096 }
00097
00098 //-----
00099 void KernelTimer::Start(void)
00100 {
00101     #if !KERNEL_TIMERS_TICKLESS
00102         TCCR1B = ((1 << WGM12) | (1 << CS11) | (1 << CS10));
00103         OCR1A = ((PORT_SYSTEM_FREQ / 1000) / 64);
00104     #else
00105         TCCR1B |= (1 << CS12);
00106     #endif
00107     TCNT1 = 0;
00108     TIFR1 &= ~TIMER_IFR;
00109     TIMSK1 |= TIMER_IMSK;
00110 }
00111
00112 //-----
00113 void KernelTimer::Stop(void)
00114 {
00115     #if KERNEL_TIMERS_TICKLESS
00116         TIFR1 &= ~TIMER_IFR;
00117         TIMSK1 &= ~TIMER_IMSK;
00118         TCCR1B &= ~(1 << CS12); // Disable count...
00119         TCNT1 = 0;
00120         OCR1A = 0;
00121     #endif
00122 }
00123
00124 //-----
00125 PORT_TIMER_COUNT_TYPE KernelTimer::Read(void)
00126 {
00127     #if KERNEL_TIMERS_TICKLESS
00128         volatile uint16_t u16Read1;
00129         volatile uint16_t u16Read2;
00130         do {
00131             u16Read1 = TCNT1;
00132             u16Read2 = TCNT1;
00133         } while (u16Read1 != u16Read2);
00134         return u16Read1;
00135     #else
00136         return 0;
00137     #endif
00138 }
00139
00140 //-----
00141 PORT_TIMER_COUNT_TYPE KernelTimer::SubtractExpiry(
00142     PORT_TIMER_COUNT_TYPE uInterval)
00143 {
00144     #if KERNEL_TIMERS_TICKLESS
00145         OCR1A -= uInterval;
00146         return OCR1A;
00147     #else
00148         return 0;
00149     #endif
00150 }
00151
00152 //-----
00153 PORT_TIMER_COUNT_TYPE KernelTimer::TimeToExpiry(void)
00154 {
00155     #if KERNEL_TIMERS_TICKLESS
00156         uint16_t u16Read = KernelTimer::Read();
00157         uint16_t u16OCR1A = OCR1A;
00158         if (u16Read >= u16OCR1A) {
00159             return 0;
00160         } else {
00161             return (u16OCR1A - u16Read);
00162         }
00163     #else
00164         return 0;
00165     #endif
00166 }
00167
00168 //-----
00169
00170
00171

```

```

00172 PORT_TIMER_COUNT_TYPE KernelTimer::GetOvertime(void)
00173 {
00174     return KernelTimer::Read();
00175 }
00176
00177 //-----
00178 PORT_TIMER_COUNT_TYPE KernelTimer::SetExpiry(uint32_t
u32Interval_)
00179 {
00180     #if KERNEL_TIMERS_TICKLESS
00181         uint16_t ul6SetInterval;
00182         if (u32Interval_ > 65535) {
00183             ul6SetInterval = 65535;
00184         } else {
00185             ul6SetInterval = (uint16_t)u32Interval_;
00186         }
00187         OCR1A = ul6SetInterval;
00188         return ul6SetInterval;
00189     #else
00190         return 0;
00191     #endif
00192 }
00193
00194 //-----
00196 void KernelTimer::ClearExpiry(void)
00197 {
00198     #if KERNEL_TIMERS_TICKLESS
00199         OCR1A = 65535; // Clear the compare value
00200     #endif
00201 }
00202
00203 //-----
00204 uint8_t KernelTimer::DI(void)
00205 {
00206     #if KERNEL_TIMERS_TICKLESS
00207         bool bEnabled = ((TIMSK1 & (TIMER_IMSK)) != 0);
00208         TIFR1 &= ~TIMER_IFR; // Clear interrupt flags
00209         TIMSK1 &= ~TIMER_IMSK; // Disable interrupt
00210         return bEnabled;
00211     #else
00212         return 0;
00213     #endif
00214 }
00215
00216 //-----
00217 void KernelTimer::EI(void)
00218 {
00219     KernelTimer::RI(0);
00220 }
00221
00222 //-----
00223 void KernelTimer::RI(bool bEnable_)
00224 {
00225     #if KERNEL_TIMERS_TICKLESS
00226         if (bEnable_) {
00227             TIMSK1 |= (1 << OCIE1A); // Enable interrupt
00228         } else {
00229             TIMSK1 &= ~(1 << OCIE1A);
00230         }
00231     #endif
00232 }

```

19.13 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kernelprofile.h

File Reference

Profiling timer hardware interface.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"

```

19.13.1 Detailed Description

Profiling timer hardware interface.

Definition in file [kernelprofile.h](#).

19.14 kernelprofile.h

```

00001  /*=====
00002
00003  _____
00004  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00005  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00006  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00007  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00020  #include "kerneltypes.h"
00021  #include "mark3cfg.h"
00022  #include "ll.h"
00023
00024  #ifndef __KPROFILE_H__
00025  #define __KPROFILE_H__
00026
00027  #if KERNEL_USE_PROFILER
00028
00029  //-----
00030  #define TICKS_PER_OVERFLOW (256)
00031  #define CLOCK_DIVIDE (8)
00032
00033  //-----
00037  class Profiler
00038  {
00039  public:
00046      static void Init();
00047
00053      static void Start();
00054
00060      static void Stop();
00061
00067      static PORT_TIMER_COUNT_TYPE Read();
00068
00074      static void Process();
00075
00081      static uint32_t GetEpoch() { return m_u32Epoch; }
00082  private:
00083      static uint32_t m_u32Epoch;
00084  };
00085
00086  #endif // KERNEL_USE_PROFILER
00087
00088  #endif

```

19.15 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kernelswi.h File Reference

[Kernel](#) Software interrupt declarations.

```
#include "kerneltypes.h"
```

Classes

- class [KernelSWI](#)

Class providing the software-interrupt required for context-switching in the kernel.

19.15.1 Detailed Description

[Kernel](#) Software interrupt declarations.

Definition in file [kernelswi.h](#).

19.16 kernelswi.h

```

00001  /*=====
00002
00003  00004  00005  00006  00007
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00022  #include "kerneltypes.h"
00023  #ifndef __KERNELSWI_H_
00024  #define __KERNELSWI_H_
00025
00026  //-----
00031  class KernelSWI
00032  {
00033  public:
00040      static void Config(void);
00041
00047      static void Start(void);
00048
00054      static void Stop(void);
00055
00061      static void Clear(void);
00062
00069      static void Trigger(void);
00070
00078      static uint8_t DI();
00079
00087      static void RI(bool bEnable_);
00088  };
00089
00090  #endif // __KERNELSWI_H_

```

19.17 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/kerneltimer.h

File Reference

[Kernel Timer](#) Class declaration.

```

#include "kerneltypes.h"
#include "mark3cfg.h"

```

Classes

- class [KernelTimer](#)
Hardware timer interface, used by all scheduling/timer subsystems.

19.17.1 Detailed Description

[Kernel Timer](#) Class declaration.

Definition in file [kerneltimer.h](#).

19.18 kerneltimer.h

```

00001  /*=====

```

```

00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023
00024 #ifndef __KERNELTIMER_H_
00025 #define __KERNELTIMER_H_
00026
00027 //-----
00031 class KernelTimer
00032 {
00033 public:
00039     static void Config(void);
00040
00046     static void Start(void);
00047
00053     static void Stop(void);
00054
00060     static uint8_t DI(void);
00061
00069     static void RI(bool bEnable_);
00070
00076     static void EI(void);
00077
00088     static PORT_TIMER_COUNT_TYPE SubtractExpiry(
PORT_TIMER_COUNT_TYPE uInterval_);
00089
00098     static PORT_TIMER_COUNT_TYPE TimeToExpiry(void);
00099
00108     static PORT_TIMER_COUNT_TYPE SetExpiry(uint32_t u32Interval_);
00109
00118     static PORT_TIMER_COUNT_TYPE GetOvertime(void);
00119
00125     static void ClearExpiry(void);
00126
00134     static PORT_TIMER_COUNT_TYPE Read(void);
00135 };
00136
00137 #endif //__KERNELTIMER_H_

```

19.19 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/portcfg.h File Reference

Mark3 Port Configuration.

Macros

- **#define AVR (1)**
Define a macro indicating the CPU architecture for which this port belongs.
- **#define K_WORD uint8_t**
Define types that map to the CPU Architecture's default data-word and address size.
- **#define K_ADDR uint16_t**
Size of an address (pointer size)
- **#define PORT_PRIO_TYPE uint8_t**
Set a base datatype used to represent each element of the scheduler's priority bitmap.
- **#define PORT_PRIO_MAP_WORD_SIZE (1)**
size of PORT_PRIO_TYPE in bytes
- **#define PORT_SYSTEM_FREQ ((uint32_t)16000000)**
Define the running CPU frequency.

- `#define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))`
Set the timer frequency.
- `#define PORT_KERNEL_TIMERS_THREAD_STACK ((K_ADDR)256)`
Define the size of the kernel-timer thread stack (if one is configured)
- `#define PORT_TIMER_COUNT_TYPE uint16_t`
Define the native type corresponding to the kernel timer hardware's counter register.
- `#define PORT_MIN_TIMER_TICKS (0)`
Minimum number of timer ticks for any delay or sleep, required to ensure that a timer cannot be initialized to a negative value.

19.19.1 Detailed Description

Mark3 Port Configuration.

This file is used to configure the kernel for your specific target CPU in order to provide the optimal set of features for a given use case.

!! NOTE: This file must ONLY be included from [mark3cfg.h](#)

Definition in file [portcfg.h](#).

19.19.2 Macro Definition Documentation

19.19.2.1 `#define AVR (1)`

Define a macro indicating the CPU architecture for which this port belongs.

This may also be set by the toolchain, but that's not guaranteed.

Definition at line 34 of file [portcfg.h](#).

19.19.2.2 `#define K_WORD uint8_t`

Define types that map to the CPU Architecture's default data-word and address size.

Size of a data word

Examples:

[buffalogger/main.cpp](#), [lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 41 of file [portcfg.h](#).

19.19.2.3 `#define PORT_PRIO_TYPE uint8_t`

Set a base datatype used to represent each element of the scheduler's priority bitmap.

`PORT_PRIO_MAP_WORD_SIZE` should map to the *size* of an element of type `PORT_PRIO_TYPE`. Type used for bitmap in the [PriorityMap](#) class

Definition at line 51 of file [portcfg.h](#).

19.19.2.4 #define PORT_SYSTEM_FREQ ((uint32_t)16000000)

Define the running CPU frequency.

This may be an integer constant, or an alias for another variable which holds the CPU's current running frequency.↔
CPU Frequency in Hz

Definition at line 59 of file [portcfg.h](#).

19.19.2.5 #define PORT_TIMER_COUNT_TYPE uint16_t

Define the native type corresponding to the kernel timer hardware's counter register.

[Timer](#) counter type

Definition at line 83 of file [portcfg.h](#).

19.19.2.6 #define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))

Set the timer frequency.

If running in tickless mode, this is simply the frequency at which the free-running kernel timer increments.

In tick-based mode, this is the frequency at which the fixed-frequency kernel tick interrupt occurs. Fixed timer interrupt frequency

Definition at line 72 of file [portcfg.h](#).

19.20 portcfg.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00025 #ifndef __PORTCFG_H__
00026 #define __PORTCFG_H__
00027
00033 #ifndef AVR
00034 # define AVR (1)
00035 #endif
00036
00041 #define K_WORD uint8_t
00042 #define K_ADDR uint16_t
00043
00044
00051 #define PORT_PRIO_TYPE uint8_t
00052 #define PORT_PRIO_MAP_WORD_SIZE (1)
00053
00054
00058 #if !defined(PORT_SYSTEM_FREQ)
00059 #define PORT_SYSTEM_FREQ ((uint32_t)16000000)
00060 #endif
00061
00069 #if KERNEL_TIMERS_TICKLESS
00070 #define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 256))
00071 #else
00072 #define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))
00073 #endif
00074
00078 #define PORT_KERNEL_TIMERS_THREAD_STACK ((K_ADDR)256)
00079
00083 #define PORT_TIMER_COUNT_TYPE uint16_t
00084
00085

```

```
00089 #define PORT_MIN_TIMER_TICKS          (0)
00090
00091 #endif // __PORTCFG_H__
```

19.21 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/public/threadport.h File Reference

ATMega328p Multithreading support.

```
#include "kerneltypes.h"
#include "thread.h"
#include <avr/builtins.h>
#include <avr/io.h>
#include <avr/interrupt.h>
```

Classes

- class [ThreadPort](#)
Class defining the architecture specific functions required by the kernel.

Macros

- #define [ASM](#)(x) asm volatile(x);
ASM Macro - simplify the use of ASM directive in C.
- #define [SR_](#) 0x3F
Status register define - map to 0x003F.
- #define [SPH_](#) 0x3E
Stack pointer define.
- #define [TOP_OF_STACK](#)(x, y) (uint8_t*)((uint16_t)x + (y - 1))
Macro to find the top of a stack given its size and top address.
- #define [PUSH_TO_STACK](#)(x, y)
Push a value y to the stack pointer x and decrement the stack pointer.
- #define [Thread_SaveContext](#)()
Save the context of the [Thread](#).
- #define [Thread_RestoreContext](#)()
Restore the context of the [Thread](#).
- #define [CS_ENTER](#)()
These macros must be used in pairs !
- #define [CS_EXIT](#)()
Exit critical section (restore status register)
- #define [ENABLE_INTS](#)() [ASM](#)("sei");
Initiate a contex switch without using the SWI.

Functions

- uint8_t [__mark3_clz8](#) (uint8_t in_)

19.21.1 Detailed Description

ATMega328p Multithreading support.

Definition in file [threadport.h](#).

19.21.2 Macro Definition Documentation

19.21.2.1 #define CS_ENTER()

Value:

```
{
    uint8_t __x = _SFR_IO8(SR_);
    ASM("cli");
}
```

These macros *must* be used in pairs !

Enter critical section (copy status register, disable interrupts)

Examples:

[buffalogger/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 162 of file [threadport.h](#).

19.21.3 Function Documentation

19.21.3.1 uint8_t __mark3_clz8 (uint8_t in_) [inline]

Lookup table based count-leading zeros implementation, used by scheduler by way of [PriorityMap](#)

Definition at line 52 of file [threadport.h](#).

19.22 threadport.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00021 #ifndef __THREADPORT_H_
00022 #define __THREADPORT_H_
00023
00024 #include "kerneltypes.h"
00025 #include "thread.h"
00026
00027 #include <avr/builtins.h>
00028 #include <avr/io.h>
00029 #include <avr/interrupt.h>
00030
00031 //-----
00033 #define ASM(x) asm volatile(x);
00034 #define SR_ 0x3F
00036 #define SPH_ 0x3E
00038 #define SPL_ 0x3D
00039
00040 //-----
00042 #define TOP_OF_STACK(x, y) (uint8_t*)((uint16_t)x) + (y - 1)
00043 #define PUSH_TO_STACK(x, y)
00045     *x = y;
00046     x--;
00047 #define STACK_GROWS_DOWN (1)
00048
00049 //-----
00052 inline uint8_t __mark3_clz8(uint8_t in_)
00053 {
```

```

00054     static const uint8_t u8Lookup[] = {4, 3, 2, 2, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0};
00055     uint8_t hi = __builtin_avr_swap(in_) & 0x0F;
00056     if (hi) {
00057         return u8Lookup[hi];
00058     }
00059     return 4 + u8Lookup[in_];
00060 }
00061
00062 //-----
00063 #define HW_CLZ    (1)
00064 #define CLZ(x)    __mark3_clz8(x)
00065
00066 //-----
00067 #define Thread_SaveContext()
00068 ASM("push r0");
00069 ASM("in r0, __SREG__");
00070 ASM("cli");
00071 ASM("push r0");
00072 ASM("push r1");
00073 ASM("push r2");
00074 ASM("push r3");
00075 ASM("push r4");
00076 ASM("push r5");
00077 ASM("push r6");
00078 ASM("push r7");
00079 ASM("push r8");
00080 ASM("push r9");
00081 ASM("push r10");
00082 ASM("push r11");
00083 ASM("push r12");
00084 ASM("push r13");
00085 ASM("push r14");
00086 ASM("push r15");
00087 ASM("push r16");
00088 ASM("push r17");
00089 ASM("push r18");
00090 ASM("push r19");
00091 ASM("push r20");
00092 ASM("push r21");
00093 ASM("push r22");
00094 ASM("push r23");
00095 ASM("push r24");
00096 ASM("push r25");
00097 ASM("push r26");
00098 ASM("push r27");
00099 ASM("push r28");
00100 ASM("push r29");
00101 ASM("push r30");
00102 ASM("push r31");
00103 ASM("lds r26, g_pclCurrent");
00104 ASM("lds r27, g_pclCurrent + 1");
00105 ASM("adiw r26, 4");
00106 ASM("in r0, 0x3D");
00107 ASM("st x+, r0");
00108 ASM("in r0, 0x3E");
00109 ASM("st x+, r0");
00110
00111 //-----
00112 #define Thread_RestoreContext()
00113 ASM("lds r26, g_pclCurrent");
00114 ASM("lds r27, g_pclCurrent + 1");
00115 ASM("adiw r26, 4");
00116 ASM("ld r28, x+");
00117 ASM("out 0x3D, r28");
00118 ASM("ld r29, x+");
00119 ASM("out 0x3E, r29");
00120 ASM("pop r31");
00121 ASM("pop r30");
00122 ASM("pop r29");
00123 ASM("pop r28");
00124 ASM("pop r27");
00125 ASM("pop r26");
00126 ASM("pop r25");
00127 ASM("pop r24");
00128 ASM("pop r23");
00129 ASM("pop r22");
00130 ASM("pop r21");
00131 ASM("pop r20");
00132 ASM("pop r19");
00133 ASM("pop r18");
00134 ASM("pop r17");
00135 ASM("pop r16");
00136 ASM("pop r15");
00137 ASM("pop r14");
00138 ASM("pop r13");
00139 ASM("pop r12");

```

```

00143 ASM("pop r11");
00144 ASM("pop r10");
00145 ASM("pop r9");
00146 ASM("pop r8");
00147 ASM("pop r7");
00148 ASM("pop r6");
00149 ASM("pop r5");
00150 ASM("pop r4");
00151 ASM("pop r3");
00152 ASM("pop r2");
00153 ASM("pop r1");
00154 ASM("pop r0");
00155 ASM("out __SREG__, r0");
00156 ASM("pop r0");
00157
00158 //-----
00160 //-----
00162 #define CS_ENTER()
00163 {
00164     uint8_t __x = _SFR_IO8(SR_);
00165     ASM("cli");
00166 //-----
00167 #define CS_EXIT()
00169     _SFR_IO8(SR_) = __x;
00170 }
00171
00172 //-----
00174 #define ENABLE_INTS() ASM("sei");
00175 #define DISABLE_INTS() ASM("cli");
00176
00177 //-----
00178 class Thread;
00186 class ThreadPort
00187 {
00188 public:
00194     static void StartThreads();
00195     friend class Thread;
00196
00197 private:
00205     static void InitStack(Thread* pstThread_);
00206 };
00207
00208 #endif //__ThreadPORT_H_

```

19.23 /media/usb/project/github/Mark3/kernel/cpu/avr/atmega328p/gcc/threadport.cpp File Reference

ATMega328p Multithreading.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "threadport.h"
#include "kernelswi.h"
#include "kerneltimer.h"
#include "timerlist.h"
#include "quantum.h"
#include "kernel.h"
#include "kernelaware.h"
#include <avr/io.h>
#include <avr/interrupt.h>

```

Functions

- [ISR](#) (INT0_vect) __attribute__((signal
[ISR\(INT0_vect\)](#) SWI using INT0 - used to trigger a context switch.


```

00081     PUSH_TO_STACK(pu8Stack, i);
00082 }
00083
00084 // Set the top o' the stack.
00085 pclThread->m_pwStackTop = (uint8_t*)pu8Stack;
00086
00087 // That's it! the thread is ready to run now.
00088 }
00089
00090 //-----
00091 static void Thread_Switch(void)
00092 {
00093     #if KERNEL_USE_IDLE_FUNC
00094         // If there's no next-thread-to-run...
00095         if (g_pclNext == Kernel::GetIdleThread()) {
00096             g_pclCurrent = Kernel::GetIdleThread();
00097
00098             // Disable the SWI, and re-enable interrupts -- enter nested interrupt
00099             // mode.
00100             KernelSWI::DI();
00101
00102             uint8_t u8SR = _SFR_IO8(SR_);
00103
00104             // So long as there's no "next-to-run" thread, keep executing the Idle
00105             // function to conclusion...
00106
00107             while (g_pclNext == Kernel::GetIdleThread()) {
00108                 // Ensure that we run this block in an interrupt enabled context (but
00109                 // with the rest of the checks being performed in an interrupt disabled
00110                 // context).
00111                 ASM("sei");
00112                 Kernel::IdleFunc();
00113                 ASM("cli");
00114             }
00115
00116             // Progress has been achieved -- an interrupt-triggered event has caused
00117             // the scheduler to run, and choose a new thread. Since we've already
00118             // saved the context of the thread we've hijacked to run idle, we can
00119             // proceed to disable the nested interrupt context and switch to the
00120             // new thread.
00121
00122             _SFR_IO8(SR_) = u8SR;
00123             KernelSWI::RI(true);
00124         }
00125     #endif
00126     g_pclCurrent = (Thread*)g_pclNext;
00127 }
00128
00129 //-----
00130 void ThreadPort::StartThreads()
00131 {
00132     KernelSWI::Config(); // configure the task switch SWI
00133     KernelTimer::Config(); // configure the kernel timer
00134
00135     Scheduler::SetScheduler(1); // enable the scheduler
00136     Scheduler::Schedule(); // run the scheduler - determine the first thread to run
00137
00138     Thread_Switch(); // Set the next scheduled thread to the current thread
00139
00140     KernelTimer::Start(); // enable the kernel timer
00141     KernelSWI::Start(); // enable the task switch SWI
00142
00143     #if KERNEL_USE_QUANTUM
00144         // Restart the thread quantum timer, as any value held prior to starting
00145         // the kernel will be invalid. This fixes a bug where multiple threads
00146         // started with the highest priority before starting the kernel causes problems
00147         // until the running thread voluntarily blocks.
00148         Quantum::RemoveThread();
00149         Quantum::AddThread(g_pclCurrent);
00150     #endif
00151
00152     // Restore the context...
00153     Thread_RestoreContext(); // restore the context of the first running thread
00154     ASM("reti"); // return from interrupt - will return to the first scheduled thread
00155 }
00156
00157 //-----
00162 //-----
00163 ISR(INT0_vect) __attribute__((signal, naked));
00164 ISR(INT0_vect)
00165 {
00166     Thread_SaveContext(); // Push the context (registers) of the current task
00167     Thread_Switch(); // Switch to the next task
00168     Thread_RestoreContext(); // Pop the context (registers) of the next task
00169     ASM("reti"); // Return to the next task
00170 }

```

19.25 /media/usb/project/github/Mark3/kernel/driver.cpp File Reference

Device driver/hardware abstraction layer.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "driver.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

19.25.1 Detailed Description

Device driver/hardware abstraction layer.

Definition in file [driver.cpp](#).

19.26 driver.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023 #include "driver.h"
00024
00025 #define _CAN_HAS_DEBUG
00026 //--[Autogenerated - Do Not Modify]-----
00027 #include "dbg_file_list.h"
00028 #include "buffalogger.h"
00029 #if defined(DBG_FILE)
00030 #error "Debug logging file token already defined! Bailing."
00031 #else
00032 #define DBG_FILE _DBG__KERNEL_DRIVER_CPP
00033 #endif
00034 //--[End Autogenerated content]-----
00035
00036 #include "kerneldebug.h"
00037
00038 //-----
00039 #if KERNEL_USE_DRIVER
00040
00041 DoubleLinkedList DriverList::m_clDriverList;
00042
00046 class DevNull : public Driver
00047 {
00048 public:
00049     virtual void Init() { SetName("/dev/null"); };
00050     virtual uint8_t Open() { return 0; }
00051     virtual uint8_t Close() { return 0; }
00052     virtual uint16_t Read(uint16_t /*u16Bytes*/, uint8_t* /*pu8Data*/) { return 0; }
00053     virtual uint16_t Write(uint16_t /*u16Bytes*/, uint8_t* /*pu8Data*/) { return 0; }
00054     virtual uint16_t
00055     Control(uint16_t /*u16Event*/, void* /*pvDataIn*/, uint16_t /*u16SizeIn*/, void* /*
00056     pvDataOut*/, uint16_t /*u16SizeOut*/)
00057     {
00058         return 0;
00059     };
00060
00061 //-----
00062 static DevNull clDevNull;
00063
00064 //-----
```



```

00075 static uint8_t DrvCmp(const char* szStr1_, const char* szStr2_)
00076 {
00077     char* szTmp1 = (char*)szStr1_;
00078     char* szTmp2 = (char*)szStr2_;
00079
00080     while ((*szTmp1 != 0) && (*szTmp2 != 0)) {
00081         if (*szTmp1++ != *szTmp2++) {
00082             return 0;
00083         }
00084     }
00085
00086     // Both terminate at the same length
00087     if (((*szTmp1 == 0) && (*szTmp2 == 0)) {
00088         return 1;
00089     }
00090
00091     return 0;
00092 }
00093
00094 //-----
00095 void DriverList::Init()
00096 {
00097     // Ensure we always have at least one entry - a default in case no match
00098     // is found (/dev/null)
00099     clDevNull.Init();
00100     Add(&clDevNull);
00101 }
00102
00103 //-----
00104 Driver* DriverList::FindByPath(const char* m_pcPath)
00105 {
00106     KERNEL_ASSERT(m_pcPath);
00107     Driver* pclTemp = static_cast<Driver*>(m_clDriverList.
GetHead());
00108
00109     // Iterate through the list of drivers until we find a match, or we
00110     // exhaust our list of installed drivers
00111     while (pclTemp != 0) {
00112         if (DrvCmp(m_pcPath, pclTemp->GetPath()) != 0u) {
00113             return pclTemp;
00114         }
00115         pclTemp = static_cast<Driver*>(pclTemp->GetNext());
00116     }
00117     // No matching driver found - return a pointer to our /dev/null driver
00118     return &clDevNull;
00119 }
00120
00121 #endif

```

19.27 /media/usb/project/github/Mark3/kernel/eventflag.cpp File Reference

Event Flag Blocking Object/IPC-Object implementation.

```

#include "mark3cfg.h"
#include "blocking.h"
#include "kernel.h"
#include "thread.h"
#include "eventflag.h"
#include "kernelaware.h"
#include "kerneldebug.h"
#include "dbg_file_list.h"
#include "buffalogger.h"

```

19.27.1 Detailed Description

Event Flag Blocking Object/IPC-Object implementation.

Definition in file [eventflag.cpp](#).


```

00100 #if KERNEL_USE_TIMEOUTS
00101 uint16_t EventFlag::Wait_i(uint16_t ul6Mask_,
    EventFlagOperation_t eMode_, uint32_t u32TimeMS_)
00102 #else
00103 uint16_t EventFlag::Wait_i(uint16_t ul6Mask_,
    EventFlagOperation_t eMode_)
00104 #endif
00105 {
00106     #if KERNEL_EXTRA_CHECKS
00107         KERNEL_ASSERT(IsInitialized());
00108     #endif
00109
00110     bool bThreadYield = false;
00111     bool bMatch       = false;
00112
00113     #if KERNEL_USE_TIMEOUTS
00114         Timer clEventTimer;
00115         bool bUseTimer = false;
00116     #endif
00117
00118     // Ensure we're operating in a critical section while we determine
00119     // whether or not we need to block the current thread on this object.
00120     CS_ENTER();
00121
00122     // Check to see whether or not the current mask matches any of the
00123     // desired bits.
00124     g_pclCurrent->SetEventFlagMask(ul6Mask_);
00125
00126     if ((eMode_ == EVENT_FLAG_ALL) || (eMode_ ==
EVENT_FLAG_ALL_CLEAR)) {
00127         // Check to see if the flags in their current state match all of
00128         // the set flags in the event flag group, with this mask.
00129         if ((m_ul6SetMask & ul6Mask_) == ul6Mask_) {
00130             bMatch = true;
00131             g_pclCurrent->SetEventFlagMask(ul6Mask_);
00132         }
00133     } else if ((eMode_ == EVENT_FLAG_ANY) || (eMode_ ==
EVENT_FLAG_ANY_CLEAR)) {
00134         // Check to see if the existing flags match any of the set flags in
00135         // the event flag group with this mask
00136         if ((m_ul6SetMask & ul6Mask_) != 0) {
00137             bMatch = true;
00138             g_pclCurrent->SetEventFlagMask(m_ul6SetMask & ul6Mask_);
00139         }
00140     }
00141
00142     // We're unable to match this pattern as-is, so we must block.
00143     if (!bMatch) {
00144         // Reset the current thread's event flag mask & mode
00145         g_pclCurrent->SetEventFlagMask(ul6Mask_);
00146         g_pclCurrent->SetEventFlagMode(eMode_);
00147
00148     #if KERNEL_USE_TIMEOUTS
00149         if (u32TimeMS_ != 0u) {
00150             g_pclCurrent->SetExpired(false);
00151             clEventTimer.Init();
00152             clEventTimer.Start(false, u32TimeMS_, TimedEventFlag_Callback, (void*)this);
00153             bUseTimer = true;
00154         }
00155     #endif
00156
00157     // Add the thread to the object's block-list.
00158     BlockPriority(g_pclCurrent);
00159
00160     // Trigger that
00161     bThreadYield = true;
00162 }
00163
00164 // If bThreadYield is set, it means that we've blocked the current thread,
00165 // and must therefore rerun the scheduler to determine what thread to
00166 // switch to.
00167 if (bThreadYield) {
00168     // Switch threads immediately
00169     Thread::Yield();
00170 }
00171
00172 // Exit the critical section and return back to normal execution
00173 CS_EXIT();
00174
00175 #if KERNEL_USE_TIMEOUTS
00176 if (bUseTimer && bThreadYield) {
00177     clEventTimer.Stop();
00178 }
00179 #endif
00180
00181 return g_pclCurrent->GetEventFlagMask();
00182 }

```

```

00187
00188 //-----
00189 uint16_t EventFlag::Wait(uint16_t ul6Mask_, EventFlagOperation_t eMode_)
00190 {
00191     #if KERNEL_USE_TIMEOUTS
00192         return Wait_i(ul6Mask_, eMode_, 0);
00193     #else
00194         return Wait_i(ul6Mask_, eMode_);
00195     #endif
00196 }
00197
00198 #if KERNEL_USE_TIMEOUTS
00199 //-----
00200 uint16_t EventFlag::Wait(uint16_t ul6Mask_, EventFlagOperation_t eMode_,
00201                          uint32_t u32TimeMS_)
00202 {
00203     return Wait_i(ul6Mask_, eMode_, u32TimeMS_);
00204 }
00205 #endif
00206 //-----
00207 void EventFlag::Set(uint16_t ul6Mask_)
00208 {
00209     #if KERNEL_EXTRA_CHECKS
00210         KERNEL_ASSERT(IsInitialized());
00211     #endif
00212
00213     Thread* pclPrev;
00214     Thread* pclCurrent;
00215     bool bReschedule = false;
00216     uint16_t ul6NewMask;
00217
00218     CS_ENTER();
00219
00220     // Walk through the whole block list, checking to see whether or not
00221     // the current flag set now matches any/all of the masks and modes of
00222     // the threads involved.
00223
00224     m_ul6SetMask |= ul6Mask_;
00225     ul6NewMask = m_ul6SetMask;
00226
00227     // Start at the head of the list, and iterate through until we hit the
00228     // "head" element in the list again. Ensure that we handle the case where
00229     // we remove the first or last elements in the list, or if there's only
00230     // one element in the list.
00231     pclCurrent = static_cast<Thread*>(m_clBlockList.GetHead());
00232
00233     // Do nothing when there are no objects blocking.
00234     if (pclCurrent != 0) {
00235         // First loop - process every thread in the block-list and check to
00236         // see whether or not the current flags match the event-flag conditions
00237         // on the thread.
00238         do {
00239             pclPrev = pclCurrent;
00240             pclCurrent = static_cast<Thread*>(pclCurrent->GetNext());
00241
00242             // Read the thread's event mask/mode
00243             uint16_t ul6ThreadMask = pclPrev->GetEventFlagMask();
00244             EventFlagOperation_t eThreadMode = pclPrev->
GetEventFlagMode();
00245
00246             // For the "any" mode - unblock the blocked threads if one or more bits
00247             // in the thread's bitmask match the object's bitmask
00248             if ((EVENT_FLAG_ANY == eThreadMode) || (
EVENT_FLAG_ANY_CLEAR == eThreadMode)) {
00249                 if ((ul6ThreadMask & m_ul6SetMask) != 0) {
00250                     pclPrev->SetEventFlagMode(
EVENT_FLAG_PENDING_UNBLOCK);
00251                     pclPrev->SetEventFlagMask(m_ul6SetMask & ul6ThreadMask);
00252                     bReschedule = true;
00253
00254                     // If the "clear" variant is set, then clear the bits in the mask
00255                     // that caused the thread to unblock.
00256                     if (EVENT_FLAG_ANY_CLEAR == eThreadMode) {
00257                         ul6NewMask &= ~(ul6ThreadMask & ul6Mask_);
00258                     }
00259                 }
00260             }
00261             // For the "all" mode, every set bit in the thread's requested bitmask must
00262             // match the object's flag mask.
00263             else if ((EVENT_FLAG_ALL == eThreadMode) || (
EVENT_FLAG_ALL_CLEAR == eThreadMode)) {
00264                 if ((ul6ThreadMask & m_ul6SetMask) == ul6ThreadMask) {
00265                     pclPrev->SetEventFlagMode(
EVENT_FLAG_PENDING_UNBLOCK);
00266                     pclPrev->SetEventFlagMask(ul6ThreadMask);
00267                     bReschedule = true;

```

```

00268
00269         // If the "clear" variant is set, then clear the bits in the mask
00270         // that caused the thread to unblock.
00271         if (EVENT_FLAG_ALL_CLEAR == eThreadMode) {
00272             ul6NewMask &= ~(ul6ThreadMask & ul6Mask_);
00273         }
00274     }
00275 }
00276
00277 // To keep looping, ensure that there's something in the list, and
00278 // that the next item isn't the head of the list.
00279 while (pclPrev != m_clBlockList.GetTail());
00280
00281 // Second loop - go through and unblock all of the threads that
00282 // were tagged for unblocking.
00283 pclCurrent = static_cast<Thread*>(m_clBlockList.
GetHead());
00284 bool bIsTail = false;
00285 do {
00286     pclPrev = pclCurrent;
00287     pclCurrent = static_cast<Thread*>(pclCurrent->GetNext());
00288
00289     // Check to see if this is the condition to terminate the loop
00290     if (pclPrev == m_clBlockList.GetTail()) {
00291         bIsTail = true;
00292     }
00293
00294     // If the first pass indicated that this thread should be
00295     // unblocked, then unblock the thread
00296     if (pclPrev->GetEventFlagMode() ==
EVENT_FLAG_PENDING_UNBLOCK) {
00297         Unblock(pclPrev);
00298     }
00299     } while (!bIsTail);
00300 }
00301
00302 // If we awoke any threads, re-run the scheduler
00303 if (bReschedule) {
00304     Thread::Yield();
00305 }
00306
00307 // Update the bitmask based on any "clear" operations performed along
00308 // the way
00309 m_ul6SetMask = ul6NewMask;
00310
00311 // Restore interrupts - will potentially cause a context switch if a
00312 // thread is unblocked.
00313 CS_EXIT();
00314 }
00315
00316 //-----
00317 void EventFlag::Clear(uint16_t ul6Mask_)
00318 {
00319     #if KERNEL_EXTRA_CHECKS
00320         KERNEL_ASSERT(IsInitialized());
00321     #endif
00322
00323     // Just clear the bitfields in the local object.
00324     CS_ENTER();
00325     m_ul6SetMask &= ~ul6Mask_;
00326     CS_EXIT();
00327 }
00328
00329 //-----
00330 uint16_t EventFlag::GetMask()
00331 {
00332     #if KERNEL_EXTRA_CHECKS
00333         KERNEL_ASSERT(IsInitialized());
00334     #endif
00335
00336     // Return the presently held event flag values in this object. Ensure
00337     // we get this within a critical section to guarantee atomicity.
00338     uint16_t ul6Return;
00339     CS_ENTER();
00340     ul6Return = m_ul6SetMask;
00341     CS_EXIT();
00342     return ul6Return;
00343 }
00344
00345 #endif // KERNEL_USE_EVENTFLAG

```

19.29 /media/usb/project/github/Mark3/kernel/kernel.cpp File Reference

[Kernel](#) initialization and startup code.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "kernel.h"
#include "scheduler.h"
#include "thread.h"
#include "threadport.h"
#include "timerlist.h"
#include "message.h"
#include "driver.h"
#include "profile.h"
#include "kernelprofile.h"
#include "autoalloc.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
#include "tracebuffer.h"
```

19.29.1 Detailed Description

[Kernel](#) initialization and startup code.

Definition in file [kernel.cpp](#).

19.30 kernel.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023
00024 #include "kernel.h"
00025 #include "scheduler.h"
00026 #include "thread.h"
00027 #include "threadport.h"
00028 #include "timerlist.h"
00029 #include "message.h"
00030 #include "driver.h"
00031 #include "profile.h"
00032 #include "kernelprofile.h"
00033 #include "autoalloc.h"
00034
00035 #define _CAN_HAS_DEBUG
00036 //--[Autogenerated - Do Not Modify]-----
00037 #include "dbg_file_list.h"
00038 #include "buffalogger.h"
00039 #if defined(DBG_FILE)
00040 #error "Debug logging file token already defined! Bailing."
00041 #else
00042 #define DBG_FILE _DBG__KERNEL_KERNEL_CPP
00043 #endif
00044 //--[End Autogenerated content]-----
00045 #include "kerneldebug.h"
00046 #include "tracebuffer.h"
00047
```

```

00048 bool          Kernel::m_bIsStarted;
00049 bool          Kernel::m_bIsPanic;
00050 PanicFunc_t  Kernel::m_pfPanic;
00051
00052 #if KERNEL_USE_STACK_GUARD
00053 uint16_t      Kernel::m_u16GuardThreshold;
00054 #endif
00055
00056 #if KERNEL_USE_IDLE_FUNC
00057 IdleFunc_t    Kernel::m_pfIdle;
00058 FakeThread_t  Kernel::m_clIdle;
00059 #endif
00060
00061 #if KERNEL_USE_THREAD_CALLOUTS
00062 ThreadCreateCallout_t  Kernel::m_pfThreadCreateCallout;
00063 ThreadExitCallout_t    Kernel::m_pfThreadExitCallout;
00064 ThreadContextCallout_t Kernel::m_pfThreadContextCallout;
00065 #endif
00066 //-----
00067 void Kernel::Init(void)
00068 {
00069     #if KERNEL_USE_AUTO_ALLOC
00070         AutoAlloc::Init();
00071     #endif
00072     #if KERNEL_USE_IDLE_FUNC
00073         ((Thread*)&m_clIdle)->InitIdle();
00074     #endif
00075     #if KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION
00076         TraceBuffer::Init();
00077     #endif
00078     KERNEL_TRACE("Initializing Mark3 Kernel");
00079
00080     // Initialize the global kernel data - scheduler, timer-scheduler, and
00081     // the global message pool.
00082     Scheduler::Init();
00083     #if KERNEL_USE_DRIVER
00084         DriverList::Init();
00085     #endif
00086     #if KERNEL_USE_TIMERS
00087         TimerScheduler::Init();
00088     #endif
00089     #if KERNEL_USE_MESSAGE
00090         GlobalMessagePool::Init();
00091     #endif
00092     #if KERNEL_USE_PROFILER
00093         Profiler::Init();
00094     #endif
00095     #if KERNEL_USE_STACK_GUARD
00096         m_u16GuardThreshold = KERNEL_STACK_GUARD_DEFAULT;
00097     #endif
00098 }
00099
00100 //-----
00101 void Kernel::Start(void)
00102 {
00103     KERNEL_TRACE("Starting Mark3 Scheduler");
00104     m_bIsStarted = true;
00105     ThreadPort::StartThreads();
00106     KERNEL_TRACE("Error starting Mark3 Scheduler");
00107 }
00108
00109 //-----
00110 void Kernel::Panic(uint16_t u16Cause_)
00111 {
00112     m_bIsPanic = true;
00113     if (m_pfPanic != 0) {
00114         m_pfPanic(u16Cause_);
00115     } else {
00116         #if KERNEL_AWARE_SIMULATION
00117             KernelAware::Print("Panic\n");
00118             KernelAware::Trace(0, 0, u16Cause_, g_pclCurrent->
GetID());
00119             KernelAware::ExitSimulator();
00120         #endif
00121         while (true) {
00122             ;
00123         }
00124     }
00125 }

```

19.31 /media/usb/project/github/Mark3/kernel/kernelaware.cpp File Reference

[Kernel](#) aware simulation support.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "kernelaware.h"
#include "threadport.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
```

19.31.1 Detailed Description

[Kernel](#) aware simulation support.

Definition in file [kernelaware.cpp](#).

19.32 kernelaware.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023 #include "kernelaware.h"
00024 #include "threadport.h"
00025
00026 #define _CAN_HAS_DEBUG
00027 //--[Autogenerated - Do Not Modify]-----
00028 #include "dbg_file_list.h"
00029 #include "buffalogger.h"
00030 #if defined(DBG_FILE)
00031 #error "Debug logging file token already defined! Bailing."
00032 #else
00033 #define DBG_FILE _DBG__KERNEL_KERNELAWARE_CPP
00034 #endif
00035 //--[End Autogenerated content]-----
00036
00037 #if KERNEL_AWARE_SIMULATION
00038
00039 //-----
00048 typedef union {
00049     volatile uint16_t au16Buffer[5];
00050
00051     struct {
00052         volatile const char* szName;
00053     } Profiler;
00054     struct {
00055         volatile uint16_t u16File;
00056         volatile uint16_t u16Line;
00057         volatile uint16_t u16Arg1;
00058         volatile uint16_t u16Arg2;
00059     } Trace;
00060     struct {
00061         volatile const char* szString;
00062     } Print;
00063 } KernelAwareData_t;
00064
00065 //-----
00077 volatile bool g_bIsKernelAware;
00078 volatile uint8_t g_u8KACommand;
00079 KernelAwareData_t g_stKADData;
00080
00081 //-----
00082 void KernelAware::ProfileInit(const char* szStr_)
00083 {
00084     CS_ENTER();
00085     g_stKADData.Profiler.szName = szStr_;
00086     g_u8KACommand = KA_COMMAND_PROFILE_INIT;
```



```

00087     CS_EXIT();
00088 }
00089
00090 //-----
00091 void KernelAware::ProfileStart(void)
00092 {
00093     g_u8KACommand = KA_COMMAND_PROFILE_START;
00094 }
00095
00096 //-----
00097 void KernelAware::ProfileStop(void)
00098 {
00099     g_u8KACommand = KA_COMMAND_PROFILE_STOP;
00100 }
00101
00102 //-----
00103 void KernelAware::ProfileReport(void)
00104 {
00105     g_u8KACommand = KA_COMMAND_PROFILE_REPORT;
00106 }
00107
00108 //-----
00109 void KernelAware::ExitSimulator(void)
00110 {
00111     g_u8KACommand = KA_COMMAND_EXIT_SIMULATOR;
00112 }
00113
00114 //-----
00115 void KernelAware::Trace(uint16_t ul6File_, uint16_t ul6Line_)
00116 {
00117     Trace_i(ul6File_, ul6Line_, 0, 0, KA_COMMAND_TRACE_0);
00118 }
00119
00120 //-----
00121 void KernelAware::Trace(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_)
00122 {
00123     Trace_i(ul6File_, ul6Line_, ul6Arg1_, 0, KA_COMMAND_TRACE_1);
00124 }
00125 //-----
00126 void KernelAware::Trace(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_, uint16_t
    ul6Arg2_)
00127 {
00128     Trace_i(ul6File_, ul6Line_, ul6Arg1_, ul6Arg2_, KA_COMMAND_TRACE_2);
00129 }
00130
00131 //-----
00132 void KernelAware::Trace_i(
00133     uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_, uint16_t ul6Arg2_,
    KernelAwareCommand_t eCmd_)
00134 {
00135     CS_ENTER();
00136     g_stKADData.Trace.ul6File = ul6File_;
00137     g_stKADData.Trace.ul6Line = ul6Line_;
00138     g_stKADData.Trace.ul6Arg1 = ul6Arg1_;
00139     g_stKADData.Trace.ul6Arg2 = ul6Arg2_;
00140     g_u8KACommand = eCmd_;
00141     CS_EXIT();
00142 }
00143
00144 //-----
00145 void KernelAware::Print(const char* szStr_)
00146 {
00147     CS_ENTER();
00148     g_stKADData.Print.szString = szStr_;
00149     g_u8KACommand = KA_COMMAND_PRINT;
00150     CS_EXIT();
00151 }
00152
00153 //-----
00154 bool KernelAware::IsSimulatorAware(void)
00155 {
00156     return g_bIsKernelAware;
00157 }
00158
00159 #endif

```

19.33 /media/usb/project/github/Mark3/kernel/ksemaphore.cpp File Reference

[Semaphore](#) Blocking-Object Implemenation.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ksemaphore.h"
#include "blocking.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

19.33.1 Detailed Description

[Semaphore](#) Blocking-Object Implemenation.

Definition in file [ksemaphore.cpp](#).

19.34 ksemaphore.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "mark3cfg.h"
00024
00025 #include "ksemaphore.h"
00026 #include "blocking.h"
00027
00028 #define _CAN_HAS_DEBUG
00029 //--[Autogenerated - Do Not Modify]-----
00030 #include "dbg_file_list.h"
00031 #include "buffalogger.h"
00032 #if defined(DBG_FILE)
00033 #error "Debug logging file token already defined! Bailing."
00034 #else
00035 #define DBG_FILE _DBG__KERNEL_KSEMAPHORE_CPP
00036 #endif
00037 //--[End Autogenerated content]-----
00038 #include "kerneldebug.h"
00039
00040 #if KERNEL_USE_SEMAPHORE
00041
00042 #if KERNEL_USE_TIMEOUTS
00043 #include "timerlist.h"
00044
00045 //-----
00056 void TimedSemaphore_Callback(Thread* pclOwner_, void* pvData_)
00057 {
00058     Semaphore* pclSemaphore = static_cast<Semaphore*>(pvData_);
00059
00060     // Indicate that the semaphore has expired on the thread
00061     pclOwner_->SetExpired(true);
00062
00063     // Wake up the thread that was blocked on this semaphore.
00064     pclSemaphore->WakeMe(pclOwner_);
00065
00066     if (pclOwner_->GetCurPriority() >= Scheduler::GetCurrentThread
00067         ()->GetCurPriority()) {
00068         Thread::Yield();
00069     }
00070 }
00071 //-----
00072 Semaphore::~Semaphore()
00073 {
00074     // If there are any threads waiting on this object when it goes out
00075     // of scope, set a kernel panic.
00076     if (m_clBlockList.GetHead() != 0) {
```

```

00077         Kernel::Panic(PANIC_ACTIVE_SEMAPHORE_DESCOPE);
00078     }
00079 }
00080
00081 //-----
00082 void Semaphore::WakeMe(Thread* pClChosenOne_)
00083 {
00084     #if KERNEL_EXTRA_CHECKS
00085         KERNEL_ASSERT(IsInitialized());
00086     #endif
00087
00088     // Remove from the semaphore waitlist and back to its ready list.
00089     Unblock(pClChosenOne_);
00090 }
00091
00092 #endif // KERNEL_USE_TIMEOUTS
00093
00094 //-----
00095 uint8_t Semaphore::WakeNext()
00096 {
00097     Thread* pClChosenOne;
00098
00099     pClChosenOne = m_clBlockList.HighestWaiter();
00100
00101     // Remove from the semaphore waitlist and back to its ready list.
00102     Unblock(pClChosenOne);
00103
00104     // Call a task switch if higher or equal priority thread
00105     if (pClChosenOne->GetCurPriority() >=
        Scheduler::GetCurrentThread()->GetCurPriority()) {
00106         return 1;
00107     }
00108     return 0;
00109 }
00110
00111 //-----
00112 void Semaphore::Init(uint16_t ul6InitVal_, uint16_t ul6MaxVal_)
00113 {
00114     #if KERNEL_EXTRA_CHECKS
00115         KERNEL_ASSERT(!m_clBlockList.GetHead());
00116     #endif
00117
00118     // Copy the paramters into the object - set the maximum value for this
00119     // semaphore to implement either binary or counting semaphores, and set
00120     // the initial count. Clear the wait list for this object.
00121     m_ul6Value = ul6InitVal_;
00122     m_ul6MaxValue = ul6MaxVal_;
00123
00124     #if KERNEL_EXTRA_CHECKS
00125         SetInitialized();
00126     #endif
00127 }
00128
00129 //-----
00130
00131 bool Semaphore::Post()
00132 {
00133     #if KERNEL_EXTRA_CHECKS
00134         KERNEL_ASSERT(IsInitialized());
00135     #endif
00136
00137     KERNEL_TRACE_1("Posting semaphore, Thread %d", (uint16_t)
        g_pClCurrent->GetID());
00138
00139     bool bThreadWake = false;
00140     bool bBail = false;
00141     // Increment the semaphore count - we can mess with threads so ensure this
00142     // is in a critical section. We don't just disable the scheduler since
00143     // we want to be able to do this from within an interrupt context as well.
00144     CS_ENTER();
00145
00146     // If nothing is waiting for the semaphore
00147     if (m_clBlockList.GetHead() == NULL) {
00148         // Check so see if we've reached the maximum value in the semaphore
00149         if (m_ul6Value < m_ul6MaxValue) {
00150             // Increment the count value
00151             m_ul6Value++;
00152         } else {
00153             // Maximum value has been reached, bail out.
00154             bBail = true;
00155         }
00156     } else {
00157         // Otherwise, there are threads waiting for the semaphore to be
00158         // posted, so wake the next one (highest priority goes first).
00159         bThreadWake = (WakeNext() != 0);
00160     }
00161 }

```

```

00162     CS_EXIT();
00163
00164     // If we weren't able to increment the semaphore count, fail out.
00165     if (bBail) {
00166         return false;
00167     }
00168
00169     // if bThreadWake was set, it means that a higher-priority thread was
00170     // woken. Trigger a context switch to ensure that this thread gets
00171     // to execute next.
00172     if (bThreadWake) {
00173         Thread::Yield();
00174     }
00175     return true;
00176 }
00177
00178 //-----
00179 #if KERNEL_USE_TIMEOUTS
00180 bool Semaphore::Pend_i(uint32_t u32WaitTimeMS_)
00181 #else
00182 void Semaphore::Pend_i(void)
00183 #endif
00184 {
00185     #if KERNEL_EXTRA_CHECKS
00186         KERNEL_ASSERT(IsInitialized());
00187     #endif
00188
00189     KERNEL_TRACE_1("Pending semaphore, Thread %d", (uint16_t)
00190         g_pclCurrent->GetID());
00191
00192     #if KERNEL_USE_TIMEOUTS
00193         Timer clSemTimer;
00194         bool bUseTimer = false;
00195     #endif
00196
00197     // Once again, messing with thread data - ensure
00198     // we're doing all of these operations from within a thread-safe context.
00199     CS_ENTER();
00200
00201     // Check to see if we need to take any action based on the semaphore count
00202     if (m_ul6Value != 0) {
00203         // The semaphore count is non-zero, we can just decrement the count
00204         // and go along our merry way.
00205         m_ul6Value--;
00206     } else {
00207         // The semaphore count is zero - we need to block the current thread
00208         // and wait until the semaphore is posted from elsewhere.
00209         #if KERNEL_USE_TIMEOUTS
00210             if (u32WaitTimeMS_ != 0u) {
00211                 g_pclCurrent->SetExpired(false);
00212                 clSemTimer.Init();
00213                 clSemTimer.Start(false, u32WaitTimeMS_, TimedSemaphore_Callback, (void*)this);
00214                 bUseTimer = true;
00215             }
00216         #endif
00217         BlockPriority(g_pclCurrent);
00218
00219         // Switch Threads immediately
00220         Thread::Yield();
00221     }
00222     CS_EXIT();
00223
00224     #if KERNEL_USE_TIMEOUTS
00225         if (bUseTimer) {
00226             clSemTimer.Stop();
00227             return (static_cast<int>(g_pclCurrent->GetExpired()) == 0);
00228         }
00229         return true;
00230     #endif
00231 }
00232
00233 //-----
00234 // Redirect the untimed pend API to the timed pend, with a null timeout.
00235 void Semaphore::Pend()
00236 {
00237     #if KERNEL_USE_TIMEOUTS
00238         Pend_i(0);
00239     #else
00240         Pend_i();
00241     #endif
00242 }
00243
00244 #if KERNEL_USE_TIMEOUTS
00245 //-----
00246 bool Semaphore::Pend(uint32_t u32WaitTimeMS_)
00247 {

```

```
00248         return Pend_i(u32WaitTimeMS_);
00249     }
00250     #endif
00251
00252     //-----
00253     uint16_t Semaphore::GetCount()
00254     {
00255         #if KERNEL_EXTRA_CHECKS
00256             KERNEL_ASSERT(IsInitialized());
00257         #endif
00258         uint16_t u16Ret;
00259         CS_ENTER();
00260         u16Ret = m_u16Value;
00261         CS_EXIT();
00262         return u16Ret;
00263     }
00264
00265     #endif
```

19.35 /media/usb/project/github/Mark3/kernel/ll.cpp File Reference

Core Linked-List implementation, from which all kernel objects are derived.

```
#include "kerneltypes.h"
#include "kernel.h"
#include "ll.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

19.35.1 Detailed Description

Core Linked-List implementation, from which all kernel objects are derived.

Definition in file [ll.cpp](#).

19.36 ll.cpp

```
00001 /*=====*/
00002
00003 |-----|-----|-----|-----|
00004 |_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/__|
00005 |_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/__|
00006 |_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/__|
00007 |_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/__|
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "kernel.h"
00024 #include "ll.h"
00025
00026 #define _CAN_HAS_DEBUG
00027 //--[Autogenerated - Do Not Modify]-----
00028 #include "dbg_file_list.h"
00029 #include "buffalogger.h"
00030 #if defined(DBG_FILE)
00031 #error "Debug logging file token already defined! Bailing."
00032 #else
00033 #define DBG_FILE _DBG___KERNEL_LL_CPP
00034 #endif
00035 //--[End Autogenerated content]-----
00036
00037 #include "kerneldebug.h"
00038
00039 //-----
00040 void LinkListNode::ClearNode()
00041 {
```

```

00042     next = NULL;
00043     prev = NULL;
00044 }
00045
00046 //-----
00047 void DoubleLinkedList::Add(LinkListNode* node_)
00048 {
00049     KERNEL_ASSERT(node_);
00050
00051     node_>prev = m_pstTail;
00052     node_>next = NULL;
00053
00054     // If the list is empty, initilize the head
00055     if (m_pstHead == 0) {
00056         m_pstHead = node_;
00057     }
00058     // Otherwise, adjust the tail's next pointer
00059     else {
00060         m_pstTail->next = node_;
00061     }
00062
00063     // Move the tail node, and assign it to the new node just passed in
00064     m_pstTail = node_;
00065 }
00066
00067 //-----
00068 void DoubleLinkedList::Remove(LinkListNode* node_)
00069 {
00070     KERNEL_ASSERT(node_);
00071
00072     if (node_>prev != 0) {
00073 #if SAFE_UNLINK
00074         if (node_>prev->next != node_) {
00075             Kernel::Panic(PANIC_LIST_UNLINK_FAILED);
00076         }
00077 #endif
00078         node_>prev->next = node_>next;
00079     }
00080     if (node_>next != 0) {
00081 #if SAFE_UNLINK
00082         if (node_>next->prev != node_) {
00083             Kernel::Panic(PANIC_LIST_UNLINK_FAILED);
00084         }
00085 #endif
00086         node_>next->prev = node_>prev;
00087     }
00088     if (node_ == m_pstHead) {
00089         m_pstHead = node_>next;
00090     }
00091     if (node_ == m_pstTail) {
00092         m_pstTail = node_>prev;
00093     }
00094     node_>ClearNode();
00095 }
00096
00097 //-----
00098 void CircularLinkedList::Add(LinkListNode* node_)
00099 {
00100     KERNEL_ASSERT(node_);
00101
00102     if (m_pstHead == 0) {
00103         // If the list is empty, initilize the nodes
00104         m_pstHead = node_;
00105         m_pstTail = node_;
00106     } else {
00107         // Move the tail node, and assign it to the new node just passed in
00108         m_pstTail->next = node_;
00109     }
00110
00111     // Add a node to the end of the linked list.
00112     node_>prev = m_pstTail;
00113     node_>next = m_pstHead;
00114
00115     m_pstTail = node_;
00116     m_pstHead->prev = node_;
00117 }
00118
00119 //-----
00120 void CircularLinkedList::Remove(LinkListNode* node_)
00121 {
00122     KERNEL_ASSERT(node_);
00123
00124     // Check to see if this is the head of the list...
00125     if ((node_ == m_pstHead) && (m_pstHead == m_pstTail)) {
00126         // Clear the head and tail pointers - nothing else left.
00127         m_pstHead = NULL;
00128         m_pstTail = NULL;

```

```

00129         return;
00130     }
00131
00132     #if SAFE_UNLINK
00133     // Verify that all nodes are properly connected
00134     if ((node_->prev->next != node_) || (node_->next->prev != node_)) {
00135         Kernel::Panic(PANIC_LIST_UNLINK_FAILED);
00136     }
00137     #endif
00138
00139     // This is a circularly linked list - no need to check for connection,
00140     // just remove the node.
00141     node_->next->prev = node_->prev;
00142     node_->prev->next = node_->next;
00143
00144     if (node_ == m_pstHead) {
00145         m_pstHead = m_pstHead->next;
00146     }
00147     if (node_ == m_pstTail) {
00148         m_pstTail = m_pstTail->prev;
00149     }
00150     node_->ClearNode();
00151 }
00152
00153 //-----
00154 void CircularLinkedList::PivotForward()
00155 {
00156     if (m_pstHead != 0) {
00157         m_pstHead = m_pstHead->next;
00158         m_pstTail = m_pstTail->next;
00159     }
00160 }
00161
00162 //-----
00163 void CircularLinkedList::PivotBackward()
00164 {
00165     if (m_pstHead != 0) {
00166         m_pstHead = m_pstHead->prev;
00167         m_pstTail = m_pstTail->prev;
00168     }
00169 }
00170
00171 //-----
00172 void CircularLinkedList::InsertNodeBefore(
    LinkListNode* node_, LinkListNode* insert_)
00173 {
00174     KERNEL_ASSERT(node_);
00175
00176     node_->next = insert_;
00177     node_->prev = insert_->prev;
00178
00179     if (insert_->prev != 0) {
00180         insert_->prev->next = node_;
00181     }
00182     insert_->prev = node_;
00183 }

```

19.37 /media/usb/project/github/Mark3/kernel/mailbox.cpp File Reference

[Mailbox](#) + Envelope IPC mechanism.

```

#include "mark3cfg.h"
#include "kerneltypes.h"
#include "ksemaphore.h"
#include "mailbox.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

19.37.1 Detailed Description

[Mailbox](#) + Envelope IPC mechanism.

Definition in file [mailbox.cpp](#).


```

00091     KERNEL_ASSERT(pvData_);
00092
00093 #if KERNEL_USE_TIMEOUTS
00094     Receive_i(pvData_, false, 0);
00095 #else
00096     Receive_i(pvData_, false);
00097 #endif
00098 }
00099
00100 #if KERNEL_USE_TIMEOUTS
00101 //-----
00102 bool Mailbox::Receive(void* pvData_, uint32_t u32TimeoutMS_)
00103 {
00104     KERNEL_ASSERT(pvData_);
00105     return Receive_i(pvData_, false, u32TimeoutMS_);
00106 }
00107 #endif
00108
00109 //-----
00110 void Mailbox::ReceiveTail(void* pvData_)
00111 {
00112     KERNEL_ASSERT(pvData_);
00113
00114 #if KERNEL_USE_TIMEOUTS
00115     Receive_i(pvData_, true, 0);
00116 #else
00117     Receive_i(pvData_, true);
00118 #endif
00119 }
00120
00121 #if KERNEL_USE_TIMEOUTS
00122 //-----
00123 bool Mailbox::ReceiveTail(void* pvData_, uint32_t u32TimeoutMS_)
00124 {
00125     KERNEL_ASSERT(pvData_);
00126     return Receive_i(pvData_, true, u32TimeoutMS_);
00127 }
00128 #endif
00129
00130 //-----
00131 bool Mailbox::Send(void* pvData_)
00132 {
00133     KERNEL_ASSERT(pvData_);
00134
00135 #if KERNEL_USE_TIMEOUTS
00136     return Send_i(pvData_, false, 0);
00137 #else
00138     return Send_i(pvData_, false);
00139 #endif
00140 }
00141
00142 //-----
00143 bool Mailbox::SendTail(void* pvData_)
00144 {
00145     KERNEL_ASSERT(pvData_);
00146
00147 #if KERNEL_USE_TIMEOUTS
00148     return Send_i(pvData_, true, 0);
00149 #else
00150     return Send_i(pvData_, true);
00151 #endif
00152 }
00153
00154 #if KERNEL_USE_TIMEOUTS
00155 //-----
00156 bool Mailbox::Send(void* pvData_, uint32_t u32TimeoutMS_)
00157 {
00158     KERNEL_ASSERT(pvData_);
00159     return Send_i(pvData_, false, u32TimeoutMS_);
00160 }
00161
00162 //-----
00163 bool Mailbox::SendTail(void* pvData_, uint32_t u32TimeoutMS_)
00164 {
00165     KERNEL_ASSERT(pvData_);
00166     return Send_i(pvData_, true, u32TimeoutMS_);
00167 }
00168 #endif
00169
00170 //-----
00171 #if KERNEL_USE_TIMEOUTS
00172 bool Mailbox::Send_i(const void* pvData_, bool bTail_, uint32_t u32TimeoutMS_)
00173 #else
00174 bool Mailbox::Send_i(const void* pvData_, bool bTail_)
00175 #endif
00176 #endif

```

```

00178 {
00179     const void* pvDst = NULL;
00180
00181     bool bRet          = false;
00182     bool bSchedState = Scheduler::SetScheduler(false);
00183
00184 #if KERNEL_USE_TIMEOUTS
00185     bool bBlock = false;
00186     bool bDone  = false;
00187     while (!bDone) {
00188         // Try to claim a slot first before resorting to blocking.
00189         if (bBlock) {
00190             bDone = true;
00191             Scheduler::SetScheduler(bSchedState);
00192             m_clSendSem.Pend(u32TimeoutMS_);
00193             Scheduler::SetScheduler(false);
00194         }
00195 #endif
00196
00197         CS_ENTER();
00198         // Ensure we have a free slot before we attempt to write data
00199         if (m_ul6Free != 0u) {
00200             m_ul6Free--;
00201
00202             if (bTail_) {
00203                 pvDst = GetTailPointer();
00204                 MoveTailBackward();
00205             } else {
00206                 MoveHeadForward();
00207                 pvDst = GetHeadPointer();
00208             }
00209             bRet = true;
00210 #if KERNEL_USE_TIMEOUTS
00211             bDone = true;
00212 #endif
00213         }
00214 #if KERNEL_USE_TIMEOUTS
00215         else if (u32TimeoutMS_ != 0u) {
00216             bBlock = true;
00217         } else {
00218             bDone = true;
00219         }
00220 #endif
00221
00222         CS_EXIT();
00223
00224 #if KERNEL_USE_TIMEOUTS
00225     }
00226 #endif
00227
00228     // Copy data to the claimed slot, and post the counting semaphore
00229     if (bRet) {
00230         CopyData(pvData_, pvDst, m_ul6ElementSize);
00231     }
00232
00233     Scheduler::SetScheduler(bSchedState);
00234
00235     if (bRet) {
00236         m_clRecvSem.Post();
00237     }
00238
00239     return bRet;
00240 }
00241
00242 //-----
00243 #if KERNEL_USE_TIMEOUTS
00244 bool Mailbox::Receive_i(const void* pvData_, bool bTail_, uint32_t u32WaitTimeMS_)
00245 {
00246     void Mailbox::Receive_i(const void* pvData_, bool bTail_)
00247 #endif
00248 {
00249     const void* pvSrc;
00250
00251 #if KERNEL_USE_TIMEOUTS
00252     if (!m_clRecvSem.Pend(u32WaitTimeMS_)) {
00253         // Failed to get the notification from the counting semaphore in the
00254         // time allotted. Bail.
00255         return false;
00256     }
00257 #else
00258     m_clRecvSem.Pend();
00259 #endif
00260
00261     // Disable the scheduler while we do this -- this ensures we don't have
00262     // multiple concurrent readers off the same queue, which could be problematic
00263     // if multiple writes occur during reads, etc.
00264     bool bSchedState = Scheduler::SetScheduler(false);

```

```

00265
00266     // Update the head/tail indexes, and get the associated data pointer for
00267     // the read operation.
00268     CS_ENTER();
00269
00270     m_ul6Free++;
00271     if (bTail_) {
00272         MoveTailForward();
00273         pvSrc = GetTailPointer();
00274     } else {
00275         pvSrc = GetHeadPointer();
00276         MoveHeadBackward();
00277     }
00278
00279     CS_EXIT();
00280
00281     CopyData(pvSrc, pvData_, m_ul6ElementSize);
00282
00283     Scheduler::SetScheduler(bSchedState);
00284
00285 #if KERNEL_USE_TIMEOUTS
00286     // Unblock a thread waiting for a free slot to send to
00287     m_clSendSem.Post();
00288
00289     return true;
00290 #endif
00291 }
00292
00293 #endif

```

19.39 /media/usb/project/github/Mark3/kernel/message.cpp File Reference

Inter-thread communications via message passing.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "message.h"
#include "threadport.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

19.39.1 Detailed Description

Inter-thread communications via message passing.

Definition in file [message.cpp](#).

19.40 message.cpp

```

00001 /*=====
00002
00003 |-----|-----|-----|-----|-----|-----|-----|-----|
00004 | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / |
00005 |  /  |  /  |  /  |  /  |  /  |  /  |  /  |  /  |  /  |  /  |
00006 |-----|-----|-----|-----|-----|-----|-----|-----|
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "mark3cfg.h"
00024
00025 #include "message.h"
00026 #include "threadport.h"
00027
00028 #define _CAN_HAS_DEBUG

```

```

00029 //--[Autogenerated - Do Not Modify]-----
00030 #include "dbg_file_list.h"
00031 #include "buffalogger.h"
00032 #if defined(DBG_FILE)
00033 #error "Debug logging file token already defined! Bailing."
00034 #else
00035 #define DBG_FILE _DBG__KERNEL_MESSAGE_CPP
00036 #endif
00037 //--[End Autogenerated content]-----
00038 #include "kerneldebug.h"
00039
00040 #if KERNEL_USE_MESSAGE
00041
00042 #if KERNEL_USE_TIMEOUTS
00043 #include "timerlist.h"
00044 #endif
00045
00046 Message      GlobalMessagePool::m_aclMessagePool[
GLOBAL_MESSAGE_POOL_SIZE];
MessagePool GlobalMessagePool::m_clPool;
00048
00049 //-----
00050 void MessagePool::Init()
00051 {
00052     m_clList.Init();
00053 }
00054
00055 //-----
00056 void MessagePool::Push(Message* pclMessage_)
00057 {
00058     KERNEL_ASSERT(pclMessage_);
00059
00060     CS_ENTER();
00061
00062     m_clList.Add(pclMessage_);
00063
00064     CS_EXIT();
00065 }
00066
00067 //-----
00068 Message* MessagePool::Pop()
00069 {
00070     Message* pclRet;
00071     CS_ENTER();
00072
00073     pclRet = static_cast<Message*>(m_clList.GetHead());
00074     if (0 != pclRet) {
00075         m_clList.Remove(static_cast<LinkListNode*>(pclRet));
00076     }
00077
00078     CS_EXIT();
00079     return pclRet;
00080 }
00081
00082 //-----
00083 Message* MessagePool::GetHead()
00084 {
00085     return static_cast<Message*>(m_clList.GetHead());
00086 }
00087
00088 //-----
00089 void GlobalMessagePool::Init()
00090 {
00091     uint8_t i;
00092     GlobalMessagePool::m_clPool.Init();
00093     for (i = 0; i < GLOBAL_MESSAGE_POOL_SIZE; i++) {
00094         GlobalMessagePool::m_aclMessagePool[i].Init();
00095         GlobalMessagePool::m_clPool.Push(&(GlobalMessagePool::m_aclMessagePool[i]));
00096     }
00097 }
00098
00099 //-----
00100 void GlobalMessagePool::Push(Message* pclMessage_)
00101 {
00102     m_clPool.Push(pclMessage_);
00103 }
00104
00105 //-----
00106 Message* GlobalMessagePool::Pop()
00107 {
00108     return m_clPool.Pop();
00109 }
00110
00111 //-----
00112 Message* GlobalMessagePool::GetHead()
00113 {
00114     return m_clPool.GetHead();

```

```

00115 }
00116
00117 //-----
00118 MessagePool* GlobalMessagePool::GetPool()
00119 {
00120     return &m_clPool;
00121 }
00122
00123 //-----
00124 void MessageQueue::Init()
00125 {
00126     m_clSemaphore.Init(0, GLOBAL_MESSAGE_POOL_SIZE);
00127 }
00128
00129 //-----
00130 Message* MessageQueue::Receive()
00131 {
00132     #if KERNEL_USE_TIMEOUTS
00133         return Receive_i(0);
00134     #else
00135         return Receive_i();
00136     #endif
00137 }
00138
00139 //-----
00140 #if KERNEL_USE_TIMEOUTS
00141 Message* MessageQueue::Receive(uint32_t u32TimeWaitMS_)
00142 {
00143     return Receive_i(u32TimeWaitMS_);
00144 }
00145 #endif
00146
00147 //-----
00148 #if KERNEL_USE_TIMEOUTS
00149 Message* MessageQueue::Receive_i(uint32_t u32TimeWaitMS_)
00150 #else
00151 Message* MessageQueue::Receive_i(void)
00152 #endif
00153 {
00154     Message* pClRet;
00155
00156     // Block the current thread on the counting semaphore
00157     #if KERNEL_USE_TIMEOUTS
00158         if (!m_clSemaphore.Pend(u32TimeWaitMS_)) {
00159             return NULL;
00160         }
00161     #else
00162         m_clSemaphore.Pend();
00163     #endif
00164
00165     CS_ENTER();
00166
00167     // Pop the head of the message queue and return it
00168     pClRet = static_cast<Message*>(m_clLinkList.GetHead());
00169     m_clLinkList.Remove(static_cast<Message*>(pClRet));
00170
00171     CS_EXIT();
00172
00173     return pClRet;
00174 }
00175
00176 //-----
00177 void MessageQueue::Send(Message* pClSrc_)
00178 {
00179     KERNEL_ASSERT(pClSrc_);
00180
00181     CS_ENTER();
00182
00183     // Add the message to the head of the linked list
00184     m_clLinkList.Add(pClSrc_);
00185
00186     // Post the semaphore, waking the blocking thread for the queue.
00187     m_clSemaphore.Post();
00188
00189     CS_EXIT();
00190 }
00191
00192 //-----
00193 uint16_t MessageQueue::GetCount()
00194 {
00195     return m_clSemaphore.GetCount();
00196 }
00197 #endif // KERNEL_USE_MESSAGE

```

19.41 /media/usb/project/github/Mark3/kernel/mutex.cpp File Reference

Mutual-exclusion object.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "blocking.h"
#include "mutex.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

19.41.1 Detailed Description

Mutual-exclusion object.

Definition in file [mutex.cpp](#).

19.42 mutex.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #include "kerneltypes.h"
00021 #include "mark3cfg.h"
00022
00023 #include "blocking.h"
00024 #include "mutex.h"
00025
00026 #define _CAN_HAS_DEBUG
00027 //--[Autogenerated - Do Not Modify]-----
00028 #include "dbg_file_list.h"
00029 #include "buffalogger.h"
00030 #if defined(DBG_FILE)
00031 #error "Debug logging file token already defined! Bailing."
00032 #else
00033 #define DBG_FILE _DBG__KERNEL_MUTEX_CPP
00034 #endif
00035 //--[End Autogenerated content]-----
00036
00037 #include "kerneldebug.h"
00038
00039 #if KERNEL_USE_MUTEX
00040
00041 #if KERNEL_USE_TIMEOUTS
00042
00043 //-----
00054 void TimedMutex_Callback(Thread* pclOwner_, void* pvData_)
00055 {
00056     Mutex* pclMutex = static_cast<Mutex*>(pvData_);
00057
00058     // Indicate that the semaphore has expired on the thread
00059     pclOwner_>SetExpired(true);
00060
00061     // Wake up the thread that was blocked on this semaphore.
00062     pclMutex->WakeMe(pclOwner_);
00063
00064     if (pclOwner_>GetCurPriority() >= Scheduler::GetCurrentThread
00065         ()>GetCurPriority()) {
00066         Thread::Yield();
00067     }
00068 }
00069 }
00070 //-----
```

```

00069 Mutex::~Mutex()
00070 {
00071     // If there are any threads waiting on this object when it goes out
00072     // of scope, set a kernel panic.
00073     if (m_clBlockList.GetHead() != 0) {
00074         Kernel::Panic(PANIC_ACTIVE_MUTEX_DESCOPED);
00075     }
00076 }
00077
00078 //-----
00079 void Mutex::WakeMe(Thread* pclOwner_)
00080 {
00081     // Remove from the semaphore waitlist and back to its ready list.
00082     Unblock(pclOwner_);
00083 }
00084
00085 #endif
00086
00087 //-----
00088 uint8_t Mutex::WakeNext()
00089 {
00090     Thread* pclChosenOne = NULL;
00091
00092     // Get the highest priority waiter thread
00093     pclChosenOne = m_clBlockList.HighestWaiter();
00094
00095     // Unblock the thread
00096     Unblock(pclChosenOne);
00097
00098     // The chosen one now owns the mutex
00099     m_pclOwner = pclChosenOne;
00100
00101     // Signal a context switch if it's a greater than or equal to the current priority
00102     if (pclChosenOne->GetCurPriority() >=
Scheduler::GetCurrentThread()->GetCurPriority()) {
00103         return 1;
00104     }
00105     return 0;
00106 }
00107
00108 //-----
00109 void Mutex::Init()
00110 {
00111     // Cannot re-init a mutex which has threads blocked on it
00112     #if KERNEL_EXTRA_CHECKS
00113         KERNEL_ASSERT(!m_clBlockList.GetHead());
00114     #endif
00115
00116     // Reset the data in the mutex
00117     m_bReady = true; // The mutex is free.
00118     m_u8MaxPri = 0; // Set the maximum priority inheritance state
00119     m_pclOwner = NULL; // Clear the mutex owner
00120     m_u8Recurse = 0; // Reset recurse count
00121
00122     #if KERNEL_EXTRA_CHECKS
00123         SetInitialized();
00124     #endif
00125 }
00126
00127 //-----
00128 #if KERNEL_USE_TIMEOUTS
00129 bool Mutex::Claim_i(uint32_t u32WaitTimeMS_)
00130 #else
00131 void Mutex::Claim_i(void)
00132 #endif
00133 {
00134     #if KERNEL_EXTRA_CHECKS
00135         KERNEL_ASSERT(IsInitialized());
00136     #endif
00137
00138     KERNEL_TRACE_1("Claiming Mutex, Thread %d", (uint16_t)
g_pclCurrent->GetID());
00139
00140     #if KERNEL_USE_TIMEOUTS
00141         Timer clTimer;
00142         bool bUseTimer = false;
00143     #endif
00144
00145     // Disable the scheduler while claiming the mutex - we're dealing with all
00146     // sorts of private thread data, can't have a thread switch while messing
00147     // with internal data structures.
00148     Scheduler::SetScheduler(false);
00149
00150     // Check to see if the mutex is claimed or not
00151     if (static_cast<int>(m_bReady) != 0) {
00152         // Mutex isn't claimed, claim it.
00153         m_bReady = false;

```

```

00154         m_u8Recurse = 0;
00155         m_u8MaxPri = g_pclCurrent->GetPriority();
00156         m_pclOwner = g_pclCurrent;
00157
00158         Scheduler::SetScheduler(true);
00159
00160 #if KERNEL_USE_TIMEOUTS
00161         return true;
00162 #else
00163         return;
00164 #endif
00165     }
00166
00167     // If the mutex is already claimed, check to see if this is the owner thread,
00168     // since we allow the mutex to be claimed recursively.
00169     if (g_pclCurrent == m_pclOwner) {
00170         // Ensure that we haven't exceeded the maximum recursive-lock count
00171         KERNEL_ASSERT((m_u8Recurse < 255));
00172         m_u8Recurse++;
00173
00174         // Increment the lock count and bail
00175         Scheduler::SetScheduler(true);
00176 #if KERNEL_USE_TIMEOUTS
00177         return true;
00178 #else
00179         return;
00180 #endif
00181     }
00182
00183     // The mutex is claimed already - we have to block now. Move the
00184     // current thread to the list of threads waiting on the mutex.
00185     #if KERNEL_USE_TIMEOUTS
00186         if (u32WaitTimeMS_ != 0u) {
00187             g_pclCurrent->SetExpired(false);
00188             clTimer.Init();
00189             clTimer.Start(false, u32WaitTimeMS_, (TimerCallback_t)TimedMutex_Callback, (void*)
this);
00190             bUseTimer = true;
00191         }
00192     #endif
00193     BlockPriority(g_pclCurrent);
00194
00195     // Check if priority inheritance is necessary. We do this in order
00196     // to ensure that we don't end up with priority inversions in case
00197     // multiple threads are waiting on the same resource.
00198     if (m_u8MaxPri <= g_pclCurrent->GetPriority()) {
00199         m_u8MaxPri = g_pclCurrent->GetPriority();
00200
00201         Thread* pclTemp = static_cast<Thread*>(m_clBlockList.GetHead());
00202         while (pclTemp != 0) {
00203             pclTemp->InheritPriority(m_u8MaxPri);
00204             if (pclTemp == static_cast<Thread*>(m_clBlockList.GetTail())) {
00205                 break;
00206             }
00207             pclTemp = static_cast<Thread*>(pclTemp->GetNext());
00208         }
00209         m_pclOwner->InheritPriority(m_u8MaxPri);
00210     }
00211
00212     // Done with thread data -reenable the scheduler
00213     Scheduler::SetScheduler(true);
00214
00215     // Switch threads if this thread acquired the mutex
00216     Thread::Yield();
00217
00218 #if KERNEL_USE_TIMEOUTS
00219     if (bUseTimer) {
00220         clTimer.Stop();
00221         return (static_cast<int>(g_pclCurrent->GetExpired()) == 0);
00222     }
00223     return true;
00224 #endif
00225 }
00226
00227 //-----
00228 void Mutex::Claim(void)
00229 {
00230     #if KERNEL_USE_TIMEOUTS
00231         Claim_i(0);
00232     #else
00233         Claim_i();
00234     #endif
00235 }
00236
00237 //-----
00238 #if KERNEL_USE_TIMEOUTS
00239 bool Mutex::Claim(uint32_t u32WaitTimeMS_)

```



```

00240 {
00241     return Claim_i(u32WaitTimeMS_);
00242 }
00243 #endif
00244
00245 //-----
00246 void Mutex::Release()
00247 {
00248     #if KERNEL_EXTRA_CHECKS
00249         KERNEL_ASSERT(IsInitialized());
00250     #endif
00251
00252     KERNEL_TRACE_1("Releasing Mutex, Thread %d", (uint16_t)
g_pclCurrent->GetID());
00253
00254     bool bSchedule = false;
00255
00256     // Disable the scheduler while we deal with internal data structures.
00257     Scheduler::SetScheduler(false);
00258
00259     // This thread had better be the one that owns the mutex currently...
00260     KERNEL_ASSERT((g_pclCurrent == m_pclOwner));
00261
00262     // If the owner had claimed the lock multiple times, decrease the lock
00263     // count and return immediately.
00264     if (m_u8Recurse != 0u) {
00265         m_u8Recurse--;
00266         Scheduler::SetScheduler(true);
00267         return;
00268     }
00269
00270     // Restore the thread's original priority
00271     if (g_pclCurrent->GetCurPriority() != g_pclCurrent->
GetPriority()) {
00272         g_pclCurrent->SetPriority(g_pclCurrent->
GetPriority());
00273
00274         // In this case, we want to reschedule
00275         bSchedule = true;
00276     }
00277
00278     // No threads are waiting on this semaphore?
00279     if (m_clBlockList.GetHead() == NULL) {
00280         // Re-initialize the mutex to its default values
00281         m_bReady = true;
00282         m_u8MaxPri = 0;
00283         m_pclOwner = NULL;
00284     } else {
00285         // Wake the highest priority Thread pending on the mutex
00286         if (WakeNext() != 0u) {
00287             // Switch threads if it's higher or equal priority than the current thread
00288             bSchedule = true;
00289         }
00290     }
00291
00292     // Must enable the scheduler again in order to switch threads.
00293     Scheduler::SetScheduler(true);
00294     if (bSchedule) {
00295         // Switch threads if a higher-priority thread was woken
00296         Thread::Yield();
00297     }
00298 }
00299
00300 #endif // KERNEL_USE_MUTEX

```

19.43 /media/usb/project/github/Mark3/kernel/notify.cpp File Reference

Lightweight thread notification - blocking object.

```

#include "mark3cfg.h"
#include "notify.h"
#include "mark3.h"
#include "kerneldebug.h"
#include "dbg_file_list.h"
#include "buffalogger.h"

```



```

00082 #endif
00083
00084     bool bReschedule = false;
00085
00086     CS_ENTER();
00087     Thread* pclCurrent = (Thread*)m_clBlockList.GetHead();
00088     if (pclCurrent == 0) {
00089         m_bPending = true;
00090     } else {
00091         while (pclCurrent != NULL) {
00092             Unblock(pclCurrent);
00093             if (!bReschedule && (pclCurrent->GetCurPriority() >=
Scheduler::GetCurrentThread()->GetCurPriority())) {
00094                 bReschedule = true;
00095             }
00096             pclCurrent = (Thread*)m_clBlockList.GetHead();
00097         }
00098         m_bPending = false;
00099     }
00100     CS_EXIT();
00101
00102     if (bReschedule) {
00103         Thread::Yield();
00104     }
00105 }
00106
00107 //-----
00108 void Notify::Wait(bool* pbFlag_)
00109 {
00110     #if KERNEL_EXTRA_CHECKS
00111         KERNEL_ASSERT(IsInitialized());
00112     #endif
00113
00114     bool bEarlyExit = false;
00115     CS_ENTER();
00116     if (!m_bPending) {
00117         Block(g_pclCurrent);
00118         if (pbFlag_ != 0) {
00119             *pbFlag_ = false;
00120         }
00121     } else {
00122         m_bPending = false;
00123         bEarlyExit = true;
00124     }
00125     CS_EXIT();
00126
00127     if (bEarlyExit) {
00128         return;
00129     }
00130
00131     Thread::Yield();
00132     if (pbFlag_ != 0) {
00133         *pbFlag_ = true;
00134     }
00135 }
00136
00137 //-----
00138 #if KERNEL_USE_TIMEOUTS
00139 bool Notify::Wait(uint32_t u32WaitTimeMS_, bool* pbFlag_)
00140 {
00141     #if KERNEL_EXTRA_CHECKS
00142         KERNEL_ASSERT(IsInitialized());
00143     #endif
00144     bool bUseTimer = false;
00145     bool bEarlyExit = false;
00146     Timer clNotifyTimer;
00147
00148     CS_ENTER();
00149     if (!m_bPending) {
00150         if (u32WaitTimeMS_ != 0u) {
00151             bUseTimer = true;
00152             g_pclCurrent->SetExpired(false);
00153
00154             clNotifyTimer.Init();
00155             clNotifyTimer.Start(false, u32WaitTimeMS_, TimedNotify_Callback, (void*)this);
00156         }
00157
00158         Block(g_pclCurrent);
00159
00160         if (pbFlag_ != 0) {
00161             *pbFlag_ = false;
00162         }
00163     } else {
00164         m_bPending = false;
00165         bEarlyExit = true;
00166     }
00167     CS_EXIT();

```



```

00029 #if HW_CLZ
00030 // Support hardware-accelerated Count-leading-zeros instruction
00031 uint8_t rc = PRIO_MAP_BITS - CLZ(uXPrio_);
00032 return rc;
00033 #else
00034 // Default un-optimized count-leading zeros operation
00035 PORT_PRIO_TYPE uXMask = (1 << (PRIO_MAP_BITS - 1));
00036 uint8_t u8Zeros = 0;
00037
00038 while (uXMask) {
00039     if (uXMask & uXPrio_) {
00040         return (PRIO_MAP_BITS - u8Zeros);
00041     }
00042     uXMask >>= 1;
00043     u8Zeros++;
00044 }
00045 return 0;
00046 #endif
00047 }
00048
00049 //-----
00050
00051 PriorityMap::PriorityMap()
00052 {
00053     #if PRIO_MAP_MULTI_LEVEL
00054         m_uXPriorityMapL2 = 0;
00055         for (int i = 0; i < PRIO_MAP_NUM_WORDS; i++) {
00056             m_auXPriorityMap[i] = 0;
00057         }
00058     #else
00059         m_uXPriorityMap = 0;
00060     #endif
00061 }
00062
00063 //-----
00064 void PriorityMap::Set(PORT_PRIO_TYPE uXPrio_)
00065 {
00066     PORT_PRIO_TYPE uXPrioBit = PRIO_BIT(uXPrio_);
00067     #if PRIO_MAP_MULTI_LEVEL
00068     PORT_PRIO_TYPE uXWordIdx = PRIO_MAP_WORD_INDEX(uXPrio_);
00069     m_auXPriorityMap[uXWordIdx] |= (1 << uXPrioBit);
00070     m_uXPriorityMapL2 |= (1 << uXWordIdx);
00071     #else
00072     m_uXPriorityMap |= (1 << uXPrioBit);
00073     #endif
00074 }
00075
00076 //-----
00077 void PriorityMap::Clear(PORT_PRIO_TYPE uXPrio_)
00078 {
00079     PORT_PRIO_TYPE uXPrioBit = PRIO_BIT(uXPrio_);
00080     #if PRIO_MAP_MULTI_LEVEL
00081     PORT_PRIO_TYPE uXWordIdx = PRIO_MAP_WORD_INDEX(uXPrio_);
00082     m_auXPriorityMap[uXWordIdx] &= ~(1 << uXPrioBit);
00083     if (!m_auXPriorityMap[uXWordIdx]) {
00084         m_uXPriorityMapL2 &= ~(1 << uXWordIdx);
00085     }
00086     #else
00087     m_uXPriorityMap &= ~(1 << uXPrioBit);
00088     #endif
00089 }
00090
00091 //-----
00092
00093 PORT_PRIO_TYPE PriorityMap::HighestPriority(void)
00094 {
00095     #if PRIO_MAP_MULTI_LEVEL
00096     PORT_PRIO_TYPE uXMapIdx = priority_from_bitmap(m_uXPriorityMapL2);
00097     if (!uXMapIdx) {
00098         return 0;
00099     }
00100     uXMapIdx--;
00101     PORT_PRIO_TYPE uXPrio = priority_from_bitmap(m_auXPriorityMap[uXMapIdx]);
00102     uXPrio += (uXMapIdx * PRIO_MAP_BITS);
00103     #else
00104     PORT_PRIO_TYPE uXPrio = priority_from_bitmap(m_uXPriorityMap);
00105     #endif
00106     return uXPrio;
00107 }
00108

```

19.47 /media/usb/project/github/Mark3/kernel/profile.cpp File Reference

Code profiling utilities.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "profile.h"
#include "kernelprofile.h"
#include "threadport.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

19.47.1 Detailed Description

Code profiling utilities.

Definition in file [profile.cpp](#).

19.48 profile.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023 #include "profile.h"
00024 #include "kernelprofile.h"
00025 #include "threadport.h"
00026
00027 #define _CAN_HAS_DEBUG
00028 //--[Autogenerated - Do Not Modify]-----
00029 #include "dbg_file_list.h"
00030 #include "buffalogger.h"
00031 #if defined(DBG_FILE)
00032 #error "Debug logging file token already defined! Bailing."
00033 #else
00034 #define DBG_FILE _DBG__KERNEL_PROFILE_CPP
00035 #endif
00036 //--[End Autogenerated content]-----
00037
00038 #include "kerneldebug.h"
00039
00040 #if KERNEL_USE_PROFILER
00041
00042 //-----
00043 void ProfileTimer::Init()
00044 {
00045     m_u32Cumulative      = 0;
00046     m_u32CurrentIteration = 0;
00047     m_ul6Iterations      = 0;
00048     m_bActive            = false;
00049 }
00050
00051 //-----
00052 void ProfileTimer::Start()
00053 {
00054     if (!m_bActive) {
00055         CS_ENTER();
00056         m_u32CurrentIteration = 0;
00057         m_u32InitialEpoch    = Profiler::GetEpoch();
00058         m_ul6Initial          = Profiler::Read();
00059         CS_EXIT();
00060     }
00061 }
```

```

00060         m_bActive = true;
00061     }
00062 }
00063
00064 //-----
00065 void ProfileTimer::Stop()
00066 {
00067     if (m_bActive) {
00068         uint16_t u16Final;
00069         uint32_t u32Epoch;
00070         CS_ENTER();
00071         u16Final = Profiler::Read();
00072         u32Epoch = Profiler::GetEpoch();
00073         // Compute total for current iteration...
00074         m_u32CurrentIteration = ComputeCurrentTicks(u16Final,
00075             u32Epoch);
00076         m_u32Cumulative += m_u32CurrentIteration;
00077         m_u16Iterations++;
00078         CS_EXIT();
00079         m_bActive = false;
00080     }
00081 }
00082 //-----
00083 uint32_t ProfileTimer::GetAverage()
00084 {
00085     if (m_u16Iterations != 0u) {
00086         return m_u32Cumulative / (uint32_t)m_u16Iterations;
00087     }
00088     return 0;
00089 }
00090
00091 //-----
00092 uint32_t ProfileTimer::GetCurrent()
00093 {
00094     if (m_bActive) {
00095         uint16_t u16Current;
00096         uint32_t u32Epoch;
00097         CS_ENTER();
00098         u16Current = Profiler::Read();
00099         u32Epoch = Profiler::GetEpoch();
00100         CS_EXIT();
00101         return ComputeCurrentTicks(u16Current, u32Epoch);
00102     }
00103     return m_u32CurrentIteration;
00104 }
00105
00106 //-----
00107 uint32_t ProfileTimer::ComputeCurrentTicks(uint16_t u16Current_, uint32_t
00108     u32Epoch_)
00109 {
00110     uint32_t u32Total;
00111     uint32_t u32Overflows;
00112     u32Overflows = u32Epoch_ - m_u32InitialEpoch;
00113     // More than one overflow...
00114     if (u32Overflows > 1) {
00115         u32Total = ((uint32_t)(u32Overflows - 1) * TICKS_PER_OVERFLOW) + (uint32_t)(TICKS_PER_OVERFLOW -
00116             m_u16Initial)
00117             + (uint32_t)u16Current_;
00118     }
00119     // Only one overflow, or one overflow that has yet to be processed
00120     else if ((u32Overflows != 0u) || (u16Current_ < m_u16Initial)) {
00121         u32Total = (uint32_t)(TICKS_PER_OVERFLOW - m_u16Initial) + (uint32_t)u16Current_;
00122     }
00123     // No overflows, none pending.
00124     else {
00125         u32Total = (uint32_t)(u16Current_ - m_u16Initial);
00126     }
00127     return u32Total;
00128 }
00129 }
00130
00131 #endif

```

19.49 /media/usb/project/github/Mark3/kernel/public/atomic.h File Reference

Basic Atomic Operations.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "threadport.h"
```

19.49.1 Detailed Description

Basic Atomic Operations.

Definition in file [atomic.h](#).

19.50 atomic.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #ifndef __ATOMIC_H__
00022 #define __ATOMIC_H__
00023
00024 #include "kerneltypes.h"
00025 #include "mark3cfg.h"
00026 #include "threadport.h"
00027
00028 #if KERNEL_USE_ATOMIC
00029
00039 class Atomic
00040 {
00041 public:
00048     static uint8_t Set(uint8_t* pu8Source_, uint8_t u8Val_);
00049     static uint16_t Set(uint16_t* pu16Source_, uint16_t u16Val_);
00050     static uint32_t Set(uint32_t* pu32Source_, uint32_t u32Val_);
00051
00058     static uint8_t Add(uint8_t* pu8Source_, uint8_t u8Val_);
00059     static uint16_t Add(uint16_t* pu16Source_, uint16_t u16Val_);
00060     static uint32_t Add(uint32_t* pu32Source_, uint32_t u32Val_);
00061
00068     static uint8_t Sub(uint8_t* pu8Source_, uint8_t u8Val_);
00069     static uint16_t Sub(uint16_t* pu16Source_, uint16_t u16Val_);
00070     static uint32_t Sub(uint32_t* pu32Source_, uint32_t u32Val_);
00071
00086     static bool TestAndSet(bool* pbLock);
00087 };
00088
00089 #endif // KERNEL_USE_ATOMIC
00090
00091 #endif //__ATOMIC_H__
```

19.51 /media/usb/project/github/Mark3/kernel/public/autoalloc.h File Reference

Automatic memory allocation for kernel objects.

```
#include <stdint.h>
#include <stdbool.h>
#include "mark3cfg.h"
```

19.51.1 Detailed Description

Automatic memory allocation for kernel objects.

Definition in file [autoalloc.h](#).

19.52 autoalloc.h

```

00001  /*
00002  |
00003  |
00004  |
00005  |
00006  |
00007  |
00008  |
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00020  #ifndef __AUTO_ALLOC_H__
00021  #define __AUTO_ALLOC_H__
00022
00023  #include <stdint.h>
00024  #include <stdbool.h>
00025  #include "mark3cfg.h"
00026
00027  #if KERNEL_USE_AUTO_ALLOC
00028  // Forward declaration of kernel objects that can be automatically allocated.
00029
00030  #if KERNEL_USE_EVENTFLAG
00031  class EventFlag;
00032  #endif
00033
00034  #if KERNEL_USE_MAILBOX
00035  class Mailbox;
00036  #endif
00037
00038  #if KERNEL_USE_MESSAGE
00039  class Message;
00040  class MessageQueue;
00041  #endif
00042
00043  #if KERNEL_USE_MUTEX
00044  class Mutex;
00045  #endif
00046
00047  #if KERNEL_USE_NOTIFY
00048  class Notify;
00049  #endif
00050
00051  #if KERNEL_USE_SEMAPHORE
00052  class Semaphore;
00053  #endif
00054
00055  class Thread;
00056
00057  #if KERNEL_USE_TIMERS
00058  class Timer;
00059  #endif
00060
00061  class AutoAlloc
00062  {
00063  public:
00070      static void Init(void);
00071
00082      static void* Allocate(uint16_t ul6Size_);
00083
00084  #if KERNEL_USE_SEMAPHORE
00085      static Semaphore* NewSemaphore(void);
00086  #endif
00087
00088  #if KERNEL_USE_MUTEX
00089      static Mutex* NewMutex(void);
00090  #endif
00091
00092  #if KERNEL_USE_EVENTFLAG
00093      static EventFlag* NewEventFlag(void);
00094  #endif
00095
00096  #if KERNEL_USE_MESSAGE
00097      static Message* NewMessage(void);
00098      static MessageQueue* NewMessageQueue(void);
00099  #endif
00100
00101  #if KERNEL_USE_NOTIFY

```

```

00102     static Notify* NewNotify(void);
00103 #endif
00104
00105 #if KERNEL_USE_MAILBOX
00106     static Mailbox* NewMailbox(void);
00107 #endif
00108
00109     static Thread* NewThread(void);
00110
00111 #if KERNEL_USE_TIMERS
00112     static Timer* NewTimer(void);
00113 #endif
00114
00115 private:
00116     static uint8_t m_au8AutoHeap[AUTO_ALLOC_SIZE]; // Heap memory
00117     static K_ADDR m_aHeapTop;                      // Top of the heap
00118 };
00119 #endif
00120
00121 #endif

```

19.53 /media/usb/project/github/Mark3/kernel/public/blocking.h File Reference

Blocking object base class declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "threadlist.h"

```

Classes

- class [BlockingObject](#)

Class implementing thread-blocking primitives.

19.53.1 Detailed Description

Blocking object base class declarations.

A Blocking object in Mark3 is essentially a thread list. Any blocking object implementation (being a semaphore, mutex, event flag, etc.) can be built on top of this class, utilizing the provided functions to manipulate thread location within the [Kernel](#).

Blocking a thread results in that thread becoming de-scheduled, placed in the blocking object's own private list of threads which are waiting on the object.

Unblocking a thread results in the reverse: The thread is moved back to its original location from the blocking list.

The only difference between a blocking object based on this class is the logic used to determine what constitutes a Block or Unblock condition.

For instance, a semaphore Pend operation may result in a call to the Block() method with the currently-executing thread in order to make that thread wait for a semaphore Post. That operation would then invoke the Unblock() method, removing the blocking thread from the semaphore's list, and back into the the appropriate thread inside the scheduler.

Care must be taken when implementing blocking objects to ensure that critical sections are used judiciously, otherwise asynchronous events like timers and interrupts could result in non-deterministic and often catastrophic behavior.

Definition in file [blocking.h](#).

19.54 blocking.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00047 #ifndef __BLOCKING_H__
00048 #define __BLOCKING_H__
00049
00050 #include "kerneltypes.h"
00051 #include "mark3cfg.h"
00052
00053 #include "ll.h"
00054 #include "threadlist.h"
00055
00056 #if KERNEL_USE_MUTEX || KERNEL_USE_SEMAPHORE || KERNEL_USE_EVENTFLAG
00057
00058 //-----
00059 // Cookies used to determine whether or not an object has been initialized
00060 #define BLOCKING_INVALID_COOKIE (0x3C)
00061 #define BLOCKING_INIT_COOKIE (0xC3)
00062
00063 class Thread;
00064
00065 //-----
00071 class BlockingObject
00072 {
00073 public:
00074 #if KERNEL_EXTRA_CHECKS
00075     BlockingObject() { m_u8Initialized = BLOCKING_INVALID_COOKIE; }
00076     ~BlockingObject() { m_u8Initialized = BLOCKING_INVALID_COOKIE; }
00077 #endif
00078
00079 protected:
00100     void Block(Thread* pclThread_);
00101
00110     void BlockPriority(Thread* pclThread_);
00111
00123     void Unblock(Thread* pclThread_);
00124
00129     ThreadList m_clBlockList;
00130
00131 #if KERNEL_EXTRA_CHECKS
00132
00136     uint8_t m_u8Initialized;
00137
00141     void SetInitialized(void) { m_u8Initialized = BLOCKING_INIT_COOKIE; }
00142
00147     bool IsInitialized(void) { return (m_u8Initialized == BLOCKING_INIT_COOKIE); }
00148
00149 #endif
00150
00151 };
00152
00153 #endif
00154
00155 #endif

```

19.55 /media/usb/project/github/Mark3/kernel/public/buffalogger.h File Reference

Super-efficient, super-secure logging routines.

```
#include <stdint.h>
```

19.55.1 Detailed Description

Super-efficient, super-secure logging routines.

Uses offline processing to ensure performance.

Definition in file [buffalogger.h](#).

19.56 buffalogger.h

```

00001  /*=====
00002
00003  _____
00004  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00005  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00006  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00007  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00020  #pragma once
00021  #include <stdint.h>
00022
00023  //-----
00024  #define STR1(s) #s
00025  #define STR(s) STR1(s)
00026
00027  //-----
00028  #define EMIT_DBG_STRING(str)
00029      do {
00030          const static volatile char    log_str[] __attribute__((section(".logger")))
00031          __attribute__((unused)) = str; \
00032          const static volatile uint16_t line_id __attribute__((section(".logger"))) __attribute__((unused))
00033          = __LINE__; \
00034          const static volatile uint16_t file_id __attribute__((section(".logger"))) __attribute__((unused))
00035          = DBG_FILE; \
00036          const static volatile uint16_t sync __attribute__((section(".logger"))) __attribute__((unused))
00037          = 0xCAFE; \
00038      } while (0);

```

19.57 /media/usb/project/github/Mark3/kernel/public/driver.h File Reference

[Driver](#) abstraction framework.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"

```

Classes

- class [Driver](#)

Base device-driver class used in hardware abstraction.

- class [DriverList](#)

List of [Driver](#) objects used to keep track of all device drivers in the system.

19.57.1 Detailed Description

[Driver](#) abstraction framework.

[Driver](#) abstraction framework for Mark3C.

19.57.2 Intro

This is the basis of the driver framework. In the context of Mark3, drivers don't necessarily have to be based on physical hardware peripherals. They can be used to represent algorithms (such as random number generators), files, or protocol stacks. Unlike FunkOS, where driver IO is protected automatically by a mutex, we do not use this kind of protection - we leave it up to the driver implementor to do what's right in its own context. This also frees up the driver to implement all sorts of other neat stuff, like sending messages to threads associated with the driver. Drivers are implemented as character devices, with the standard array of posix-style accessor methods for reading, writing, and general driver control.

A global driver list is provided as a convenient and minimal "filesystem" structure, in which devices can be accessed by name.

19.57.3 Driver Design

A device driver needs to be able to perform the following operations: -Initialize a peripheral -Start/stop a peripheral -Handle I/O control operations -Perform various read/write operations

At the end of the day, that's pretty much all a device driver has to do, and all of the functionality that needs to be presented to the developer.

We abstract all device drivers using a base-class which implements the following methods: -Start/Open -Stop/Close -Control -Read -Write

A basic driver framework and API can thus be implemented in five function calls - that's it! You could even reduce that further by handling the initialize, start, and stop operations inside the "control" operation.

19.57.4 Driver API

In C++, we can implement this as a class to abstract these event handlers, with virtual void functions in the base class overridden by the inherited objects.

To add and remove device drivers from the global table, we use the following methods:

```
void DriverList::Add( Driver *pclDriver_ );
void DriverList::Remove( Driver *pclDriver_ );
```

`DriverList::Add()/Remove()` takes a single arguments the pointer to the object to operate on.

Once a driver has been added to the table, drivers are opened by NAME using `DriverList::FindBy←Name("/dev/name")`. This function returns a pointer to the specified driver if successful, or to a built in /dev/null device if the path name is invalid. After a driver is open, that pointer is used for all other driver access functions.

This abstraction is incredibly useful any peripheral or service can be accessed through a consistent set of APIs, that make it easy to substitute implementations from one platform to another. Portability is ensured, the overhead is negligible, and it emphasizes the reuse of both driver and application code as separate entities.

Consider a system with drivers for I2C, SPI, and UART peripherals - under our driver framework, an application can initialize these peripherals and write a greeting to each using the same simple API functions for all drivers:

```
pclI2C = DriverList::FindByName("/dev/i2c");
pclUART = DriverList::FindByName("/dev/tty0");
pclSPI = DriverList::FindByName("/dev/spi");

pclI2C->Write(12, "Hello World!");
pclUART->Write(12, "Hello World!");
pclSPI->Write(12, "Hello World!");
```

Definition in file [driver.h](#).

19.58 driver.h

```
00001 /*=====
```

```

00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00105 #include "kerneltypes.h"
00106 #include "mark3cfg.h"
00107
00108 #include "ll.h"
00109
00110 #ifndef __DRIVER_H__
00111 #define __DRIVER_H__
00112
00113 #if KERNEL_USE_DRIVER
00114
00115 class DriverList;
00116 //-----
00121 class Driver : public LinkListNode
00122 {
00123 public:
00124     void* operator new(size_t sz, void* pv) { return (Driver*)pv; };
00130     virtual void Init() = 0;
00131
00139     virtual uint8_t Open() = 0;
00140
00148     virtual uint8_t Close() = 0;
00149
00164     virtual uint16_t Read(uint16_t ul6Bytes_, uint8_t* pu8Data_) = 0;
00165
00181     virtual uint16_t Write(uint16_t ul6Bytes_, uint8_t* pu8Data_) = 0;
00182
00201     virtual uint16_t
00202     Control(uint16_t ul6Event_, void* pvDataIn_, uint16_t ul6SizeIn_, void* pvDataOut_, uint16_t
00203             ul6SizeOut_)
00204             = 0;
00213     void SetName(const char* pcName_) { m_pcPath = pcName_; }
00221     const char* GetPath() { return m_pcPath; }
00222 private:
00224     const char* m_pcPath;
00225 };
00226
00227 //-----
00232 class DriverList
00233 {
00234 public:
00242     static void Init();
00243
00252     static void Add(Driver* pclDriver_) { m_clDriverList.
00261     Add(pclDriver_); }
00261     static void Remove(Driver* pclDriver_) { m_clDriverList.
00270     Remove(pclDriver_); }
00270     static Driver* FindByPath(const char* m_pcPath);
00271
00272 private:
00274     static DoubleLinkedList m_clDriverList;
00275 };
00276
00277 #endif // KERNEL_USE_DRIVER
00278
00279 #endif

```

19.59 /media/usb/project/github/Mark3/kernel/public/eventflag.h File Reference

Event Flag Blocking Object/IPC-Object definition.

```

#include "mark3cfg.h"
#include "kernel.h"
#include "kerneltypes.h"
#include "blocking.h"
#include "thread.h"

```

Classes

- class [EventFlag](#)

The [EventFlag](#) class is a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system.

19.59.1 Detailed Description

Event Flag Blocking Object/IPC-Object definition.

Definition in file [eventflag.h](#).

19.60 eventflag.h

```

00001  /*=====
00002
00003
00004
00005
00006
00007
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00019  #ifndef __EVENTFLAG_H__
00020  #define __EVENTFLAG_H__
00021
00022  #include "mark3cfg.h"
00023  #include "kernel.h"
00024  #include "kerneltypes.h"
00025  #include "blocking.h"
00026  #include "thread.h"
00027
00028  #if KERNEL_USE_EVENTFLAG
00029
00030  //-----
00046  class EventFlag : public BlockingObject
00047  {
00048  public:
00049      void* operator new(size_t sz, void* pv) { return (EventFlag*)pv; };
00050      ~EventFlag();
00051
00055      void Init();
00056
00064      uint16_t Wait(uint16_t ul6Mask_, EventFlagOperation_t eMode_);
00065
00066  #if KERNEL_USE_TIMEOUTS
00067
00075      uint16_t Wait(uint16_t ul6Mask_, EventFlagOperation_t eMode_, uint32_t
u32TimeMS_);
00076
00084      void WakeMe(Thread* p1ChosenOne_);
00085
00086  #endif
00087
00093      void Set(uint16_t ul6Mask_);
00094
00099      void Clear(uint16_t ul6Mask_);
00100
00105      uint16_t GetMask();
00106
00107  private:
00108  #if KERNEL_USE_TIMEOUTS
00109
00121      uint16_t Wait_i(uint16_t ul6Mask_, EventFlagOperation_t eMode_, uint32_t
u32TimeMS_);
00122  #else
00123
00133      uint16_t Wait_i(uint16_t ul6Mask_, EventFlagOperation_t eMode_);
00134  #endif
00135
00136      uint16_t m_ul6SetMask;
00137  };
00138

```



```

00096
00097 #if KERNEL_USE_IDLE_FUNC
00098
00103     static void SetIdleFunc(IdleFunc_t pfIdle_) { m_pfIdle = pfIdle_; }
00108     static void IdleFunc(void)
00109     {
00110         if (m_pfIdle != 0) {
00111             m_pfIdle();
00112         }
00113     }
00114
00122     static Thread* GetIdleThread(void) { return (Thread*)&
m_clIdle; }
00123 #endif
00124
00125 #if KERNEL_USE_THREAD_CALLOUTS
00126
00136     static void SetThreadCreateCallout(ThreadCreateCallout_t pfCreate_) {
m_pfThreadCreateCallout = pfCreate_; }
00148     static void SetThreadExitCallout(ThreadExitCallout_t pfExit_) {
m_pfThreadExitCallout = pfExit_; }
00159     static void SetThreadContextSwitchCallout(ThreadContextCallout_t
pfContext_)
00160     {
00161         m_pfThreadContextCallout = pfContext_;
00162     }
00163
00172     static ThreadCreateCallout_t GetThreadCreateCallout(void) { return
m_pfThreadCreateCallout; }
00181     static ThreadExitCallout_t GetThreadExitCallout(void) { return
m_pfThreadExitCallout; }
00190     static ThreadContextCallout_t GetThreadContextSwitchCallout(void) { return
m_pfThreadContextCallout; }
00191 #endif
00192
00193 #if KERNEL_USE_STACK_GUARD
00194     static void SetStackGuardThreshold(uint16_t ul6Threshold_) { m_ul6GuardThreshold = ul6Threshold_; }
00195     static uint16_t
        GetStackGuardThreshold(void) { return m_ul6GuardThreshold; }
00196 #endif
00197
00198 private:
00199     static bool        m_bIsStarted;
00200     static bool        m_bIsPanic;
00201     static PanicFunc_t m_pfPanic;
00202 #if KERNEL_USE_IDLE_FUNC
00203     static IdleFunc_t  m_pfIdle;
00204     static FakeThread_t m_clIdle;
00205 #endif
00206
00207 #if KERNEL_USE_THREAD_CALLOUTS
00208     static ThreadCreateCallout_t m_pfThreadCreateCallout;
00209     static ThreadExitCallout_t m_pfThreadExitCallout;
00210     static ThreadContextCallout_t m_pfThreadContextCallout;
00211 #endif
00212
00213 #if KERNEL_USE_STACK_GUARD
00214     static uint16_t m_ul6GuardThreshold;
00215 #endif
00216 };
00217
00218 #endif

```

19.63 /media/usb/project/github/Mark3/kernel/public/kernelaware.h File Reference

Kernel aware simulation support.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
```

Classes

- class [KernelAware](#)

The *KernelAware* class.

Enumerations

- enum `KernelAwareCommand_t` {
`KA_COMMAND_IDLE = 0, KA_COMMAND_PROFILE_INIT, KA_COMMAND_PROFILE_START, KA_COMMAND_PROFILE_STOP,`
`KA_COMMAND_PROFILE_REPORT, KA_COMMAND_EXIT_SIMULATOR, KA_COMMAND_TRACE_0,`
`KA_COMMAND_TRACE_1,`
`KA_COMMAND_TRACE_2, KA_COMMAND_PRINT }`

This enumeration contains a list of supported commands that can be executed to invoke a response from a kernel aware host.

19.63.1 Detailed Description

Kernel aware simulation support.

Definition in file `kernelaware.h`.

19.63.2 Enumeration Type Documentation

19.63.2.1 enum `KernelAwareCommand_t`

This enumeration contains a list of supported commands that can be executed to invoke a response from a kernel aware host.

Enumerator

`KA_COMMAND_IDLE` Null command, does nothing.
`KA_COMMAND_PROFILE_INIT` Initialize a new profiling session.
`KA_COMMAND_PROFILE_START` Begin a profiling sample.
`KA_COMMAND_PROFILE_STOP` End a profiling sample.
`KA_COMMAND_PROFILE_REPORT` Report current profiling session.
`KA_COMMAND_EXIT_SIMULATOR` Terminate the host simulator.
`KA_COMMAND_TRACE_0` 0-argument kernel trace
`KA_COMMAND_TRACE_1` 1-argument kernel trace
`KA_COMMAND_TRACE_2` 2-argument kernel trace
`KA_COMMAND_PRINT` Print an arbitrary string of data.

Definition at line 33 of file `kernelaware.h`.

19.64 kernelaware.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #ifndef __KERNEL_AWARE_H__
00022 #define __KERNEL_AWARE_H__
00023
00024 #include "kerneltypes.h"
00025 #include "mark3cfg.h"
```

```

00026
00027 #if KERNEL_AWARE_SIMULATION
00028 //-----
00033 typedef enum {
00034     KA_COMMAND_IDLE = 0,
00035     KA_COMMAND_PROFILE_INIT,
00036     KA_COMMAND_PROFILE_START,
00037     KA_COMMAND_PROFILE_STOP,
00038     KA_COMMAND_PROFILE_REPORT,
00039     KA_COMMAND_EXIT_SIMULATOR,
00040     KA_COMMAND_TRACE_0,
00041     KA_COMMAND_TRACE_1,
00042     KA_COMMAND_TRACE_2,
00043     KA_COMMAND_PRINT
00044 } KernelAwareCommand_t;
00045
00046 //-----
00064 class KernelAware
00065 {
00066 public:
00067     //-----
00078     static void ProfileInit(const char* szStr_);
00079
00080     //-----
00088     static void ProfileStart(void);
00089
00090     //-----
00097     static void ProfileStop(void);
00098
00099     //-----
00107     static void ProfileReport(void);
00108
00109     //-----
00117     static void ExitSimulator(void);
00118
00119     //-----
00127     static void Print(const char* szStr_);
00128
00129     //-----
00139     static void Trace(uint16_t ul6File_, uint16_t ul6Line_);
00140
00141     //-----
00152     static void Trace(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_);
00153
00154     //-----
00166     static void Trace(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_, uint16_t ul6Arg2_);
00167
00168     //-----
00178     static bool IsSimulatorAware(void);
00179
00180 private:
00181     //-----
00194     static void
00195     Trace_i(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_, uint16_t ul6Arg2_,
00196             KernelAwareCommand_t eCmd_);
00197
00198 #endif
00199
00200 #endif

```

19.65 /media/usb/project/github/Mark3/kernel/public/kerneldebug.h File Reference

Macros and functions used for assertions, kernel traces, etc.

```

#include "mark3cfg.h"
#include "tracebuffer.h"
#include "kernelaware.h"
#include "paniccodes.h"
#include "kernel.h"
#include "buffalogger.h"
#include "dbg_file_list.h"

```



```

00042         \
        aul6Msg__[2] = __LINE__;
00043         \
        aul6Msg__[3] = TraceBuffer::Increment();
00044         \
        TraceBuffer::Write(aul6Msg__, 4);
00045         \
00046     };
00047
00048     //-----
00049     #define KERNEL_TRACE_1(x, arg1)
00050         \
00051     {
00052         \
        EMIT_DBG_STRING(x);
00053         \
        uint16_t aul6Msg__[5];
00054         \
        aul6Msg__[0] = 0xACDC;
00055         \
        aul6Msg__[1] = DBG_FILE;
00056         \
        aul6Msg__[2] = __LINE__;
00057         \
        aul6Msg__[3] = TraceBuffer::Increment();
00058         \
        aul6Msg__[4] = arg1;
00059         \
        TraceBuffer::Write(aul6Msg__, 5);
00060         \
00061     }
00062
00063     //-----
00064     #define KERNEL_TRACE_2(x, arg1, arg2)
00065         \
00066     {
00067         \
        EMIT_DBG_STRING(x);
00068         \
        uint16_t aul6Msg__[6];
00069         \
        aul6Msg__[0] = 0xACDC;
00070         \
        aul6Msg__[1] = DBG_FILE;
00071         \
        aul6Msg__[2] = __LINE__;
00072         \
        aul6Msg__[3] = TraceBuffer::Increment();
00073         \
        aul6Msg__[4] = arg1;
00074         \
        aul6Msg__[5] = arg2;
00075         \
        TraceBuffer::Write(aul6Msg__, 6);
00076         \
00077     }
00078
00079     //-----
00080     #define KERNEL_ASSERT(x)
00081         \
00082     {
00083         \
        if ((x) == false) {
00084             \
            EMIT_DBG_STRING("ASSERT FAILED");
00085             \
            uint16_t aul6Msg__[4];
00086             \
            aul6Msg__[0] = 0xACDC;
00087             \
            aul6Msg__[1] = DBG_FILE;
00088             \
            aul6Msg__[2] = __LINE__;
00089             \
            aul6Msg__[3] = TraceBuffer::Increment();
00090             \
            TraceBuffer::Write(aul6Msg__, 4);
00091             \
            Kernel::Panic(PANIC_ASSERT_FAILED);
00092             \
        }

```

```

00093     \
00094 }
00095 #elif (KERNEL_USE_DEBUG && KERNEL_AWARE_SIMULATION && KERNEL_ENABLE_LOGGING)
00096
00097 //-----
00098 #define KERNEL_TRACE(x)
00099     \
00100 {
00101     \
00102     EMIT_DBG_STRING(x);
00103     \
00104     KernelAware::Trace(DBG_FILE, __LINE__);
00105 }
00106 //-----
00107 #define KERNEL_TRACE_1(x, arg1)
00108     \
00109 {
00110     \
00111     EMIT_DBG_STRING(x);
00112     \
00113     KernelAware::Trace(DBG_FILE, __LINE__, arg1);
00114 }
00115 //-----
00116 #define KERNEL_TRACE_2(x, arg1, arg2)
00117     \
00118 {
00119     \
00120     EMIT_DBG_STRING(x);
00121     \
00122     KernelAware::Trace(DBG_FILE, __LINE__, arg1, arg2);
00123 }
00124 //-----
00125 #define KERNEL_ASSERT(x)
00126     \
00127 {
00128     \
00129     if ((x) == false) {
00130         \
00131         EMIT_DBG_STRING("ASSERT FAILED");
00132         \
00133         KernelAware::Trace(DBG_FILE, __LINE__);
00134         \
00135         Kernel::Panic(PANIC_ASSERT_FAILED);
00136     }
00137 }
00138 //-----
00139 #define KERNEL_TRACE(x)
00140 //-----
00141 #define KERNEL_TRACE_1(x, arg1)
00142 //-----
00143 #define KERNEL_TRACE_2(x, arg1, arg2)
00144 //-----
00145 #define KERNEL_ASSERT(x) \
00146 {
00147     \
00148     if ((x) == false) {
00149         \
00150         Kernel::Panic(PANIC_ASSERT_FAILED);
00151     }
00152 }
00153 //-----
00154 // Note -- when kernel-debugging is disabled, we still have to define the
00155 // macros to ensure that the expressions compile (albeit, by elimination
00156 // during pre-processing).

```

```

00156 //-----
00157 #define KERNEL_TRACE(x)
00158 //-----
00159 #define KERNEL_TRACE_1(x, arg1)
00160 //-----
00161 #define KERNEL_TRACE_2(x, arg1, arg2)
00162 //-----
00163 #define KERNEL_ASSERT(x)
00164
00165 #endif // KERNEL_USE_DEBUG
00166
00167 //-----
00168 #if (KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION && KERNEL_ENABLE_USER_LOGGING)
00169
00170 //-----
00171 #define USER_TRACE(x)
00172     \
00173 {
00174     \
00175     EMIT_DBG_STRING(x);
00176     \
00177     uint16_t aul6Msg__[4];
00178     \
00179     aul6Msg__[0] = 0xACDC;
00180     \
00181     aul6Msg__[1] = DBG_FILE;
00182     \
00183     aul6Msg__[2] = __LINE__;
00184     \
00185     aul6Msg__[3] = TraceBuffer::Increment();
00186     \
00187     TraceBuffer::Write(aul6Msg__, 4);
00188 }
00189
00190 //-----
00191 #define USER_TRACE_1(x, arg1)
00192     \
00193 {
00194     \
00195     EMIT_DBG_STRING(x);
00196     \
00197     uint16_t aul6Msg__[5];
00198     \
00199     aul6Msg__[0] = 0xACDC;
00200     \
00201     aul6Msg__[1] = DBG_FILE;
00202     \
00203     aul6Msg__[2] = __LINE__;
00204     \
00205     aul6Msg__[3] = TraceBuffer::Increment();
00206     \
00207     aul6Msg__[4] = arg1;
00208     \
00209     TraceBuffer::Write(aul6Msg__, 5);
00210 }
00211
00212 //-----
00213 #define USER_TRACE_2(x, arg1, arg2)
00214     \
00215 {
00216     \
00217     EMIT_DBG_STRING(x);
00218     \
00219     uint16_t aul6Msg__[6];
00220     \
00221     aul6Msg__[0] = 0xACDC;
00222     \
00223     aul6Msg__[1] = DBG_FILE;
00224     \
00225     aul6Msg__[2] = __LINE__;
00226     \
00227     aul6Msg__[3] = TraceBuffer::Increment();
00228     \
00229     aul6Msg__[4] = arg1;
00230     \
00231     aul6Msg__[5] = arg2;
00232     \
00233     TraceBuffer::Write(aul6Msg__, 6);
00234 }
00235
00236 //-----

```

```

00213 }
00214
00215 //-----
00216 #define USER_ASSERT(x)
00217     \
00218 {
00219     \
00219     if ((x) == false) {
00220         \
00220         EMIT_DBG_STRING("ASSERT FAILED");
00221         \
00221         uint16_t au16Msg__[4];
00222         \
00222         au16Msg__[0] = 0xACDC;
00223         \
00223         au16Msg__[1] = DBG_FILE;
00224         \
00224         au16Msg__[2] = __LINE__;
00225         \
00225         au16Msg__[3] = TraceBuffer::Increment();
00226         \
00226         TraceBuffer::Write(au16Msg__, 4);
00227         \
00227         Kernel::Panic(PANIC_ASSERT_FAILED);
00228     }
00229     \
00230 }
00231 #elif (KERNEL_USE_DEBUG && KERNEL_AWARE_SIMULATION && KERNEL_ENABLE_USER_LOGGING)
00232
00233 //-----
00234 #define USER_TRACE(x)
00235     \
00236 {
00237     \
00237     EMIT_DBG_STRING(x);
00238     \
00238     KernelAware::Trace(DBG_FILE, __LINE__);
00239     \
00240 };
00241
00242 //-----
00243 #define USER_TRACE_1(x, arg1)
00244     \
00245 {
00246     \
00246     EMIT_DBG_STRING(x);
00247     \
00247     KernelAware::Trace(DBG_FILE, __LINE__, arg1);
00248     \
00249 }
00250
00251 //-----
00252 #define USER_TRACE_2(x, arg1, arg2)
00253     \
00254 {
00255     \
00255     EMIT_DBG_STRING(x);
00256     \
00256     KernelAware::Trace(DBG_FILE, __LINE__, arg1, arg2);
00257     \
00258 }
00259
00260 //-----
00261 #define USER_ASSERT(x)
00262     \
00263 {
00264     \
00264     if ((x) == false) {
00265         \
00265         EMIT_DBG_STRING("ASSERT FAILED");
00266         \
00266         KernelAware::Trace(DBG_FILE, __LINE__);
00267         \
00267         Kernel::Panic(PANIC_ASSERT_FAILED);
00268     }
00269     \

```



```

00269     \
00270 }
00271
00272 #else
00273 //-----
00274 // Note -- when kernel-debugging is disabled, we still have to define the
00275 // macros to ensure that the expressions compile (albeit, by elimination
00276 // during pre-processing).
00277 //-----
00278 #define USER_TRACE(x)
00279 //-----
00280 #define USER_TRACE_1(x, arg1)
00281 //-----
00282 #define USER_TRACE_2(x, arg1, arg2)
00283 //-----
00284 #define USER_ASSERT(x)
00285
00286 #endif // KERNEL_USE_DEBUG
00287
00288 #endif

```

19.67 /media/usb/project/github/Mark3/kernel/public/kerneltypes.h File Reference

Basic data type primitives used throughout the OS.

```

#include <stdint.h>
#include <stdbool.h>
#include <stddef.h>

```

Typedefs

- typedef void(* [PanicFunc_t](#))(uint16_t u16PanicCode_)
Function pointer type used to implement kernel-panic handlers.
- typedef void(* [IdleFunc_t](#))(void)
Function pointer type used to implement the idle function, where support for an idle function (as opposed to an idle thread) exists.
- typedef void(* [ThreadEntry_t](#))(void *pvArg_)
Function pointer type used for thread entrypoint functions.

Enumerations

- enum [EventFlagOperation_t](#) {
 [EVENT_FLAG_ALL](#), [EVENT_FLAG_ANY](#), [EVENT_FLAG_ALL_CLEAR](#), [EVENT_FLAG_ANY_CLEAR](#),
 [EVENT_FLAG_MODES](#), [EVENT_FLAG_PENDING_UNBLOCK](#) }
This enumeration describes the different operations supported by the event flag blocking object.
- enum [ThreadState_t](#)
Enumeration representing the different states a thread can exist in.

19.67.1 Detailed Description

Basic data type primitives used throughout the OS.

Definition in file [kerneltypes.h](#).

19.67.2 Enumeration Type Documentation

19.67.2.1 enum EventFlagOperation_t

This enumeration describes the different operations supported by the event flag blocking object.

Enumerator

EVENT_FLAG_ALL Block until all bits in the specified bitmask are set.

EVENT_FLAG_ANY Block until any bits in the specified bitmask are set.

EVENT_FLAG_ALL_CLEAR Block until all bits in the specified bitmask are cleared.

EVENT_FLAG_ANY_CLEAR Block until any bits in the specified bitmask are cleared.

EVENT_FLAG_MODES Count of event-flag modes. Not used by user

EVENT_FLAG_PENDING_UNBLOCK Special code. Not used by user

Definition at line 50 of file [kerneltypes.h](#).

19.68 kerneltypes.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00019 #include <stdint.h>
00020 #include <stdbool.h>
00021 #include <stddef.h>
00022
00023 #ifndef __KERNELTYPES_H__
00024 #define __KERNELTYPES_H__
00025
00026 //-----
00030 typedef void (*PanicFunc_t)(uint16_t ul6PanicCode_);
00031
00032 //-----
00037 typedef void (*IdleFunc_t)(void);
00038
00039 //-----
00043 typedef void (*ThreadEntry_t)(void* pvArg_);
00044
00045 //-----
00050 typedef enum {
00051     EVENT_FLAG_ALL,
00052     EVENT_FLAG_ANY,
00053     EVENT_FLAG_ALL_CLEAR,
00054     EVENT_FLAG_ANY_CLEAR,
00055
00056     EVENT_FLAG_MODES,          //---
00057     EVENT_FLAG_PENDING_UNBLOCK
00058 } EventFlagOperation_t;
00059
00060 //-----
00064 typedef enum {
00065     THREAD_STATE_EXIT = 0,
00066     THREAD_STATE_READY,
00067     THREAD_STATE_BLOCKED,
00068     THREAD_STATE_STOP,
00069     //---
00070     THREAD_STATES,
00071     THREAD_STATE_INVALID
00072 } ThreadState_t;
00073
00074 #endif

```

19.69 /media/usb/project/github/Mark3/kernel/public/ksemaphore.h File Reference

[Semaphore](#) Blocking Object class declarations.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "blocking.h"
#include "threadlist.h"
```

Classes

- class [Semaphore](#)

Binary & Counting semaphores, based on [BlockingObject](#) base class.

19.69.1 Detailed Description

[Semaphore](#) Blocking Object class declarations.

Definition in file [ksemaphore.h](#).

19.70 ksemaphore.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00022 #ifndef __KSEMAPHORE_H__
00023 #define __KSEMAPHORE_H__
00024
00025 #include "kerneltypes.h"
00026 #include "mark3cfg.h"
00027
00028 #include "blocking.h"
00029 #include "threadlist.h"
00030
00031 #if KERNEL_USE_SEMAPHORE
00032
00033 //-----
00037 class Semaphore : public BlockingObject
00038 {
00039 public:
00040     void* operator new(size_t sz, void* pv) { return (Semaphore*)pv; };
00041     ~Semaphore();
00042
00043     void Init(uint16_t u16InitVal_, uint16_t u16MaxVal_);
00044
00045     bool Post();
00046
00047     void Pend();
00048
00049     uint16_t GetCount();
00050
00051 #if KERNEL_USE_TIMEOUTS
00052     bool Pend(uint32_t u32WaitTimeMS_);
00053     void WakeMe(Thread* pclChosenOne_);
00054 #endif
00055 private:
00056     uint8_t WakeNext();
00057
00058 #if KERNEL_USE_TIMEOUTS
00059     bool Pend_i(uint32_t u32WaitTimeMS_);
00060 #else
00061
00062 
```

```

00156     void Pend_i(void);
00157 #endif
00158
00159     uint16_t m_u16Value;
00160     uint16_t m_u16MaxValue;
00161 };
00162
00163 #endif // KERNEL_USE_SEMAPHORE
00164
00165 #endif

```

19.71 /media/usb/project/github/Mark3/kernel/public/ll.h File Reference

Core linked-list declarations, used by all kernel list types.

```
#include "kerneltypes.h"
```

Classes

- class [LinkedListNode](#)
Basic linked-list node data structure.
- class [LinkedList](#)
Abstract-data-type from which all other linked-lists are derived.
- class [DoubleLinkedList](#)
Doubly-linked-list data type, inherited from the base [LinkedList](#) type.
- class [CircularLinkedList](#)
Circular-linked-list data type, inherited from the base [LinkedList](#) type.

19.71.1 Detailed Description

Core linked-list declarations, used by all kernel list types.

At the heart of RTOS data structures are linked lists. Having a robust and efficient set of linked-list types that we can use as a foundation for building the rest of our kernel types allows u16 to keep our RTOS code efficient and logically-separated.

So what data types rely on these linked-list classes?

-Threads -ThreadLists -The [Scheduler](#) -Timers, -The [Timer Scheduler](#) -Blocking objects (Semaphores, Mutexes, etc...)

Pretty much everything in the kernel uses these linked lists. By having objects inherit from the base linked-list node type, we're able to leverage the double and circular linked-list classes to manager virtually every object type in the system without duplicating code. These functions are very efficient as well, allowing for very deterministic behavior in our code.

Definition in file [ll.h](#).

19.72 ll.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.

```

```

00012 See license.txt for more information
00013 ===== */
00043 #ifndef __LL_H__
00044 #define __LL_H__
00045
00046 #include "kerneltypes.h"
00047
00048 //-----
00049 #ifndef NULL
00050 #define NULL (0)
00051 #endif
00052
00053 //-----
00059 class LinkList;
00060 class DoubleLinkList;
00061 class CircularLinkList;
00062
00063 //-----
00068 class LinkListNode
00069 {
00070 protected:
00071     LinkListNode* next;
00072     LinkListNode* prev;
00073
00074     LinkListNode() {}
00080     void ClearNode();
00081
00082 public:
00090     LinkListNode* GetNext(void) { return next; }
00098     LinkListNode* GetPrev(void) { return prev; }
00099     friend class LinkList;
00100     friend class DoubleLinkList;
00101     friend class CircularLinkList;
00102     friend class ThreadList;
00103 };
00104
00105 //-----
00109 class LinkList
00110 {
00111 protected:
00112     LinkListNode* m_pstHead;
00113     LinkListNode* m_pstTail;
00114
00115 public:
00121     void Init()
00122     {
00123         m_pstHead = NULL;
00124         m_pstTail = NULL;
00125     }
00126
00134     LinkListNode* GetHead() { return m_pstHead; }
00142     LinkListNode* GetTail() { return m_pstTail; }
00143 };
00144
00145 //-----
00149 class DoubleLinkList : public LinkList
00150 {
00151 public:
00152     void* operator new(size_t sz, void* pv) { return (DoubleLinkList*)pv; };
00158     DoubleLinkList()
00159     {
00160         m_pstHead = NULL;
00161         m_pstTail = NULL;
00162     }
00163
00171     void Add(LinkListNode* node_);
00172
00180     void Remove(LinkListNode* node_);
00181 };
00182
00183 //-----
00187 class CircularLinkList : public LinkList
00188 {
00189 public:
00190     void* operator new(size_t sz, void* pv) { return (CircularLinkList*)pv; };
00191     CircularLinkList()
00192     {
00193         m_pstHead = NULL;
00194         m_pstTail = NULL;
00195     }
00196
00204     void Add(LinkListNode* node_);
00205
00213     void Remove(LinkListNode* node_);
00214
00221     void PivotForward();
00222

```

```

00229     void PivotBackward();
00230
00240     void InsertNodeBefore(LinkListNode* node_,
        LinkListNode* insert_);
00241 };
00242
00243 #endif

```

19.73 /media/usb/project/github/Mark3/kernel/public/mailbox.h File Reference

Mailbox + Envelope IPC Mechanism.

```

#include "mark3cfg.h"
#include "kerneltypes.h"
#include "threadport.h"
#include "ksemaphore.h"

```

Classes

- class **Mailbox**

*The **Mailbox** class implements an IPC mechnism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user.*

19.73.1 Detailed Description

Mailbox + Envelope IPC Mechanism.

Definition in file [mailbox.h](#).

19.74 mailbox.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #ifndef __MAILBOX_H__
00022 #define __MAILBOX_H__
00023
00024 #include "mark3cfg.h"
00025 #include "kerneltypes.h"
00026 #include "threadport.h"
00027 #include "ksemaphore.h"
00028
00029 #if KERNEL_USE_MAILBOX
00030
00036 class Mailbox
00037 {
00038 public:
00039     void* operator new(size_t sz, void* pv) { return (Mailbox*)pv; };
00040     ~Mailbox();
00041
00052     void Init(void* pvBuffer_, uint16_t ul6BufferSize_, uint16_t ul6ElementSize_);
00053
00054 #if KERNEL_USE_AUTO_ALLOC
00055
00068     static Mailbox* Init(uint16_t ul6BufferSize_, uint16_t ul6ElementSize_);
00069
00070 #endif

```

```

00071
00085     bool Send(void* pvData_);
00086
00100     bool SendTail(void* pvData_);
00101
00102 #if KERNEL_USE_TIMEOUTS
00103
00117     bool Send(void* pvData_, uint32_t u32TimeoutMS_);
00118
00133     bool SendTail(void* pvData_, uint32_t u32TimeoutMS_);
00134 #endif
00135
00145     void Receive(void* pvData_);
00146
00156     void ReceiveTail(void* pvData_);
00157
00158 #if KERNEL_USE_TIMEOUTS
00159
00171     bool Receive(void* pvData_, uint32_t u32TimeoutMS_);
00172
00185     bool ReceiveTail(void* pvData_, uint32_t u32TimeoutMS_);
00186 #endif
00187
00188     uint16_t GetFreeSlots(void)
00189     {
00190         uint16_t rc;
00191         CS_ENTER();
00192         rc = m_ul6Free;
00193         CS_EXIT();
00194         return rc;
00195     }
00196
00197     bool IsFull(void) { return (GetFreeSlots() == 0); }
00198     bool IsEmpty(void) { return (GetFreeSlots() == m_ul6Count); }
00199 private:
00208     void* GetHeadPointer(void)
00209     {
00210         K_ADDR uAddr = (K_ADDR)m_pvBuffer;
00211         uAddr += (K_ADDR)(m_ul6ElementSize) * (K_ADDR)(
m_ul6Head);
00212         return (void*)uAddr;
00213     }
00214
00223     void* GetTailPointer(void)
00224     {
00225         K_ADDR uAddr = (K_ADDR)m_pvBuffer;
00226         uAddr += (K_ADDR)(m_ul6ElementSize) * (K_ADDR)(
m_ul6Tail);
00227         return (void*)uAddr;
00228     }
00229
00239     void CopyData(const void* src_, const void* dst_, uint16_t len_)
00240     {
00241         uint8_t* u8Src = (uint8_t*)src_;
00242         uint8_t* u8Dst = (uint8_t*)dst_;
00243         while (len_-- > 0) {
00244             *u8Dst++ = *u8Src++;
00245         }
00246     }
00247
00253     void MoveTailForward(void)
00254     {
00255         m_ul6Tail++;
00256         if (m_ul6Tail == m_ul6Count) {
00257             m_ul6Tail = 0;
00258         }
00259     }
00260
00266     void MoveHeadForward(void)
00267     {
00268         m_ul6Head++;
00269         if (m_ul6Head == m_ul6Count) {
00270             m_ul6Head = 0;
00271         }
00272     }
00273
00279     void MoveTailBackward(void)
00280     {
00281         if (m_ul6Tail == 0) {
00282             m_ul6Tail = m_ul6Count;
00283         }
00284         m_ul6Tail--;
00285     }
00286
00292     void MoveHeadBackward(void)
00293     {
00294         if (m_ul6Head == 0) {

```


19.77 /media/usb/project/github/Mark3/kernel/public/mark3.h File Reference

Single include file given to users of the Mark3 [Kernel](#) API.

```
#include "mark3cfg.h"
#include "kerneltypes.h"
#include "threadport.h"
#include "kernelswi.h"
#include "kerneltimer.h"
#include "kernelprofile.h"
#include "kernel.h"
#include "thread.h"
#include "timerlist.h"
#include "ksemaphore.h"
#include "mutex.h"
#include "eventflag.h"
#include "message.h"
#include "notify.h"
#include "mailbox.h"
#include "atomic.h"
#include "driver.h"
#include "kernelaware.h"
#include "profile.h"
#include "autoalloc.h"
#include "priomap.h"
```

19.77.1 Detailed Description

Single include file given to users of the Mark3 [Kernel](#) API.

Definition in file [mark3.h](#).

19.78 mark3.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #ifndef __MARK3_H__
00022 #define __MARK3_H__
00023
00024 #include "mark3cfg.h"
00025 #include "kerneltypes.h"
00026
00027 #include "threadport.h"
00028 #include "kernelswi.h"
00029 #include "kerneltimer.h"
00030 #include "kernelprofile.h"
00031
00032 #include "kernel.h"
00033 #include "thread.h"
00034 #include "timerlist.h"
00035
00036 #include "ksemaphore.h"
00037 #include "mutex.h"
00038 #include "eventflag.h"
00039 #include "message.h"
```

```

00040 #include "notify.h"
00041 #include "mailbox.h"
00042
00043 #include "atomic.h"
00044 #include "driver.h"
00045
00046 #include "kernelaware.h"
00047
00048 #include "profile.h"
00049 #include "autoalloc.h"
00050 #include "priomap.h"
00051
00052 #endif

```

19.79 /media/usb/project/github/Mark3/kernel/public/mark3cfg.h File Reference

Mark3 [Kernel](#) Configuration.

```
#include "portcfg.h"
```

Macros

- `#define` [KERNEL_NUM_PRIORITIES](#) (8)
Define the number of thread priorities that the kernel's scheduler will support.
- `#define` [KERNEL_USE_TIMERS](#) (1)
The following options is related to all kernel time-tracking.
- `#define` [KERNEL_TIMERS_TICKLESS](#) (1)
If you've opted to use the kernel timers module, you have an option as to which timer implementation to use: Tick-based or Tick-less.
- `#define` [KERNEL_TIMERS_MINIMUM_DELAY_US](#) (25)
When using tickless timers, it is useful to define a minimum sleep value.
- `#define` [KERNEL_TIMERS_THREADED](#) (0)
When timers are enabled, configure whether or not a dedicated thread is used to service timer maintenance.
- `#define` [KERNEL_USE_TIMEOUTS](#) (1)
Set the priority of the timer thread, if the kernel is configured to use a dedicated timer thread.
- `#define` [KERNEL_USE_QUANTUM](#) (1)
Do you want to enable time quanta? This is useful when you want to have tasks in the same priority group share time in a controlled way.
- `#define` [THREAD_QUANTUM_DEFAULT](#) (4)
This value defines the default thread quantum when [KERNEL_USE_QUANTUM](#) is enabled.
- `#define` [KERNEL_USE_NOTIFY](#) (1)
This is a simple blocking object, where a thread (or threads) are guaranteed to block until an asynchronous event signals the object.
- `#define` [KERNEL_USE_SEMAPHORE](#) (1)
Do you want the ability to use counting/binary semaphores for thread synchronization? Enabling this features provides fully-blocking semaphores and enables all API functions declared in [semaphore.h](#).
- `#define` [KERNEL_USE_MUTEX](#) (1)
Do you want the ability to use mutual exclusion semaphores (mutex) for resource/block protection? Enabling this feature provides mutexes, with priority inheritance, as declared in [mutex.h](#).
- `#define` [KERNEL_USE_EVENTFLAG](#) (1)
Provides additional event-flag based blocking.
- `#define` [KERNEL_USE_MESSAGE](#) (1)
Enable inter-thread messaging using message queues.
- `#define` [GLOBAL_MESSAGE_POOL_SIZE](#) (8)
If Messages are enabled, define the size of the default kernel message pool.

- `#define KERNEL_USE_MAILBOX (1)`
Enable inter-thread messaging using mailboxes.
- `#define KERNEL_USE_SLEEP (1)`
Do you want to be able to set threads to sleep for a specified time? This enables the [Thread::Sleep\(\)](#) API.
- `#define KERNEL_USE_DRIVER (1)`
Enabling device drivers provides a posix-like filesystem interface for peripheral device drivers.
- `#define KERNEL_USE_THREADNAME (0)`
Provide [Thread](#) method to allow the user to set a name for each thread in the system.
- `#define KERNEL_USE_EXTENDED_CONTEXT (1)`
Allocate an extra pointer's worth of storage within a [Thread](#) object (and corresponding accessor methods) to provide the user with a means to implement arbitrary Thread-local storage.
- `#define KERNEL_USE_DYNAMIC_THREADS (1)`
Provide extra [Thread](#) methods to allow the application to create (and more importantly destroy) threads at runtime.
- `#define KERNEL_USE_PROFILER (1)`
Provides extra classes for profiling the performance of code.
- `#define KERNEL_USE_DEBUG (0)`
Provides extra logic for kernel debugging, and instruments the kernel with extra asserts, and kernel trace functionality.
- `#define KERNEL_USE_ATOMIC (0)`
Provides support for atomic operations, including addition, subtraction, set, and test-and-set.
- `#define SAFE_UNLINK (0)`
"Safe unlinking" performs extra checks on data to make sure that there are no consistencies when performing operations on linked lists.
- `#define KERNEL_AWARE_SIMULATION (1)`
Include support for kernel-aware simulation.
- `#define KERNEL_USE_IDLE_FUNC (1)`
Enabling this feature removes the necessity for the user to dedicate a complete thread for idle functionality.
- `#define KERNEL_USE_AUTO_ALLOC (0)`
This feature enables an additional set of APIs that allow for objects to be created on-the-fly out of a special heap, without having to explicitly allocate them (from stack, heap, or static memory).
- `#define KERNEL_USE_THREAD_CALLOUTS (1)`
This feature provides additional kernel APIs to register callout functions that are activated when threads are created or exited.
- `#define KERNEL_USE_STACK_GUARD (1)`
This feature, when enabled, tells the kernel to check whether any [Thread](#)'s stack has been exhausted (or slack falls below a certain safety threshold) before executing each context switch.
- `#define KERNEL_EXTRA_CHECKS (1)`
This option provides extra safety checks within the kernel APIs in order to minimize the potential for unsafe operations.

19.79.1 Detailed Description

Mark3 [Kernel](#) Configuration.

This file is used to configure the kernel for your specific application in order to provide the optimal set of features for a given use case.

Since you only pay the price (code space/RAM) for the features you use, you can usually find a sweet spot between features and resource usage by picking and choosing features a-la-carte. This config file is written in an "interactive" way, in order to minimize confusion about what each option provides, and to make dependencies obvious.

Definition in file [mark3cfg.h](#).

19.79.2 Macro Definition Documentation

19.79.2.1 `#define GLOBAL_MESSAGE_POOL_SIZE` (8)

If Messages are enabled, define the size of the default kernel message pool.

Messages can be manually added to the message pool, but this mechanism is more convenient and automatic. All message queues share their message objects from this global pool to maximize efficiency and simplify data management.

Definition at line 191 of file [mark3cfg.h](#).

19.79.2.2 `#define KERNEL_AWARE_SIMULATION` (1)

Include support for kernel-aware simulation.

Enabling this feature adds advanced profiling, trace, and environment-aware debugging and diagnostic functionality when Mark3-based applications are run on the flavr AVR simulator.

Definition at line 301 of file [mark3cfg.h](#).

19.79.2.3 `#define KERNEL_EXTRA_CHECKS` (1)

This option provides extra safety checks within the kernel APIs in order to minimize the potential for unsafe operations.

This is especially helpful during development, and can help catch problems at development time, instead of in the field. include CPU/Port specific configuration options

Definition at line 356 of file [mark3cfg.h](#).

19.79.2.4 `#define KERNEL_NUM_PRIORITIES` (8)

Define the number of thread priorities that the kernel's scheduler will support.

The number of thread priorities is limited only by the memory of the host CPU, as a [ThreadList](#) object is statically-allocated for each thread priority.

In practice, systems rarely need more than 32 priority levels, with the most complex having the capacity for 256.

Definition at line 41 of file [mark3cfg.h](#).

19.79.2.5 `#define KERNEL_TIMERS_MINIMUM_DELAY_US` (25)

When using tickless timers, it is useful to define a minimum sleep value.

In the event that a delay/sleep/timeout value lower than this is provided to a timer-based API, the minimum value will be substituted.

Definition at line 86 of file [mark3cfg.h](#).

19.79.2.6 `#define KERNEL_TIMERS_THREADED` (0)

When timers are enabled, configure whether or not a dedicated thread is used to service timer maintenance.

If set to 0, timer handlers are executed from a nested interrupt context.

Definition at line 95 of file [mark3cfg.h](#).

19.79.2.7 `#define KERNEL_TIMERS_TICKLESS (1)`

If you've opted to use the kernel timers module, you have an option as to which timer implementation to use: Tick-based or Tick-less.

Tick-based timers provide a "traditional" RTOS timer implementation based on a fixed-frequency timer interrupt. While this provides very accurate, reliable timing, it also means that the CPU is being interrupted far more often than may be necessary (as not all timer ticks result in "real work" being done).

Tick-less timers still rely on a hardware timer interrupt, but uses a dynamic expiry interval to ensure that the interrupt is only called when the next timer expires. This increases the complexity of the timer interrupt handler, but reduces the number and frequency.

Note that the CPU port ([kerneltimer.cpp](#)) must be implemented for the particular timer variant desired.

Definition at line 77 of file [mark3cfg.h](#).

19.79.2.8 `#define KERNEL_USE_ATOMIC (0)`

Provides support for atomic operations, including addition, subtraction, set, and test-and-set.

Add/Sub/Set contain 8, 16, and 32-bit variants.

Definition at line 285 of file [mark3cfg.h](#).

19.79.2.9 `#define KERNEL_USE_AUTO_ALLOC (0)`

This feature enables an additional set of APIs that allow for objects to be created on-the-fly out of a special heap, without having to explicitly allocate them (from stack, heap, or static memory).

Note that auto-alloc memory cannot be reclaimed.

Definition at line 322 of file [mark3cfg.h](#).

19.79.2.10 `#define KERNEL_USE_DYNAMIC_THREADS (1)`

Provide extra [Thread](#) methods to allow the application to create (and more importantly destroy) threads at runtime. useful for designs implementing worker threads, or threads that can be restarted after encountering error conditions.

Definition at line 245 of file [mark3cfg.h](#).

19.79.2.11 `#define KERNEL_USE_EVENTFLAG (1)`

Provides additional event-flag based blocking.

This relies on an additional per-thread flag-mask to be allocated, which adds 2 bytes to the size of each thread object.

Definition at line 170 of file [mark3cfg.h](#).

19.79.2.12 `#define KERNEL_USE_IDLE_FUNC (1)`

Enabling this feature removes the necessity for the user to dedicate a complete thread for idle functionality.

This saves a full thread stack, but also requires a bit extra static data. This also adds a slight overhead to the context switch and scheduler, as a special case has to be taken into account.

Definition at line 311 of file [mark3cfg.h](#).

19.79.2.13 #define KERNEL_USE_MAILBOX (1)

Enable inter-thread messaging using mailboxes.

A mailbox manages a blob of data provided by the user, that is partitioned into fixed-size blocks called envelopes. The size of an envelope is set by the user when the mailbox is initialized. Any number of threads can read-from and write-to the mailbox. Envelopes can be sent-to or received-from the mailbox at the head or tail. In this way, mailboxes essentially act as a circular buffer that can be used as a blocking FIFO or LIFO queue.

Definition at line 204 of file [mark3cfg.h](#).

19.79.2.14 #define KERNEL_USE_MESSAGE (1)

Enable inter-thread messaging using message queues.

This is the preferred mechanism for IPC for serious multi-threaded communications; generally anywhere a semaphore or event-flag is insufficient.

Definition at line 178 of file [mark3cfg.h](#).

19.79.2.15 #define KERNEL_USE_PROFILER (1)

Provides extra classes for profiling the performance of code.

useful for debugging and development, but uses an additional hardware timer.

Definition at line 251 of file [mark3cfg.h](#).

19.79.2.16 #define KERNEL_USE_QUANTUM (1)

Do you want to enable time quanta? This is useful when you want to have tasks in the same priority group share time in a controlled way.

This allows equal tasks to use unequal amounts of the CPU, which is a great way to set up CPU budgets per thread in a round-robin scheduling system. If enabled, you can specify a number of ticks that serves as the default time period (quantum). Unless otherwise specified, every thread in a priority will get the default quantum.

Definition at line 133 of file [mark3cfg.h](#).

19.79.2.17 #define KERNEL_USE_SEMAPHORE (1)

Do you want the ability to use counting/binary semaphores for thread synchronization? Enabling this features provides fully-blocking semaphores and enables all API functions declared in semaphore.h.

If you have to pick one blocking mechanism, this is the one to choose.

Definition at line 156 of file [mark3cfg.h](#).

19.79.2.18 #define KERNEL_USE_STACK_GUARD (1)

This feature, when enabled, tells the kernel to check whether any [Thread](#)'s stack has been exhausted (or slack falls below a certain safety threshold) before executing each context switch.

Enabling this is the most effective means to guard against stack corruption and stack overflow in the kernel, at the cost of increased context switch latency.

Definition at line 344 of file [mark3cfg.h](#).

19.79.2.19 #define KERNEL_USE_THREAD_CALLOUTS (1)

This feature provides additional kernel APIs to register callout functions that are activated when threads are created or exited.

This is useful for implementing low-level instrumentation based on information held in the threads.

Definition at line 334 of file [mark3cfg.h](#).

19.79.2.20 #define KERNEL_USE_THREADNAME (0)

Provide [Thread](#) method to allow the user to set a name for each thread in the system.

Adds a const char* pointer to the size of the thread object.

Definition at line 230 of file [mark3cfg.h](#).

19.79.2.21 #define KERNEL_USE_TIMEOUTS (1)

Set the priority of the timer thread, if the kernel is configured to use a dedicated timer thread.

By default, if you opt to enable kernel timers, you also get timeout-enabled versions of the blocking object APIs along with it. This support comes at a small cost to code size, but a slightly larger cost to realtime performance - as checking for the use of timers in the underlying internal code costs some cycles.

As a result, the option is given to the user here to manually disable these timeout-based APIs if desired by the user for performance and code-size reasons.

Definition at line 118 of file [mark3cfg.h](#).

19.79.2.22 #define KERNEL_USE_TIMERS (1)

The following options is related to all kernel time-tracking.

-timers provide a way for events to be periodically triggered in a lightweight manner. These can be periodic, or one-shot.

-Thread [Quantum](#) (used for round-robin scheduling) is dependent on this module, as is [Thread](#) Sleep functionality.

Definition at line 56 of file [mark3cfg.h](#).

19.79.2.23 #define SAFE_UNLINK (0)

"Safe unlinking" performs extra checks on data to make sure that there are no inconsistencies when performing operations on linked lists.

This goes beyond pointer checks, adding a layer of structural and metadata validation to help detect system corruption early.

Definition at line 293 of file [mark3cfg.h](#).

19.79.2.24 #define THREAD_QUANTUM_DEFAULT (4)

This value defines the default thread quantum when KERNEL_USE_QUANTUM is enabled.

The thread quantum value is in milliseconds

Definition at line 142 of file [mark3cfg.h](#).


```

00218
00223 #define KERNEL_USE_DRIVER (1)
00224
00230 #define KERNEL_USE_THREADNAME (0)
00231
00237 #define KERNEL_USE_EXTENDED_CONTEXT (1)
00238
00245 #define KERNEL_USE_DYNAMIC_THREADS (1)
00246
00251 #define KERNEL_USE_PROFILER (1)
00252
00257 #define KERNEL_USE_DEBUG (0)
00258
00259 #if KERNEL_USE_DEBUG
00260
00266 #define KERNEL_ENABLE_LOGGING (0)
00267
00275 #define KERNEL_ENABLE_USER_LOGGING (0)
00276 #else
00277 #define KERNEL_ENABLE_LOGGING (0)
00278 #define KERNEL_ENABLE_USER_LOGGING (0)
00279 #endif
00280
00285 #define KERNEL_USE_ATOMIC (0)
00286
00293 #define SAFE_UNLINK (0)
00294
00301 #define KERNEL_AWARE_SIMULATION (1)
00302
00310 #if !defined(ARM)
00311 #define KERNEL_USE_IDLE_FUNC (1) // Supported everywhere but ARM
00312 #else
00313 #define KERNEL_USE_IDLE_FUNC (0) // Not currently supported on ARM
00314 #endif
00315
00322 #define KERNEL_USE_AUTO_ALLOC (0)
00323
00324 #if KERNEL_USE_AUTO_ALLOC
00325 #define AUTO_ALLOC_SIZE (512)
00326 #endif
00327
00334 #define KERNEL_USE_THREAD_CALLOUTS (1)
00335
00344 #define KERNEL_USE_STACK_GUARD (1)
00345
00346 #if KERNEL_USE_STACK_GUARD
00347 #define KERNEL_STACK_GUARD_DEFAULT (32) // words
00348 #endif
00349
00356 #define KERNEL_EXTRA_CHECKS (1)
00357
00358 #include "portcfg.h"
00359
00360 #endif

```

19.81 /media/usb/project/github/Mark3/kernel/public/message.h File Reference

Inter-thread communication via message-passing.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "ksemaphore.h"
#include "timerlist.h"

```

Classes

- class [Message](#)
Class to provide message-based IPC services in the kernel.
- class [MessagePool](#)
Implements a list of message objects.
- class [GlobalMessagePool](#)

Implements a list of message objects shared between all threads.

- class [MessageQueue](#)

List of messages, used as the channel for sending and receiving messages between threads.

19.81.1 Detailed Description

Inter-thread communication via message-passing.

Embedded systems guru Jack Ganssle once said that without a robust form of interprocess communications (IPC), an RTOS is just a toy. Mark3 implements a form of IPC to provide safe and flexible messaging between threads.

using kernel-managed IPC offers significant benefits over other forms of data sharing (i.e. Global variables) in that it avoids synchronization issues and race conditions common to the practice. using IPC also enforces a more disciplined coding style that keeps threads decoupled from one another and minimizes global data, preventing careless and hard-to-debug errors.

19.81.2 using Messages, Queues, and the Global Message Pool

```
// Declare a message queue shared between two threads
MessageQueue my_queue;

int main()
{
    ...
    // Initialize the message queue
    my_queue.init();
    ...
}

void Thread1()
{
    // Example TX thread - sends a message every 10ms
    while(1)
    {
        // Grab a message from the global message pool
        Message *tx_message = GlobalMessagePool::Pop();

        // Set the message data/parameters
        tx_message->SetCode( 1234 );
        tx_message->SetData( NULL );

        // Send the message on the queue.
        my_queue.Send( tx_message );
        Thread::Sleep(10);
    }
}

void Thread2()
{
    while()
    {
        // Blocking receive - wait until we have messages to process
        Message *rx_message = my_queue.Recv();

        // Do something with the message data...

        // Return back into the pool when done
        GlobalMessagePool::Push(rx_message);
    }
}
```

Definition in file [message.h](#).

19.82 message.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
```

```

00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00080 #ifndef __MESSAGE_H__
00081 #define __MESSAGE_H__
00082
00083 #include "kerneltypes.h"
00084 #include "mark3cfg.h"
00085
00086 #include "ll.h"
00087 #include "ksemaphore.h"
00088
00089 #if KERNEL_USE_MESSAGE
00090
00091 #if KERNEL_USE_TIMEOUTS
00092 #include "timerlist.h"
00093 #endif
00094
00095 //-----
00099 class Message : public LinkListNode
00100 {
00101 public:
00102     void* operator new(size_t sz, void* pv) { return (Message*)pv; };
00108     void Init()
00109     {
00110         ClearNode();
00111         m_pvData = NULL;
00112         m_ul6Code = 0;
00113     }
00114
00122     void SetData(void* pvData_) { m_pvData = pvData_; }
00130     void* GetData() { return m_pvData; }
00138     void SetCode(uint16_t ul6Code_) { m_ul6Code = ul6Code_; }
00146     uint16_t GetCode() { return m_ul6Code; }
00147 private:
00149     void* m_pvData;
00150
00152     uint16_t m_ul6Code;
00153 };
00154
00155 //-----
00159 class MessagePool
00160 {
00161 public:
00167     void Init();
00168
00178     void Push(Message* pclMessage_);
00179
00188     Message* Pop();
00189
00197     Message* GetHead();
00198
00199 private:
00201     DoubleLinkedList m_clList;
00202 };
00203
00204 //-----
00208 class GlobalMessagePool
00209 {
00210 public:
00216     static void Init();
00217
00227     static void Push(Message* pclMessage_);
00228
00237     static Message* Pop();
00238
00246     static Message* GetHead();
00247
00255     static MessagePool* GetPool();
00256
00257 private:
00259     static Message m_aclMessagePool[
GLOBAL_MESSAGE_POOL_SIZE];
00260
00261     static MessagePool m_clPool;
00262 };
00263
00264 //-----
00269 class MessageQueue
00270 {
00271 public:
00272     void* operator new(size_t sz, void* pv) { return (MessageQueue*)pv; };
00278     void Init();
00279

```

```

00288     Message* Receive();
00289
00290 #if KERNEL_USE_TIMEOUTS
00291
00305     Message* Receive(uint32_t u32TimeWaitMS);
00306 #endif
00307
00316     void Send(Message* pclSrc_);
00317
00325     uint16_t GetCount();
00326
00327 private:
00328 #if KERNEL_USE_TIMEOUTS
00329
00338     Message* Receive_i(uint32_t u32TimeWaitMS);
00339 #else
00340
00347     Message* Receive_i(void);
00348 #endif
00349
00351     Semaphore m_clSemaphore;
00352
00354     DoubleLinkedList m_clLinkList;
00355 };
00356
00357 #endif // KERNEL_USE_MESSAGE
00358
00359 #endif

```

19.83 /media/usb/project/github/Mark3/kernel/public/mutex.h File Reference

Mutual exclusion class declaration.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "blocking.h"

```

Classes

- class [Mutex](#)

Mutual-exclusion locks, based on [BlockingObject](#).

19.83.1 Detailed Description

Mutual exclusion class declaration.

Resource locks are implemented using mutual exclusion semaphores ([Mutex_t](#)). Protected blocks can be placed around any resource that may only be accessed by one thread at a time. If additional threads attempt to access the protected resource, they will be placed in a wait queue until the resource becomes available. When the resource becomes available, the thread with the highest original priority claims the resource and is activated. Priority inheritance is included in the implementation to prevent priority inversion. Always ensure that you claim and release your mutex objects consistently, otherwise you may end up with a deadlock scenario that's hard to debug.

19.83.2 Initializing

Initializing a mutex object by calling:

```
clMutex.Init();
```

19.83.3 Resource protection example

```
clMutex.Claim();
```

```
...
<resource protected block>
...
clMutex.Release();
```

Definition in file [mutex.h](#).

19.84 mutex.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00050 #ifndef __MUTEX_H_
00051 #define __MUTEX_H_
00052
00053 #include "kerneltypes.h"
00054 #include "mark3cfg.h"
00055
00056 #include "blocking.h"
00057
00058 #if KERNEL_USE_MUTEX
00059
00060 //-----
00064 class Mutex : public BlockingObject
00065 {
00066 public:
00067     void* operator new(size_t sz, void* pv) { return (Mutex*)pv; };
00068     ~Mutex();
00069
00076     void Init();
00077
00095     void Claim();
00096
00097 #if KERNEL_USE_TIMEOUTS
00098     bool Claim(uint32_t u32WaitTimeMS_);
00109
00110     void WakeMe(Thread* pclOwner_);
00123
00124 #endif
00125
00126     void Release();
00147
00148 private:
00155     uint8_t WakeNext();
00156
00157 #if KERNEL_USE_TIMEOUTS
00158     bool Claim_i(uint32_t u32WaitTimeMS_);
00166 #else
00167     void Claim_i(void);
00174 #endif
00175 #endif
00176
00177     uint8_t m_u8Recurse;
00178     bool m_bReady;
00179     uint8_t m_u8MaxPri;
00180     Thread* m_pclOwner;
00181 };
00182
00183 #endif // KERNEL_USE_MUTEX
00184
00185 #endif // __MUTEX_H_
```

19.85 /media/usb/project/github/Mark3/kernel/public/notify.h File Reference

Lightweight thread notification - blocking object.

```
#include "mark3cfg.h"
#include "blocking.h"
```

Classes

- class [Notify](#)

The [Notify](#) class is a blocking object type, that allows one or more threads to wait for an event to occur before resuming operation.

19.85.1 Detailed Description

Lightweight thread notification - blocking object.

Definition in file [notify.h](#).

19.86 notify.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #ifndef __NOTIFY_H__
00022 #define __NOTIFY_H__
00023
00024 #include "mark3cfg.h"
00025 #include "blocking.h"
00026
00027 #if KERNEL_USE_NOTIFY
00028
00033 class Notify : public BlockingObject
00034 {
00035 public:
00036     void* operator new(size_t sz, void* pv) { return (Notify*)pv; };
00037     ~Notify();
00038
00044     void Init(void);
00045
00055     void Signal(void);
00056
00066     void Wait(bool* pbFlag_);
00067
00068 #if KERNEL_USE_TIMEOUTS
00069
00081     bool Wait(uint32_t u32WaitTimeMS_, bool* pbFlag_);
00082 #endif
00083
00093     void WakeMe(Thread* pclChosenOne_);
00094
00095 private:
00096
00097     bool m_bPending;
00098 };
00099
00100 #endif
00101
00102 #endif
```

19.87 /media/usb/project/github/Mark3/kernel/public/paniccodes.h File Reference

Defines the reason codes thrown when a kernel panic occurs.


```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
```

Classes

- class ProfileTimer
Profiling timer.

19.91.1 Detailed Description

High-precision profiling timers.

Enables the profiling and instrumentation of performance-critical code. Multiple timers can be used simultaneously to enable system-wide performance metrics to be computed in a lightweight manner.

Usage:

```

ProfileTimer clMyTimer;
int i;

clMyTimer.Init();

// Profile the same block of code ten times
for (i = 0; i < 10; i++)
{
    clMyTimer.Start();
    ...
    //Block of code to profile
    ...
    clMyTimer.Stop();
}

// Get the average execution time of all iterations
u32LastTimer = clMyTimer.GetAverage();

// Get the execution time from the last iteration
u32LastTimer = clMyTimer.GetCurrent();

```

Definition in file [profile.h](#).

19.92 profile.h

[illegible]

```

00086     void Start();
00087
00094     void Stop();
00095
00103     uint32_t GetAverage();
00104
00113     uint32_t GetCurrent();
00114
00115 private:
00126     uint32_t ComputeCurrentTicks(uint16_t u16Current_, uint32_t u32Epoch_);
00127
00128     uint32_t m_u32Cumulative;
00129     uint32_t m_u32CurrentIteration;
00130     uint16_t m_u16Initial;
00131     uint32_t m_u32InitialEpoch;
00132     uint16_t m_u16Iterations;
00133     bool     m_bActive;
00134 };
00135
00136 #endif // KERNEL_USE_PROFILE
00137
00138 #endif

```

19.93 /media/usb/project/github/Mark3/kernel/public/quantum.h File Reference

[Thread Quantum](#) declarations for Round-Robin Scheduling.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "timer.h"
#include "timerlist.h"
#include "timerscheduler.h"

```

Classes

- class [Quantum](#)

Static-class used to implement [Thread](#) quantum functionality, which is a key part of round-robin scheduling.

19.93.1 Detailed Description

[Thread Quantum](#) declarations for Round-Robin Scheduling.

Definition in file [quantum.h](#).

19.94 quantum.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00022 #ifndef __KQUANTUM_H__
00023 #define __KQUANTUM_H__
00024
00025 #include "kerneltypes.h"
00026 #include "mark3cfg.h"
00027
00028 #include "thread.h"

```

```

00029 #include "timer.h"
00030 #include "timerlist.h"
00031 #include "timerscheduler.h"
00032
00033 #if KERNEL_USE_QUANTUM
00034 class Timer;
00035
00041 class Quantum
00042 {
00043 public:
00052     static void UpdateTimer();
00053
00060     static void AddThread(Thread* pclThread_);
00061
00067     static void RemoveThread();
00068
00077     static void SetInTimer(void) { m_bInTimer = true; }
00083     static void ClearInTimer(void) { m_bInTimer = false; }
00084
00085 #if KERNEL_TIMERS_THREADED
00086     static void SetTimerThread(Thread* pclTimerThread_);
00087
00088     static Thread* GetTimerThread();
00089 #endif
00090
00091 private:
00103     static void SetTimer(Thread* pclThread_);
00104
00105 #if KERNEL_TIMERS_THREADED
00106     static Thread* m_pclTimerThread;
00107 #endif
00108
00109     static Timer m_clQuantumTimer;
00110     static bool m_bActive;
00111     static bool m_bInTimer;
00112 };
00113
00114 #endif // KERNEL_USE_QUANTUM
00115
00116 #endif

```

19.95 /media/usb/project/github/Mark3/kernel/public/scheduler.h File Reference

[Thread](#) scheduler function declarations.

```

#include "kerneltypes.h"
#include "thread.h"
#include "threadport.h"
#include "priomap.h"

```

Classes

- class [Scheduler](#)

Priority-based round-robin [Thread](#) scheduling, using [ThreadLists](#) for housekeeping.

Variables

- volatile [Thread](#) * [g_pclNext](#)
Pointer to the currently-chosen next-running thread.
- [Thread](#) * [g_pclCurrent](#)
Pointer to the currently-running thread.

19.95.1 Detailed Description

[Thread](#) scheduler function declarations.


```
00179
00181     static PriorityMap m_clPrioMap;
00182 };
00183 #endif
```

19.97 /media/usb/project/github/Mark3/kernel/public/thread.h File Reference

Platform independent thread class declarations.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "threadlist.h"
#include "scheduler.h"
#include "threadport.h"
#include "quantum.h"
#include "autoalloc.h"
#include "priomap.h"
```

Classes

- class Thread

Object providing fundamental multitasking support in the kernel.

- struct FakeThread t

If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system.

19.97.1 Detailed Description

Platform independent thread class declarations.

Threads are an atomic unit of execution, and each instance of the thread class represents an instance of a program running on the processor. The `Thread` is the fundamental user-facing object in the kernel - it is what makes multiprocessing possible from application code.

In Mark3, threads each have their own context - consisting of a stack, and all of the registers required to multiplex a processor between multiple threads.

The `Thread` class inherits directly from the `LinkedListNode` class to facilitate efficient thread management using Double, or Double-Circular linked lists.

Definition in file [thread.h](#).

19.98 thread.h

```
00001 /*=====
```

```
00002
```

```
00003      |-----|-----|-----|-----|-----|-----|
```

```
00004      | \    / | \    / | \    / | \    / | \    / | \    / |
```

```
00005      |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
```

```
00006      |___\_/___|___\_/___|___\_/___|___\_/___|___\_/___|___\_/___|
```

```
00007      |_____|   |_____|   |_____|   |_____|   |_____|   |_____|   |
```

```
00008
```

```
00009 --[Mark3 Realtime Platform]-----
```

```
00010
```

```
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
```

```
00012 See license.txt for more information
```

```
00013 ===== */
```

```
00035 #ifndef __THREAD_H__
```

```
00036 #define __THREAD_H__
```

```
00037
```

```

00038 #include "kerneltypes.h"
00039 #include "mark3cfg.h"
00040
00041 #include "ll.h"
00042 #include "threadlist.h"
00043 #include "scheduler.h"
00044 #include "threadport.h"
00045 #include "quantum.h"
00046 #include "autoalloc.h"
00047 #include "priomap.h"
00048
00049 class Thread;
00050
00051 //-----
00052 typedef void (*ThreadCreateCallout_t)(Thread* pclThread_);
00053 typedef void (*ThreadExitCallout_t)(Thread* pclThread_);
00054 typedef void (*ThreadContextCallout_t)(Thread* pclThread_);
00055
00056 //-----
00060 class Thread : public LinkListNode
00061 {
00062 public:
00063     void* operator new(size_t sz, void* pv) { return (Thread*)pv; };
00064     ~Thread();
00065
00066     #if KERNEL_EXTRA_CHECKS
00067     Thread() { m_eState = THREAD_STATE_INVALID; }
00068
00069     bool IsInitialized() { return (m_eState != THREAD_STATE_INVALID); }
00070     #endif
00071
00072     void
00073     Init(K_WORD* pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_,
00074         ThreadEntry_t pfEntryPoint_, void* pvArg_);
00075
00076     #if KERNEL_USE_AUTO_ALLOC
00077     static Thread* Init(uint16_t u16StackSize_, uint8_t uXPriority_,
00078         ThreadEntry_t pfEntryPoint_, void* pvArg_);
00079     #endif
00080
00081     void Start();
00082     void Stop();
00083
00084     #if KERNEL_USE_THREADNAME
00085     void SetName(const char* szName_) { m_szName = szName_; }
00086     const char* GetName() { return m_szName; }
00087     #endif
00088
00089     ThreadList* GetOwner(void) { return m_pclOwner; }
00090     ThreadList* GetCurrent(void) { return m_pclCurrent; }
00091     PORT_PRIO_TYPE GetPriority(void) { return
00092     m_uXPriority; }
00093     PORT_PRIO_TYPE GetCurPriority(void) { return
00094     m_uXCurPriority; }
00095
00096     #if KERNEL_USE_QUANTUM
00097     void SetQuantum(uint16_t u16Quantum_) { m_u16Quantum = u16Quantum_; }
00098     uint16_t GetQuantum(void) { return m_u16Quantum; }
00099     #endif
00100
00101     void SetCurrent(ThreadList* pclNewList_) { m_pclCurrent = pclNewList_; }
00102
00103     void SetOwner(ThreadList* pclNewList_) { m_pclOwner = pclNewList_; }
00104     void SetPriority(PORT_PRIO_TYPE uXPriority_);
00105
00106     void InheritPriority(PORT_PRIO_TYPE uXPriority_);
00107
00108     #if KERNEL_USE_DYNAMIC_THREADS
00109     void Exit();
00110     #endif
00111
00112     #if KERNEL_USE_SLEEP
00113     static void Sleep(uint32_t u32TimeMs_);
00114     static void USleep(uint32_t u32TimeUs_);
00115     #endif
00116
00117     static void Yield(void);
00118
00119     void SetID(uint8_t u8ID_) { m_u8ThreadID = u8ID_; }
00120     uint8_t GetID() { return m_u8ThreadID; }
00121     uint16_t GetStackSlack();

```

```

00318
00319 #if KERNEL_USE_EVENTFLAG
00320
00327     uint16_t GetEventFlagMask() { return m_ul6FlagMask; }
00332     void SetEventFlagMask(uint16_t ul6Mask_) { m_ul6FlagMask = ul6Mask_; }
00338     void SetEventFlagMode(EventFlagOperation_t eMode_) {
00343         m_eFlagMode = eMode_; }
00344 #endif
00345
00346 #if KERNEL_USE_TIMEOUTS || KERNEL_USE_SLEEP
00347     Timer* GetTimer();
00351 #endif
00352 #if KERNEL_USE_TIMEOUTS
00353     void SetExpired(bool bExpired_);
00362     bool GetExpired();
00370 #endif
00371
00372 #if KERNEL_USE_IDLE_FUNC
00373     void InitIdle();
00379 #endif
00380
00381 #if KERNEL_USE_EXTENDED_CONTEXT
00382     void* GetExtendedContext() { return m_pvExtendedContext; }
00391     void SetExtendedContext(void* pvData_) {
00404         m_pvExtendedContext = pvData_; }
00405 #endif
00406
00413     ThreadState_t GetState() { return m_eState; }
00421     void SetState(ThreadState_t eState_) { m_eState = eState_; }
00422
00427     K_WORD* GetStack() { return m_pwStack; }
00428
00433     uint16_t GetStackSize() { return m_ul6StackSize; }
00434
00435     friend class ThreadPort;
00436 private:
00444     static void ContextSwitchSWI(void);
00445
00451     void SetPriorityBase(PORT_PRIO_TYPE uXPriority_);
00452
00454     K_WORD* m_pwStackTop;
00455
00457     K_WORD* m_pwStack;
00458
00460     uint8_t m_u8ThreadID;
00461
00463     PORT_PRIO_TYPE m_uXPriority;
00464
00466     PORT_PRIO_TYPE m_uXCurPriority;
00467
00469     ThreadState_t m_eState;
00470
00471 #if KERNEL_USE_EXTENDED_CONTEXT
00472     void* m_pvExtendedContext;
00474 #endif
00475
00476 #if KERNEL_USE_THREADNAME
00477     const char* m_szName;
00479 #endif
00480
00482     uint16_t m_ul6StackSize;
00483
00485     ThreadList* m_pclCurrent;
00486
00488     ThreadList* m_pclOwner;
00489
00491     ThreadEntry_t m_pfEntryPoint;
00492
00494     void* m_pvArg;
00495
00496 #if KERNEL_USE_QUANTUM
00497     uint16_t m_ul6Quantum;
00499 #endif
00500
00501 #if KERNEL_USE_EVENTFLAG
00502     uint16_t m_ul6FlagMask;
00504
00506     EventFlagOperation_t m_eFlagMode;
00507 #endif

```

```

00508
00509 #if KERNEL_USE_TIMEOUTS || KERNEL_USE_SLEEP
00510     Timer m_clTimer;
00512 #endif
00513
00514 #if KERNEL_USE_TIMEOUTS
00515     bool m_bExpired;
00517 #endif
00518 };
00519
00520 #if KERNEL_USE_IDLE_FUNC
00521 //-----
00533 typedef struct {
00534     LinkListNode* next;
00535     LinkListNode* prev;
00536
00538     K_WORD* m_pwStackTop;
00539
00541     K_WORD* m_pwStack;
00542
00544     uint8_t m_u8ThreadID;
00545
00547     PORT_PRIO_TYPE m_uXPriority;
00548
00550     PORT_PRIO_TYPE m_uXCurPriority;
00551
00553     ThreadState_t m_eState;
00554
00555 #if KERNEL_USE_EXTENDED_CONTEXT
00556     void* m_pvExtendedContext;
00558 #endif
00559
00560 #if KERNEL_USE_THREADNAME
00561     const char* m_szName;
00563 #endif
00564
00565 } FakeThread_t;
00566 #endif
00567
00568 #endif

```

19.99 /media/usb/project/github/Mark3/kernel/public/threadlist.h File Reference

[Thread](#) linked-list declarations.

```

#include "kerneltypes.h"
#include "priomap.h"
#include "ll.h"

```

Classes

- class [ThreadList](#)

This class is used for building thread-management facilities, such as schedulers, and blocking objects.

19.99.1 Detailed Description

[Thread](#) linked-list declarations.

Definition in file [threadlist.h](#).

19.100 threadlist.h

```

00001 /*=====
00002
00003
00004
00005
00006

```



```

00007      |_____|      |_____|      |_____|      |_____|
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00022  #ifndef __THREADLIST_H__
00023  #define __THREADLIST_H__
00024
00025  #include "kerneltypes.h"
00026  #include "priomap.h"
00027  #include "ll.h"
00028
00029  class Thread;
00030
00035  class ThreadList : public CircularLinkedList
00036  {
00037  public:
00038      void* operator new(size_t sz, void* pv) { return (ThreadList*)pv; };
00044      ThreadList() : m_uXPriority(0), m_pclMap(NULL)
00045      { }
00046
00054      void SetPriority(PORT_PRIO_TYPE uXPriority_);
00055
00065      void SetMapPointer(PriorityMap* pclMap_);
00066
00074      void Add(LinkListNode* node_);
00075
00087      void Add(LinkListNode* node_, PriorityMap* pclMap_,
PORT_PRIO_TYPE uXPriority_);
00088
00097      void AddPriority(LinkListNode* node_);
00098
00106      void Remove(LinkListNode* node_);
00107
00115      Thread* HighestWaiter();
00116
00117  private:
00119      PORT_PRIO_TYPE m_uXPriority;
00120
00122      PriorityMap* m_pclMap;
00123  };
00124
00125  #endif

```

19.101 /media/usb/project/github/Mark3/kernel/public/timer.h File Reference

Timer object declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"

```

Classes

- class **Timer**
Kernel-managed software timers.

Macros

- #define **TIMERLIST_FLAG_ONE_SHOT** (0x01)
Timer is one-shot.
- #define **TIMERLIST_FLAG_ACTIVE** (0x02)
Timer is currently active.
- #define **TIMERLIST_FLAG_CALLBACK** (0x04)
Timer is pending a callback.
- #define **TIMERLIST_FLAG_EXPIRED** (0x08)


```

00028
00029 #if KERNEL_USE_TIMERS
00030 class Thread;
00031
00032 //-----
00033 #define TIMERLIST_FLAG_ONE_SHOT (0x01)
00034 #define TIMERLIST_FLAG_ACTIVE (0x02)
00035 #define TIMERLIST_FLAG_CALLBACK (0x04)
00036 #define TIMERLIST_FLAG_EXPIRED (0x08)
00037
00038 //-----
00039 #define TIMER_INVALID_COOKIE (0x3C)
00040 #define TIMER_INIT_COOKIE (0xC3)
00041
00042 //-----
00043 #define MAX_TIMER_TICKS (0x7FFFFFFF)
00044 #define TIMER_TICKS_INVALID (0x80000000)
00045 //-----
00046 #if KERNEL_TIMERS_TICKLESS
00047
00048 //-----
00049 /*
00050     Ugly macros to support a wide resolution of delays.
00051     Given a 16-bit timer @ 16MHz & 256 cycle prescaler, this gives u16...
00052     Max time, SECONDS_TO_TICKS: 68719s
00053     Max time, MSECONDS_TO_TICKS: 6871.9s
00054     Max time, USECONDS_TO_TICKS: 6.8719s
00055
00056     ...With a 16us tick resolution.
00057
00058     Depending on the system frequency and timer resolution, you may want to
00059     customize these values to suit your system more appropriately.
00060 */
00061 //-----
00062 #define SECONDS_TO_TICKS(x) (((uint32_t)x) * PORT_TIMER_FREQ)
00063 #define MSECONDS_TO_TICKS(x) (((((uint32_t)x) * (PORT_TIMER_FREQ / 100)) + 5) / 10))
00064 #define USECONDS_TO_TICKS(x) (((((uint32_t)x) * PORT_TIMER_FREQ) + 50000) / 100000))
00065
00066 //-----
00067 #define MIN_TICKS (3)
00068 //-----
00069
00070 #else
00071
00072 //-----
00073 // add time because we don't know how far in an epoch we are when a call is made.
00074 #define SECONDS_TO_TICKS(x) (((uint32_t)(x)*1000) + 1)
00075 #define MSECONDS_TO_TICKS(x) ((uint32_t)(x + 1))
00076 #define USECONDS_TO_TICKS(x) (((uint32_t)(x + 999)) / 1000)
00077
00078 //-----
00079 #define MIN_TICKS (1)
00080 //-----
00081
00082 #endif // KERNEL_TIMERS_TICKLESS
00083
00084 //-----
00095 typedef void (*TimerCallback_t)(Thread* pOwner_, void* pvData_);
00096
00097 //-----
00098 class TimerList;
00099 class TimerScheduler;
00100 class Quantum;
00101
00102 //-----
00112 class Timer : public LinkListNode
00113 {
00114 public:
00115     void* operator new(size_t sz, void* pv) { return (Timer*)pv; };
00122     Timer();
00123
00129     void Init();
00130
00142     void Start(bool bRepeat_, uint32_t u32IntervalMs_, TimerCallback_t pfCallback_,
void* pvData_);
00143
00157     void
00158     Start(bool bRepeat_, uint32_t u32IntervalMs_, uint32_t u32ToleranceMs_,
TimerCallback_t pfCallback_, void* pvData_);
00159
00168     void Start();
00169
00176     void Stop();
00177
00187     void SetFlags(uint8_t u8Flags_) { m_u8Flags = u8Flags_; }
00195     void SetCallback(TimerCallback_t pfCallback_) {
m_pfCallback = pfCallback_; }

```

```

00203     void SetData(void* pvData_) { m_pvData = pvData_; }
00212     void SetOwner(Thread* pclOwner_) { m_pclOwner = pclOwner_; }
00220     void SetIntervalTicks(uint32_t u32Ticks_);
00221
00229     void SetIntervalSeconds(uint32_t u32Seconds_);
00230
00238     uint32_t GetInterval() { return m_u32Interval; }
00246     void SetIntervalMSeconds(uint32_t u32MSeconds_);
00247
00255     void SetIntervalUSeconds(uint32_t u32USeconds_);
00256
00265     void SetTolerance(uint32_t u32Ticks_);
00266
00267 private:
00268     friend class TimerList;
00269
00270 #if KERNEL_EXTRA_CHECKS
00271
00274     void SetInitialized() { m_u8Initialized = TIMER_INIT_COOKIE; }
00275
00280     bool IsInitialized(void) { return (m_u8Initialized == TIMER_INIT_COOKIE); }
00281
00283     uint8_t m_u8Initialized;
00284 #endif
00285
00287     uint8_t m_u8Flags;
00288
00290     TimerCallback_t m_pfCallback;
00291
00293     uint32_t m_u32Interval;
00294
00296     uint32_t m_u32TimeLeft;
00297
00299     uint32_t m_u32TimerTolerance;
00300
00302     Thread* m_pclOwner;
00303
00305     void* m_pvData;
00306 };
00307
00308 #endif // KERNEL_USE_TIMERS
00309
00310 #endif

```

19.103 /media/usb/project/github/Mark3/kernel/public/timerlist.h File Reference

[Timer](#) list declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "mutex.h"

```

Classes

- class [TimerList](#)

[TimerList](#) class - a doubly-linked-list of timer objects.

19.103.1 Detailed Description

[Timer](#) list declarations.

These classes implements a linked list of timer objects attached to the global kernel timer scheduler.

Definition in file [timerlist.h](#).

19.104 timerlist.h

```

00001 /*=====

```

```

00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00024 #ifndef __TIMERLIST_H__
00025 #define __TIMERLIST_H__
00026
00027 #include "kerneltypes.h"
00028 #include "mark3cfg.h"
00029
00030 #include "mutex.h"
00031 #if KERNEL_USE_TIMERS
00032
00033 class Timer;
00034
00035 //-----
00039 class TimerList : public DoubleLinkedList
00040 {
00041 public:
00048     void Init();
00049
00057     void Add(Timer* pclListNode_);
00058
00066     void Remove(Timer* pclLinkedListNode_);
00067
00074     void Process();
00075
00076 private:
00078     uint32_t m_u32NextWakeup;
00079
00081     bool m_bTimerActive;
00082
00083 #if KERNEL_TIMERS_THREADED
00084     Mutex m_clMutex;
00086 #endif
00087
00088 };
00089
00090 #endif // KERNEL_USE_TIMERS
00091
00092 #endif

```

19.105 /media/usb/project/github/Mark3/kernel/public/timerscheduler.h File Reference

[Timer](#) scheduler declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "timer.h"
#include "timerlist.h"

```

Classes

- class [TimerScheduler](#)

"Static" Class used to interface a global [TimerList](#) with the rest of the kernel.

19.105.1 Detailed Description

[Timer](#) scheduler declarations.

Definition in file [timerscheduler.h](#).

[illegible]

19.109 /media/usb/project/github/Mark3/kernel/quantum.cpp File Reference

Thread Quantum Implementation for Round-Robin Scheduling.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "timerlist.h"
#include "quantum.h"
#include "kernelaware.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

19.109.1 Detailed Description

Thread Quantum Implementation for Round-Robin Scheduling.

Definition in file [quantum.cpp](#).

19.110 quantum.cpp

00001 / *-----
00002
00003
00004
00005
00006
00007

```

00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "mark3cfg.h"
00024
00025 #include "thread.h"
00026 #include "timerlist.h"
00027 #include "quantum.h"
00028 #include "kernelaware.h"
00029
00030 #define _CAN_HAS_DEBUG
00031 //--[Autogenerated - Do Not Modify]-----
00032 #include "dbg_file_list.h"
00033 #include "buffalogger.h"
00034 #if defined(DBG_FILE)
00035 #error "Debug logging file token already defined! Bailing."
00036 #else
00037 #define DBG_FILE _DBG__KERNEL_QUANTUM_CPP
00038 #endif
00039 //--[End Autogenerated content]-----
00040 #include "kerneldebug.h"
00041
00042 #if KERNEL_USE_QUANTUM
00043
00044 //-----
00045 static volatile bool bAddQuantumTimer; // Indicates that a timer add is pending
00046
00047 //-----
00048 #if KERNEL_TIMERS_THREADED
00049 Thread* Quantum::m_pclTimerThread;
00050 #endif // KERNEL_TIMERS_THREADED
00051
00052 Timer Quantum::m_clQuantumTimer; // The global timernodelist_t object
00053 bool Quantum::m_bActive;
00054 bool Quantum::m_bInTimer;
00055 //-----
00066 static void QuantumCallback(Thread* pclThread_, void* /*pvData_*/)
00067 {
00068     // Validate thread pointer, check that source/destination match (it's
00069     // in its real priority list).
00070     if (pclThread_>GetCurrent()->GetHead() != pclThread_>
        GetCurrent()->GetTail()) {
00071         bAddQuantumTimer = true;
00072         pclThread_>GetCurrent()->PivotForward();
00073     }
00074 }
00075
00076 //-----
00077 void Quantum::SetTimer(Thread* pclThread_)
00078 {
00079     m_clQuantumTimer.SetIntervalMSeconds(pclThread_>
        GetQuantum());
00080     m_clQuantumTimer.SetFlags(TIMERLIST_FLAG_ONE_SHOT);
00081     m_clQuantumTimer.SetData(NULL);
00082     m_clQuantumTimer.SetCallback((TimerCallback_t)QuantumCallback);
00083     m_clQuantumTimer.SetOwner(pclThread_);
00084 }
00085
00086 //-----
00087 void Quantum::AddThread(Thread* pclThread_)
00088 {
00089     if (m_bActive
00090 #if KERNEL_USE_IDLE_FUNC
00091         || (pclThread_ == Kernel::GetIdleThread())
00092 #endif
00093     ) {
00094         return;
00095     }
00096
00097     // If this is called from the timer callback, queue a timer add...
00098     if (m_bInTimer) {
00099         bAddQuantumTimer = true;
00100         return;
00101     }
00102
00103     // If this isn't the only thread in the list.
00104     if (pclThread_>GetCurrent()->GetHead() != pclThread_>
        GetCurrent()->GetTail()) {
00105 #if KERNEL_EXTRA_CHECKS
00106         m_clQuantumTimer.Init();
00107 #endif
00108         Quantum::SetTimer(pclThread_);
00109         TimerScheduler::Add(&m_clQuantumTimer);

```



```

00110         m_bActive = true;
00111     }
00112 }
00113
00114 //-----
00115 void Quantum::RemoveThread(void)
00116 {
00117     if (!m_bActive) {
00118         return;
00119     }
00120
00121     // Cancel the current timer
00122     TimerScheduler::Remove(&m_clQuantumTimer);
00123     m_bActive = false;
00124 }
00125
00126 //-----
00127 void Quantum::UpdateTimer(void)
00128 {
00129     // If we have to re-add the quantum timer (more than 2 threads at the
00130     // high-priority level...)
00131     if (bAddQuantumTimer) {
00132         // Trigger a thread yield - this will also re-schedule the
00133         // thread *and* reset the round-robin scheduler.
00134         Thread::Yield();
00135         bAddQuantumTimer = false;
00136     }
00137 }
00138
00139 //-----
00140 #if KERNEL_TIMERS_THREADED
00141 void Quantum::SetTimerThread(Thread* pclThread_)
00142 {
00143     m_pclTimerThread = pclThread_;
00144 }
00145
00146 Thread* Quantum::GetTimerThread()
00147 {
00148     return m_pclTimerThread;
00149 }
00150
00151 #endif // KERNEL_TIMERS_THREADED
00152
00153 #endif // KERNEL_USE_QUANTUM

```

19.111 /media/usb/project/github/Mark3/kernel/scheduler.cpp File Reference

Strict-Priority + Round-Robin thread scheduler implementation.

```

#include "kerneltypes.h"
#include "ll.h"
#include "scheduler.h"
#include "thread.h"
#include "threadport.h"
#include "kernel.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

Variables

- volatile Thread * g_pclNext
Pointer to the currently-chosen next-running thread.
- Thread * g_pclCurrent
Pointer to the currently-running thread.

19.111.1 Detailed Description

Strict-Priority + Round-Robin thread scheduler implementation.

Definition in file [scheduler.cpp](#).

19.112 scheduler.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "ll.h"
00024 #include "scheduler.h"
00025 #include "thread.h"
00026 #include "threadport.h"
00027 #include "kernel.h"
00028
00029 #define _CAN_HAS_DEBUG
00030 //--[Autogenerated - Do Not Modify]-----
00031 #include "dbg_file_list.h"
00032 #include "buffalogger.h"
00033 #if defined(DBG_FILE)
00034 #error "Debug logging file token already defined! Bailing."
00035 #else
00036 #define DBG_FILE _DBG__KERNEL_SCHEDULER_CPP
00037 #endif
00038 //--[End Autogenerated content]-----
00039
00040 #include "kerneldebug.h"
00041 volatile Thread* g_pclNext;
00042 Thread* g_pclCurrent;
00043
00044 //-----
00045 bool Scheduler::m_bEnabled;
00046 bool Scheduler::m_bQueuedSchedule;
00047
00048 //-----
00049 ThreadList Scheduler::m_clStopList;
00050 ThreadList Scheduler::m_aclPriorities[
    KERNEL_NUM_PRIORITIES];
00051 PriorityMap Scheduler::m_clPrioMap;
00052
00053 //-----
00054 void Scheduler::Init()
00055 {
00056     for (int i = 0; i < KERNEL_NUM_PRIORITIES; i++) {
00057         m_aclPriorities[i].SetPriority(i);
00058         m_aclPriorities[i].SetMapPointer(&
00059             m_clPrioMap);
00059     }
00060 }
00061
00062 //-----
00063 void Scheduler::Schedule()
00064 {
00065     PORT_PRIO_TYPE uXPrio;
00066
00067     uXPrio = m_clPrioMap.HighestPriority();
00068
00069 #if KERNEL_USE_IDLE_FUNC
00070     if (uXPrio == 0) {
00071         // There aren't any active threads at all - set g_pclNext to IDLE
00072         g_pclNext = Kernel::GetIdleThread();
00073     } else
00074 #endif
00075     {
00076         if (uXPrio == 0) {
00077             Kernel::Panic(PANIC_NO_READY_THREADS);
00078         }
00079         // Priorities are one-indexed
00080         uXPrio--;
00081
00082         // Get the thread node at this priority.
00083         g_pclNext = (Thread*)(m_aclPriorities[uXPrio].GetHead());
00084     }
00085     KERNEL_TRACE_1("Next Thread: %d\n", (uint16_t)((Thread*)g_pclNext)->GetID());

```

```

00086 }
00087
00088 //-----
00089 void Scheduler::Add(Thread* pClThread_)
00090 {
00091     m_aclPriorities[pClThread_->GetPriority()].Add(pClThread_);
00092 }
00093
00094 //-----
00095 void Scheduler::Remove(Thread* pClThread_)
00096 {
00097     m_aclPriorities[pClThread_->GetPriority()].Remove(pClThread_);
00098 }
00099
00100 //-----
00101 bool Scheduler::SetScheduler(bool bEnable_)
00102 {
00103     bool bRet;
00104     CS_ENTER();
00105     bRet = m_bEnabled;
00106     m_bEnabled = bEnable_;
00107     // If there was a queued scheduler event, dequeue and trigger an
00108     // immediate Yield
00109     if (m_bEnabled && m_bQueuedSchedule) {
00110         m_bQueuedSchedule = false;
00111         Thread::Yield();
00112     }
00113     CS_EXIT();
00114     return bRet;
00115 }

```

19.113 /media/usb/project/github/Mark3/kernel/thread.cpp File Reference

Platform-Independent thread class Definition.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "scheduler.h"
#include "kernelswi.h"
#include "timerlist.h"
#include "ksemaphore.h"
#include "quantum.h"
#include "kernel.h"
#include "priomap.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

19.113.1 Detailed Description

Platform-Independent thread class Definition.

Definition in file [thread.cpp](#).

19.114 thread.cpp

[illegible]

```

00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "mark3cfg.h"
00024
00025 #include "thread.h"
00026 #include "scheduler.h"
00027 #include "kernelswi.h"
00028 #include "timerlist.h"
00029 #include "ksemaphore.h"
00030 #include "quantum.h"
00031 #include "kernel.h"
00032 #include "priomap.h"
00033
00034 #define _CAN_HAS_DEBUG
00035 //--[Autogenerated - Do Not Modify]-----
00036 #include "dbg_file_list.h"
00037 #include "buffalogger.h"
00038 #if defined(DBG_FILE)
00039 #error "Debug logging file token already defined! Bailing."
00040 #else
00041 #define DBG_FILE _DBG__KERNEL_THREAD_CPP
00042 #endif
00043 //--[End Autogenerated content]-----
00044
00045 #include "kerneldebug.h"
00046 //-----
00047 Thread::~Thread()
00048 {
00049     // On destruction of a thread located on a stack,
00050     // ensure that the thread is either stopped, or exited.
00051     // If the thread is stopped, move it to the exit state.
00052     // If not in the exit state, kernel panic -- it's catastrophic to have
00053     // running threads on stack suddenly disappear.
00054     if (m_eState == THREAD_STATE_STOP) {
00055         CS_ENTER();
00056         m_pclCurrent->Remove(this);
00057         m_pclCurrent = 0;
00058         m_pclOwner = 0;
00059         m_eState = THREAD_STATE_EXIT;
00060         CS_EXIT();
00061     } else if (m_eState != THREAD_STATE_EXIT) {
00062         #if KERNEL_AWARE_SIMULATION
00063             KernelAware::Trace(0, 0, m_u8ThreadID,
00064                               m_eState);
00065         #endif
00066         Kernel::Panic(PANIC_RUNNING_THREAD_DESCOPE);
00067     }
00068 }
00069 //-----
00070 void Thread::Init(
00071     K_WORD* pwStack_, uint16_t ul6StackSize_, PORT_PRIO_TYPE uXPriority_,
00072     ThreadEntry_t pfEntryPoint_, void* pvArg_)
00073 {
00074     static uint8_t u8ThreadID = 0;
00075     KERNEL_ASSERT(pwStack_);
00076     KERNEL_ASSERT(pfEntryPoint_);
00077     ClearNode();
00078     m_u8ThreadID = u8ThreadID++;
00079     #if KERNEL_USE_IDLE_FUNC
00080         if (u8ThreadID == 255) {
00081             u8ThreadID = 0;
00082         }
00083     #endif
00084     #if KERNEL_TRACE_1
00085         KERNEL_TRACE_1("Stack Size: %d", ul6StackSize_);
00086         KERNEL_TRACE_1("Thread Pri: %d", (uint8_t)uXPriority_);
00087         KERNEL_TRACE_1("Thread Id: %d", (uint16_t)m_u8ThreadID);
00088         KERNEL_TRACE_1("Entrypoint: %x", (uint16_t)pfEntryPoint_);
00089     #endif
00090     // Initialize the thread parameters to their initial values.
00091     m_pwStack = pwStack_;
00092     m_pwStackTop = TOP_OF_STACK(pwStack_, ul6StackSize_);
00093     m_ul6StackSize = ul6StackSize_;
00094     #if KERNEL_USE_QUANTUM
00095         m_ul6Quantum = THREAD_QUANTUM_DEFAULT;
00096     #endif
00097     m_uXPriority = uXPriority_;
00098     m_uXCurPriority = m_uXPriority;

```

```

00104     m_pfEntryPoint = pfEntryPoint_;
00105     m_pvArg        = pvArg_;
00106
00107     #if KERNEL_USE_THREADNAME
00108         m_szName = NULL;
00109     #endif
00110     #if KERNEL_USE_TIMERS
00111         m_clTimer.Init();
00112     #endif
00113
00114     // Call CPU-specific stack initialization
00115     ThreadPort::InitStack(this);
00116
00117     // Add to the global "stop" list.
00118     CS_ENTER();
00119     m_pclOwner = Scheduler::GetThreadList(
00120         m_uXPriority);
00121     m_pclCurrent = Scheduler::GetStopList();
00122     m_eState     = THREAD_STATE_STOP;
00123     m_pclCurrent->Add(this);
00124     CS_EXIT();
00125
00126     #if KERNEL_USE_THREAD_CALLOUTS
00127     ThreadCreateCallout_t pfCallout = Kernel::GetThreadCreateCallout();
00128     if (pfCallout != 0) {
00129         pfCallout(this);
00130     }
00131     #endif
00132 }
00133
00134 #if KERNEL_USE_AUTO_ALLOC
00135 //-----
00136 Thread* Thread::Init(uint16_t ul6StackSize_, PORT_PRIO_TYPE uXPriority_,
00137     ThreadEntry_t pfEntryPoint_, void* pvArg_)
00138 {
00139     Thread* pclNew = (Thread*)AutoAlloc::Allocate(sizeof(Thread));
00140     K_WORD* pwStack = (K_WORD*)AutoAlloc::Allocate(ul6StackSize_);
00141     pclNew->Init(pwStack, ul6StackSize_, uXPriority_, pfEntryPoint_, pvArg_);
00142     return pclNew;
00143 }
00144 #endif
00145 //-----
00146 void Thread::Start(void)
00147 {
00148     #if KERNEL_EXTRA_CHECKS
00149     KERNEL_ASSERT(IsInitialized());
00150     #endif
00151
00152     // Remove the thread from the scheduler's "stopped" list, and add it
00153     // to the scheduler's ready list at the proper priority.
00154     KERNEL_TRACE_1("Starting Thread %d", (uint16_t)m_u8ThreadID);
00155
00156     CS_ENTER();
00157     Scheduler::GetStopList()->Remove(this);
00158     Scheduler::Add(this);
00159     m_pclOwner = Scheduler::GetThreadList(
00160         m_uXPriority);
00161     m_pclCurrent = m_pclOwner;
00162     m_eState     = THREAD_STATE_READY;
00163
00164     #if KERNEL_USE_QUANTUM
00165     if (Kernel::IsStarted()) {
00166         if (GetCurPriority() >= Scheduler::GetCurrentThread()->
00167             GetCurPriority()) {
00168             // Deal with the thread Quantum
00169             #if KERNEL_TIMERS_THREADED
00170             if (Quantum::GetTimerThread() != this) {
00171                 Quantum::RemoveThread();
00172                 Quantum::AddThread(this);
00173             }
00174             #endif
00175         }
00176     }
00177     #endif
00178
00179     if (Kernel::IsStarted()) {
00180         if (GetCurPriority() >= Scheduler::GetCurrentThread()->
00181             GetCurPriority()) {
00182             Thread::Yield();
00183         }
00184     }
00185     CS_EXIT();
00186 }

```

```

00186
00187 //-----
00188 void Thread::Stop()
00189 {
00190     #if KERNEL_EXTRA_CHECKS
00191         KERNEL_ASSERT(IsInitialized());
00192     #endif
00193
00194     bool bReschedule = false;
00195     if (m_eState == THREAD_STATE_STOP) {
00196         return;
00197     }
00198
00199     CS_ENTER();
00200
00201     // If a thread is attempting to stop itself, ensure we call the scheduler
00202     if (this == Scheduler::GetCurrentThread()) {
00203         bReschedule = true;
00204     }
00205
00206     // Add this thread to the stop-list (removing it from active scheduling)
00207     // Remove the thread from scheduling
00208     if (m_eState == THREAD_STATE_READY) {
00209         Scheduler::Remove(this);
00210     } else if (m_eState == THREAD_STATE_BLOCKED) {
00211         m_pclCurrent->Remove(this);
00212     }
00213
00214     m_pclOwner = Scheduler::GetStopList();
00215     m_pclCurrent = m_pclOwner;
00216     m_pclOwner->Add(this);
00217     m_eState = THREAD_STATE_STOP;
00218
00219     #if KERNEL_USE_TIMERS
00220         // Just to be safe - attempt to remove the thread's timer
00221         // from the timer-scheduler (does no harm if it isn't
00222         // in the timer-list)
00223         TimerScheduler::Remove(&m_clTimer);
00224     #endif
00225
00226     CS_EXIT();
00227
00228     if (bReschedule) {
00229         Thread::Yield();
00230     }
00231 }
00232
00233 #if KERNEL_USE_DYNAMIC_THREADS
00234 //-----
00235 void Thread::Exit()
00236 {
00237     #if KERNEL_EXTRA_CHECKS
00238         KERNEL_ASSERT(IsInitialized());
00239     #endif
00240
00241     bool bReschedule = false;
00242
00243     KERNEL_TRACE_1("Exit Thread %d", m_u8ThreadID);
00244     if (m_eState == THREAD_STATE_EXIT) {
00245         return;
00246     }
00247
00248     CS_ENTER();
00249
00250     // If this thread is the actively-running thread, make sure we run the
00251     // scheduler again.
00252     if (this == Scheduler::GetCurrentThread()) {
00253         bReschedule = true;
00254     }
00255
00256     // Remove the thread from scheduling
00257     if (m_eState == THREAD_STATE_READY) {
00258         Scheduler::Remove(this);
00259     } else if ((m_eState == THREAD_STATE_BLOCKED) || (m_eState == THREAD_STATE_STOP)) {
00260         m_pclCurrent->Remove(this);
00261     }
00262
00263     m_pclCurrent = 0;
00264     m_pclOwner = 0;
00265     m_eState = THREAD_STATE_EXIT;
00266
00267     // We've removed the thread from scheduling, but interrupts might
00268     // trigger checks against this thread's currently priority before
00269     // we get around to scheduling new threads. As a result, set the
00270     // priority to idle to ensure that we always wind up scheduling
00271     // new threads.
00272     m_uXCurPriority = 0;
00273     m_uXPriority = 0;

```

```

00273
00274 #if KERNEL_USE_TIMERS
00275     // Just to be safe - attempt to remove the thread's timer
00276     // from the timer-scheduler (does no harm if it isn't
00277     // in the timer-list)
00278     TimerScheduler::Remove(&m_clTimer);
00279 #endif
00280     CS_EXIT();
00281
00282 #if KERNEL_USE_THREAD_CALLOUTS
00283     ThreadExitCallout_t pfCallout = Kernel::GetThreadExitCallout();
00284     if (pfCallout != 0) {
00285         pfCallout(this);
00286     }
00287 #endif
00288
00289     if (bReschedule) {
00290         // Choose a new "next" thread if we must
00291         Thread::Yield();
00292     }
00293 }
00294 #endif
00295
00296 #if KERNEL_USE_SLEEP
00297 //-----
00298 static void ThreadSleepCallback(Thread* /*pclOwner*/, void* pvData_)
00299 {
00300     Semaphore* pclSemaphore = static_cast<Semaphore*>(pvData_);
00301     // Post the semaphore, which will wake the sleeping thread.
00302     pclSemaphore->Post();
00303 }
00304 //-----
00305
00306 void Thread::Sleep(uint32_t u32TimeMs_)
00307 {
00308     Semaphore clSemaphore;
00309     Timer* pclTimer = g_pclCurrent->GetTimer();
00310
00311     // Create a semaphore that this thread will block on
00312     clSemaphore.Init(0, 1);
00313
00314     // Create a one-shot timer that will call a callback that posts the
00315     // semaphore, waking our thread.
00316     pclTimer->Init();
00317     pclTimer->SetIntervalMSeconds(u32TimeMs_);
00318     pclTimer->SetCallback(ThreadSleepCallback);
00319     pclTimer->SetData((void*)&clSemaphore);
00320     pclTimer->SetFlags(TIMERLIST_FLAG_ONE_SHOT);
00321
00322     // Add the new timer to the timer scheduler, and block the thread
00323     TimerScheduler::Add(pclTimer);
00324     clSemaphore.Pend();
00325 }
00326
00327 //-----
00328 void Thread::USleep(uint32_t u32TimeUs_)
00329 {
00330     Semaphore clSemaphore;
00331     Timer* pclTimer = g_pclCurrent->GetTimer();
00332
00333     // Create a semaphore that this thread will block on
00334     clSemaphore.Init(0, 1);
00335
00336     // Create a one-shot timer that will call a callback that posts the
00337     // semaphore, waking our thread.
00338     pclTimer->Init();
00339     pclTimer->SetIntervalUSeconds(u32TimeUs_);
00340     pclTimer->SetCallback(ThreadSleepCallback);
00341     pclTimer->SetData((void*)&clSemaphore);
00342     pclTimer->SetFlags(TIMERLIST_FLAG_ONE_SHOT);
00343
00344     // Add the new timer to the timer scheduler, and block the thread
00345     TimerScheduler::Add(pclTimer);
00346     clSemaphore.Pend();
00347 }
00348 #endif // KERNEL_USE_SLEEP
00349
00350 //-----
00351 uint16_t Thread::GetStackSlack()
00352 {
00353     #if KERNEL_EXTRA_CHECKS
00354         KERNEL_ASSERT(IsInitialized());
00355     #endif
00356
00357     K_ADDR wTop = (K_ADDR)m_u16StackSize - 1;
00358     K_ADDR wBottom = (K_ADDR)0;
00359     K_ADDR wMid = ((wTop + wBottom) + 1) / 2;

```

```

00361
00362     CS_ENTER();
00363
00364     // Logarithmic bisection - find the point where the contents of the
00365     // stack go from 0xFF's to non 0xFF. Not Definitive, but accurate enough
00366     while ((wTop - wBottom) > 1) {
00367 #if STACK_GROWS_DOWN
00368         if (m_pwStack[wMid] != (K_WORD) (-1))
00369 #else
00370         if (m_pwStack[wMid] == (K_WORD) (-1))
00371 #endif
00372         {
00373             wTop = wMid;
00374         } else {
00375             wBottom = wMid;
00376         }
00377         wMid = (wTop + wBottom + 1) / 2;
00378     }
00379
00380     CS_EXIT();
00381
00382     return wMid;
00383 }
00384
00385
00386 //-----
00387 void Thread::Yield()
00388 {
00389
00390     CS_ENTER();
00391     // Run the scheduler
00392     if (Scheduler::IsEnabled()) {
00393         Scheduler::Schedule();
00394
00395         // Only switch contexts if the new task is different than the old task
00396         if (Scheduler::GetCurrentThread() !=
00397             Scheduler::GetNextThread()) {
00398 #if KERNEL_USE_QUANTUM
00399 #if KERNEL_TIMERS_THREADED
00400             if (Quantum::GetTimerThread() != g_pclNext) {
00401 #endif
00402                 // new thread scheduled. Stop current quantum timer (if it exists),
00403                 // and restart it for the new thread (if required).
00404                 Quantum::RemoveThread();
00405                 Quantum::AddThread((Thread*)g_pclNext);
00406 #if KERNEL_TIMERS_THREADED
00407             }
00408 #endif
00409 #endif
00410             Thread::ContextSwitchSWI();
00411         } else {
00412             Scheduler::QueueScheduler();
00413         }
00414     }
00415     CS_EXIT();
00416 }
00417
00418 //-----
00419 void Thread::SetPriorityBase(PORT_PRIO_TYPE /*uXPriority_*/)
00420 {
00421 #if KERNEL_EXTRA_CHECKS
00422     KERNEL_ASSERT(IsInitialized());
00423 #endif
00424
00425     GetCurrent()->Remove(this);
00426
00427     SetCurrent(Scheduler::GetThreadList(
00428         m_uXPriority));
00429
00430     GetCurrent()->Add(this);
00431 }
00432
00433 //-----
00434 void Thread::SetPriority(PORT_PRIO_TYPE uXPriority_)
00435 {
00436 #if KERNEL_EXTRA_CHECKS
00437     KERNEL_ASSERT(IsInitialized());
00438 #endif
00439
00440     bool bSchedule = false;
00441
00442     CS_ENTER();
00443     // If this is the currently running thread, it's a good idea to reschedule
00444     // Or, if the new priority is a higher priority than the current thread's.
00445     if ((g_pclCurrent == this) || (uXPriority_ > g_pclCurrent->
00446         GetPriority())) {
00447         bSchedule = true;
00448     }

```



```

00446     }
00447     Scheduler::Remove(this);
00448     CS_EXIT();
00449
00450     m_uXCurPriority = uXPriority_;
00451     m_uXPriority     = uXPriority_;
00452
00453     CS_ENTER();
00454     Scheduler::Add(this);
00455     CS_EXIT();
00456
00457     if (bSchedule) {
00458         if (Scheduler::IsEnabled()) {
00459             CS_ENTER();
00460             Scheduler::Schedule();
00461 #if KERNEL_USE_QUANTUM
00462 #if KERNEL_TIMERS_THREADED
00463             if (Quantum::GetTimerThread() != g_pclNext) {
00464 #endif
00465                 // new thread scheduled. Stop current quantum timer (if it exists),
00466                 // and restart it for the new thread (if required).
00467                 Quantum::RemoveThread();
00468                 Quantum::AddThread((Thread*)g_pclNext);
00469 #if KERNEL_TIMERS_THREADED
00470             }
00471 #endif
00472 #endif
00473             CS_EXIT();
00474             Thread::ContextSwitchSWI();
00475         } else {
00476             Scheduler::QueueScheduler();
00477         }
00478     }
00479 }
00480
00481 //-----
00482 void Thread::InheritPriority(PORT_PRIO_TYPE uXPriority_)
00483 {
00484     #if KERNEL_EXTRA_CHECKS
00485         KERNEL_ASSERT(IsInitialized());
00486     #endif
00487
00488     SetOwner(Scheduler::GetThreadList(uXPriority_));
00489     m_uXCurPriority = uXPriority_;
00490 }
00491
00492 //-----
00493 void Thread::ContextSwitchSWI()
00494 {
00495     // Call the context switch interrupt if the scheduler is enabled.
00496     if (static_cast<int>(Scheduler::IsEnabled()) == 1) {
00497         KERNEL_TRACE_1("Context switch to Thread %d", (uint16_t)((
00498             Thread*)g_pclNext->GetID());
00499 #if KERNEL_USE_STACK_GUARD
00500 #if KERNEL_USE_IDLE_FUNC
00501         if ((g_pclCurrent != 0) && (g_pclCurrent->GetID() != 255)) {
00502 #endif
00503             if (g_pclCurrent->GetStackSlack() <= Kernel::GetStackGuardThreshold())
00504             {
00505 #if KERNEL_AWARE_SIMULATION
00506                 KernelAware::Trace(DBG_FILE, __LINE__,
00507                     g_pclCurrent->GetID(), g_pclCurrent->GetStackSlack());
00508 #endif
00509                 Kernel::Panic(PANIC_STACK_SLACK_VIOLATED);
00510             }
00511 #if KERNEL_USE_IDLE_FUNC
00512         }
00513 #endif
00514 #if KERNEL_USE_THREAD_CALLOUTS
00515         ThreadContextCallout_t pfCallout = Kernel::GetThreadContextSwitchCallout
00516         ();
00517         if (pfCallout != 0) {
00518             pfCallout(g_pclCurrent);
00519         }
00520 #endif
00521         KernelSWI::Trigger();
00522     }
00523 }
00524 //-----
00525 Timer* Thread::GetTimer()
00526 {
00527     #if KERNEL_EXTRA_CHECKS
00528         KERNEL_ASSERT(IsInitialized());
00529     #endif

```


Generated on Fri Jul 7 2017 01:40:39 for Mark3 Realtime Kernel by Doxygen


```

00024
00025 #include "timer.h"
00026 #include "timerlist.h"
00027 #include "timerscheduler.h"
00028 #include "kerneltimer.h"
00029 #include "threadport.h"
00030 #include "quantum.h"
00031
00032 #define _CAN_HAS_DEBUG
00033 //--[Autogenerated - Do Not Modify]-----
00034 #include "dbg_file_list.h"
00035 #include "buffalogger.h"
00036 #if defined(DBG_FILE)
00037 #error "Debug logging file token already defined! Bailing."
00038 #else
00039 #define DBG_FILE _DBG__KERNEL_TIMER_CPP
00040 #endif
00041 //--[End Autogenerated content]-----
00042
00043 #include "kerneldebug.h"
00044
00045 #if KERNEL_USE_TIMERS
00046
00047 //-----
00048 Timer::Timer()
00049 {
00050     #if KERNEL_EXTRA_CHECKS
00051         m_u8Initialized = TIMER_INVALID_COOKIE;
00052     #endif
00053     m_u8Flags = 0;
00054 }
00055
00056 //-----
00057 void Timer::Init()
00058 {
00059     #if KERNEL_EXTRA_CHECKS
00060         if (IsInitialized()) {
00061             KERNEL_ASSERT((m_u8Flags & TIMERLIST_FLAG_ACTIVE) == 0);
00062         }
00063     #endif
00064
00065     ClearNode();
00066     m_u32Interval = 0;
00067     m_u32TimerTolerance = 0;
00068     m_u32TimeLeft = 0;
00069     m_u8Flags = 0;
00070
00071     #if KERNEL_EXTRA_CHECKS
00072         SetInitialized();
00073     #endif
00074 }
00075
00076 //-----
00077 void Timer::Start(bool bRepeat_, uint32_t u32IntervalMs_,
00078     TimerCallback_t pfCallback_, void* pvData_)
00079 {
00080     #if KERNEL_EXTRA_CHECKS
00081         KERNEL_ASSERT(IsInitialized());
00082     #endif
00083
00084     if ((m_u8Flags & TIMERLIST_FLAG_ACTIVE) != 0) {
00085         return;
00086     }
00087
00088     SetIntervalMSeconds(u32IntervalMs_);
00089     m_u32TimerTolerance = 0;
00090     m_pfCallback = pfCallback_;
00091     m_pvData = pvData_;
00092
00093     if (!bRepeat_) {
00094         m_u8Flags = TIMERLIST_FLAG_ONE_SHOT;
00095     } else {
00096         m_u8Flags = 0;
00097     }
00098
00099     Start();
00100 }
00101
00102 //-----
00103 void Timer::Start(
00104     bool bRepeat_, uint32_t u32IntervalMs_, uint32_t u32ToleranceMs_,
00105     TimerCallback_t pfCallback_, void* pvData_)
00106 {
00107     #if KERNEL_EXTRA_CHECKS
00108         KERNEL_ASSERT(IsInitialized());
00109     #endif
00110 }

```

```

00109     if ((m_u8Flags & TIMERLIST_FLAG_ACTIVE) != 0) {
00110         return;
00111     }
00112
00113     m_u32TimerTolerance = MSECONDS_TO_TICKS(u32ToleranceMs_);
00114     Start(bRepeat_, u32IntervalMs_, pfCallback_, pvData_);
00115 }
00116
00117 //-----
00118 void Timer::Start()
00119 {
00120     #if KERNEL_EXTRA_CHECKS
00121         KERNEL_ASSERT(IsInitialized());
00122     #endif
00123
00124     if ((m_u8Flags & TIMERLIST_FLAG_ACTIVE) != 0) {
00125         return;
00126     }
00127
00128     m_pclOwner = Scheduler::GetCurrentThread();
00129     TimerScheduler::Add(this);
00130 }
00131
00132 //-----
00133 void Timer::Stop()
00134 {
00135     #if KERNEL_EXTRA_CHECKS
00136         KERNEL_ASSERT(IsInitialized());
00137     #endif
00138
00139     if ((m_u8Flags & TIMERLIST_FLAG_ACTIVE) == 0) {
00140         return;
00141     }
00142     TimerScheduler::Remove(this);
00143 }
00144
00145 //-----
00146 void Timer::SetIntervalTicks(uint32_t u32Ticks_)
00147 {
00148     #if KERNEL_EXTRA_CHECKS
00149         KERNEL_ASSERT(IsInitialized());
00150     #endif
00151
00152     m_u32Interval = u32Ticks_;
00153 }
00154
00155 //-----
00156 //-----
00157 void Timer::SetIntervalSeconds(uint32_t u32Seconds_)
00158 {
00159     #if KERNEL_EXTRA_CHECKS
00160         KERNEL_ASSERT(IsInitialized());
00161     #endif
00162
00163     m_u32Interval = SECONDS_TO_TICKS(u32Seconds_);
00164 }
00165
00166 //-----
00167 void Timer::SetIntervalMSeconds(uint32_t u32MSeconds_)
00168 {
00169     #if KERNEL_EXTRA_CHECKS
00170         KERNEL_ASSERT(IsInitialized());
00171     #endif
00172
00173     m_u32Interval = MSECONDS_TO_TICKS(u32MSeconds_);
00174 }
00175
00176 //-----
00177 void Timer::SetIntervalUSeconds(uint32_t u32USeconds_)
00178 {
00179     #if KERNEL_EXTRA_CHECKS
00180         KERNEL_ASSERT(IsInitialized());
00181     #endif
00182
00183     #if KERNEL_TIMERS_TICKLESS
00184         if (u32USeconds_ < KERNEL_TIMERS_MINIMUM_DELAY_US) {
00185             u32USeconds_ = KERNEL_TIMERS_MINIMUM_DELAY_US;
00186         }
00187     #endif
00188
00189     m_u32Interval = USECONDS_TO_TICKS(u32USeconds_);
00190 }
00191
00192 //-----
00193 void Timer::SetTolerance(uint32_t u32Ticks_)
00194 {
00195     #if KERNEL_EXTRA_CHECKS
00196

```

19.119 /media/usb/project/github/Mark3/kernel/timerlist.cpp File Reference

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "timerlist.h"
#include "kerneltimer.h"
#include "threadport.h"
#include "quantum.h"
#include "mutex.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

Definition in file [timerlist.cpp](#).

```

00001 /*
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00023 #include "kerneltypes.h"
00024 #include "mark3cfg.h"
00025
00026 #include "timerlist.h"
00027 #include "kerneltimer.h"
00028 #include "threadport.h"
00029 #include "quantum.h"
00030 #include "mutex.h"
00031
00032 #define _CAN_HAS_DEBUG
00033 //--[Autogenerated - Do Not Modify]-----
00034 #include "dbg_file_list.h"
00035 #include "buffallogger.h"
00036 #if defined(DBG_FILE)
00037 #error "Debug logging file token already defined! Bailing."
00038 #else
00039 #define DBG_FILE _DBG__KERNEL_TIMERLIST_CPP
00040 #endif
00041 //--[End Autogenerated content]-----
00042
00043 #include "kerneldebug.h"
00044
00045 #if KERNEL_USE_TIMERS
00046 //-----
00047 TimerList TimerScheduler::m_clTimerList;
00048

```

```

00049 #if KERNEL_TIMERS_THREADED
00050 # define TIMERLIST_LOCK() (m_clMutex.Claim())
00051 # define TIMERLIST_UNLOCK() (m_clMutex.Release())
00052 #else
00053 # define TIMERLIST_LOCK() CS_ENTER()
00054 # define TIMERLIST_UNLOCK() CS_EXIT()
00055 #endif
00056
00057 //-----
00058 void TimerList::Init(void)
00059 {
00060     m_bTimerActive = false;
00061     m_u32NextWakeup = 0;
00062     #if KERNEL_TIMERS_THREADED
00063         m_clMutex.Init();
00064     #endif
00065 }
00066
00067 //-----
00068 void TimerList::Add(Timer* pclListNode_)
00069 {
00070     #if KERNEL_TIMERS_TICKLESS
00071         bool bStart = false;
00072         int32_t lDelta;
00073     #endif
00074     TIMERLIST_LOCK();
00075     #if KERNEL_TIMERS_TICKLESS
00076     if (GetHead() == NULL) {
00077         bStart = true;
00078     }
00079     if (pclListNode_>m_u32Interval < PORT_MIN_TIMER_TICKS) {
00080         pclListNode_>m_u32Interval = PORT_MIN_TIMER_TICKS;
00081     }
00082     #endif
00083     pclListNode_>ClearNode();
00084     DoubleLinkedList::Add(pclListNode_);
00085     // Set the initial timer value
00086     pclListNode_>m_u32TimeLeft = pclListNode_>m_u32Interval;
00087     #if KERNEL_TIMERS_TICKLESS
00088     if (!bStart) {
00089         // If the new interval is less than the amount of time remaining...
00090         lDelta = (int32_t)((uint32_t)KernelTimer::TimeToExpiry() - pclListNode_>
00091             m_u32Interval);
00092         if (lDelta > 0) {
00093             // Set the new expiry time on the timer.
00094             m_u32NextWakeup = (uint32_t)
00095                 KernelTimer::SubtractExpiry((uint32_t)lDelta);
00096         }
00097     } else {
00098         m_u32NextWakeup = pclListNode_>m_u32Interval;
00099         KernelTimer::SetExpiry(m_u32NextWakeup);
00100         KernelTimer::Start();
00101     }
00102     #endif
00103     // Set the timer as active.
00104     pclListNode_>m_u8Flags |= TIMERLIST_FLAG_ACTIVE;
00105     TIMERLIST_UNLOCK();
00106 }
00107
00108 //-----
00109 void TimerList::Remove(Timer* pclLinkListNode_)
00110 {
00111     TIMERLIST_LOCK();
00112     DoubleLinkedList::Remove(pclLinkListNode_);
00113     pclLinkListNode_>m_u8Flags &= ~TIMERLIST_FLAG_ACTIVE;
00114     #if KERNEL_TIMERS_TICKLESS
00115     if (this->GetHead() == NULL) {
00116         KernelTimer::Stop();
00117     }
00118     #endif
00119     TIMERLIST_UNLOCK();
00120 }
00121
00122 //-----
00123 void TimerList::Process(void)
00124 {

```



```

00134 #if KERNEL_TIMERS_TICKLESS
00135     uint32_t u32NewExpiry;
00136     uint32_t u32Overtime;
00137     bool     bContinue;
00138 #endif
00139
00140     Timer* pclNode;
00141     Timer* pclPrev;
00142
00143     TIMERLIST_LOCK();
00144
00145 #if KERNEL_USE_QUANTUM
00146     Quantum::SetInTimer();
00147 #endif
00148 #if KERNEL_TIMERS_TICKLESS
00149 #if !KERNEL_TIMERS_THREADED
00150     // Clear the timer and its expiry time - keep it running though
00151     KernelTimer::ClearExpiry();
00152 #endif
00153     do {
00154 #endif
00155         pclNode = static_cast<Timer*>(GetHead());
00156         pclPrev = NULL;
00157
00158 #if KERNEL_TIMERS_TICKLESS
00159         bContinue = false;
00160         u32NewExpiry = MAX_TIMER_TICKS; // Used to indicate that no timers are pending
00161 #endif
00162
00163         // Subtract the elapsed time interval from each active timer.
00164         while (pclNode != 0) {
00165             // Active timers only...
00166             if ((pclNode->m_u8Flags & TIMERLIST_FLAG_ACTIVE) != 0) {
00167                 // Did the timer expire?
00168                 #if KERNEL_TIMERS_TICKLESS
00169                     if (pclNode->m_u32TimeLeft <= m_u32NextWakeup)
00170                 #else
00171                     pclNode->m_u32TimeLeft--;
00172                     if (0 == pclNode->m_u32TimeLeft)
00173                 #endif
00174                     {
00175                         // Yes - set the "callback" flag - we'll execute the callbacks later
00176                         pclNode->m_u8Flags |= TIMERLIST_FLAG_CALLBACK;
00177
00178                         if ((pclNode->m_u8Flags & TIMERLIST_FLAG_ONE_SHOT) != 0
00179 ) {
00180                             // If this was a one-shot timer, deactivate the timer.
00181                             pclNode->m_u8Flags |= TIMERLIST_FLAG_EXPIRED;
00182                             pclNode->m_u8Flags &= ~TIMERLIST_FLAG_ACTIVE;
00183                         } else {
00184                             // Reset the interval timer.
00185                             // I think we're good though...
00186                             pclNode->m_u32TimeLeft = pclNode->
m_u32Interval;
00187
00188 #if KERNEL_TIMERS_TICKLESS
00189                             // If the time remaining (plus the length of the tolerance interval)
00190                             // is less than the next expiry interval, set the next expiry interval.
00191                             uint32_t u32Tmp = pclNode->m_u32TimeLeft + pclNode->
m_u32TimerTolerance;
00192
00193                             if (u32Tmp < u32NewExpiry) {
00194                                 u32NewExpiry = u32Tmp;
00195                             }
00196 #endif
00197                         }
00198                     }
00199 #if KERNEL_TIMERS_TICKLESS
00200                     else {
00201                         // Not expiring, but determine how long to run the next timer interval for.
00202                         pclNode->m_u32TimeLeft -= m_u32NextWakeup;
00203                         if (pclNode->m_u32TimeLeft < u32NewExpiry) {
00204                             u32NewExpiry = pclNode->m_u32TimeLeft;
00205                         }
00206                     }
00207 #endif
00208                 }
00209                 pclNode = static_cast<Timer*>(pclNode->GetNext());
00210             }
00211
00212             // Process the expired timers callbacks.
00213             pclNode = static_cast<Timer*>(GetHead());
00214             while (pclNode != 0) {
00215                 pclPrev = pclNode;
00216                 pclNode = static_cast<Timer*>(pclNode->GetNext());
00217
00218                 // If the timer expired, run the callbacks now.

```

```

00219         if ((pclPrev->m_u8Flags & TIMERLIST_FLAG_CALLBACK) != 0) {
00220             bool bRemove = false;
00221             // If this was a one-shot timer, tag it for removal
00222             if ((pclPrev->m_u8Flags & TIMERLIST_FLAG_ONE_SHOT) != 0) {
00223                 bRemove = true;
00224             }
00225             // Run the callback. these callbacks must be very fast...
00226             pclPrev->m_pfCallback(pclPrev->m_pclOwner, pclPrev->
m_pvData);
00228             pclPrev->m_u8Flags &= ~TIMERLIST_FLAG_CALLBACK;
00229
00230             // Remove one-shot-timers
00231             if (bRemove) {
00232                 Remove(pclPrev);
00233             }
00234         }
00235     }
00236
00237 #if KERNEL_TIMERS_TICKLESS
00238     // Check to see how much time has elapsed since the time we
00239     // acknowledged the interrupt...
00240     u32Overtime = (uint32_t)KernelTimer::GetOvertime();
00241
00242     if (u32Overtime >= u32NewExpiry) {
00243         m_u32NextWakeup = u32Overtime;
00244         bContinue = true;
00245     }
00246
00247     // If it's taken longer to go through this loop than would take us to
00248     // the next expiry, re-run the timing loop
00249 } while (bContinue);
00250
00251 // This timer elapsed, but there's nothing more to do...
00252 // Turn the timer off.
00253 if (u32NewExpiry >= MAX_TIMER_TICKS) {
00254     KernelTimer::Stop();
00255 } else {
00256     // Update the timer with the new "Next Wakeup" value, plus whatever
00257     // overtime has accumulated since the last time we called this handler
00258     m_u32NextWakeup = (uint32_t)KernelTimer::SetExpiry(
u32NewExpiry + u32Overtime);
00261 }
00262 #endif
00263 #if KERNEL_USE_QUANTUM
00264     Quantum::ClearInTimer();
00265 #endif
00266     TIMERLIST_UNLOCK();
00267 }
00268 }
00269
00270 #endif // KERNEL_USE_TIMERS

```

19.121 /media/usb/project/github/Mark3/kernel/tracebuffer.cpp File Reference

[Kernel](#) trace buffer class definition.

```

#include "kerneltypes.h"
#include "tracebuffer.h"
#include "mark3cfg.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

19.121.1 Detailed Description

[Kernel](#) trace buffer class definition.

Definition in file [tracebuffer.cpp](#).

19.122 tracebuffer.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00019 #include "kerneltypes.h"
00020 #include "tracebuffer.h"
00021 #include "mark3cfg.h"
00022
00023 #define _CAN_HAS_DEBUG
00024 //--[Autogenerated - Do Not Modify]-----
00025 #include "dbg_file_list.h"
00026 #include "buffallogger.h"
00027 #if defined(DBG_FILE)
00028 #error "Debug logging file token already defined! Bailing."
00029 #else
00030 #define DBG_FILE _DBG__KERNEL_TRACEBUFFER_CPP
00031 #endif
00032
00033 #include "kerneldebug.h"
00034
00035 //--[End Autogenerated content]-----
00036
00037 #if KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION
00038 //--
00039 TraceBufferCallback_t TraceBuffer::m_pfCallback;
00040 uint16_t TraceBuffer::m_ul6Index;
00041 uint16_t TraceBuffer::m_ul6SyncNumber;
00042 uint16_t TraceBuffer::m_aul6Buffer[(TRACE_BUFFER_SIZE / sizeof(uint16_t))];
00043
00044 //--
00045 void TraceBuffer::Init()
00046 {
00047 }
00048
00049 //--
00050 void TraceBuffer::Write(uint16_t* pul6Data_, uint16_t ul6Size_)
00051 {
00052     // Pipe the data directly to the circular buffer
00053     uint16_t ul6Start;
00054
00055     // Update the circular buffer index in a critical section. The
00056     // rest of the operations can take place in any context.
00057     CS_ENTER();
00058     uint16_t ul6NextIndex;
00059     ul6Start = m_ul6Index;
00060     ul6NextIndex = m_ul6Index + ul6Size_;
00061     if (ul6NextIndex >= (sizeof(m_aul6Buffer) / sizeof(uint16_t))) {
00062         ul6NextIndex -= (sizeof(m_aul6Buffer) / sizeof(uint16_t));
00063     }
00064     m_ul6Index = ul6NextIndex;
00065     CS_EXIT();
00066
00067     // Write the data into the circular buffer.
00068     uint16_t i;
00069     bool bCallback = false;
00070     bool bPingPong = false;
00071     for (i = 0; i < ul6Size_; i++) {
00072         m_aul6Buffer[ul6Start++] = pul6Data_[i];
00073         if (ul6Start >= (sizeof(m_aul6Buffer) / sizeof(uint16_t))) {
00074             ul6Start = 0;
00075             bCallback = true;
00076         } else if (ul6Start == ((sizeof(m_aul6Buffer) / sizeof(uint16_t)) / 2)) {
00077             bPingPong = true;
00078             bCallback = true;
00079         }
00080     }
00081
00082     // Done writing - see if there's a 50% or rollover callback
00083     if (bCallback && m_pfCallback) {
00084         uint16_t ul6Size = (sizeof(m_aul6Buffer) / sizeof(uint16_t)) / 2;
00085         if (bPingPong) {
00086             m_pfCallback(m_aul6Buffer, ul6Size, bPingPong);
00087         } else {
00088             m_pfCallback(m_aul6Buffer + ul6Size, ul6Size, bPingPong);
00089         }
00090     }
00091 }

```

```

00090     }
00091 }
00092
00093 #endif

```

19.123 /media/usb/project/github/Mark3/libs/mark3c/public/fake_types.h File Reference

C-struct definitions that mirror.

```

#include "kerneltypes.h"
#include "mark3cfg.h"

```

19.123.1 Detailed Description

C-struct definitions that mirror.

This header contains a set of "fake" structures that have the same memory layout as the kernel objects in C++ (taking into account inheritance, etc.). These are used for sizing the opaque data blobs that are declared in C, which then become instantiated as C++ kernel objects via the bindings provided.

Definition in file [fake_types.h](#).

19.124 fake_types.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00026 #include "kerneltypes.h"
00027 #include "mark3cfg.h"
00028
00029 #ifndef __FAKE_TYPES_H__
00030 #define __FAKE_TYPES_H__
00031
00032 #if defined(__cplusplus)
00033 extern "C" {
00034 #endif
00035
00036 //-----
00037 typedef struct {
00038     void* prev;
00039     void* next;
00040 } Fake_LinkedListNode;
00041
00042 //-----
00043 typedef struct {
00044     void* head;
00045     void* tail;
00046 } Fake_LinkedList;
00047
00048 //-----
00049 typedef struct {
00050     Fake_LinkedList fake_list;
00051     PORT_PRIO_TYPE m_uXPriority;
00052     void* m_pclMap;
00053 } Fake_ThreadList;
00054
00055 //-----
00056 typedef struct {
00057     Fake_LinkedListNode m_ll_node;
00058 #if KERNEL_EXTRA_CHECKS
00059     uint8_t m_u8Initialized;

```

```

00060 #endif
00061     uint8_t          m_u8Flags;
00062     void*            m_pfCallback;
00063     uint32_t         m_u32Interval;
00064     uint32_t         m_u32TimeLeft;
00065     uint32_t         m_u32TimerTolerance;
00066     void*            m_pclOwner;
00067     void*            m_pvData;
00068 } Fake_Timer;
00069
00070 //-----
00071 typedef struct {
00072     Fake_LinkedListNode m_ll_node;
00073     K_WORD*             m_pwStackTop;
00074     K_WORD*             m_pwStack;
00075     uint8_t             m_u8ThreadID;
00076     PORT_PRIO_TYPE      m_uXPriority;
00077     PORT_PRIO_TYPE      m_uXCurPriority;
00078     uint8_t             m_eState;
00079 #if KERNEL_USE_EXTENDED_CONTEXT
00080     void*               m_pvExtendedContext;
00081 #endif
00082 #if KERNEL_USE_THREADNAME
00083     const char*         m_szName;
00084 #endif
00085     uint16_t            m_u16StackSize;
00086     void*               m_pclCurrent;
00087     void*               m_pclOwner;
00088     void*               m_pfEntryPoint;
00089     void*               m_pvArg;
00090 #if KERNEL_USE_QUANTUM
00091     uint16_t            m_u16Quantum;
00092 #endif
00093 #if KERNEL_USE_EVENTFLAG
00094     uint16_t            m_u16FlagMask;
00095     uint8_t             m_eFlagMode;
00096 #endif
00097 #if KERNEL_USE_TIMEOUTS || KERNEL_USE_SLEEP
00098     Fake_Timer           m_clTimer;
00099 #endif
00100 #if KERNEL_USE_TIMEOUTS
00101     bool                m_bExpired;
00102 #endif
00103 } Fake_Thread;
00104
00105 //-----
00106 typedef struct {
00107     Fake_ThreadList      thread_list;
00108 #if KERNEL_EXTRA_CHECKS
00109     uint8_t              m_u8Initialized;
00110 #endif
00111     uint16_t             m_u16Value;
00112     uint16_t             m_u16MaxValue;
00113 } Fake_Semaphore;
00114
00115 //-----
00116 typedef struct {
00117     Fake_ThreadList      thread_list;
00118 #if KERNEL_EXTRA_CHECKS
00119     uint8_t              m_u8Initialized;
00120 #endif
00121     uint8_t              m_u8Recurse;
00122     bool                 m_bReady;
00123     uint8_t              m_u8MaxPri;
00124     void*                m_pclOwner;
00125 } Fake_Mutex;
00126
00127 //-----
00128 typedef struct {
00129     Fake_LinkedListNode  list_node;
00130     void*                m_pvData;
00131     uint16_t             m_u16Code;
00132 } Fake_Message;
00133
00134 //-----
00135 typedef struct {
00136     Fake_Semaphore       m_clSemaphore;
00137     Fake_LinkedList       m_clLinkList;
00138 } Fake_MessageQueue;
00139
00140 //-----
00141 typedef struct {
00142     uint16_t             m_u16Head;
00143     uint16_t             m_u16Tail;
00144     uint16_t             m_u16Count;
00145     uint16_t             m_u16Free;
00146     uint16_t             m_u16ElementSize;

```

```

00147     void*          m_pvBuffer;
00148     Fake_Semaphore m_clRecvSem;
00149 #if KERNEL_USE_TIMEOUTS
00150     Fake_Semaphore m_clSendSem;
00151 #endif
00152 } Fake_Mailbox;
00153
00154 //-----
00155 typedef struct {
00156     Fake_ThreadList thread_list;
00157 #if KERNEL_EXTRA_CHECKS
00158     uint8_t          m_u8Initialized;
00159 #endif
00160     bool             m_bPending;
00161 } Fake_Notify;
00162
00163 //-----
00164 typedef struct {
00165     Fake_ThreadList thread_list;
00166 #if KERNEL_EXTRA_CHECKS
00167     uint8_t          m_u8Initialized;
00168 #endif
00169     uint16_t         m_ul6EventFlag;
00170 } Fake_EventFlag;
00171
00172 #if defined(__cplusplus)
00173 }
00174 #endif
00175
00176 #endif // __FAKE_TYPES_H__

```

19.125 /media/usb/project/github/Mark3/libs/mark3c/public/mark3c.h File Reference

Header providing C-language API bindings for the Mark3 kernel.

```

#include "mark3cfg.h"
#include "kerneltypes.h"
#include "fake_types.h"
#include "driver3c.h"
#include <stdint.h>
#include <stdbool.h>

```

Typedefs

- typedef void * [EventFlag_t](#)
EventFlag opaque handle data type.
- typedef void * [Mailbox_t](#)
Mailbox opaque handle data type.
- typedef void * [Message_t](#)
Message opaque handle data type.
- typedef void * [MessageQueue_t](#)
MessageQueue opaque handle data type.
- typedef void * [Mutex_t](#)
Mutex opaque handle data type.
- typedef void * [Notify_t](#)
Notification object opaque handle data type.
- typedef void * [Semaphore_t](#)
Semaphore opaque handle data type.
- typedef void * [Thread_t](#)
Thread opaque handle data type.
- typedef void * [Timer_t](#)
Timer opaque handle data type.

Functions

- void [Kernel_Init](#) (void)
Kernel_Init.
- void [Kernel_Start](#) (void)
Kernel_Start.
- bool [Kernel_IsStarted](#) (void)
Kernel_IsStarted.
- void [Kernel_SetPanic](#) ([PanicFunc_t](#) pfPanic_)
Kernel_SetPanic.
- bool [Kernel_IsPanic](#) (void)
Kernel_IsPanic.
- void [Kernel_Panic](#) (uint16_t u16Cause_)
Kernel_Panic.
- void [Scheduler_Enable](#) (bool bEnable_)
Scheduler_Enable.
- bool [Scheduler_IsEnabled](#) (void)
Scheduler_IsEnabled.
- [Thread_t](#) [Scheduler_GetCurrentThread](#) (void)
Scheduler_GetCurrentThread.
- void [Thread_Init](#) ([Thread_t](#) handle, [K_WORD](#) *pwStack_, uint16_t u16StackSize_, [PORT_PRIO_TYPE](#) uXPriority_, [ThreadEntry_t](#) pfEntryPoint_, void *pvArg_)
Thread_Init.
- void [Thread_Start](#) ([Thread_t](#) handle)
Thread_Start.
- void [Thread_Stop](#) ([Thread_t](#) handle)
Thread_Stop.
- [PORT_PRIO_TYPE](#) [Thread_GetPriority](#) ([Thread_t](#) handle)
Thread_GetPriority.
- [PORT_PRIO_TYPE](#) [Thread_GetCurPriority](#) ([Thread_t](#) handle)
Thread_GetCurPriority.
- void [Thread_SetPriority](#) ([Thread_t](#) handle, [PORT_PRIO_TYPE](#) uXPriority_)
Thread_SetPriority.
- void [Thread_Yield](#) (void)
Thread_Yield.
- void [Thread_SetID](#) ([Thread_t](#) handle, uint8_t u8ID_)
Thread_SetID.
- uint8_t [Thread_GetID](#) ([Thread_t](#) handle)
Thread_GetID.
- uint16_t [Thread_GetStackSlack](#) ([Thread_t](#) handle)
Thread_GetStackSlack.
- [ThreadState_t](#) [Thread_GetState](#) ([Thread_t](#) handle)
Thread_GetState.

19.125.1 Detailed Description

Header providing C-language API bindings for the Mark3 kernel.

Definition in file [mark3c.h](#).

19.125.2 Function Documentation

19.125.2.1 void Kernel_Init (void)

Kernel_Init.

See also

void [Kernel::Init\(\)](#)

19.125.2.2 bool Kernel_IsPanic (void)

Kernel_IsPanic.

See also

bool [Kernel::IsPanic\(\)](#)

Returns

Whether or not the kernel is in a panic state

19.125.2.3 bool Kernel_IsStarted (void)

Kernel_IsStarted.

See also

bool [Kernel::IsStarted\(\)](#)

Returns

Whether or not the kernel has started - true = running, false = not started

19.125.2.4 void Kernel_Panic (uint16_t u16Cause_)

Kernel_Panic.

See also

void [Kernel::Panic\(uint16_t u16Cause_\)](#)

Parameters

<i>u16Cause_</i>	Reason for the kernel panic
------------------	-----------------------------

19.125.2.5 void Kernel_SetPanic (PanicFunc_t pfPanic_)

Kernel_SetPanic.

See also

void [Kernel::SetPanic\(PanicFunc_t pfPanic_\)](#)

Parameters

<i>pfPanic_</i>	Panic function pointer
-----------------	------------------------

19.125.2.6 void Kernel_Start (void)

Kernel_Start.

See also

void [Kernel::Start\(\)](#)

19.125.2.7 void Scheduler_Enable (bool *bEnable_*)

Scheduler_Enable.

See also

void [Scheduler::SetScheduler\(bool bEnable_\)](#)

Parameters

<i>bEnable_true</i>	to enable, false to disable the scheduler
---------------------	---

19.125.2.8 Thread_t Scheduler_GetCurrentThread (void)

Scheduler_GetCurrentThread.

See also

Thread* [Scheduler::GetCurrentThread\(\)](#)

Returns

Handle of the currently-running thread

19.125.2.9 bool Scheduler_IsEnabled (void)

Scheduler_IsEnabled.

See also

bool [Scheduler::IsEnabled\(\)](#)

Returns

true - scheduler enabled, false - disabled

19.125.2.10 PORT_PRIO_TYPE Thread_GetCurPriority (Thread_t *handle*)

Thread_GetCurPriority.

See also

[PORT_PRIO_TYPE Thread::GetCurPriority\(\)](#)

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Current priority of the thread considering priority inheritance

19.125.2.11 `uint8_t Thread_GetID (Thread_t handle)`

Thread_GetID.

See also

`uint8_t Thread::GetID()`

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Return ID assigned to the thread

19.125.2.12 `PORT_PRIO_TYPE Thread_GetPriority (Thread_t handle)`

Thread_GetPriority.

See also

`PORT_PRIO_TYPE Thread::GetPriority()`

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Current priority of the thread not considering priority inheritance

19.125.2.13 `uint16_t Thread_GetStackSlack (Thread_t handle)`

Thread_GetStackSlack.

See also

`uint16_t Thread::GetStackSlack()`

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Return the amount of unused stack on the given thread

19.125.2.14 ThreadState_t Thread_GetState (Thread_t handle)

Thread_GetState.

See also

[ThreadState_t Thread::GetState\(\)](#)

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

The thread's current execution state

19.125.2.15 void Thread_Init (Thread_t handle, K_WORD * pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_, ThreadEntry_t pfEntryPoint_, void * pvArg_)

Thread_Init.

See also

void [Thread::Init](#)(K_WORD *pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_, [ThreadEntry_t](#) pfEntryPoint_, void *pvArg_)

Parameters

<i>handle</i>	Handle of the thread to initialize
<i>pwStack_</i>	Pointer to the stack to use for the thread
<i>u16StackSize_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint_</i>	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entrypoint function.

19.125.2.16 void Thread_SetID (Thread_t handle, uint8_t u8ID_)

Thread_SetID.

See also

void [Thread::SetID](#)(uint8_t u8ID_)

Parameters

<i>handle</i>	Handle of the thread
<i>u8ID_</i>	ID To assign to the thread

19.125.2.17 void Thread_SetPriority (Thread_t handle, PORT_PRIO_TYPE uXPriority_)

Thread_SetPriority.

See also

void [Thread::SetPriority](#)(PORT_PRIO_TYPE uXPriority_)

Parameters

<i>handle</i>	Handle of the thread
<i>uXPriority_</i>	New priority level

19.125.2.18 void Thread_Start (Thread_t handle)

Thread_Start.

See also

void [Thread::Start\(\)](#)

Parameters

<i>handle</i>	Handle of the thread to start
---------------	-------------------------------

19.125.2.19 void Thread_Stop (Thread_t handle)

Thread_Stop.

See also

void [Thread::Stop\(\)](#)

Parameters

<i>handle</i>	Handle of the thread to stop
---------------	------------------------------

19.125.2.20 void Thread_Yield (void)

Thread_Yield.

See also

void [Thread::Yield\(\)](#)

19.126 mark3c.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "mark3cfg.h"
00022 #include "kerneltypes.h"
00023 #include "fake_types.h"
00024 #include "driver3c.h"
00025
00026 #include <stdint.h>
00027 #include <stdbool.h>
00028
00029 #ifndef __MARK3C_H__
00030 #define __MARK3C_H__

```

```

00031
00032 #if defined(__cplusplus)
00033 extern "C" {
00034 #endif
00035
00036 //-----
00037 // Define a series of handle types to be used in place of the underlying classes
00038 // of Mark3.
00039 typedef void* EventFlag_t;
00040 typedef void* Mailbox_t;
00041 typedef void* Message_t;
00042 typedef void* MessageQueue_t;
00043 typedef void* Mutex_t;
00044 typedef void* Notify_t;
00045 typedef void* Semaphore_t;
00046 typedef void* Thread_t;
00047 typedef void* Timer_t;
00048
00049 //-----
00050 // Function pointer types used by Kernel APIs
00051 typedef void (*thread_create_callout_t)(Thread_t hThread_);
00052 typedef void (*thread_exit_callout_t)(Thread_t hThread_);
00053 typedef void (*thread_context_callout_t)(Thread_t hThread_);
00054
00055 //-----
00056 // Use the sizes of the structs in fake_types.h to generate opaque object-blobs
00057 // that get instantiated as kernel objects (from the C++ code) later.
00058 #define THREAD_SIZE (sizeof(Fake_Thread))
00059 #define TIMER_SIZE (sizeof(Fake_Timer))
00060 #define SEMAPHORE_SIZE (sizeof(Fake_Semaphore))
00061 #define MUTEX_SIZE (sizeof(Fake_Mutex))
00062 #define MESSAGE_SIZE (sizeof(Fake_Message))
00063 #define MESSAGEQUEUE_SIZE (sizeof(Fake_MessageQueue))
00064 #define MAILBOX_SIZE (sizeof(Fake_Mailbox))
00065 #define NOTIFY_SIZE (sizeof(Fake_Notify))
00066 #define EVENTFLAG_SIZE (sizeof(Fake_EventFlag))
00067
00068 //-----
00069 // Macros for declaring opaque buffers of an appropriate size for the given
00070 // kernel objects
00071 #define TOKEN_1(x, y) x##y
00072 #define TOKEN_2(x, y) TOKEN_1(x, y)
00073
00074 // Ensure that opaque buffers are sized to the nearest word - which is
00075 // a platform-dependent value.
00076 #define WORD_ROUND(x) ((x) + (sizeof(K_WORD) - 1)) / sizeof(K_WORD)
00077
00078 #define DECLARE_THREAD(name)
00079 \
    K_WORD TOKEN_2(__thread_, name)[WORD_ROUND(THREAD_SIZE)];
00080 \
    Thread_t name = (Thread_t)TOKEN_2(__thread_, name);
00081
00082 #define DECLARE_TIMER(name)
00083 \
    K_WORD TOKEN_2(__timer_, name)[WORD_ROUND(TIMER_SIZE)];
00084 \
    Timer_t name = (Timer_t)TOKEN_2(__timer_, name);
00085
00086 #define DECLARE_SEMAPHORE(name)
00087 \
    K_WORD TOKEN_2(__semaphore_, name)[WORD_ROUND(SEMAPHORE_SIZE)];
00088 \
    Semaphore_t name = (Semaphore_t)TOKEN_2(__semaphore_, name);
00089
00090 #define DECLARE_MUTEX(name)
00091 \
    K_WORD TOKEN_2(__mutex_, name)[WORD_ROUND(MUTEX_SIZE)];
00092 \
    Mutex_t name = (Mutex_t)TOKEN_2(__mutex_, name);
00093
00094 #define DECLARE_MESSAGE(name)
00095 \
    K_WORD TOKEN_2(__message_, name)[WORD_ROUND(MESSAGE_SIZE)];
00096 \
    Message_t name = (Message_t)TOKEN_2(__message_, name);
00097
00098 #define DECLARE_MESSAGEQUEUE(name)
00099 \
    K_WORD TOKEN_2(__messagequeue_, name)[WORD_ROUND(MESSAGEQUEUE_SIZE)];
00100 \
    MessageQueue_t name = (MessageQueue_t)TOKEN_2(__messagequeue_, name);
00101
00102 #define DECLARE_MAILBOX(name)
00103 \
    K_WORD TOKEN_2(__mailbox_, name)[WORD_ROUND(MAILBOX_SIZE)];

```

```

00104     Mailbox_t name = (Mailbox_t)TOKEN_2(__mailbox_, name);
00105
00106 #define DECLARE_NOTIFY(name)
00107     \
00108     K_WORD     TOKEN_2(__notify_, name) [WORD_ROUND(NOTIFY_SIZE)];
00109
00108     Notify_t name = (Notify_t)TOKEN_2(__notify_, name);
00109
00110 #define DECLARE_EVENTFLAG(name)
00111     \
00112     K_WORD     TOKEN_2(__eventflag_, name) [WORD_ROUND(EVENTFLAG_SIZE)];
00113
00112     EventFlag_t name = (EventFlag_t)TOKEN_2(__eventflag_, name);
00113
00114 //-----
00115 // Allocate-once Memory managment APIs
00116 #if defined KERNEL_USE_AUTO_ALLOC
00117
00123 void* AutoAlloc(uint16_t ul6Size_);
00124 #if KERNEL_USE_SEMAPHORE
00125
00130 Semaphore_t Alloc_Semaphore(void);
00131 #endif
00132 #if KERNEL_USE_MUTEX
00133
00138 Mutex_t Alloc_Mutex(void);
00139 #endif
00140 #if KERNEL_USE_EVENTFLAG
00141
00146 EventFlag_t Alloc_EventFlag(void);
00147 #endif
00148 #if KERNEL_USE_MESSAGE
00149
00154 Message_t Alloc_Message(void);
00160 MessageQueue_t Alloc_MessageQueue(void);
00161 #endif
00162 #if KERNEL_USE_NOTIFY
00163
00168 Notify_t Alloc_Notify(void);
00169 #endif
00170 #if KERNEL_USE_MAILBOX
00171
00176 Mailbox_t Alloc_Mailbox(void);
00177 #endif
00178
00183 Thread_t Alloc_Thread(void);
00184 #if KERNEL_USE_TIMERS
00185
00190 Timer_t Alloc_Timer(void);
00191 #endif
00192 #endif
00193
00194 //-----
00195 // Kernel APIs
00200 void Kernel_Init(void);
00205 void Kernel_Start(void);
00212 bool Kernel_IsStarted(void);
00218 void Kernel_SetPanic(PanicFunc_t pfPanic_);
00224 bool Kernel_IsPanic(void);
00230 void Kernel_Panic(uint16_t ul6Cause_);
00231 #if KERNEL_USE_IDLE_FUNC
00232
00237 void Kernel_SetIdleFunc(IdleFunc_t pfIdle_);
00238 #endif
00239
00240 #if KERNEL_USE_THREAD_CALLOUTS
00241
00246 void Kernel_SetThreadCreateCallout(thread_create_callout_t pfCreate_);
00252 void Kernel_SetThreadExitCallout(thread_exit_callout_t pfExit_);
00253
00259 void Kernel_SetThreadContextSwitchCallout(thread_context_callout_t pfContext_);
00260
00266 thread_create_callout_t Kernel_GetThreadCreateCallout(void);
00267
00273 thread_exit_callout_t Kernel_GetThreadExitCallout(void);
00274
00280 thread_context_callout_t Kernel_GetThreadContextSwitchCallout(void);
00281 #endif
00282
00283 #if KERNEL_USE_STACK_GUARD
00284
00290 void Kernel_SetStackGuardThreshold(uint16_t ul6Threshold_);
00291
00297 uint16_t Kernel_GetStackGuardThreshold(void);
00298 #endif
00299 //-----
00300 // Scheduler APIs

```

```

00306 void Scheduler_Enable(bool bEnable_);
00312 bool Scheduler_IsEnabled(void);
00318 Thread_t Scheduler_GetCurrentThread(void);
00319
00320 //-----
00321 // Thread APIs
00335 void Thread_Init(Thread_t handle,
00336                  K_WORD* pwStack_,
00337                  uint16_t ul6StackSize_,
00338                  PORT_PRIO_TYPE uXPriority_,
00339                  ThreadEntry_t pfEntryPoint_,
00340                  void* pvArg_);
00346 void Thread_Start(Thread_t handle);
00352 void Thread_Stop(Thread_t handle);
00353 #if KERNEL_USE_THREADNAME
00354
00360 void Thread_SetName(Thread_t handle, const char* szName_);
00367 const char* Thread_GetName(Thread_t handle);
00368 #endif
00369
00375 PORT_PRIO_TYPE Thread_GetPriority(Thread_t handle);
00382 PORT_PRIO_TYPE Thread_GetCurPriority(Thread_t handle);
00383 #if KERNEL_USE_QUANTUM
00384
00390 void Thread_SetQuantum(Thread_t handle, uint16_t ul6Quantum_);
00397 uint16_t Thread_GetQuantum(Thread_t handle);
00398 #endif
00399
00405 void Thread_SetPriority(Thread_t handle, PORT_PRIO_TYPE uXPriority_);
00406 #if KERNEL_USE_DYNAMIC_THREADS
00407
00412 void Thread_Exit(Thread_t handle);
00413 #endif
00414 #if KERNEL_USE_SLEEP
00415
00420 void Thread_Sleep(uint32_t u32TimeMs_);
00426 void Thread_USleep(uint32_t u32TimeUs_);
00427 #endif
00428 #if KERNEL_USE_EXTENDED_CONTEXT
00429
00435 void* Thread_GetExtendedContext(Thread_t handle);
00436
00443 void Thread_SetExtendedContext(Thread_t handle, void* pvData_);
00444
00445 #endif
00446
00450 void Thread_Yield(void);
00457 void Thread_SetID(Thread_t handle, uint8_t u8ID_);
00464 uint8_t Thread_GetID(Thread_t handle);
00471 uint16_t Thread_GetStackSlack(Thread_t handle);
00478 ThreadState_t Thread_GetState(Thread_t handle);
00479
00480 //-----
00481 // Timer APIs
00482 #if KERNEL_USE_TIMERS
00483 typedef void (*TimerCallbackC_t)(Thread_t hOwner_, void* pvData_);
00489 void Timer_Init(Timer_t handle);
00501 void Timer_Start(Timer_t handle,
00502                  bool bRepeat_,
00503                  uint32_t u32IntervalMs_,
00504                  uint32_t u32ToleranceMs_,
00505                  TimerCallbackC_t pfCallback_,
00506                  void* pvData_);
00507
00513 void Timer_Restart(Timer_t handle);
00514
00520 void Timer_Stop(Timer_t handle);
00521 #endif
00522
00523 //-----
00524 // Semaphore APIs
00525 #if KERNEL_USE_SEMAPHORE
00526
00533 void Semaphore_Init(Semaphore_t handle, uint16_t ul6InitVal_, uint16_t ul6MaxVal_);
00539 void Semaphore_Post(Semaphore_t handle);
00545 void Semaphore_Pend(Semaphore_t handle);
00546 #if KERNEL_USE_TIMEOUTS
00547
00554 bool Semaphore_TimedPend(Semaphore_t handle, uint32_t u32WaitTimeMS_);
00555 #endif
00556 #endif
00557
00558 //-----
00559 // Mutex APIs
00560 #if KERNEL_USE_MUTEX
00561
00566 void Mutex_Init(Mutex_t handle);

```

```

00572 void Mutex_Claim(Mutex_t handle);
00578 void Mutex_Release(Mutex_t handle);
00579 #if KERNEL_USE_TIMEOUTS
00580
00587 bool Mutex_TimedClaim(Mutex_t handle, uint32_t u32WaitTimeMS_);
00588 #endif
00589 #endif
00590
00591 //-----
00592 // EventFlag APIs
00593 #if KERNEL_USE_EVENTFLAG
00594
00599 void EventFlag_Init(EventFlag_t handle);
00608 uint16_t EventFlag_Wait(EventFlag_t handle, uint16_t u16Mask_,
    EventFlagOperation_t eMode_);
00609 #if KERNEL_USE_TIMEOUTS
00610
00619 uint16_t EventFlag_TimedWait(EventFlag_t handle, uint16_t u16Mask_,
    EventFlagOperation_t eMode_, uint32_t u32TimeMS_);
00620 #endif
00621
00627 void EventFlag_Set(EventFlag_t handle, uint16_t u16Mask_);
00634 void EventFlag_Clear(EventFlag_t handle, uint16_t u16Mask_);
00641 uint16_t EventFlag_GetMask(EventFlag_t handle);
00642 #endif
00643
00644 //-----
00645 // Notification APIs
00646 #if KERNEL_USE_NOTIFY
00647
00652 void Notify_Init(Notify_t handle);
00658 void Notify_Signal(Notify_t handle);
00665 void Notify_Wait(Notify_t handle, bool* pbFlag_);
00666 #if KERNEL_USE_TIMEOUTS
00667
00675 bool Notify_TimedWait(Notify_t handle, uint32_t u32WaitTimeMS_, bool* pbFlag_);
00676 #endif
00677 #endif
00678
00679 //-----
00680 // Atomic Functions
00681 #if KERNEL_USE_ATOMIC
00682
00689 uint8_t Atomic_Set8(uint8_t* pu8Source_, uint8_t u8Val_);
00697 uint16_t Atomic_Set16(uint16_t* pu16Source_, uint16_t u16Val_);
00705 uint32_t Atomic_Set32(uint32_t* pu32Source_, uint32_t u32Val_);
00713 uint8_t Atomic_Add8(uint8_t* pu8Source_, uint8_t u8Val_);
00721 uint16_t Atomic_Add16(uint16_t* pu16Source_, uint16_t u16Val_);
00729 uint32_t Atomic_Add32(uint32_t* pu32Source_, uint32_t u32Val_);
00737 uint8_t Atomic_Sub8(uint8_t* pu8Source_, uint8_t u8Val_);
00745 uint16_t Atomic_Sub16(uint16_t* pu16Source_, uint16_t u16Val_);
00753 uint32_t Atomic_Sub32(uint32_t* pu32Source_, uint32_t u32Val_);
00762 bool Atomic_TestAndSet(bool* pbLock);
00763 #endif
00764
00765 //-----
00766 // Message/Message Queue APIs
00767 #if KERNEL_USE_MESSAGE
00768
00773 void Message_Init(Message_t handle);
00780 void Message_SetData(Message_t handle, void* pvData_);
00787 void* Message_GetData(Message_t handle);
00794 void Message_SetCode(Message_t handle, uint16_t u16Code_);
00801 uint16_t Message_GetCode(Message_t handle);
00807 void GlobalMessagePool_Push(Message_t handle);
00813 Message_t GlobalMessagePool_Pop(void);
00819 void MessageQueue_Init(MessageQueue_t handle);
00826 Message_t MessageQueue_Receive(MessageQueue_t handle);
00827 #if KERNEL_USE_TIMEOUTS
00828
00838 Message_t MessageQueue_TimedReceive(MessageQueue_t handle, uint32_t u32TimeWaitMS_);
00839 #endif
00840
00847 void MessageQueue_Send(MessageQueue_t handle, Message_t hMessage_);
00848
00854 uint16_t MessageQueue_GetCount(void);
00855 #endif
00856
00857 //-----
00858 // Mailbox APIs
00859 #if KERNEL_USE_MAILBOX
00860
00869 void Mailbox_Init(Mailbox_t handle, void* pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_);
00870
00878 bool Mailbox_Send(Mailbox_t handle, void* pvData_);
00879
00887 bool Mailbox_SendTail(Mailbox_t handle, void* pvData_);

```



```
00888
00897 bool Mailbox_TimedSend(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
00898
00907 bool Mailbox_TimedSendTail(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
00908
00916 void Mailbox_Receive(Mailbox_t handle, void* pvData_);
00917
00925 void Mailbox_ReceiveTail(Mailbox_t handle, void* pvData_);
00926 #if KERNEL_USE_TIMEOUTS
00927
00937 bool Mailbox_TimedReceive(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
00938
00948 bool Mailbox_TimedReceiveTail(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
00949
00956 uint16_t Mailbox_GetFreeSlots(Mailbox_t handle);
00957
00964 bool Mailbox_IsFull(Mailbox_t handle);
00965
00972 bool Mailbox_IsEmpty(Mailbox_t handle);
00973 #endif
00974 #endif
00975
00976 //-----
00977 // Kernel-Aware Simulation APIs
00978 #if KERNEL_AWARE_SIMULATION
00979
00985 void KernelAware_ProfileInit(const char* szStr_);
00986
00991 void KernelAware_ProfileStart(void);
00992
00997 void KernelAware_ProfileStop(void);
00998
01003 void KernelAware_ProfileReport(void);
01004
01010 void KernelAware_ExitSimulator(void);
01011
01017 void KernelAware_Print(const char* szStr_);
01018
01025 void KernelAware_Trace(uint16_t ul6File_, uint16_t ul6Line_);
01026
01034 void KernelAware_Trace1(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_);
01043 void KernelAware_Trace2(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_, uint16_t ul6Arg2_);
01053 bool KernelAware_IsSimulatorAware(void);
01054 #endif
01055
01056 #if defined(__cplusplus)
01057 }
01058 #endif
01059
01060 #endif // __MARK3C_H__
```


Example Documentation

This example demonstrates how low-overhead logging can be implemented using `buffalogger`.

```
/*=====*/  
#include "mark3.h"  
#include "kerneldebug.h"  
#include "drvATMegaUART.h"  
#include "tracebuffer.h"  
#include "ksemaphore.h"  
  
/*=====*/  
  
Example - Logging data via buffalogger/debug APIs.  
  
/*=====*/  
  
#if !(KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION)  
# error "Buffalogger demo requires tracebuffer support"  
#endif  
  
#if KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION  
  
#define _CAN_HAS_DEBUG  
//--[Autogenerated - Do Not Modify]-----  
#include "dbg_file_list.h"  
#include "buffalogger.h"  
#if defined(DBG_FILE)  
#error "Debug logging file token already defined! Bailing."  
#else  
#define DBG_FILE _DBG__EXAMPLES AVR_BUFFALOGGER_MAIN_CPP  
#endif  
//--[End Autogenerated content]-----  
  
//-----  
// This block declares the thread data for the main application thread. It  
// defines a thread object, stack (in word-array form), and the entry-point  
// function used by the application thread.  
#define APP_STACK_SIZE (192 / sizeof(K_WORD))  
static Thread clAppThread;  
static K_WORD awAppStack[APP_STACK_SIZE];  
static void AppMain(void* unused_);  
  
#define IDLE_STACK_SIZE (192 / sizeof(K_WORD))  
static Thread clIdleThread;  
static K_WORD awIdleStack[APP_STACK_SIZE];  
static void IdleMain(void* unused_);
```

```

#define LOGGER_STACK_SIZE (192 / sizeof(K_WORD))
static Thread clLoggerThread;
static K_WORD awLoggerStack[APP_STACK_SIZE];
static void LoggerMain(void* unused_);
static volatile bool bPingPong;
static Semaphore clSem;

//-----
static ATmegaUART clUART;

//-----
#define UART_SIZE_TX (32)
#define UART_SIZE_RX (8)

static uint8_t aucTxBuffer[UART_SIZE_TX];
static uint8_t aucRxBuffer[UART_SIZE_RX];

static volatile uint16_t* pul6Log;
static volatile uint16_t ul6LogLen;

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

void IdleMain(void* unused_)
{
    while (1) {
    }
}

//-----
void LoggerCallback(uint16_t* pul6Data_, uint16_t ul6Len_, bool bPingPong_)
{
    CS_ENTER();
    bPingPong = bPingPong_;
    pul6Log = pul6Data_;
    ul6LogLen = ul6Len_;
    CS_EXIT();

    clSem.Post();
}

//-----
void LoggerMain(void* unused_)
{
    while (1) {
        uint8_t* src;
        uint16_t len;

        clSem.Pend();

        CS_ENTER();
        src = (uint8_t*)pul6Log;
        len = ul6LogLen * sizeof(uint16_t);
        CS_EXIT();

        uint16_t written = 0;
        while (len != written) {
            written += clUART.Write(len - written, src + written);
        }
    }
}

//-----
int main(void)
{
    Kernel::Init();

    // Example assumes use of built-in idle.
    clAppThread.Init(awAppStack, APP_STACK_SIZE, 2, AppMain, 0);
    clAppThread.Start();

    clLoggerThread.Init(awLoggerStack, LOGGER_STACK_SIZE, 1, LoggerMain, 0);
    clLoggerThread.Start();

    clIdleThread.Init(awIdleStack, IDLE_STACK_SIZE, 0, IdleMain, 0);
    clIdleThread.Start();

    clUART.SetName("/dev/tty");
    clUART.Init();
    clUART.Open();

    DriverList::Add(&clUART);
}

```

```

    Kernel::Start();

    return 0;
}

//-----
void AppMain(void* unused_)
{
    {
        uint32_t u32Baud = 57600 * 4;
        clUART.Control(CMD_SET_BAUDRATE, &u32Baud, 0, 0, 0);
    }
    clUART.Control(CMD_SET_BUFFERS, (void*)aucRxBuffer, UART_SIZE_RX, (void*)aucTxBuffer, UART_SIZE_TX);

    clSem.Init(0, 1);

    TraceBuffer::SetCallback(LoggerCallback);
    volatile uint16_t ul6Iteration = 0;
    while (1) {
        Thread::Sleep(100);
        USER_TRACE("Beginning of the main application loop!");

        Thread::Sleep(100);
        USER_TRACE_1(" Iteration: %d", ul6Iteration++);

        Thread::Sleep(100);
        USER_TRACE("End of the main application loop!");
    }
}

#endif //if KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION

```

20.2 lab10_notifications/main.cpp

This examples demonstrates how to use notification objects as a thread synchronization mechanism.

[illegible]

Takeaway:

- Mailboxes are a powerful IPC mechanism used to pass messages of a fixed-size between threads.

```

=====*/
#ifdef !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

//-----
#define APP_STACK_SIZE (256 / sizeof(K_WORD))
static Thread clApp1Thread;
static K_WORD awApp1Stack[APP_STACK_SIZE];
static void App1Main(void* unused_);

//-----
static Thread clApp2Thread;
static K_WORD awApp2Stack[APP_STACK_SIZE];
static void App2Main(void* unused_);

//-----
static Mailbox clMailbox;
static uint8_t au8MBData[100];

typedef struct {
    uint8_t au8Buffer[10];
} MBType_t;

//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    // Initialize the threads used in this example
    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp1Thread.Start();

    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 2, App2Main, 0);
    clApp2Thread.Start();

    // Initialize the mailbox used in this example
    clMailbox.Init(au8MBData, 100, sizeof(MBType_t));

    Kernel::Start();

    return 0;
}

//-----
void App1Main(void* unused_)
{
    while (1) {
        MBType_t stMsg;

        // Wait until there is an envelope available in the shared mailbox, and
        // then log a trace message.
        clMailbox.Receive(&stMsg);
        KernelAware::Trace(0, __LINE__, stMsg.au8Buffer[0], stMsg.au8Buffer[9]);
    }
}

//-----
void App2Main(void* unused_)
{
    while (1) {
        MBType_t stMsg;

        // Place a bunch of envelopes in the mailbox, and then wait for a
        // while. Note that this thread has a higher priority than the other
        // thread, so it will keep pushing envelopes to the other thread until
        // it gets to the sleep, at which point the other thread will be allowed
        // to execute.

        KernelAware::Print("Messages Begin\n");

        for (uint8_t i = 0; i < 10; i++) {
            for (uint8_t j = 0; j < 10; j++) {
                stMsg.au8Buffer[j] = (i * 10) + j;
            }
            clMailbox.Send(&stMsg);
        }
    }
}

```



```
//-----
int main(void)
{
    // Before any Mark3 RTOS APIs can be called, the user must call Kernel::Init().
    // Note that if you have any hardware-specific init code, it can be called
    // before Kernel::Init, so long as it does not enable interrupts, or
    // rely on hardware peripherals (timer, software interrupt, etc.) used by the
    // kernel.
    Kernel::Init();

    // Once the kernel initialization has been complete, the user can add their
    // application thread(s) and idle thread. Threads added before the kernel
    // is started are referred to as the "static threads" in the system, as they
    // are the default working-set of threads that make up the application on
    // kernel startup.

    // Initialize the application thread to use a specified word-array as its stack.
    // The thread will run at priority level "1", and start execution the
    // "AppMain" function when it's started.
    clAppThread.Init(awAppStack, sizeof(awAppStack), 1, AppMain, 0);

    // Initialize the idle thread to use a specific word-array as its stack.
    // The thread will run at priority level "0", which is reserved for the idle
    // priority thread. IdleMain will be run when the thread is started.
    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);

    // Once the static threads have been added, the user must then ensure that the
    // threads are ready to execute. By default, creating a thread is created
    // in a STOPPED state. All threads must manually be started using the
    // Start() API before they will be scheduled by the system. Here, we are
    // starting the application and idle threads before starting the kernel - and
    // that's OK. When the kernel is started, it will choose which thread to run
    // first from the pool of ready threads.

    clAppThread.Start();
    clIdleThread.Start();

    // All threads have been initialized and made ready. The kernel will now
    // select the first thread to run, enable the hardware required to run the
    // kernel (Timers, software interrupts, etc.), and then do whatever is
    // necessary to maneuver control of thread execution to the kernel. At this
    // point, execution will transition to the highest-priority ready thread.
    // This function will not return.

    Kernel::Start();

    // As Kernel::Start() results in the operating system being executed, control
    // will not be relinquished back to main(). The "return 0" is simply to
    // avoid warnings.

    return 0;
}

//-----
void AppMain(void* unused_)
{
    // This function is run from within the application thread. Here, we
    // simply print a friendly greeting and allow the thread to sleep for a
    // while before repeating the message. Note that while the thread is
    // sleeping, CPU execution will transition to the Idle thread.

    while (1) {
        KernelAware::Print("Hello World!\n");
        Thread::Sleep(1000);
    }
}

//-----
void IdleMain(void* unused_)
{
    while (1) {
        // Low priority task + power management routines go here.
        // The actions taken in this context must *not* cause the thread
        // to block, as the kernel requires that at least one thread is
        // schedulable at all times when not using an idle thread.

        // Note that if you have no special power-management code or idle
        // tasks, an empty while(1){} loop is sufficient to guarantee that
        // condition.
    }
}

```



```
// whenever there are no ready threads in the system. Note that if no
// Idle function is specified, a default will be used. Note that this default
// function is essentially a null operation.
Kernel::SetIdleFunc(IdleMain);

Kernel::Start();

return 0;
}

//-----
void AppMain(void* unused_)
{
    // Same as in lab1.
    while (1) {
        KernelAware::Print("Hello World!\n");
        Thread::Sleep(1000);
    }
}

//-----
void IdleMain(void)
{
    // Low priority task + power management routines go here.
    // The actions taken in this context must *not* cause a blocking call,
    // similar to the requirements for an idle thread.

    // Note that unlike an idle thread, the idle function must run to
    // completion. As this is also called from a nested interrupt context,
    // it's worthwhile keeping this function brief, limited to absolutely
    // necessary functionality, and with minimal stack use.
}
```

20.6 lab3_round_robin/main.cpp

This example demonstrates how to use round-robin thread scheduling with multiple threads of the same priority.

```

=====
/*
   |_____|_____|_____|_____|_____|_____|_____|_____|_____|_____|
   |  /  \  /  \  /  \  /  \  /  \  /  \  /  \  /  \  /  \  /  \  |
   |_____|_____|_____|_____|_____|_____|_____|_____|_____|_____|
   |  /  \  /  \  /  \  /  \  /  \  /  \  /  \  /  \  /  \  /  \  |
   |_____|_____|_____|_____|_____|_____|_____|_____|_____|_____|

--[Mark3 Realtime Platform]-----

Copyright (c) 2012 - 2017 Funkenstein Software Consulting, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"

/*=====

Lab Example 3:  using round-robin scheduling to time-slice the CPU.

Lessons covered in this example include:
- Threads at the same priority get timesliced automatically
- The Thread::SetQuantum() API can be used to set the maximum amount of CPU
  time a thread can take before being swapped for another task at that
  priority level.

Takeaway:

- CPU Scheduling can be achieved using not just strict Thread priority, but
  also with round-robin time-slicing between threads at the same priority.

=====*/
#if !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

//-----
// This block declares the thread data for one main application thread.  It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.

```

```

#define APP1_STACK_SIZE (320 / sizeof(K_WORD))
static Thread clApp1Thread;
static K_WORD awApp1Stack[APP1_STACK_SIZE];
static void App1Main(void* unused_);

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (320 / sizeof(K_WORD))
static Thread clApp2Thread;
static K_WORD awApp2Stack[APP2_STACK_SIZE];
static void App2Main(void* unused_);

//-----
int main(void)
{
    // See the annotations in lab1.
    Kernel::Init();

    // In this exercise, we create two threads at the same priority level.
    // As a result, the CPU will automatically swap between these threads
    // at runtime to ensure that each get a chance to execute.

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    // Set the threads up so that Thread 1 can get 4ms of CPU time uninterrupted,
    // but Thread 2 can get 8ms of CPU time uninterrupted. This means that
    // in an ideal situation, Thread 2 will get to do twice as much work as
    // Thread 1 - even though they share the same scheduling priority.

    // Note that if SetQuantum() isn't called on a thread, a default value
    // is set such that each thread gets equal timeslicing in the same
    // priority group by default. You can play around with these values and
    // observe how it affects the execution of both threads.

    clApp1Thread.SetQuantum(4);
    clApp2Thread.SetQuantum(8);

    clApp1Thread.Start();
    clApp2Thread.Start();

    Kernel::Start();

    return 0;
}

//-----
void App1Main(void* unused_)
{
    // Simple loop that increments a volatile counter to 1000000 then resets
    // it while printing a message.
    volatile uint32_t u32Counter = 0;
    while (1) {
        u32Counter++;
        if (u32Counter == 1000000) {
            u32Counter = 0;
            KernelAware::Print("Thread 1 - Did some work\n");
        }
    }
}

//-----
void App2Main(void* unused_)
{
    // Same as App1Main. However, as this thread gets twice as much CPU time
    // as Thread 1, you should see its message printed twice as often as the
    // above function.
    volatile uint32_t u32Counter = 0;
    while (1) {
        u32Counter++;
        if (u32Counter == 1000000) {
            u32Counter = 0;
            KernelAware::Print("Thread 2 - Did some work\n");
        }
    }
}

```

20.7 lab4_semaphores/main.cpp

This example demonstrates how to use semaphores for [Thread](#) synchronization.

Generated on Fri Jul 7 2017 01:40:39 for Mark3 Realtime Kernel by Doxygen


```

void __cxa_pure_virtual(void)
{
}

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (320 / sizeof(K_WORD))
static Thread clApp1Thread;
static K_WORD awApp1Stack[APP1_STACK_SIZE];
static void App1Main(void* unused_);

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (320 / sizeof(K_WORD))
static Thread clApp2Thread;
static K_WORD awApp2Stack[APP2_STACK_SIZE];
static void App2Main(void* unused_);

//-----
// This is the mutex that we'll use to synchronize two threads in this
// demo application.
static Mutex clMyMutex;

// This counter variable is the "shared resource" in the example, protected
// by the mutex. Only one thread should be given access to the counter at
// any time.
static volatile uint32_t u32Counter = 0;

//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    clApp1Thread.Start();
    clApp2Thread.Start();

    // Initialize the mutex used in this example.
    clMyMutex.Init();

    Kernel::Start();

    return 0;
}

//-----
void App1Main(void* unused_)
{
    while (1) {
        // Claim the mutex. This will prevent any other thread from claiming
        // this lock simulatenously. As a result, the other thread has to
        // wait until we're done before it can do its work. You will notice
        // that the Start/Done prints for the thread will come as a pair (i.e.
        // you won't see "Thread2: Start" then "Thread1: Start").

        clMyMutex.Claim();

        // Start our work (incrementing a counter). Notice that the Start and
        // Done prints wind up as a pair when simulated with flAVR.

        KernelAware::Print("Thread1: Start\n");
        u32Counter++;
        while (u32Counter <= 1000000) {
            u32Counter++;
        }
        u32Counter = 0;
        KernelAware::Print("Thread1: Done\n");

        // Release the lock, allowing the other thread to do its thing.
        clMyMutex.Release();
    }
}

//-----
void App2Main(void* unused_)
{
    while (1) {
        // Claim the mutex. This will prevent any other thread from claiming
        // this lock simulatenously. As a result, the other thread has to

```



```

Kernel::Start();

return 0;
}

//-----
void PeriodicCallback(Thread* owner, void* pvData_)
{
    // Timer callback function used to post a semaphore. Posting the semaphore
    // will wake up a thread that's pending on that semaphore.
    Semaphore* pclSem = (Semaphore*)pvData_;
    pclSem->Post();
}

//-----
void OneShotCallback(Thread* owner, void* pvData_)
{
    KernelAware::Print("One-shot timer expired.\n");
}

//-----
void ApplMain(void* unused_)
{
    Timer clMyTimer; // Periodic timer object
    Timer clOneShot; // One-shot timer object

    Semaphore clMySem; // Semaphore used to wake this thread

    // Initialize a binary semaphore (maximum value of one, initial value of
    // zero).
    clMySem.Init(0, 1);

    // Start a timer that triggers every 500ms that will call PeriodicCallback.
    // This timer simulates an external stimulus or event that would require
    // an action to be taken by this thread, but would be serviced by an
    // interrupt or other high-priority context.

    // PeriodicCallback will post the semaphore which wakes the thread
    // up to perform an action. Here that action consists of a trivial message
    // print.
    clMyTimer.Start(true, 500, PeriodicCallback, (void*)&clMySem);

    // Set up a one-shot timer to print a message after 2.5 seconds, asynchronously
    // from the execution of this thread.
    clOneShot.Start(false, 2500, OneShotCallback, 0);

    while (1) {
        // Wait until the semaphore is posted from the timer expiry
        clMySem.Pend();

        // Take some action after the timer posts the semaphore to wake this
        // thread.
        KernelAware::Print("Thread Triggered.\n");
    }
}

```

20.10 lab7_events/main.cpp

This example demonstrates how to create and use event groups

[illegible]

-Explore the behavior of the `EVENT_FLAG_ANY` and `EVENT_FLAG_ALL`, and the event-mask bitfield.

Takeaway:

Like Semaphores and Mutexes, `EventFlag` objects can be used to synchronize the execution of threads in a system. The `EventFlag` class allows for many threads to share the same object, blocking on different event combinations. This provides an efficient, robust way for threads to process asynchronous system events that occur with a unified interface.

```
=====*/
#if !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (320 / sizeof(K_WORD))
static Thread clApp1Thread;
static K_WORD awApp1Stack[APP1_STACK_SIZE];
static void App1Main(void* unused_);

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (320 / sizeof(K_WORD))
static Thread clApp2Thread;
static K_WORD awApp2Stack[APP2_STACK_SIZE];
static void App2Main(void* unused_);

//-----
//
static EventFlag clFlags;

//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    clApp1Thread.Start();
    clApp2Thread.Start();

    clFlags.Init();

    Kernel::Start();

    return 0;
}

//-----
void App1Main(void* unused_)
{
    while (1) {
        uint16_t ul6Flags;

        // Block this thread until any of the event flags have been set by
        // some outside force (here, we use Thread 2). As an exercise to the
        // user, try playing around with the event mask to see the effect it
        // has on which events get processed. Different threads can block on
        // different bitmasks - this allows events with different real-time
        // priorities to be handled in different threads, while still using
        // the same event-flag object.

        // Also note that EVENT_FLAG_ANY indicates that the thread will be
        // unblocked whenever any of the flags in the mask are selected. If
        // you wanted to trigger an action that only takes place once multiple
        // bits are set, you could block the thread waiting for a specific
        // event bitmask with EVENT_FLAG_ALL specified.
        ul6Flags = clFlags.Wait(0xFFFF, EVENT_FLAG_ANY);

        // Print a message indicating which bit was set this time.
        switch (ul6Flags) {
            case 0x0001: KernelAware::Print("Event1\n"); break;

```

```

        case 0x0002: KernelAware::Print("Event2\n"); break;
        case 0x0004: KernelAware::Print("Event3\n"); break;
        case 0x0008: KernelAware::Print("Event4\n"); break;
        case 0x0010: KernelAware::Print("Event5\n"); break;
        case 0x0020: KernelAware::Print("Event6\n"); break;
        case 0x0040: KernelAware::Print("Event7\n"); break;
        case 0x0080: KernelAware::Print("Event8\n"); break;
        case 0x0100: KernelAware::Print("Event9\n"); break;
        case 0x0200: KernelAware::Print("Event10\n"); break;
        case 0x0400: KernelAware::Print("Event11\n"); break;
        case 0x0800: KernelAware::Print("Event12\n"); break;
        case 0x1000: KernelAware::Print("Event13\n"); break;
        case 0x2000: KernelAware::Print("Event14\n"); break;
        case 0x4000: KernelAware::Print("Event15\n"); break;
        case 0x8000: KernelAware::Print("Event16\n"); break;
        default: break;
    }

    // Clear the event-flag that we just printed a message about. This
    // will allow u16 to acknowledge further events in that bit in the future.
    clFlags.Clear(u16Flags);
}

//-----
void App2Main(void* unused_)
{
    uint16_t u16Flag = 1;
    while (1) {
        Thread::Sleep(100);

        // Event flags essentially map events to bits in a bitmap. Here we
        // set one bit each 100ms. In this loop, we cycle through bits 0-15
        // repeatedly. Note that this will wake the other thread, which is
        // blocked, waiting for *any* of the flags in the bitmap to be set.
        clFlags.Set(u16Flag);

        // Bitshift the flag value to the left. This will be the flag we set
        // the next time this thread runs through its loop.
        if (u16Flag != 0x8000) {
            u16Flag <<= 1;
        } else {
            u16Flag = 1;
        }
    }
}

```

20.11 lab8_messages/main.cpp

This example demonstrates how to pass data between threads using message passing.

[illegible]

"data" members. This mechanism can be used to pass data between threads extremely efficiently, with a simple and flexible API. Any number of threads can write to/block on a single message queue, which give this method of IPC even more flexibility.

```

=====*/
#if !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (320 / sizeof(K_WORD))
static Thread c1App1Thread;
static K_WORD awApp1Stack[APP1_STACK_SIZE];
static void App1Main(void* unused_);

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (320 / sizeof(K_WORD))
static Thread c1App2Thread;
static K_WORD awApp2Stack[APP2_STACK_SIZE];
static void App2Main(void* unused_);

//-----
static MessageQueue c1MsgQ;

//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    c1App1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    c1App2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    c1App1Thread.Start();
    c1App2Thread.Start();

    c1MsgQ.Init();

    Kernel::Start();

    return 0;
}

//-----
void App1Main(void* unused_)
{
    uint16_t ul6Data = 0;
    while (1) {
        // This thread grabs a message from the global message pool, sets a
        // code-value and the message data pointer, then sends the message to
        // a message queue object. Another thread (Thread2) is blocked, waiting
        // for a message to arrive in the queue.

        // Get the message object
        Message* pclMsg = GlobalMessagePool::Pop();

        // Set the message object's data (contrived in this example)
        pclMsg->SetCode(0x1337);
        ul6Data++;
        pclMsg->SetData(&ul6Data);

        // Send the message to the shared message queue
        c1MsgQ.Send(pclMsg);

        // Wait before sending another message.
        Thread::Sleep(200);
    }
}

//-----
void App2Main(void* unused_)
{
    while (1) {
        // This thread waits until it receives a message on the shared global

```

```
// message queue. When it gets the message, it prints out information
// about the message's code and data, before returning the message object
// back to the global message pool. In a more practical application,
// the user would typically use the code to tell the receiving thread
// what kind of message was sent, and what type of data to expect in the
// data field.

// Wait for a message to arrive on the specified queue. Note that once
// this thread receives the message, it is "owned" by the thread, and
// must be returned back to its source message pool when it is no longer
// needed.
Message* pclMsg = clMsgQ.Receive();

// We received a message, now print out its information
KernelAware::Print("Received Message\n");
KernelAware::Trace(0, __LINE__, pclMsg->GetCode(), *((uint16_t*)pclMsg->
GetData()));

// Done with the message, return it back to the global message queue.
GlobalMessagePool::Push(pclMsg);
}
}
```

20.12 lab9_dynamic_threads/main.cpp

This example demonstrates how to create and destroy threads dynamically at runtime.

[illegible]

```

static void ApplMain(void* unused_);

//-----
// This block declares the thread stack data for a thread that we'll create
// dynamically.
#define APP2_STACK_SIZE (400 / sizeof(K_WORD))
static K_WORD awApp2Stack[APP2_STACK_SIZE];

#if KERNEL_USE_THREAD_CALLOUTS
#define MAX_THREADS (10)
static Thread* apclActiveThreads[10];
static uint32_t aul6ActiveTime[10];

static void PrintThreadSlack(void)
{
    KernelAware::Print("Stack Slack");
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] != 0) {
            char szStr[10];

            uint16_t ul6Slack = apclActiveThreads[i]->GetStackSlack();
            MemUtil::DecimalToHex((K_ADDR)apclActiveThreads[i], szStr);
            KernelAware::Print(szStr);
            KernelAware::Print(" ");
            MemUtil::DecimalToString(ul6Slack, szStr);
            KernelAware::Print(szStr);
            KernelAware::Print("\n");
        }
    }
}

static void PrintCPUUsage(void)
{
    KernelAware::Print("Cpu usage\n");
    for (int i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] != 0) {
            KernelAware::Trace(0, __LINE__, (K_ADDR)apclActiveThreads[i],
                aul6ActiveTime[i]);
        }
    }
}

static void ThreadCreateCallout(Thread* pclThread_)
{
    KernelAware::Print("TC\n");
    CS_ENTER();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == 0) {
            apclActiveThreads[i] = pclThread_;
            break;
        }
    }
    CS_EXIT();

    PrintThreadSlack();
    PrintCPUUsage();
}

static void ThreadExitCallout(Thread* pclThread_)
{
    KernelAware::Print("TX\n");
    CS_ENTER();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == pclThread_) {
            apclActiveThreads[i] = 0;
            aul6ActiveTime[i] = 0;
            break;
        }
    }
    CS_EXIT();

    PrintThreadSlack();
    PrintCPUUsage();
}

static void ThreadContextSwitchCallback(Thread* pclThread_)
{
    KernelAware::Print("CS\n");
    static uint16_t ul6LastTick = 0;
    uint16_t ul6Ticks = KernelTimer::Read();

    CS_ENTER();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == pclThread_) {
            aul6ActiveTime[i] += ul6Ticks - ul6LastTick;
            break;
        }
    }
}

```

```

    }
    CS_EXIT();

    u16LastTick = u16Ticks;
}

#endif

//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    Kernel::SetThreadCreateCallout(ThreadCreateCallout);
    Kernel::SetThreadExitCallout(ThreadExitCallout);
    Kernel::SetThreadContextSwitchCallout(ThreadContextSwitchCallback);
    ;

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp1Thread.Start();
    Kernel::Start();

    return 0;
}

//-----
static void WorkerMain1(void* arg_)
{
    Semaphore* pClSem = (Semaphore*)arg_;
    uint32_t u32Count = 0;

    // Do some work. Post a semaphore to notify the other thread that the
    // work has been completed.
    while (u32Count < 1000000) {
        u32Count++;
    }

    KernelAware::Print("Worker1 -- Done Work\n");
    pClSem->Post();

    // Work is completed, just spin now. Let another thread destroy u16.
    while (1) {
    }
}

//-----
static void WorkerMain2(void* arg_)
{
    uint32_t u32Count = 0;
    while (u32Count < 1000000) {
        u32Count++;
    }

    KernelAware::Print("Worker2 -- Done Work\n");

    // A dynamic thread can self-terminate as well:
    Scheduler::GetCurrentThread()->Exit();
}

//-----
void App1Main(void* unused_)
{
    Thread clMyThread;
    Semaphore clMySem;

    clMySem.Init(0, 1);
    while (1) {
        // Example 1 - create a worker thread at our current priority in order to
        // parallelize some work.
        clMyThread.Init(awApp2Stack, sizeof(awApp2Stack), 1, WorkerMain1, (void*)&clMySem);
        clMyThread.Start();

        // Do some work of our own in parallel, while the other thread works on its project.
        uint32_t u32Count = 0;
        while (u32Count < 100000) {
            u32Count++;
        }

        KernelAware::Print("Thread -- Done Work\n");

        PrintThreadSlack();

        // Wait for the other thread to finish its job.
        clMySem.Pend();

        // Once the thread has signalled u16, we can safely call "Exit" on the thread to
        // remove it from scheduling and recycle it later.
    }
}

```

```
clMyThread.Exit();

// Spin the thread up again to do something else in parallel. This time, the thread
// will run completely asynchronously to this thread.
clMyThread.Init(awApp2Stack, sizeof(awApp2Stack), 1, WorkerMain2, 0);
clMyThread.Start();

u32Count = 0;
while (u32Count < 1000000) {
    u32Count++;
}

KernelAware::Print("Thread -- Done Work\n");

// Check that we're sure the worker thread has terminated before we try running the
// test loop again.
while (clMyThread.GetState() != THREAD_STATE_EXIT) {
}

KernelAware::Print("  Test Done\n");
Thread::Sleep(1000);
PrintThreadSlack();
}
}
```


Index

- Add
 - Scheduler, [141](#)
- Claim
 - Mutex, [130](#)
- Close
 - Driver, [91](#)
- Control
 - Driver, [91](#)
- Driver, [90](#)
 - Close, [91](#)
 - Control, [91](#)
 - Init, [91](#)
 - Open, [91](#)
 - Read, [92](#)
 - Write, [92](#)
- EVENT_FLAG_ALL
 - kerneltypes.h, [244](#)
- EVENT_FLAG_ALL_CLEAR
 - kerneltypes.h, [244](#)
- EVENT_FLAG_ANY
 - kerneltypes.h, [244](#)
- EVENT_FLAG_ANY_CLEAR
 - kerneltypes.h, [244](#)
- EVENT_FLAG_MODES
 - kerneltypes.h, [244](#)
- EVENT_FLAG_PENDING_UNBLOCK
 - kerneltypes.h, [244](#)
- Exit
 - Thread, [149](#)
- Init
 - Driver, [91](#)
 - Kernel, [101](#)
 - Mailbox, [118](#)
 - Message, [125](#)
 - Mutex, [131](#)
 - Notify, [134](#)
 - Scheduler, [142](#)
 - Semaphore, [144](#)
 - Thread, [152](#)
 - Timer, [162](#)
- KA_COMMAND_EXIT_SIMULATOR
 - kernelaware.h, [236](#)
- KA_COMMAND_IDLE
 - kernelaware.h, [236](#)
- KA_COMMAND_PRINT
 - kernelaware.h, [236](#)
- KA_COMMAND_PROFILE_INIT
 - kernelaware.h, [236](#)
- KA_COMMAND_PROFILE_REPORT
 - kernelaware.h, [236](#)
- KA_COMMAND_PROFILE_START
 - kernelaware.h, [236](#)
- KA_COMMAND_PROFILE_STOP
 - kernelaware.h, [236](#)
- KA_COMMAND_TRACE_0
 - kernelaware.h, [236](#)
- KA_COMMAND_TRACE_1
 - kernelaware.h, [236](#)
- KA_COMMAND_TRACE_2
 - kernelaware.h, [236](#)
- Kernel, [99](#)
 - Init, [101](#)
 - Panic, [102](#)
 - Start, [104](#)
- kernelaware.h
 - KA_COMMAND_EXIT_SIMULATOR, [236](#)
 - KA_COMMAND_IDLE, [236](#)
 - KA_COMMAND_PRINT, [236](#)
 - KA_COMMAND_PROFILE_INIT, [236](#)
 - KA_COMMAND_PROFILE_REPORT, [236](#)
 - KA_COMMAND_PROFILE_START, [236](#)
 - KA_COMMAND_PROFILE_STOP, [236](#)
 - KA_COMMAND_TRACE_0, [236](#)
 - KA_COMMAND_TRACE_1, [236](#)
 - KA_COMMAND_TRACE_2, [236](#)
- kerneltypes.h
 - EVENT_FLAG_ALL, [244](#)
 - EVENT_FLAG_ALL_CLEAR, [244](#)
 - EVENT_FLAG_ANY, [244](#)
 - EVENT_FLAG_ANY_CLEAR, [244](#)
 - EVENT_FLAG_MODES, [244](#)
 - EVENT_FLAG_PENDING_UNBLOCK, [244](#)
- Mailbox, [116](#)
 - Init, [118](#)
 - Receive, [119](#)
 - Send, [120](#), [122](#)
- Message, [123](#)
 - Init, [125](#)
- Mutex, [129](#)
 - Claim, [130](#)
 - Init, [131](#)
 - Release, [131](#)
- Notify, [133](#)
 - Init, [134](#)

- Signal, [134](#)
- Wait, [134](#)
- Open
 - Driver, [91](#)
- Panic
 - Kernel, [102](#)
- Pend
 - Semaphore, [144](#), [145](#)
- Post
 - Semaphore, [145](#)
- Quantum, [138](#)
- Read
 - Driver, [92](#)
- Receive
 - Mailbox, [119](#)
- Release
 - Mutex, [131](#)
- Remove
 - Scheduler, [142](#)
- Schedule
 - Scheduler, [142](#)
- Scheduler, [140](#)
 - Add, [141](#)
 - Init, [142](#)
 - Remove, [142](#)
 - Schedule, [142](#)
- Semaphore, [143](#)
 - Init, [144](#)
 - Pend, [144](#), [145](#)
 - Post, [145](#)
- Send
 - Mailbox, [120](#), [122](#)
- Signal
 - Notify, [134](#)
- Sleep
 - Thread, [155](#)
- Start
 - Kernel, [104](#)
 - Thread, [155](#)
 - Timer, [164](#), [166](#)
- Stop
 - Thread, [156](#)
 - Timer, [166](#)
- Thread, [146](#)
 - Exit, [149](#)
 - Init, [152](#)
 - Sleep, [155](#)
 - Start, [155](#)
 - Stop, [156](#)
 - Yield, [156](#)
- Timer, [160](#)
 - Init, [162](#)
 - Start, [164](#), [166](#)
 - Stop, [166](#)
- Timer, [162](#)
- Wait
 - Notify, [134](#)
- Write
 - Driver, [92](#)
- Yield
 - Thread, [156](#)